

Pure Mathematics

Chapter 1

Number Theory

1.1 Natural Numbers

1.1.1 Peano Axioms

Axiom 1. *Closure under addition:*

For all $a, b \in \mathbb{N}$ we have $a + b \in \mathbb{N}$.

Axiom 2. *Closure under multiplication:*

For all $a, b \in \mathbb{N}$ we have $a \times b \in \mathbb{N}$.

Axiom 3. *Commutative Law for addition:*

For all $a, b \in \mathbb{N}$ we have $a + b = b + a$.

Axiom 4. *Associative Law for addition:*

For all $a, b, c \in \mathbb{N}$ we have $(a + b) + c = a + (b + c)$.

Axiom 5. *Commutative Law for multiplication:*

For all $a, b \in \mathbb{N}$ we have $a \times b = b \times a$.

Axiom 6. *Associative Law for multiplication:*

For all $a, b, c \in \mathbb{N}$ we have $(a \times b) \times c = a \times (b \times c)$.

Axiom 7. *Multiplicative Identity:*

There is a special element of \mathbb{N} , denoted by 1, which has the property that for all $n \in \mathbb{N}$, $n \times 1 = n$.

Axiom 8. *Additive cancellation:*

For all $a, b, c \in \mathbb{N}$ if $a + c = b + c$ then $a = b$.

Axiom 9. *Multiplicative cancellation:*

For all $a, b, c \in \mathbb{N}$ if $a \times c = b \times c$ then $a = b$.

Axiom 10. *Distributive Law:*

For all $a, b, c \in \mathbb{N}$, $a \times (b + c) = (a \times b) + (a \times c)$.

Axiom 11. *Definition of "less than":*

For all $a, b \in \mathbb{N}$, $a < b$ if and only if there is some $c \in \mathbb{N}$ s.t. $a + c = b$.

Axiom 12. *Trichotomous property:*

For all $a, b \in \mathbb{N}$ exactly one of the following is true: $a = b$, $a < b$, $b < a$.

Notation. We also write ab for $a \times b$.

1.1.1.1 Properties following from these axioms

Proposition 1. *If $a, b \in \mathbb{N}$ satisfy $a \times b = a$, then $b = 1$.*

Proof.

$$\begin{array}{lll}
 & a \times b = a = a \times 1 & \text{by Multiplicative Identity axiom} \\
 \Longleftrightarrow & b \times a = 1 \times a & \text{by Commutative Law for multiplication} \\
 \Longleftrightarrow & b = 1 & \text{by Multiplicative cancellation}
 \end{array}$$

□

Proposition 2. *If $a, b, c \in \mathbb{N}$ and $a < b$ then $a \times c < b \times c$.*

Proof.

$$\begin{array}{lll}
 a < b & \implies a + d = b \text{ for some } d \in \mathbb{N} & \text{by Definition of "less than"} \\
 \therefore b \times c & = (a + d) \times c = (a \times c) + (d \times c) & \text{by Distributive Law} \\
 \therefore a \times c & < (a \times c) + (d \times c) = b \times c & \text{by defn. "less than" and closure}
 \end{array}$$

□

Proposition 3. *1 is the least element of \mathbb{N} .*

Proof. Assume m is the least element of \mathbb{N} . Then, also $m < 1$. So, by Proposition 2,

$$m < 1 \implies m \times m < 1 \times m = m$$

But, closure of multiplication and $m \times m < m$ together contradict the assumption that m is the least element of \mathbb{N} .

Therefore m cannot be less than 1. Since we know that $1 \in \mathbb{N}$ and that the minimum element of \mathbb{N} , m , cannot be less than 1, it follows that 1 must be the minimum element of \mathbb{N} and $m = 1$. \square

1.1.2 Integers

1.1.3 Odd and Even Numbers

Definition. An *even* number, $n \in \mathbb{Z}$, is one that satisfies,

$$\exists m \in \mathbb{Z} \cdot n = 2m$$

Definition. An *odd* number, $n \in \mathbb{Z}$, is one that satisfies,

$$\exists m \in \mathbb{Z} \cdot n = 2m + 1$$

1.1.3.1 Consequences

Sum of even numbers, $m + n$:

$$\begin{aligned} m + n &= 2k + 2l \quad \text{where } k, l \in \mathbb{Z} && \text{by defn. of even no.s } m, n \\ &= 2(k + l) \\ &= 2q \quad \text{where } q \in \mathbb{Z} \end{aligned}$$

So $m + n$ is also even. However, if $m + n$ is even:

$$\begin{aligned} m + n &= 2k \quad \text{where } k \in \mathbb{Z} && \text{by defn. of even } m + n \\ k &= \frac{m}{2} + \frac{n}{2} \end{aligned}$$

So m and n are not necessarily even. A counterexample is

$$3 + 5 = 8 \iff \frac{3}{2} + \frac{5}{2} = 4$$

To summarize:

- $m, n \text{ even} \implies m + n \text{ even}$
- $m + n \text{ even} \implies m, n \text{ even}$ **Wrong!**

1.1.4 The Fundamental Theorem of Arithmetic

Definition of prime number: An integer that is only divided cleanly by itself and one. More formally, an integer, p , is prime if it is greater than 1 and,

$$\nexists! m, n \in \mathbb{Z} \cdot \frac{p}{m} = n \wedge (m \neq p \wedge m \neq 1)$$

Primality \implies Unique Prime Factorization:

“Any number either is prime or is measured by some prime number.”

Euclid, Elements Book VII, Proposition 32

So, if an integer n is not prime then,

$$\begin{aligned} \exists a, b \in \mathbb{Z} \cdot \frac{n}{a} = b \\ \iff n = ab \end{aligned}$$

Then, for a (the same applies to b),

$$\begin{aligned} \exists c, d \in \mathbb{Z}, c, d \notin \{1, a\} \cdot \frac{n}{a} = b \\ \iff n = cd \end{aligned}$$

We can continue to descend like this until we must eventually encounter one or more primes. Furthermore, if a number, n , has a prime factorization, $p_1 p_2$ then,

$$n = p_1 p_2 = p_3 p_4 \iff \frac{p_1}{p_3} = \frac{p_4}{p_2} = n$$

But $\frac{p_1}{p_3} = n$ contradicts the definition of primeness of p_1 . Therefore prime factorizations are unique.

Proof of existence

Proof. It must be shown that every integer greater than 1 is either prime or a product of primes. First, 2 is prime. Then, by strong induction, assume this is true for all numbers greater than 1 and less than n . If n is prime, there is nothing more to prove. Otherwise, there are integers a, b where $n = ab$, and $1 < a \leq b < n$. By the induction hypothesis, $a = p_1 p_2 \dots p_j$ and $b = q_1 q_2 \dots q_k$ are products of primes. But then $n = ab = p_1 p_2 \dots p_j q_1 q_2 \dots q_k$ is a product of primes. \square

Proof of uniqueness

Proof. Suppose, to the contrary, that there is an integer that has two distinct prime factorizations. Let n be the least such integer and write $n = p_1 p_2 \dots p_j = q_1 q_2 \dots q_k$, where each p_i and q_i is prime. (Note that j and k are both at least 2.) We see that p_1 divides $q_1 q_2 \dots q_k$, so p_1 divides some q_i by Euclid's lemma. Without loss of generality, say that p_1 divides q_1 . Since p_1 and q_1 are both prime, it follows that $p_1 = q_1$. Returning to our factorizations of n , we may cancel these two terms to conclude that $p_2 \dots p_j = q_2 \dots q_k$. We now have two distinct prime factorizations of some integer strictly smaller than n , which contradicts the minimality of n . \square

1.1.5 Modular Arithmetic

1.1.5.1 Greatest Common Divisor (also called Highest Common Factor)

Definition. The **Greatest Common Divisor (gcd)** of two integers – say a and b – is an integer d that satisfies,

$$d = z_1a + z_2b$$

for some $z_1, z_2 \in \mathbb{Z}$. This means that, whenever we are adding or subtracting multiples of the two numbers a and b , the result will always be a multiple of d and, therefore also, d is the smallest such result obtainable.

The greatest common divisor of 16 and 6 can be visualized as follows:

$\begin{array}{c} \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \\ \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \\ \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \cdot \end{array}$	$16 = 6 \times 2 + 4$ $6 = 4 \times 1 + 2$ $4 = 2 \times 2 + 0$
--	---

This implies the algorithm:

```

gcd(a, b) :
  if b == 0 then
    return a
  else
    return gcd(b, a mod b)
  end if

```

The greatest common divisor of a and b is the smallest difference of multiples of a and b . This is because – for any difference, d , of multiples of a and b – we have,

$$d = ma + nb \text{ for } m, n \in \mathbb{Z}$$

and, if $g = \gcd(a, b)$ then, by definition, g divides both a and b and, therefore, also divides d . So any such sum (or difference) of integer multiples of a and

b is a multiple of g.

Proposition 4. For non-zero integers a and b , if $a = bq + r$ where $q, r \in \mathbb{Z}$, then $\gcd(a, b) = \gcd(b, r) = \gcd(b, a \bmod b)$.

Proof. $(a \bmod b) = r = a - bq$. For any m s.t. $m \mid a$ and $m \mid b$ we must also have $m \mid (a - bq)$ so the set of divisors of a and b is a subset of the set of divisors of b and $r = (a \bmod b)$. Conversely, for any m s.t. $m \mid b$ and $m \mid r = (a \bmod b)$ we have that $m \mid (bq + r) = a$ so the set of divisors of b and $r = (a \bmod b)$ is a subset of the set of divisors of a and b . So the sets are equal proving that they must have the same maximum element - the greatest common divisor. \square

Proposition 5. If $d = \gcd(a, b)$ then there is no integer linear combination of a and b that equals any positive value less than d .

Proof. Assume $d = \gcd(a, b)$ and that $\exists e < d \in \mathbb{N}, m, n \in \mathbb{Z}$ s.t. $e = am + bn$. Then,

$$\begin{aligned} e = am + bn &= dz_1m + dz_2n = d(z_1m + z_2n) \quad \text{for } z_1, z_2 \in \mathbb{Z} \\ \iff z_1m + z_2n &= \frac{e}{d} \notin \mathbb{Z} \end{aligned}$$

where we know that $\frac{e}{d} \notin \mathbb{Z}$ because $e < d$. But the field properties of the integers ensures that the integers are closed under integer linear combinations so that $z_1m + z_2n \in \mathbb{Z}$. Therefore such an e does not exist. \square

Corollary 1. If $d = \gcd(a, b)$ then every integer linear combination of a and b is a multiple of d .

Proof. Proposition 5 showed that there is no integer linear combination of a and b less than d . Suppose that we have,

$$e > d \in \mathbb{N}, m, n \in \mathbb{Z} \text{ s.t. } e = am + bn$$

then, because d divides both a and b ,

$$e = adz_1 + bdz_2 = d(az_1 + bz_2), \quad z_1, z_2 \in \mathbb{Z}.$$

Now the closure of the integer field means that $z_3 = az_1 + bz_2 \in \mathbb{Z}$ so that,

$$e = z_3d \implies d \mid e.$$

\square

Proposition 6. *A number $x \in \mathbb{Z}_m$ has a multiplicative inverse if and only if $\gcd(x, m) = 1$.*

Proof. Assume x^{-1} is a multiplicative inverse for $x \in \mathbb{Z}_m$. Then,

$$x^{-1}x = 1 \iff x^{-1}x \equiv 1 \pmod{m} \iff x^{-1}x = am + 1, \quad a \in \mathbb{Z}.$$

This means that we must have $1 = am + bx$ for some $a, b \in \mathbb{Z}$. Now if we have $d = \gcd(x, m)$ then by Corollary 1 we must have $d \mid 1$. Therefore $d = 1$.

Clearly, also, if we have $\gcd(x, m) = 1$ then we also have $1 = am + bx$ for some $a, b \in \mathbb{Z}$ and by following the previous logic in reverse we obtain that $b = x^{-1}$ is the multiplicative inverse of $x \in \mathbb{Z}_m$. \square

1.1.5.2 Lowest Common Multiple

The lowest common multiple of two numbers is formed by the multiplication of all the prime factors that occur in the two numbers where repetitions of prime factors are important. That's to say, the lowest common multiple of 4 and 8 is not 2 (which is the highest common factor/greatest common divisor) but 8 because in 8, the factor 2 occurs three times (as 2^3) and it occurs twice in 4,

$$\text{lcm}(4, 8) = \text{lcm}(2 \times 2, 2 \times 2 \times 2) = 2 \times 2 \times 2.$$

The general formula for the lowest common multiple may be expressed in terms of the gcd as follows

$$d = \gcd(a, b) \implies \text{lcm}(a, b) = d \times (a/d) \times (b/d).$$

1.1.6 Some Proofs on the Integers

Proposition 7. *For any integer m , \sqrt{m} is rational iff m is a square, i.e. $m = a^2$ for some integer a .*

To begin with we show the easier direction of implication: $(m = a^2) \implies (\sqrt{m} \text{ is rational})$.

Proof. Assume $m, a, b \in \mathbb{Z}$.

$$\begin{aligned} m &= a^2 \\ \iff \sqrt{m} &= |a| \\ &= a/b \text{ for } b = 1 \text{ or } -1. \end{aligned} \quad \square$$

Now the other (harder) direction, $(\sqrt{m} \text{ is rational}) \implies (m = a^2)$.

Proof. Assume $m, a, b \in \mathbb{Z}$. $(\sqrt{m} \text{ is rational})$ can be formalized as:

$$\exists m, a, b \in \mathbb{Z} \cdot (\sqrt{m} = \frac{a}{b}) \wedge (a \text{ and } b \text{ are coprime})$$

$$\begin{aligned} \sqrt{m} &= \frac{a}{b} \\ \implies m &= \frac{a^2}{b^2} \\ \iff mb^2 &= a^2 \end{aligned}$$

But a and b are coprime so they don't share any prime factors. This means that a^2 and b^2 also don't share any prime factors. So, if $|b| > 1$, the prime factorization of mb^2 is necessarily different from that of a^2 meaning that $mb^2 \neq a^2$ contradicting the hypothesis of coprimality. On the other hand, if $|b| = 1$, then b has no prime factors (its prime factorization is empty) and so mb^2 has the same prime factorization as m which may be equal to that of a^2 in the case that $m = a^2$. \square

Proposition 8. *For all nonnegative integers $a > b$ the difference of squares $a^2 - b^2$ does not give a remainder of 2 when divided by 4.*

Beginner's attempt - try proof by contradiction:

$$\begin{aligned} a^2 - b^2 &= 4n + 2 \\ 2k &= 4n + 2 && \text{by } a^2 - b^2 \text{ even} \\ k &= 2n + 1 \implies k \text{ is some odd number.} \end{aligned}$$

So, proof by contradiction is our first instinct but doesn't seem to get us anywhere. Instead, proceed by cases:

Case a, b are even:

$$\begin{aligned} \exists k, l \in \mathbb{Z} \cdot a &= 2k, b = 2l \\ \implies a^2 - b^2 &= 4k^2 - 4l^2 \\ &= 4(k^2 - l^2) \\ &= 4m \text{ where } m \in \mathbb{Z} \end{aligned}$$

So 4 divides $a^2 - b^2$ with 0 remainder.

Case a, b are odd:

$$\begin{aligned} \exists k, l \in \mathbb{Z} \cdot a &= 2k + 1, b = 2l + 1 \\ \implies a^2 - b^2 &= (4k^2 + 4k + 1) - (4l^2 + 4l + 1) \\ &= 4(k^2 + k - l^2 - l) \\ &= 4m \text{ where } m \in \mathbb{Z} \end{aligned}$$

So, again, 4 divides $a^2 - b^2$ with 0 remainder.

Case a even, b odd:

$$\begin{aligned}
& \exists k, l \in \mathbb{Z} \cdot a = 2k, b = 2l + 1 \\
& \implies a^2 - b^2 = 4k^2 - (4l^2 + 4l + 1) \\
& \quad = 4(k^2 - l^2 - l) - 1 \\
& \quad = 4m + 3 \text{ where } m = k^2 - l^2 - l - 1 \in \mathbb{Z}
\end{aligned}$$

So, here, 4 divides $a^2 - b^2$ with 3 remainder. So the proposition is proven as we have proven all the possible cases.

There is also another approach given in the Cambridge University Discrete Mathematics lecture notes, TODO

1.1.7 Absolute Value

$$|x| \geq x, |y| \geq y \implies |x| + |y| \geq x + y$$

$$|x + y| = \begin{cases} |x| + |y| & x, y \geq 0 \\ -|x| + |y| & x < 0, y \geq 0 \\ ||x| - |y|| & x \geq 0, y < 0 \\ |-(|x| + |y|)| & x, y < 0 \end{cases} \iff \begin{cases} ||x| + |y|| & x, y \geq 0 \text{ or } x, y < 0 \\ ||x| - |y|| & x < 0, y \geq 0 \text{ or } x \geq 0, y < 0 \end{cases}$$

Clearly, $||x| + |y|| \geq ||x| - |y||$ so that,

$$|x + y| \leq ||x| + |y|| = |x| + |y|$$

and this is known as the "triangle inequality".

Proposition 9. $|x - y| \leq |x - z| + |y - z|$

Proof.

$$|x - y| = |(x - z) + (z - y)| \leq |x - z| + |z - y| = |x - z| + |y - z|$$

□

Proposition 10. $|x - y| \geq ||x| - |y||$

Proof. Need to show $-|x - y| \leq |x| - |y| \leq |x - y|$. So, prove as two separate inequalities:

$$\begin{aligned} & |y| = |x + (y - x)| \leq |x| + |y - x| \\ \iff & -|y - x| = -|x - y| \leq |x| - |y| \end{aligned}$$

$$\begin{aligned} & |x| = |(x - y) + y| \leq |x - y| + |y| \\ \iff & |x| - |y| \leq |x - y| \end{aligned}$$

□

1.1.8 Complex Number

Definition. The ***modulus*** of a complex number, $z = a + bi$, is the quantity defined as,

$$|z| = \sqrt{a^2 + b^2}.$$

TODO: modulus is also called "absolute value" and can be calculated as product with conjugate The modulus obeys the following properties:

- $|z_1 z_2| = |z_1| |z_2|$

Chapter 2

Algebra

2.1 Group Theory

2.1.1 Groups

Definition. *A binary operation is a function,*

$$f : G \times G \mapsto G$$

which - by the definition of a function - maps a unique tuple from $G \times G$ to a unique value in the codomain G .

Definition. Let G be a set and $*$ a binary operation on G and denote this $(G, *)$. Then $(G, *)$ is a **group** if:

closure $\forall x, y \in G, x * y \in G$;

associativity $\forall x, y, z \in G, (x * y) * z = x * (y * z)$;

identity $\exists e \in G$ s.t. $\forall x \in G, e * x = x * e = x$;

inverse $\forall x \in G, \exists x^{-1} \in G$ s.t. $x * x^{-1} = x^{-1} * x = e$.

These are known as the **group axioms**.

Definition. The group is an **Abelian group** if it has the additional property:

commutativity $\forall x, y \in G, x * y = y * x \in G$.

Notation. from here on we will use juxtaposition notation for the group operation (so $xy = x * y$) and (usually) 1 for the identity element instead of e . This is known as *multiplicative notation*.

Theorem 1. Suppose an associative law of composition is given on a set S . Then there is a unique way to define a product of n elements a_1, \dots, a_n for any $n \in \mathbb{N}$.

Proof. Denote the product of n elements as $[a_1 \dots a_n]$. We show that a product can be defined with the following properties:

- (i) $[a_1] = a_1$;
- (ii) $[a_1 a_2] = a_1 * a_2$ is defined by the law of composition;
- (iii) for any integer i such that $1 \leq i \leq n$, $[a_1 \dots a_n] = [a_1 \dots a_i][a_{i+1} \dots a_n]$.

Following a proof by induction, firstly note that the product is defined for $n \leq 2$ by (i) and (ii) and that (ii) also satisfies the requirement (iii). Then, assume that the product is defined for $n \leq 2$ and that this product is the unique product satisfying (iii).

Then the induction step is to show that,

$$[a_1 \dots a_n] = [a_1 \dots a_{n-1}][a_n]$$

[TODO: complete this from Artin\[56\]](#)

□

2.1.1.1 Corollaries of the group axioms

The group operation is defined to map a unique tuple in $G \times G$ to a unique value in G so that if we have $x, y \in G$ then $f((x, y)) = f(x, y) = xy \in G$ and for $a, b, c \in G$,

$$a = b \iff (c, a) = (c, b) \implies f((c, a)) = f((c, b)) \iff ca = cb$$

$$\therefore a = b \implies ca = cb$$

Then, using all the group axioms - associativity, inverse and identity,

$$ca = cb \implies c^{-1}(ca) = c^{-1}(cb) \iff (c^{-1}c)a = (c^{-1}c)b \iff 1a = 1b \iff a = b$$

Therefore we have the principle of cancellation,

$$ca = cb \implies a = b$$

Note that, since we have used the axioms of inverse and identity and the definitions of these require these elements to exhibit these properties from both the left and the right, the principle of cancellation can also be shown from both the left and the right. So, also,

$$ac = bc \implies a = b$$

There are (at least) two approaches to finding the other consequences of the group axioms.

First approach. We begin by noticing that the law of cancellation implies that,

unique identity and inverses $\forall a, x, b \in G, ax = b$ has a unique solution because,

$$ax = ax' \iff x = x'$$

That unique solution is $a^{-1}b$. If $b = a$ we have $ax = a$ and x , by identity axiom, is an identity element. Since, the solution to this equation - x - is unique, it follows that there is a unique value that is the identity element. Then, if we let b be this unique identity element we have $ax = 1$ and the unique solution, x , is the inverse of a , i.e. a^{-1} . Therefore, the inverses of group elements are also unique.

Second approach. This approach begins by showing the uniqueness of the identity element solely using the definition of the identity. Here, for clarity, we revert to using e to denote the identity element.

unique identity Assume there are two identity elements, e, e' . Then, by the definition of the identity $ee' = e'e = e = e'$ so that there is a single value that has the property of the identity element.

Then, using the definition of the inverse we have,

unique inverses Assume there are two distinct inverses of an element a : a^{-1} and a' . Then,

$$\begin{array}{ll} aa^{-1} = 1 = aa' & \text{defn. of inverse, uniqueness of identity} \\ \iff a^{-1} = a' & \text{law of cancellation} \end{array}$$

Some Examples of Groups

- $(\mathbb{R} \setminus \{0\}, \times)$ is a group whereas (\mathbb{R}, \times) is not a group because 0 has no multiplicative inverse.

- $(\mathbb{R}, +)$ is a group.
- The set of $n \times n$ invertible matrices is called the General Linear group and denoted GL_n
- Let G denote the set of matrices

$$G = \left\{ \begin{pmatrix} a & b \\ 0 & 1 \end{pmatrix} \mid a, b \in \mathbb{Z}_7, a \neq 0 \right\}.$$

Then G is a group with respect to matrix multiplication (where all additions and multiplications are carried out in \mathbb{Z}_7). This is because closure can be shown by,

$$\begin{pmatrix} a & b \\ 0 & 1 \end{pmatrix} \begin{pmatrix} a' & b' \\ 0 & 1 \end{pmatrix} = \begin{pmatrix} aa' & ab' + b \\ 0 & 1 \end{pmatrix} \in G$$

where the result is in G because $aa' \neq 0$. Next we need to show that we have inverses. So we need to show existence in G of matrices such that,

$$\begin{pmatrix} a & b \\ 0 & 1 \end{pmatrix} \begin{pmatrix} a' & b' \\ 0 & 1 \end{pmatrix} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} \iff \begin{pmatrix} aa' & ab' + b \\ 0 & 1 \end{pmatrix} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$

which implies that $aa' = 1$ and $ab' + b = 0$. Because we are in \mathbb{Z}_7 every non-zero element has a multiplicative inverse so we have,

$$a' = a^{-1} \quad \text{and} \quad b' = -a^{-1}b$$

so that,

$$\begin{pmatrix} a & b \\ 0 & 1 \end{pmatrix}^{-1} = \begin{pmatrix} a^{-1} & -a^{-1}b \\ 0 & 1 \end{pmatrix}.$$

2.1.1.2 Permutations and Symmetric Groups

Definition. A **permutation** is a bijection from a set to itself. Since permutations are bijective, they are invertible and since they are functions, function composition defines an associative law of composition over them. As a result, they form a group.

Definition. The ***symmetric group*** defined over a set is the group whose elements are the permutations of the objects of the set and whose law of composition is the composition of functions. The name probably comes from the study of symmetries of geometric objects that were eventually realised to be equivalent to permutations of the vertices.

Definition. A ***generating set of a group*** is a subset such that every element of the group can be expressed as a combination (under the group operation) of finitely many elements of the subset and their inverses.

Notation. The symmetric group over the integers from 1 to n is denoted S_n . The symmetric group over a set G may be denoted $Sym(G)$.

S_2 The symmetric group S_2 consists of the two elements i and τ which are, respectively, the identity map and the transposition which interchanges 1 and 2. The group composition law is described by the fact that the identity map is the identity of the composition and by the relation $\tau\tau = \tau^2 = i$. Which results in the multiplication table:

$$\begin{aligned} i \cdot i &= i \\ i \cdot \tau &= \tau \\ \tau \cdot i &= \tau \\ \tau \cdot \tau &= i \end{aligned}$$

Note that the law of composition is commutative.

S_3 The symmetric group S_3 contains $3!$ elements. It is the smallest group whose law of composition is not commutative. It can be described using any two permutations of $\{1, 2, 3\}$. For example, if we take,

$$x = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \end{bmatrix}, y = \begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Then the permutations are,

$$\{1, x, x^2, y, xy, x^2y\} = \{x^i y^j \mid 0 \leq i \leq 2, 0 \leq j \leq 1\}$$

These are the elements of the group. The composition law over these elements is the function composition of these permutation functions and its multiplication table is characterized by the rules:

$$x^3 = 1, y^2 = 1, yx = x^2y$$

These are derived directly from the permutations themselves. Note that this composition law is not associative as $yx \neq xy$.

Any product of the elements x, y and of their inverses can be brought into the form $x^i y^j$ with i, j taking the ranges given above by repeated application of the above rules. To do so, we move all occurrences of y to the right side using the last relation and bring the exponents into the indicated ranges using the first two relations:

$$\begin{aligned} x^{-1}y^3x^2y &= x^2yx^2y = x^2(yx)xy = x^2(x^2y)xy = x^4(yx)y \\ &= x^4(x^2y)y = x^6y^2 = (x^3)^2y^2 = 1 \cdot 1 = 1 \end{aligned}$$

Rules like these that determine a complete multiplication table are called *defining relations* for the group.

2.1.2 Subgroups

Definition. A subset H of a group G is called a **subgroup** if it has the following properties:

- **Closure:** If $a \in H$ and $b \in H$ then $ab \in H$.
- **Identity:** $1 \in H$.
- **Inverses:** If $a \in H$ then $a^{-1} \in H$.

These conditions show that the subset H is a group with respect to the induced law of composition created by applying the law of composition of G on the members of H . Note that the associative property is not mentioned because the associativity of the composition of members of G automatically carries over to H .

Notation. If H and G are groups then we may write $H \leq G$ to indicate that H is a subgroup of G .

Note that an alternative, more compact, formulation of the definition of a subgroup is as follows.

Let G be a group and $\emptyset \neq H \subseteq G$. Then H is a subgroup if

$$x, y \in H \implies x^{-1}y \in H.$$

This is because,

$$[(x, y \in H \implies xy \in H) \wedge (x \in H \implies x^{-1} \in H)] \iff (x, y \in H \implies x^{-1}y \in H).$$

The implication,

$$[(x, y \in H \implies xy \in H) \wedge (x \in H \implies x^{-1} \in H)] \implies (x, y \in H \implies x^{-1}y \in H)$$

is obvious. In the other direction,

$$(x, y \in H \implies x^{-1}y \in H) \implies [(x, y \in H \implies xy \in H) \wedge (x \in H \implies x^{-1} \in H)]$$

is because, if we set $x = y$,

$$x^{-1}x = e \in H \implies x^{-1}e = x^{-1} \in H$$

and then, for $x \neq y$,

$$x^{-1}, y \in H \implies xy \in H.$$

Every group, at a minimum, has two trivial subgroups: the maximal subgroup - the group itself; and the minimal subgroup - the set containing just the identity. A subgroup that is neither of these is known as a *proper subgroup*.

Proposition 11. *Suppose H and K are subgroups of G such that neither $H \subseteq K$ nor $K \subseteq H$. Then $H \cup K$ is not a subgroup of G .*

*Note that it might be possible, for this proof, to just show that we have no reason to think that $H \cup K$ is a subgroup (i.e. the definitions don't require it) so, in general, it is not. But, actually, we can show something much stronger: that $H \cup K$ **cannot** be a subgroup. If we can easily show something stronger then in most cases it's going to add clarity.*

Proof. Since neither $H \subseteq K$ nor $K \subseteq H$ we can conclude that $H \setminus K$ and $K \setminus H$ are both non-empty. So, select an element from each,

$$h \in H \setminus K, k \in K \setminus H.$$

Then we have $h, k \in H \cup K$ and if $H \cup K$ were a group then the closure property of the group would require that

$$hk \in H \cup K.$$

Assume $hk \in H \cup K$. Then, $hk \in H$ or $hk \in K$. If $hk \in H$ then the group properties of H require that

$$h^{-1}hk = k \in H$$

which contradicts the selection of k . We have a similar situation if $hk \in K$. Therefore, $hk \notin H \cup K$. \square

2.1.2.1 Additive Groups of Integers

Important examples are the subgroups of the additive group of integers \mathbb{Z}^+ . Denote the subset of \mathbb{Z}^+ consisting of all multiples of a given integer b by $b\mathbb{Z}$ such that,

$$b\mathbb{Z} = \{ n \in \mathbb{Z} \mid n = bk, k \in \mathbb{Z} \}$$

Proposition 12. *For any integer b , the subset $b\mathbb{Z}$ is a subgroup of \mathbb{Z}^+ and every subgroup of \mathbb{Z}^+ is of the form $b\mathbb{Z}$ for some integer b .*

Proof. $b\mathbb{Z}$ is a subgroup of \mathbb{Z}^+ because,

- $b(0) = 0 \in b\mathbb{Z}$;
- If $a_1, a_2 \in b\mathbb{Z}$ then $a_1 = bk_1, a_2 = bk_2$ for $k_1, k_2 \in \mathbb{Z}$ and so $a_1 + a_2 = bk_1 + bk_2 = b(k_1 + k_2) \in b\mathbb{Z}$
- For any $a = bk \in b\mathbb{Z}$, $-a = b(-k) \in b\mathbb{Z}$

Now we need to prove that any subgroup of \mathbb{Z}^+ is $b\mathbb{Z}$ for some b . Let H be an arbitrary subgroup of \mathbb{Z}^+ . Then by subgroup properties,

- $0 \in H$;
- If $a_1, a_2 \in H$ then $a_1 + a_2 \in H$
- For any $a \in H$, $-a \in H$

We proceed to show that there is always some integer b such that $H = b\mathbb{Z}$. Firstly, if H is the minimal subgroup $\{0\}$ then H trivially conforms to $b\mathbb{Z}$ with $b = 0$.

Otherwise, $\exists a \in H$ s.t. $a \neq 0$ then also $\exists -a \in H$ s.t. $-a \neq 0$. One of these must be a positive non-zero integer so there is at least one such member of H . We take b to be the smallest positive non-zero integer in H . Then,

$b\mathbb{Z} \in H$

- $b \in H$ (by selection) so by subgroup properties $b+b \in H$ and $(b+b)+b \in H$ and $b + \dots + b \in H$
- By subgroup properties $b \in H \implies -b \in H$

So, $\{bk \in \mathbb{Z} \mid k \in \mathbb{Z}\}$ is in H .

$H \in b\mathbb{Z}$ Take any $n \in H$. Using division with remainder and dividing by b we get,

$$n = bq + r \quad q \in \mathbb{Z}, 0 \leq r < b$$

But, since $b\mathbb{Z} \in H$ this means that $bq \in H$ and so $-bq \in H$. Therefore $n - bq = r \in H$. But $0 \leq r < b$ and, by assumption, b is the smallest positive *non-zero* integer in H and so, $r = 0$. So, every $n \in H$ divides by b . \square

2.1.2.2 Greatest Common Divisor

If we extend this to groups which are generated by two integers a, b , then we have a subgroup of \mathbb{Z}^+ ,

$$a\mathbb{Z} + b\mathbb{Z} = \{ n \in \mathbb{Z} \mid n = ar + bs \quad r, s \in \mathbb{Z} \}$$

This is known as the subgroup *generated* by a, b because it is the smallest subgroup which contains a and b . Proposition 12 tells us that it has the form $d\mathbb{Z}$ for some integer d .

Corollary 2. *If d is the positive integer which generates the subgroup $a\mathbb{Z} + b\mathbb{Z}$ then d is the greatest common divisor of a and b and so,*

- *d can be written in the form $d = ar + bs$ for some integers r and s .*
- *d divides a and b .*
- *If an integer e divides a and b , it also divides d .*

Proof. The first property follows directly from the definition of the subgroup. The second property is a result of the fact that a, b are in the subgroup $a\mathbb{Z} + b\mathbb{Z}$ so that $d\mathbb{Z} = a$ and $d\mathbb{Z} = b$. The third property is evident because $d = ar + bs = ek_1r + ek_2s = e(k_1r + k_2s)$. \square

2.1.2.3 Deductions about Subgroups

Proposition 13. *Suppose $G = \{g_1, g_2, \dots, g_n\}$ is a finite group of order n and that $x \in G$. Then $\{xg_1, xg_2, \dots, xg_n\} = G$.*

Proof. Let $X = \{xg_1, xg_2, \dots, xg_n\}$. By closure in G , every element of X must be in G and by the inverses property in G every element of X is distinct. So, there are n distinct elements of X , each of which are members of G . Since the order of G is n we can conclude that $X = G$. \square

Proposition 14. *Suppose that G is a finite group and that H is a non-empty subset of G such that $x, y \in H \implies xy \in H$. Then H is a subgroup.*

Proof. H is non-empty so it contains at least one element, say x . H is closed under the group operation so it must also contain x^2, x^3, \dots . But G is finite so the order of x in G must be finite also and so $\exists n \in \mathbb{N}$ s.t. $x^n = e$. But also $x^n = e \iff x^{n-1} = x^{-1}$. Therefore, for every element in H , the inverse of the element is also in H . \square

Proposition 15. *Suppose that p is a prime number and we have integers $1 \leq x, g, h < p$ such that $xg \equiv xh \pmod{p}$. Then $g = h$ and $x \in \mathbb{Z}_p^*$ has an inverse.*

Proof. If $xg \equiv xh \pmod{p}$ then the difference between xg and xh is a multiple of p . But Euclid's Lemma (https://en.wikipedia.org/wiki/Euclid's_lemma) tells us that, because p is prime,

$$p \mid (xg - xh) = x(g - h) \implies (p \mid x) \wedge (p \mid (g - h)).$$

But since we have $x, (g - h) < p$ it is impossible for p to divide either of them unless they are 0. Only $g - h$ can be 0. Therefore,

$$g - h = 0 \iff g = h.$$

So, if we define a function $f : \mathbb{Z}_p^* \longrightarrow \mathbb{Z}_p^*$ such that $f(a) = xa$ for some fixed $x \in \mathbb{Z}_p^*$ then f is injective because

$$f(a) = f(b) \iff xa \equiv xb \pmod{p} \iff a \equiv b \pmod{p}.$$

This means that f maps the $p - 1$ different values of \mathbb{Z}_p^* to $p - 1$ different values in \mathbb{Z}_p^* . Therefore f is a bijection and there exists an inverse function f^{-1} such that $f^{-1}(xa) = a$. \square

2.1.2.4 Examples of Subgroups

- (1) $(\mathbb{Z}_p^*, \otimes)$ for prime p is a subgroup of the integers. This can be seen as closure of modular multiplication is clear and the existence of inverses has been shown in Proposition 15.

This is not the case however, for non-prime p . For example $(\mathbb{Z}_6^*, \otimes)$ is not a subgroup as it does not have inverses. We can see this by looking at the values generated by selecting a non-identity element and multiplying it by all the elements in \mathbb{Z}_6^* :

$$2 \otimes 1 = \mathbf{2}, 2 \otimes 2 = \mathbf{4}, 2 \otimes 3 = \mathbf{0}, 2 \otimes 4 = \mathbf{2}, 2 \otimes 5 = \mathbf{4}.$$

As can be seen, it doesn't generate all the values of \mathbb{Z}_6^* but repeats a subset of them. Compare with the same for \mathbb{Z}_7^* :

$$2 \otimes 1 = \mathbf{2}, 2 \otimes 2 = \mathbf{4}, 2 \otimes 3 = \mathbf{6}, 2 \otimes 4 = \mathbf{1}, 2 \otimes 5 = \mathbf{3}, 2 \otimes 6 = \mathbf{5}.$$


In the case of prime p all the values are generated so that multiplication by other elements is a bijective function with a corresponding inverse.

- (2) Let R be a non-square rectangle in \mathbb{R}^2 with corners having coordinates $(-1, -1), (-1, 2), (1, 2), (1, -1)$. Then there are four symmetries i, a, b, c of R , as follows:

- i is the identity
- a is reflection in the x -axis
- b is reflection in the y -axis
- c is a rotation of π radians around the origin.

These symmetry operations form a group whose group table is as follows.

	i	a	b	c
i	i	a	b	c
a	a	i	c	b
b	b	c	i	a
c	c	b	a	i



2.1.2.5 Cyclic Subgroups

Definition. If we take a single member of a group (along with its inverse and the identity), the subgroup generated by that element takes the form (using multiplicative notation),

$$H = \{x^{-(n-1)}, \dots, x^{-2}, x^{-1}, 1, x, x^2, \dots, x^{n-1}\}$$

where, either, $x^n = 1$ so that there are n distinct values in the group, or else n is infinite and the values never repeat. This is known as a **cyclic group** and also as the **subgroup generated by x** and is denoted by $\langle x \rangle$.

The cyclic subgroup, $\langle x \rangle$, generated by x is the smallest subgroup of G containing x in the sense that, if $H \leq G$ and $x \in H$ then $\langle x \rangle \subseteq H$.

Proposition 16. Every cyclic group is Abelian.

Proof. In a cyclic group every element has the form x^i for $i \in \mathbb{Z}$. So we have,

$$x^m x^n = x^{m+n} = x^n x^m$$

for all elements x^i in the group. □

Proposition 17. The set S of integers n such that $x^n = 1$ is a subgroup of \mathbb{Z}^+ .

Proof. If $x^m = 1$ and $x^n = 1$, then $x^{m+n} = x^m x^n = 1$ also so we have closure of addition. Since $x^0 = 1$, 0 is in the subgroup so we have an identity. Finally, for some n in the subgroup, $x^n = 1 \iff x^{-n} = x^n x^{-n} = x^0 = 1$ so n being in the subgroup implies that $-n$ is also in the subgroup and we have inverses. \square

Corollary 3. *It follows from S being a subgroup of \mathbb{Z}^+ and from Proposition 12 that S has the form $m\mathbb{Z}$ where m is the smallest positive integer such that $x^m = 1$. Therefore, in H , the m elements $1, x, x^2, \dots, x^{m-1}$ are all different and any element in H will simplify to one of them: for $n \in S$, $n = mq + r$ such that $x^n = (x^m)^q x^r = 1^q x^r = x^r$.*

2.1.2.6 Order

Definition. The **order** of a group G is the number of distinct elements it contains. It is typically denoted $|G|$.

An element of a group is said to have **order** m (possibly infinity) if the cyclic subgroup it generates has order m . This means that m is the smallest positive integer with the property $x^m = 1$ or, if the order is infinite, that, $x^m \neq 1$ for all $m \neq 0$.

Theorem 2. *An element and its inverse have the same order.*

Proof. Firstly we need to consider the case that an element x has infinite order. In this case, $\nexists m \in \mathbb{N}$ s.t. $x^m = e$. Now suppose that $\exists n \in \mathbb{N}$ s.t. $(x^{-1})^n = e$. Then we have,

$$(x^{-1})^n = (x^n)^{-1} = e \iff e = x^n$$

which contradicts the hypothesis that x has infinite order. Therefore x^{-1} has infinite order also. Clearly also this argument can be used in reverse to show the reverse implication also holds.

Now consider the case that x has finite order. Let $x \in G$ be an arbitrary member of an arbitrary group such that $x^m = e$. Then $x^{-1} = x^{m-1}$ and if we consider powers $i \in \mathbb{N}$ of the inverse $(x^{-1})^i = (x^{m-1})^i$ then the order is the lowest value of $i(m-1)$ such that $x^{i(m-1)} = e$. But we know that the lowest

power of x equal to e is m so we're looking for the lowest multiple of m that has the form $i(m-1)$. So we require,

$$m \mid i(m-1) = im - i \iff (m \mid im) \wedge (m \mid i)$$

which clearly requires that $m \mid i$. Also, clearly, the lowest such i is $i = m$.

Another way to show this is to say that if $x^m = e$ then,

$$(x^{-1})^m = (x^m)^{-1} = e$$

so that the order of x^{-1} is less than or equal to m . Conversely, if x^{-1} has order n then,

$$x^n = ((x^{-1})^{-1})^n = ((x^{-1})^n)^{-1} = e^{-1} = e$$

so that the order of x is less than or equal to n . Thus we have,

$$m \leq n, n \leq m \implies m = n.$$

□

Theorem 3. *An element has order 2 iff it is equal to its inverse.*

Proof. Let $x \in G$ be an arbitrary member of an arbitrary group such that $x^2 = e$. Then by Theorem 2 we have,

$$e = x^2 = (x^{-1})^2 = x^{-2} \iff x = x^{-1}.$$

Also,

$$x = x^{-1} \iff x^2 = e.$$

□

Theorem 4. *A group of finite order cannot have any element of infinite order.*

Proof. If G is a group and $x \in G$ has infinite order then,

$$\begin{aligned} & x^m = x^n \\ \iff & x^{m-n} = 1 = x^{n-m} \\ \iff & x^{|m-n|} = 1 \\ \therefore & |m-n| = 0. \end{aligned} \quad \text{because order of } x \text{ is infinite}$$

So, there are no two distinct powers of x that produce the same object so that $\langle x \rangle \leq G$, the cyclic group generated by x , is infinite. Since $\langle x \rangle \subseteq G$ this requires that G also be infinite. \square

Theorem 5. *If a group element x has finite order m then:*

1. *Let $n \in \mathbb{Z}$. If $n = km + r$ where $k, r \in \mathbb{Z}$ and $0 \leq r \leq m - 1$, then $x^n = x^r$.*
2. *For $n \in \mathbb{N}$, $x^n = 1 \iff m|n$.*
3. *$1, x, x^2, \dots, x^{m-1}$ is a complete, repetition-free, list of elements of $\langle x \rangle$.*
4. *The subgroup $\langle x \rangle$ generated by x has cardinality m .*

Theorem 6. *In an Abelian group the set of all elements of finite order forms a subgroup.*

Proof. Let S be the set of all elements of finite order,

$$S = \{ g \in G \mid \exists m \in \mathbb{N} . g^m = e \}.$$

- Firstly, S is non-empty because it contains the identity.
- Secondly, if we have $a, b \in S$ then $a^i = b^j = e$ for some $i, j \in \mathbb{N}$. Then, if we let $m = i \times j$,

$$(ab)^m = a^m b^m = (a^i)^j (b^j)^i = e$$

where $(ab)^m = a^m b^m$ is valid *only* because the group is Abelian.

- Lastly, S contains all inverses because,

$$a^i = e \iff a^{i-1} = a^{-1} \iff (a^{-1})^i = (a^i)^{i-1} = e.$$

Therefore $a^{-1} \in S$.

\square

Theorem 7. *If every non-identity element of a group has order 2 then the group is Abelian.*

Proof. Let $x, y \in G$ be two arbitrary elements of order 2 of an arbitrary group. Then by Theorem 2 we have $x = x^{-1}$, $y = y^{-1}$, $xy = xy^{-1}$ and so,

$$xy = (xy)^{-1} = y^{-1}x^{-1} = yx.$$

Note that $(xy)^{-1} = y^{-1}x^{-1}$ relies only on the associativity of the group operation and is therefore valid for all groups.

We can also show it this way,

$$(xy)^2 = e \iff xyxy = e \iff yxy = xe = x \iff yx = xy.$$

□

Theorem 8. *If a finite group has even-numbered order then it must have at least one element of order 2.*

Proof. By the group properties we know that the group contains the identity element – which has order 1 – and, for every non-identity element, the group also contains its inverse. Also, since the group is finite, every element must have finite order. Now, if every non-identity element is distinct from its inverse then the order of the group will be odd (because of the identity and then every other element is paired with its inverse). For the group's order to be even we must have at least one non-identity element that is not distinct from its inverse which, by Theorem 3, is equivalent to having order 2. □

Corollary 4. *If a finite group has even-numbered order then it must have an odd number of elements of order 2.*

Proof. Let G be a finite group with even-numbered order and M be the number of elements that are distinct from their inverse and N be the number of elements that are not distinct from their inverse (these correspond to elements with order greater than 2 and elements with order 2 respectively). Then the order of G can be expressed as,

$$|G| = 1 + 2M + N.$$

Therefore, $|G|$ is even if N is an odd natural number. □

Theorem 9. *In an infinite cyclic group all elements have infinite order.*

Proof. Let $G = \langle x \rangle$ be an infinite cyclic group. Suppose there is some non-identity element of G , x^n with finite order m . Then,

$$(x^n)^m = e \iff x^{nm} = e$$

which contradicts the hypothesis that $\langle x \rangle$ is infinite. □

Note that the elements of an infinite cyclic group having infinite order does not mean that they generate the group. For example in an infinite cyclic group $\langle x \rangle$, the element x^2 generates the cyclic subgroup,

$$\dots x^{-4}, x^{-2}, e, x^2, x^4, \dots$$

which is infinite but clearly doesn't generate the whole group $\langle x \rangle$.

Theorem 10. *An infinite cyclic group has 2 generators.*

Proof. Let $G = \langle x \rangle$ be an infinite cyclic group and suppose there is some non-identity element of G , x^n that generates the group. To show this we only need to show that x^n can generate x because, since x is a member of the group, it is obviously necessary to generate it but, also, if we generate x then we can generate all the other members of the group since they are powers of x .

So let there be an integer a such that $(x^n)^a = x^{an} = x \iff x^n = x^{1/a}$. The cyclic group $\langle x \rangle$ only contains integer powers of x so it therefore follows that $|a| = 1$ which implies that $n = 1$ or -1 and $x^n = x$ or x^{-1} .

We could also say that,

$$x^{an} = x \iff x^{an-1} = e$$

but this implies that the order of x is finite and so contradicts the hypothesis that x generates an infinite cyclic group.

□

Theorem 11. *Let $G = \langle x \rangle$ be a finite cyclic group of order n . If r is a positive integer then $G = \langle x^r \rangle$ if and only if the greatest common divisor of n and r is 1.*

Proof. Members of $\langle x^r \rangle$ have the form $(x^r)^a$ for some $a \in \mathbb{Z}$. For integers b, i ,

$$(x^r)^a = x^{ar} = x^{bn+i} = x^{bn}x^i = ex^i = x^i$$

so that the generated elements are x^i where $i = ar - bn$ is the remainder when dividing ar by n . If $d = \gcd(n, r)$ then $d \mid i$ and the generated elements are powers of x that are multiples of d . Therefore, to generate every power

of x it is necessary to have $d = 1$. Conversely, we can see – by the same argument in reverse – that it is sufficient if $d = 1$ to generate all the powers of x .

Alternatively, we can say if $d > 1$ then n/d is a positive integer less than n and r/d is a positive integer so,

$$(x^r)^{n/d} = (x^n)^{r/d} = e^{r/d} = e$$

which shows that the order of x^r is less than or equal to n/d which is less than n . Therefore the order of the cyclic group it generates is less than n and so it cannot be equal to G .

Conversely, if $d = 1$ then n and r are coprime and so we have,

$$(x^r)^m = e \iff x^{rm} = e \implies n \mid rm \implies n \mid m \implies m \geq n.$$

This says that the order of x^r in G is greater than or equal to n , the order of G . Well, clearly it cannot be greater than the order of G so it follows therefore, that the order of x^n is n . Since $|\langle x^r \rangle| = |G|$ we can conclude that $\langle x^r \rangle = G$. \square

Theorem 12. *A group G is such that G contains at least 2 elements and the only subgroups of G are $\{e\}$ and G itself. Then G is a finite cyclic group of prime order.*

Proof. G contains at least 2 elements so there is at least one non-identity element x . The only subgroups of G are the whole group and $\{e\}$ but $\langle x \rangle$ cannot equal $\{e\}$ so it must equal G . Therefore G is the cyclic group generated by x .

But if G were the infinite cyclic group generated by x then only x and x^{-1} would generate the group and all other non-identity elements – say x^n for $n > 1 \in \mathbb{N}$ – would generate subgroups $\langle x^n \rangle \neq G$. Therefore, G cannot be infinite and is therefore finite.

Now, we have a finite cyclic group where every non-identity element generates the group. Let $|G| = n$. Then, for every m s.t. $0 < m < n$, $\langle x^m \rangle = G$ and, by Theorem 11, m and n are coprime. Therefore, n is prime.

Here we could also use a proof by contradiction: Assume n is not prime and it has factors $r, s > 1 \in \mathbb{N}$. Then,

$$(x^r)^s = x^{rs} = x^n = e$$

so that the order of x^r in G is less than or equal to s which is less than n (because it is a factor of n). It follows then that $\langle x^r \rangle \neq G$, contradicting the definition of G . Therefore n is prime. \square

Proposition 18. *Suppose the elements x, y in a group G have orders m, n respectively and that the $\gcd(m, n) = 1$. Then $\langle x \rangle \cap \langle y \rangle = \{e\}$ and, if x and y commute, then the order of xy in G is mn .*

Proof. One way to approach this is to say that for $z \in \langle x \rangle \cap \langle y \rangle$ we have some $0 \leq i < m, 0 \leq j < n$ such that,

$$z = x^i = y^j \iff x^{im} = e = y^{jm}, x^{in} = y^{jn} = e$$

so that $x^{in} = y^{jm} = e \iff (m \mid in \text{ and } n \mid jm)$. Note that,

$$(m \mid in \text{ and } n \mid jm) \iff m, n \mid in + jm.$$

Now, applying the fact that $\gcd(m, n) = 1$ we see that both m and n must divide 1. But both m and n are orders of elements and so, by definition, greater than 1. The only other alternative is that both $i, j = 0$ which results in $z = x^0 = y^0 = e$.

We could also have said,

$$x^{im} = e = x^{in} \iff x^{im-in} = e \iff m \mid im - in.$$

In this case we can apply the fact that $\gcd(m, n) = 1$ to the statement that $m \mid i(m - n)$ to deduce that: either $m \mid (m - n) \iff m \mid 1$ which is impossible because m must be greater than 1; or $m \mid i$ which is also impossible because $i < m$. So, again, we are only left with the alternative that $i = 0$ which results in $z = x^0 = y^0 = e$.

Next we prove the order of $xy \in G$ and we begin by noting that, if x and y commute, then $(xy)^r = x^r y^r$. So if we assume that xy has order r then we must have $m, n \mid r$ and the lowest such r is the order. Well, the lowest common multiple of m, n is defined according to the gcd as described in the Number Theory treatment of Modular Arithmetic (1.1.5.2) as,

$$d = \gcd(m, n) \implies \text{lcm}(m, n) = d \cdot (m/d) \cdot (n/d).$$

Clearly then, if $d = \gcd(m, n) = 1$, then the lowest common multiple is mn and so the order of $xy \in G$ is mn .

Or to describe it a different way: $(xy)^{mn} = x^{mn}y^{mn} = ee = e$ so that the order of xy ,

$$|xy| \leq mn.$$

Conversely any r such that $(xy)^r = e$ must have $m, n \mid r$ and the lowest common multiple of m and n is mn so

$$r \geq mn.$$

Therefore, $|xy| = mn$. □

2.1.2.7 Examples of Cyclic Subgroups

(3) Cyclic group with order 3

$$G = \{1, x, x^2\}$$

where $x^3 = 1$ is a cyclic group of order 3 generated by the element x . Note that, since this is a group, it must also contain the inverses, x^{-1}, x^{-2} but $x^3 = 1$ so $x^{-1} = x^2$ and $x^{-2} = x$.

(4) Symmetries of an equilateral triangle

Consider an equilateral triangle with vertices labeled A, B, C :

$$\begin{array}{c} A \\ B \quad C. \end{array}$$

Every permutation of the vertices is a transformation that produces an object that occupies the same space as the original, i.e. a *symmetry*. If we take one of them, say, the clockwise rotation one place that results in,

$$\begin{array}{c} B \\ C \quad A \end{array}$$

and we name this r , then clearly – since there are 3 vertices – performing this same rotation 3 times leaves us back where we started. So, using function composition as the law of composition and multiplicative notation, $r^3 = i$ where i is the identity transformation. Also the inverse of r is r^2 . So, we have a group consisting of $\{i, r, r^2\}$ and function composition. Notice the resemblance of this group to the previous group $\{1, x, x^2\}$; this group is *isomorphic* to the cyclic group of order 3.

(5) **Group $(\mathbb{Z}_5^*, \otimes)$**

Consider the element 2 modulo 5. Using multiplicative notation we have,

$$2^2 = 4, 2^3 = 8 = 3, 2^4 = 16 = 1, 2^5 = 32 = 2.$$

So $2^1 = 2^5 \iff 1 = 2^4$ meaning that the element 2 has order 4 in the group and we see, as expected that the group it generates, $\langle 2 \rangle$ has 4 members. In this case, the members are all the members of the group – that's to say, *the element 2 generates the whole group*. If we consider the element 4 we have,

$$4^2 = 16 = 1, 4^3 = 64 = 4.$$

So this element oscillates between 1 and 4 and so, the cyclic subgroup that it generates $\langle 4 \rangle$ has order 2.

Since the group $(\mathbb{Z}_5^*, \otimes) = \langle 2 \rangle$ it can also be described as a cyclic group. This will be the case for any such group modulo a prime number – i.e. $(\mathbb{Z}_p^*, \otimes)$ where p is prime.

(6) **Cyclic group with infinite order**

$$\begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}$$

under matrix multiplication (which is commutative in this case), generates a cyclic group of infinite order because

$$\begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}^n = \begin{bmatrix} 1 & n \\ 0 & 1 \end{bmatrix}.$$

(7) **Cyclic groups in a non-Abelian group**

Consider the following two elements in $GL(2, \mathbb{R})$,

$$A = \begin{pmatrix} 0 & 1 \\ -1 & 0 \end{pmatrix}, \quad B = \begin{pmatrix} 0 & -1 \\ 1 & -1 \end{pmatrix}.$$

Both of these elements have finite order as,

$$(A^2)^2 = \begin{pmatrix} -1 & 0 \\ 0 & -1 \end{pmatrix}^2 = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix},$$

$$(B^2)B = \begin{pmatrix} -1 & 1 \\ -1 & 0 \end{pmatrix} \begin{pmatrix} 0 & -1 \\ 1 & -1 \end{pmatrix} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}.$$

But their product AB does not have finite order.

$$AB = \begin{pmatrix} 1 & -1 \\ 0 & 1 \end{pmatrix}, \quad (AB)^n = \begin{pmatrix} 1 & -1 \\ 0 & 1 \end{pmatrix}^n = \begin{pmatrix} 1 & -n \\ 0 & 1 \end{pmatrix}$$

for any $n \in \mathbb{N}$.

- (8) **The Klein Four Group**, V is the simplest group that is not cyclic (it cannot be generated by a single element). It appears in many forms but, as an example, it can be realized as the group consisting of the four matrices,

$$\begin{bmatrix} \pm 1 & 0 \\ 0 & \pm 1 \end{bmatrix}$$

Any two non-identity elements generate V .

2.1.3 Isomorphisms

Definition. An **isomorphism** is a bijection between two groups that preserves the structure of the groups by being compatible with the law of composition of both groups. More formally, two groups are **isomorphic** if there exists a bijection $\phi : G \mapsto G'$ such that,

$$\phi(ab) = \phi(a)\phi(b) \text{ for all } a, b \in G$$

where ab represents composition according to the law of composition of G and $\phi(a)\phi(b)$ represents composition according to the law of composition of G' .

An **isomorphism** is a **bijection** between two **groups**. That's to say, it is already assumed in the definition of an isomorphism that the codomain G' is a group.

Proposition 19. As a consequence of this sole property that, across the bijection, the respective laws of composition are preserved, all other properties of the groups are also preserved.

Proof. Let e be the identity in G and $e' = \phi(e) \in G'$, and $1'$ be the identity element in G' then,

- Since G' is a group, it has the inverses property that every element has an inverse so,

$$\begin{aligned} e' &= \phi(e) = \phi(ee) = \phi(e)\phi(e) = e'e' && \text{using preservation of law of composition} \\ \iff (e')^{-1}e' &= ((e')^{-1}e')e' && \text{using the inverses property of } G' \\ \iff 1' &= e' \end{aligned}$$

which implies that e' is the identity in G' so that ϕ maps the identity in G to the identity in G' .

- We can use the fact just shown that $\phi(e) = e' = 1'$ to show,

$$\begin{aligned} 1' = e' &= \phi(e) = \phi(aa^{-1}) = \phi(a)\phi(a^{-1}) && \text{using preservation of law of composition} \\ \iff \phi(a)^{-1}1' &= \phi(a)^{-1}\phi(a)\phi(a^{-1}) && \text{using the inverses property of } G' \\ \iff \phi(a)^{-1} &= \phi(a^{-1}) \end{aligned}$$

which shows that ϕ maps $a^{-1} \in G$ to $\phi(a)^{-1} \in G'$.

□

For example, if $e \in G$ is the identity of G mapped to an element $e' = \phi(e) \in G'$, then for any $a \in G$ mapped to $a' = \phi(a) \in G'$,

$$a' = \phi(a) = \phi(ea) = \phi(e)\phi(a) = e'a'$$

And $a' = e'a' = a'e'$ means that e' is the identity in G' . Furthermore, the order of elements in G and G' will also be the same as,

$$a^n = e \iff e' = \phi(e) = \phi(a^n) = \phi(a)^n = (a')^n$$

Since two isomorphic groups have the same properties, it is often convenient to identify them with each other when speaking informally. For example, the symmetric group S_n of permutations of $\{1, \dots, n\}$ is isomorphic to the group of permutation matrices, a subgroup of $GL_n(\mathbb{R})$ and we often blur the distinction between these two groups.

Notation. Sometimes when two groups are isomorphic this is indicated using the notation,

$$G \approx G'$$

2.1.3.1 Examples

- Let $C = \{\dots, a^{-2}, a^{-1}, 1, a, a^2, \dots\}$ be an infinite cyclic group. Then the map,

$$\phi : \mathbb{Z}^+ \mapsto C \text{ s.t. } \phi(n) = a^n$$

is an isomorphism where the preservation of the respective laws of composition can be seen as,

$$\phi(m+n) = a^{m+n} = a^m a^n = \phi(m)\phi(n)$$

and also $n + (-n) = 0$ and,

$$\phi(-n) = a^{-n} = (a^n)^{-1}.$$

- Let G be the set of real matrices of the form,

$$\begin{bmatrix} 1 & x \\ 0 & 1 \end{bmatrix}$$

This is a subgroup of $GL_2(\mathbb{R})$ and so, its law of composition is the same as that of $GL_2(\mathbb{R})$, i.e. matrix multiplication.

$$\begin{bmatrix} 1 & x \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & y \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} 1 & x+y \\ 0 & 1 \end{bmatrix}$$

So, G is isomorphic to \mathbb{R}^+ , the additive group of reals.

Definition. The groups isomorphic to a given group G form what is called the **isomorphism class** of G . Groups are often classified into isomorphism classes, for example, there is one isomorphism class of groups of order 3 and there are two classes of groups of order 4 and five classes of 12.

Proposition 20. *There is only one isomorphism class for each order of cyclic group.*

Proof. Any two cyclic groups of the same order are isomorphic because, if

$$G = \{1, x, x^2, \dots, x^{n-1}\}, G' = \{1, y, y^2, \dots, y^{n-1}\}$$

are two cyclic groups of order n then the map $\phi(x^i) = y^i$ is an isomorphism. \square

Proposition 21. Cayley's Theorem states that every group is isomorphic to a group of permutations of the same underlying set, or in other words, to a subgroup of the symmetric group acting on the group.

Proof. Let G be a group and $x \in G$ and define $f_x : G \rightarrow G$ as $f_x(g) = xg$. Then f_x is a bijection because it has an inverse $f_x^{-1}(g) = x^{-1}g = f_{x^{-1}}(g)$. Therefore f_x is a permutation.

Now we define a map, from G to the symmetric group of permutations of G , that maps each element x in G to the permutation defined by f_x . Let $\phi : G \rightarrow \text{Sym}(G)$ be defined as $\phi(x) = f_x$ then,

- ϕ is homomorphic because

$$(f_x \circ f_{x'})(g) = f_x(f_{x'}(g)) = x(x'g) = xx'g = f_{xx'}(g)$$

and so,

$$\phi(xx') = f_{xx'} = f_x \circ f_{x'} = \phi(x) \circ \phi(x').$$

- ϕ is injective because if $x, x' \in G$ such that $x \neq x'$ then $f_x(e) = x \neq x' = f_{x'}(e)$ is sufficient to show that $f_x \neq f_{x'}$. Or alternatively, the kernel of ϕ comprises the elements $k \in G$ such that $f_k(g) = kg = g \iff k = e$ so that the kernel is the trivial subgroup $\{e\}$.

Since ϕ is homomorphic, its image $\text{im } \phi$ is a subgroup of $\text{Sym}(G)$ and since ϕ is injective, it is in bijective correspondence with its image $\text{im } \phi \leq \text{Sym}(G)$. Therefore G is isomorphic to a subgroup of $\text{Sym}(G)$. \square

2.1.3.2 Automorphisms

Definition. The domain and codomain of an isomorphism can be the same set of objects so that $\phi : G \mapsto G$. This is known as an **automorphism**.

Example Let $G = \{1, x, x^2\}$ be a cyclic group of order 3 so that $x^3 = 1$. The transposition which interchanges x and x^2 is an automorphism of G ,

$$\begin{array}{ccc} 1 & \mapsto & 1 \\ x & \mapsto & x^2 \\ x^2 & \mapsto & x \end{array}$$

$$\begin{array}{|c|c|c|c|} \hline & 1 & x & x^2 \\ \hline 1 & 1 & x & x^2 \\ x & x & x^2 & 1 \\ x^2 & x^2 & 1 & x \\ \hline \end{array} \mapsto \begin{array}{|c|c|c|c|} \hline & 1 & x^2 & x \\ \hline 1 & 1 & x^2 & x \\ x^2 & x^2 & x & 1 \\ x & x & 1 & x^2 \\ \hline \end{array}$$

This is because the group is cyclic and x and x^2 have the same order ($x^3 = 1$ and also $(x^2)^3 = x^6 = (x^3)^2 = 1^2 = 1$). So the law of composition is preserved.

2.1.3.3 Conjugation

The most important example of automorphism is conjugation.

Definition. *Conjugation* by $b \in G$ is the map from G to itself defined by,

$$\phi(a) = bab^{-1}$$

with the result that,

$$ba = \phi(a)b$$

so that we can think of conjugation of a by b as the way that we need to change a if we want to move the multiplication by b to the other side.

This is an automorphism (known as an *inner automorphism*) because it

- is compatible with law of composition,

$$\phi(xy) = bxyb^{-1} = bxb^{-1}byb^{-1} = \phi(x)\phi(y).$$

- has an inverse so it is bijective,

$$(\phi^{-1} \circ \phi)(a) = \phi^{-1}(\phi(a)) = b^{-1}(bab^{-1})b = (b^{-1}b)a(b^{-1}b) = a.$$

Note that this is different from the inverse element of a corresponding under the mapping ϕ ,

$$\phi(a)\phi(a^{-1}) = bab^{-1}ba^{-1}b^{-1} = ba(1)a^{-1}b^{-1} = b(1)b^{-1} = 1.$$

A couple more important properties of the conjugate are as follows.

- (i) In an abelian group where the composition law is commutative, conjugation becomes the identity map.

$$ba = ab \iff bab^{-1} = a \iff \phi(a) = a.$$

- (ii) The inverse of the conjugate $bab^{-1} = b^{-1}ab$.

2.1.4 Homomorphisms

Definition. A **homomorphism** is a mapping (not necessarily bijective) between two groups, $\phi : G \mapsto G'$, such that,

$$\phi(ab) = \phi(a)\phi(b) \text{ for all } a, b \in G$$

where ab represents composition according to the law of composition of G and $\phi(a)\phi(b)$ represents composition according to the law of composition of G' .

So, the difference between a *homomorphism* and a *isomorphism* is that the latter is bijective whereas the former is not. As a result, a *homomorphism* may be one-way only.

*A **homomorphism** is a **mapping** between two **groups**. That's to say, it is already assumed in the definition of a homomorphism that the codomain G' is a group.*

Examples of homomorphisms

- (9) Let $C = \{a^{n-1}, \dots, a^{-2}, a^{-1}, 1, a, a^2, \dots, a^{n-1}\}$ be a finite cyclic group. Then the map,

$$\phi : \mathbb{Z}^+ \mapsto C \text{ s.t. } \phi(n) = a^n$$

is a homomorphism. Note that if C were an infinite cyclic group then this would be an isomorphism.

- (10) the sign of a permutation $sign : S_n \mapsto \pm 1$
(11) the determinant function $det : GL_n(\mathbb{R}) \mapsto \mathbb{R}^\times$
(12) an arguably trivial example is called the *inclusion* map $i : H \mapsto G$ of a subgroup H into a group G , defined by $i(x) = x$. It functions as the identity for elements in the subgroup H but, since it is not surjective, there is no inverse mapping.

2.1.4.1 Image of a homomorphism

Since a homomorphism is not bijective it has an image different to the codomain group,

$$\text{im } \phi = \{ x \in G' \mid \exists a \in G \text{ s.t. } \phi(a) = x \}$$

The image of a homomorphism is a subgroup of the codomain group G' because the homomorphism preserves the group structure as described in Proposition 19.

Notation. The image of the mapping ϕ with domain G is sometimes denoted $\phi(G)$.

2.1.4.2 Kernel of a homomorphism

Definition. The **kernel** of a homomorphism is the set of elements in the domain that are mapped to the identity,

$$\ker \phi = \{ a \in G \mid \phi(a) = 1' \}$$

Proposition 22. *The kernel of a homomorphism is a subgroup of the domain group G .*

Proof. If $a, b \in \ker \phi$ then,

- closure: $\phi(ab) = \phi(a)\phi(b) = 1' \cdot 1' = 1'$ which shows that

$$a, b \in \ker \phi \implies ab \in \ker \phi.$$

- identity: By Proposition 19, $1' = e' = \phi(e)$ and so $e \in \ker \phi$.
- inverses: Since $a \in \ker \phi$, then

$$\begin{aligned} 1' &= e' = \phi(e) = \phi(aa^{-1}) = \phi(a)\phi(a^{-1}) = 1'\phi(a^{-1}) \\ \iff 1' &= \phi(a^{-1}) \end{aligned}$$

so that $a \in \ker \phi \iff a^{-1} \in \ker \phi$.

□

Proposition 23. *If $\phi : G \mapsto G'$ is a group homomorphism with kernel N then, for $a, b \in G$,*

$$\phi(a) = \phi(b) \iff \exists n \in N, \text{ s.t. } b = an$$

or, equivalently, $a^{-1}b \in N$.

Proof.

$$\begin{array}{lll}
 & b = an & \\
 \implies & \phi(b) = \phi(an) & \\
 \implies & \phi(b) = \phi(a)\phi(n) & \text{by homomorphism property} \\
 \implies & \phi(b) = \phi(a)1' & \text{n is in the kernel} \\
 \implies & \phi(b) = \phi(a) &
 \end{array}$$

$$\begin{array}{lll}
 & \phi(b) = \phi(a) & \\
 \implies & \phi(a)^{-1}\phi(b) = 1' & \text{codomain is a group so has inverses} \\
 \implies & \phi(a^{-1})\phi(b) = 1' & \text{by Proposition 19} \\
 \implies & \phi(a^{-1}b) = 1' & \text{by homomorphism property} \\
 \implies & a^{-1}b = n \in N & \\
 \implies & b = an &
 \end{array}$$

□

Theorem 13. *A homomorphism is injective iff its kernel is the trivial subgroup $\{e\}$.*

When asked to prove the proposition that a homomorphism is injective iff its kernel is the trivial subgroup $\{e\}$, it's tempting to begin proving each direction of the bidirectional implication with a proof by contradiction (e.g. "Assuming there is a non-identity element in the kernel...") but the direct positive proof can be made very quick and simple with the above corollary.

Proof. Let $\phi : G \mapsto G'$ be a homomorphism.

Assume that ϕ is injective and $k \in \ker \phi$. Remembering that we always at least have $e_G \in \ker \phi$,

$$\phi(k) = e_{G'} = \phi(e_G) \iff k = e_G$$

where the last implication is by the injectivity of ϕ .

Now assume that $\ker \phi = \{e_G\}$, $a, b \in G$. Then,

$$\begin{aligned} & \phi(a) = \phi(b) \\ \iff & \phi(a)\phi(b)^{-1} = e_{G'} \\ \iff & \phi(a)\phi(b^{-1}) = \phi(ab^{-1}) = e_{G'} && \text{using homomorphism properties} \\ \iff & ab^{-1} \in \ker \phi \\ \iff & ab^{-1} = e_G && \text{by assumption } \ker \phi = \{e_G\} \\ \iff & a = b. && \square \end{aligned}$$

Corollary 5. *A homomorphism is an isomorphism if its kernel contains only the identity and its image is the whole of the codomain (i.e. it's surjective).*

2.1.4.3 Examples of Kernels of Homomorphisms

- (13) The determinant function 11, $\det : GL_n(\mathbb{R}) \mapsto \mathbb{R}^\times$, has a kernel,

$$\{\text{real } n \times n \text{ matrices } A \mid \det A = 1\},$$

which is a subgroup of $GL_n(\mathbb{R})$ known as the *special linear group* $SL_n(\mathbb{R})$.

- (14) The sign of a permutation 10 has a kernel that is the set of *even* permutations,

$$A_n = \{\text{even permutations}\},$$

which is a subgroup of the symmetric group S_n and is known as the *alternating group*, A_n .

- (15) The map from the additive group of integers to a finite cyclic group 9,

$$\phi : \mathbb{Z}^+ \mapsto C \text{ s.t. } \phi(n) = a^n$$

has the kernel,

$$\ker \phi = \{ n \in \mathbb{Z}^+ \mid a^n = 1 \}$$

which has been proven to be a subgroup in Proposition 17.

2.1.5 Equivalence Relations and Partitions

Notation. In the following treatment of equivalence relations we will use the notation $a \sim b$ to denote the equivalence of a and b ; \bar{a} to indicate the equivalence class of a ; and \bar{S} to indicate the partition of S comprised of equivalence classes such as the class $\bar{a} = \bar{b}$ which includes both a and b .

Any map of sets $\phi : S \mapsto T$ defines an equivalence relation on the domain S such that $a \sim b$ iff $\phi(a) = \phi(b)$. We will refer to this as the *equivalence relation determined by the map*. The corresponding partition is made up of the sets of elements in the domain S that are mapped to the same element in the codomain T .

Definition. Let $\phi : S \mapsto T$ be a map, then the **inverse image** of an element $t \in T$ is defined as,

$$\phi^{-1}(t) = \{ s \in S \mid \phi(s) = t \}$$

and can also be applied to a set $U \in T$ as,

$$\phi^{-1}(U) = \{ s \in S \mid \phi(s) \in U \}$$

Note that in this notation, ϕ^{-1} **does not indicate an inverse function** as the inverse of the function may not exist but the inverse image is nevertheless defined.

The inverse images - the sets $\phi^{-1}(t)$ for all $t \in T$ - may also be called the **fibres** of the map ϕ .

Clearly, the non-empty fibres of the map ϕ form a partition of S . We can express this partition of S as a bijection, that we shall call $\bar{\phi}$, between the

fibres of ϕ in S and the element of the image of S to which their members are mapped,

$$\bar{\phi} : \bar{S} \mapsto \text{im } \phi$$

so that,

$$\bar{\phi}(\bar{s}) = \phi(s).$$




Figure 2.1: A schematic diagram of a group homomorphism

2.1.5.1 Congruence

Since a homomorphism maps the identity to the identity and inverses to inverses (Proposition 19), we can deduce that the inverse image of the identity in G' is going to contain at least the identity of G and that the inverse image of an element $(a')^{-1} \in G'$ will contain at least the element $a^{-1} \in G$. So, in terms of equivalence classes we can say that, for a homomorphism ϕ ,

$$\begin{aligned} 1 = e \in \phi^{-1}(1') &\implies \bar{\phi}(\bar{1}) = 1' \\ a^{-1} \in \phi^{-1}((a')^{-1}) &\implies \bar{\phi}(\overline{a^{-1}}) = (a')^{-1} \end{aligned}$$

Definition. The equivalence relation determined by a homomorphism is known as ***congruence*** and is commonly denoted using \equiv instead of \sim . For a homomorphism ϕ ,

$$a \equiv b \iff \phi(a) = \phi(b).$$

Since ϕ is a homomorphism we also have,

$$a \equiv b \iff \phi(ac) = \phi(bc), \phi(a^{-1}) = \phi(b^{-1}).$$

More generally, a ***congruence relation*** is an equivalence relation on an algebraic structure (such as a group, ring, or vector space) that is compatible with the structure in the sense that algebraic operations done with equivalent elements will yield equivalent elements.

2.1.5.2 Congruence Examples

- (16) The modulus function of complex numbers forms a homomorphism from the multiplicative group of complex numbers to the multiplicative group of reals,

$$\phi : \mathbb{C}^\times \mapsto \mathbb{R}^\times \text{ s.t. } \phi(a) = |a|$$

and the induced equivalence relation is $a \equiv b \iff |a| = |b|$. The fibres of this map are the concentric circles about 0. They are in bijective correspondence with elements of $\text{im } \phi$, the set of positive reals.

2.1.6 Cosets

The set of elements of the form an - described in Proposition 23 - is denoted by aN and is called a *coset* of N in G .

Definition. A coset can be defined for any subgroup H of a group G . A **left coset** is a subset of the form,

$$aH = \{ ah \mid h \in H \}.$$

Cosets are not, in general, subgroups. This can be easily seen as the left coset aH does not contain the identity a , although H contains the identity, aH contains $a1 = a$.

Note that the arbitrary subgroup H could also be thought of as a coset $1H = H$ and also that the left cosets aH are equivalence classes for the congruence relation,

$$a \equiv b \iff b = ah, h \in H.$$

This is a congruence because, for some arbitrary $c \in G$,

$$1 \equiv c \iff \exists h \in H \text{ s.t. } c = 1h = h.$$

That's to say, the elements that are congruent to the identity are precisely the members of the subgroup H so that it plays a similar role to the kernel N in Proposition 23. Furthermore, since the congruence relation is an equivalence relation it forms a partition of the domain G .

Proposition 24. For a group G with a subgroup H and $x \in G$, the coset xH is equal to H iff $x \in H$.

Proof. Assume $x \in H$. Then $\forall xh \in xH . xh \in H$. Therefore $xH \subseteq H$. Conversely, $x^{-1} \in H$ so,

$$\forall h \in H . x^{-1}h \in H \implies x(x^{-1}h) = h \in xH.$$

Therefore $H \subseteq xH$ and so $H = xH$.

Now assume that $H = xH$. Since $e \in H$ then $xe = x \in xH = H$ and so $x \in H$. \square

Proposition 25. *The left cosets of a subgroup partition the group.*

Proof. The left cosets are equivalence classes and, as a result, they partition the group. \square

2.1.6.1 Examples of cosets

- (17) The coset of an element with the kernel N ,

$$aN = \{ g \in G \mid g = an, n \in N \}$$

is the set of all elements that are *congruent* to a . The *congruence classes* are precisely the cosets aN for each $a \in G$. They are also the nonempty *fibres* of the homomorphic map.

- (18) 2.1.1.2 Continuing the example of the symmetric group S_3 represented as

$$G = \{1, x, x^2, y, xy, x^2y\}$$

with group multiplication rules,

$$x^3 = 1, y^2 = 1, yx = x^2y.$$

The element xy has order 2 so it generates a cyclic subgroup $H = \{1, xy\}$ of order 2. The left cosets of H in G are the three sets,

$$\{1, xy\} = 1H = xyH, \{x, x^2y\} = xH = x^2yH, \{x^2, y\} = x^2H = yH.$$

Note that they do partition the group G . Also, notice that the cosets aH for $a \in H$ produce the subgroup H itself as should be expected as the group properties of the subgroup dictate that all products of its elements are already present in the subgroup. For this reason, the cosets aH that are distinct from H are those such that $a \notin H$.

- (19) Let $G = (\mathbb{R}^3, +)$ be the group of 3d vectors with vector addition and $\vec{w} \in G$. Then if

$$H = \{ \vec{x} \in G \mid \vec{w}^T \vec{x} = \vec{0} \}$$

then H is a subgroup, $H \leq G$. H is a vector space representing a plane through the origin in \mathbb{R}^3 and its cosets are

$$\vec{v} + H = \{ \vec{v} + \vec{h} \mid \vec{v} \in G, \vec{h} \in H \}$$

which are the affine spaces representing the translated planes, parallel to H , but not passing through the origin. Once again we see that the cosets partition the space even if there may be an infinite number of them.

2.1.6.2 The index of a subgroup

Definition. The *index* of a subgroup is the number of left cosets it forms in the parent group.

Notation. The **index** of a subgroup H in G is denoted by $[G : H]$.

In the example (18) the index of H is 3. Note that if G were to contain infinitely many elements then the index of a subgroup may also be infinite.

Proposition 26. *Each coset aH has the same number of elements as H .*

Proof. As usual, equal cardinality is demonstrated by showing the existence of a bijection. It is clear that there is a bijective map between the subgroup H and any coset aH because the map $H \mapsto aH$ is,

- injective because $ah = ah' \implies h = h'$ because by group properties a has an inverse in G ;
- surjective because every $c \in aH$ has the form ah and is therefore mapped to by some $h \in H$.

□

2.1.6.3 Lagrange's Theorem

Since the left cosets of H in G form a partition of G and their order is the same as that of H we see that the order of G is the order of H multiplied by its index in G . This results in a formula known as the *Counting Formula* as follows,

$$|G| = |H| \cdot [G : H].$$

If G is of infinite order and H is finite, then the index of H in G will be infinite.

Theorem 14. Lagrange's Theorem: *Let G be a finite group, and let H be a subgroup of G . The order of H divides the order of G .*

Corollary 6. *Let G be a finite group, and let a be an element of G . Then the order of a divides the order of G . That's to say, the order of the cyclic group generated by a , $|\langle a \rangle|$, divides $|G|$.*

Corollary 7. *If G is a group of order n , then $g^n = e$ for every element g of G .*

Proof. This is clearly a consequence of the previous corollary. If we let the order of g be m , then by the previous corollary,

$$m \mid n \iff n = km \text{ for } k \in \mathbb{N} \iff g^n = g^{km} = (g^m)^k = e^k = e. \quad \square$$

Corollary 8. *Suppose that a group G has p elements and that p is a prime integer. Let $a \in G$ be any element, not the identity. Then G is the cyclic group $\{1, a, \dots, a^{p-1}\}$ generated by a .*

Proof. Since $a \neq 1$ by selection, it has order greater than 1. Since its order must divide the order of G , which is prime, its order is equal to the order of G , p . So, the order of the nonidentity element a is the same as the order of G and so it generates the whole group. \square

Corollary 9. *All groups with some prime order, p , are in the same isomorphism class.*

Proof. Any group with prime order p is the cyclic group of order p and by Proposition 20 there is only a single isomorphism class for each cyclic group of a given order. \square

Proposition 27. *Suppose the elements x, y in a group G have orders m, n respectively and that the $\gcd(m, n) = 1$. Then $\langle x \rangle \cap \langle y \rangle = \{e\}$.*

Here we will prove, using Lagrange's Theorem, something that we previously proved here (Proposition 18) using modular arithmetic. Notice how the proofs are similar but the proof with Lagrange's Theorem allows us to remain within Group Theory.

Proof. Firstly, note that the intersection of the two cyclic groups,

$$H = \langle x \rangle \cap \langle y \rangle$$

is a subgroup both of the parent group G **and** of $\langle x \rangle$ and $\langle y \rangle$. So Lagrange's Theorem tells us that its order must divide into the order of the parent group **and** the orders of the cyclic groups of x and y . Therefore, we have,

$$|H| \mid m \quad \text{and} \quad |H| \mid n.$$

Now, applying the fact that the $\gcd(m, n) = 1$ we see that $|H| \mid 1$ and therefore $|H| = 1$. Furthermore, any group of order 1 must be the minimal group $\{e\}$. \square

Proposition 28. *Suppose that H is a subgroup of G and $x \in G$. Then there exists some $k \in \mathbb{N}$, $1 \leq k \leq [G : H]$ s.t. $x^k \in H$.*

Proof. Let $n = [G : H]$ be the index of H in G . Then there are precisely n cosets of H in G . But $x \in G \implies x^m \in G$ for any $m \in \mathbb{N}$ (we don't need to consider the negative powers of x because they are inverses of positive powers and are similar for these purposes) and so we have cosets of the form $x^m H$ for each $m \in \mathbb{N}$. Therefore, amongst the $n + 1$ cosets generated by, $x^i H$ for $i \in \{0, 1, \dots, n\}$ we must have at least one repetition of the same coset.

So, for some fixed x^i, x^j with $0 \leq i, j \leq n$ and $i \neq j$, we have,

$$\begin{aligned} & x^i H = x^j H \\ \iff & \forall h \in H . x^i h \in x^j H \\ \iff & \forall h \in H . \exists h' \in H . x^i h = x^j h' \\ \iff & \forall h \in H . \exists h' \in H . x^{i-j} = h^{-1} h' \in H \end{aligned}$$

Since necessarily we have $1 \leq i - j \leq n$ we let $k = i - j$ and then $x^k \in H$ as required. \square

2.1.6.4 Example applications of Lagrange Theorem

(20) **Fermat's Little Theorem:** *If p is a prime number then*

$$a^p \equiv a \pmod{p} \text{ for all } a \in \mathbb{Z}.$$

We need to be a little careful here. We might assume – given that we are multiplying the integer a in modulo p that the group we want to use is (\mathbb{Z}_p, \otimes) . However, this is not a group! The reason is that \mathbb{Z}_p contains 0 which has no inverse under the proposed law of composition, multiplication.

If, however, we take \mathbb{Z}_p^ where the $*$ means $\mathbb{Z}/\{0\}$ then we have a set of $p - 1$ distinct elements. Over this set we can form the multiplicative group $G = (\mathbb{Z}_p^*, \otimes)$ because the primality of p means that every element has a multiplicative inverse.*

*Note that this is **not** a group of prime order. The primality of p is essential to make sure that every element has a multiplicative inverse but, since we also have to eliminate 0 for the same reason, the order is $p - 1$ which is not necessarily prime.*

Proof. Take the set \mathbb{Z}_p^* under multiplication and some arbitrary $a \in \mathbb{Z}$.

- (i) Primality of p means that it is possible to find $1 = na + mp$ for $m, n \in \mathbb{Z}$ (see Corollary 2). This implies that there exists a multiplicative inverse of every non-zero element in modulo p . Specifically, n is the inverse of a because $na = (-m)p + 1 \iff na \pmod{p} \equiv 1$.
- (ii) Existence of the multiplicative inverses implies that we have a group $G = (\mathbb{Z}_p^*, \otimes)$.
- (iii) G being a group implies that, for any element $a \in G$, by Corollary 7 we have $a^{p-1} = 1$.
- (iv) In G , $a^{p-1} = 1 \iff a^p = a$ which translates to $a^p \equiv a \pmod{p}$.

□

2.1.6.5 Lagrange's Theorem and Homomorphisms

The Counting Formula can also be applied when a homomorphism is given. Let $\phi : G \mapsto G'$ be a homomorphism. As we saw in coset example 17, the left cosets of $\ker \phi$ are the fibres of the map ϕ . They are in bijective correspondence with the elements of the image. Therefore,

$$[G : \ker \phi] = |\operatorname{im} \phi|.$$

Which implies that,

Corollary 10. *If $\phi : G \mapsto G'$ is a homomorphism of finite groups then,*

$$|G| = |\ker \phi| \cdot |\operatorname{im} \phi|.$$

As a result, $|\ker \phi|$ divides $|G|$, and $|\operatorname{im} \phi|$ divides both $|G|$ and $|G'|$.

2.1.6.6 Restriction of a Homomorphism to a Subgroup

A useful way of understanding the structure of a complicated group is to understand its subgroups and then derive an understanding of the parent group from knowledge about the subgroups it contains. This frequently involves the application of Lagrange's Theorem. *Restriction of a Homomorphism to a subgroup* refers to studying the behaviour of a homomorphism on subgroups of the parent group.

Suppose that $\phi : G \mapsto G'$ is a homomorphism and that H is a subgroup of G . Then we may *restrict ϕ to H* to obtain a homomorphism whose domain is a subset of the original,

$$\phi|_H : H \mapsto G'.$$

This *restriction* is a homomorphism because ϕ is a homomorphism and the restriction domain is a group. Clearly, the kernel of the restricted homomorphism is the intersection of the domain H with $\ker \phi$.

2.1.6.7 Examples of using Lagrange's Theorem with a homomorphism restricted to a subgroup

- (21) Referring again to the sign of a permutation (10) $S_n \mapsto \{-1, 1\}$: the order of the codomain of this homomorphism is clearly 2. Suppose

we form the restriction of this homomorphism to a subgroup H of S_n . Then, denoting the image by $\phi|_H(H)$, by Corollary 10 we have that $|\phi|_H(H)|$ divides both 2 and $|H|$.

So, if the subgroup H has odd order then $|\phi|_H(H)| = 1$ and – since $\phi|_H(H)$ must be a group because the group structure is preserved across the homomorphism – $\phi|_H(H) = \{1\}$. This means that H is in the kernel of the sign map and that the subgroup of permutations in S_n represented by H consists of only even permutations.

Therefore, every permutation whose order in S_n is odd is an even permutation (since the cyclic group that it generates has odd order). However, we can not make any conclusions about permutations of even order; they may be even or odd permutations.

2.1.6.8 Right Cosets

Right cosets also exist and are defined as,

$$Ha = \{ g \in G \mid g = ha, h \in H \}$$

and these are equivalence classes for the *right congruence* relation,

$$a \equiv b \iff b = ha, h \in H.$$

Right cosets are not necessarily the same as left cosets. For instance, continuing the example in 18, the right cosets of the subgroup $\{1, xy\}$ of S_3 are,

$$\{1, xy\} = H1 = Hxy, \{x, y\} = Hx = H, \{x^2, x^2y\} = Hx^2 = Hx^2y.$$

Note that this generates a different partition of G than was generated by the left cosets.

2.1.7 Normal Subgroups and Centers

2.1.7.1 Normal Subgroups

Definition. A subgroup N of a group G is called a **normal subgroup** if it has the property that,

$$\forall a \in N, b \in G, bab^{-1} \in N$$

which is to say, that the conjugate by any element of G of any element in N is also in N .

Proposition 29. A subset H of a group G is normal if and only if every left coset is also a right coset. If H is normal then,

$$\forall a \in G, aH = Ha.$$

Proof. Suppose that H is normal. For any $h \in H$ and any $a \in G$,

$$ah = (aha^{-1})a.$$

Since H is normal, the conjugate by h of a is also in H , that's to say, $aha^{-1} \in H$ which implies that $(aha^{-1})a \in Ha$. Therefore, any arbitrary member of aH is also a member of Ha . Clearly, the same proof also works in the other direction so that any member of Ha is also a member of aH and the two cosets are equal. So, we have shown that $(H \text{ is normal}) \implies (\text{left and right cosets of } H \text{ are equal})$.

Now we need to show that $(\text{left and right cosets of } H \text{ are equal}) \implies (H \text{ is normal})$. Firstly, clearly the above logic doesn't apply if H is not normal; there will be at least one element whose conjugate is not in H so $aH \neq Ha$. However, it could still be the case that each left coset is also a right coset if, for every a in G , there is some b in G such that $aH = Hb$. However, this is not possible because aH and Ha both contain a which means that in a given partition of G they must be the same partition. So $aH \neq Ha$ implies that the partitions are different; Ha creates different equivalence classes. Therefore $(\text{left and right cosets of } H \text{ are equal}) \implies (H \text{ is normal})$. \square

This is really the point of normal subgroups: That their cosets contain multiplication from both sides. As a result, when the cosets themselves are used as members of groups (see: Quotient Groups 2.1.8.6), we can define a group composition operation between them such that $aHbH = abH$ and the operation is well defined (i.e. equal arguments give equal results) because,

$$aHbH = abHH = abH \quad \text{and} \quad aH = a'H \implies abH = aHb = a'Hb = a'bH$$

where $aH = Ha$ because H is normal. .

2.1.7.2 Examples of Normal Subgroups

(22) The kernel of a homomorphism is a normal subgroup because,

$$\begin{aligned} a \in \ker \phi &\iff \phi(a) = 1 \\ \implies &\phi(bab^{-1}) = \phi(b) \cdot 1 \cdot \phi(b^{-1}) \\ \iff &\phi(bab^{-1}) = \phi(b)\phi(b)^{-1} && \text{using Proposition 19} \\ \iff &\phi(bab^{-1}) = 1. \end{aligned}$$

For example,

- a. $SL_n(\mathbb{R})$ is a normal subgroup of $GL_n(\mathbb{R})$ even though it is not Abelian, which can be seen as for $M \in GL_n(\mathbb{R})$, $A, B \in SL_n(\mathbb{R})$,

$$AB \neq BA \quad \text{and} \quad \det A = 1, \det M^{-1} = 1/(\det M)$$

so that,

$$\det M^{-1}AM = (\det M^{-1}) \cdot 1 \cdot (\det M) = (\det M)/(\det M) = 1.$$

- b. A_n is a normal subgroup of the symmetric group S_n .
c. In fact, any subgroup of $GL_n(\mathbb{F})$ with some fixed determinant $d \in \mathbb{R}$ will be a normal subgroup.

$$N = \{ A \in GL_n(\mathbb{F}) \mid \det A = d \}$$

For $A \in N$, $M \in GL_n(\mathbb{F})$. $\det(MAM^{-1}) = d$ by the same logic as in example a. These conjugate matrices are known as *similar* matrices.

- (23) Any subgroup of an abelian group is normal because when the composition law is commutative, as was mentioned in the section on conjugation,

$$ba = ab \iff bab^{-1} = abb^{-1} = a$$

so that conjugation becomes the identity map and so, trivially, all conjugates of elements in a subgroup are also in the subgroup.

Subgroups of non-abelian groups, however, need not be normal. For example,

- (24) Group T of invertible upper triangular matrices is not a normal subgroup of $GL_2(\mathbb{R})$. To show this note,

$$A = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}, B = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}, BAB^{-1} = \begin{bmatrix} 1 & 0 \\ 1 & 1 \end{bmatrix}$$

where $A \in T, B \in GL_2(\mathbb{R})$ but $BAB^{-1} \notin T$.

Proposition 30. *Let $\phi : G \mapsto G'$ be a homomorphism and let H' be a subgroup of G' . Denote the inverse image $\phi^{-1}(H') = \{x \in G \mid \phi(x) \in H'\}$ by \tilde{H} . Then,*

- (i) \tilde{H} is a subgroup of G .
- (ii) If H' is a normal subgroup of G' then \tilde{H} is a normal subgroup of G .
- (iii) \tilde{H} contains $\ker \phi$.
- (iv) The restriction of ϕ to \tilde{H} defines a homomorphism $\tilde{H} \mapsto H'$ whose kernel is $\ker \phi$.

Proof. Proofs are as follows:

- (i) \tilde{H} is a subgroup of G because ϕ is a homomorphism and its image, H' , is a group (which is required for a homomorphism).
- (ii) If H' is a normal subgroup of G' then \tilde{H} is a normal subgroup of G because for every element in \tilde{H} the mapped element is in H' . Then, since H' is normal, the conjugates of the mapped element are also in H' which means that their inverse images are in \tilde{H} . Since the map is homomorphic, the inverse images of the conjugates in G' are the respective conjugates in G .

- (iii) \tilde{H} contains $\ker \phi$ because it contains every element in G that maps to an element in H' and, since H' is a group, it includes the identity of G' . Therefore \tilde{H} contains every element that maps to the identity of G' which is $\ker \phi$.
- (iv) The restriction of ϕ to \tilde{H} is clearly a homomorphism and, since it contains $\ker \phi$, its kernel is equal to the kernel of ϕ .

□

2.1.7.3 The Center of a Group

Definition. The **center** of a group G is the set of elements that commute with every element of G ,

$$Z = \{ z \in G \mid zx = xz, \forall x \in G \}.$$

We can also define,

$$C(x) = \{ g \in G \mid gx = xg \}$$

as the set of elements in G that commute with a single fixed element x .

Notation. The **center** of a group G may be denoted by Z or by $Z(G)$.

The center of a group, Z , is a subgroup of G . This can be easily seen as, first of all, Z is non-empty because the identity is in the center of any group. Then, also, the center Z is closed under the group operation,

$$\forall a, b \in Z, x \in G . (ab)x = axb = x(ab)$$

and it contains the inverses,

$$\forall a \in Z, x \in G . ax = xa \iff a^{-1}ax = x = a^{-1}xa \iff xa^{-1} = a^{-1}x.$$

2.1.7.4 Examples of group centers

- (25) Let $G = GL(2, \mathbb{R})$ be the group of invertible 2x2 matrices with real coefficients and take two elements in G ,

$$M = \begin{pmatrix} 2 & 0 \\ 0 & 1 \end{pmatrix}, \quad N = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}.$$

Then we can identify the center of M by observing that an arbitrary matrix in $C(M)$ satisfies,

$$\begin{pmatrix} a & b \\ c & d \end{pmatrix} \begin{pmatrix} 2 & 0 \\ 0 & 1 \end{pmatrix} = \begin{pmatrix} 2a & b \\ 2c & d \end{pmatrix} = \begin{pmatrix} 2 & 0 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} a & b \\ c & d \end{pmatrix} = \begin{pmatrix} 2a & 2b \\ c & d \end{pmatrix}$$

which gives $b = 2b$, $c = 2c$ implying that b and c are 0. So,

$$C(M) = \left\{ \begin{pmatrix} a & 0 \\ 0 & d \end{pmatrix} \mid a, d \in \mathbb{R} \setminus \{0\} \right\}.$$

While matrices in the center of N satisfy,

$$\begin{pmatrix} a & b \\ c & d \end{pmatrix} \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} = \begin{pmatrix} a & a+b \\ c & c+d \end{pmatrix} = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} a & b \\ c & d \end{pmatrix} = \begin{pmatrix} a+c & b+d \\ c & d \end{pmatrix}$$

which gives $a = a + c$, $c + d = d$, $a + b = b + d$ implying that $c = 0$ and $a = d$. Therefore,

$$C(N) = \left\{ \begin{pmatrix} a & b \\ 0 & a \end{pmatrix} \mid a, b \in \mathbb{R}, a \neq 0 \right\}.$$

Note that in both cases some coefficients were required to be non-zero because to be members of the general linear group they must be invertible and so their determinant must be non-zero.

- (26) The center of the general linear group $GL_n(\mathbb{R})$ is the group of *scalar matrices* of the form cI for $c \in \mathbb{R}$, i.e. matrices of the form,

$$\begin{bmatrix} c & 0 \\ 0 & c \end{bmatrix}$$

in $GL_2(\mathbb{R})$. Note that, for *diagonal* matrices whose elements on the main diagonal are all non-zero but not-necessarily equal as in *scalar* matrices, multiplication is commutative with other diagonal matrices but not generally so with other matrices in the general linear group.

2.1.8 Products Groups and Quotient Groups

Definition. *If we take the cartesian product of two sets then:*

- *if the two sets are the underlying sets of two distinct groups then we have no way to combine them (as there is no common group operation) but we can take the pairing and define a component-wise multiplication over the pairs where each component is multiplied using that group's composition operation. In this way we create a new group over the pairs.*
- *if the two sets are subsets of a common group then there is a common group operation between them and so we can multiply them using this group operation. The result is another subset of the common group (not necessarily a subgroup).*

*Both of these may at times be referred to as **Product Groups** but the first one is more specifically referred to as a **Direct Product** and the second one may be referred to as a **Product Set**.*

2.1.8.1 Direct Products

Definition. *Let G, G' be two groups. The **direct product** is the set $G \times G'$ with component-wise multiplication using the group composition operation for the group corresponding to the component. Its order is the product of the orders of G and G' .*

Notation. The **direct product** of the two groups G, G' may be denoted by $G \times G'$ or GG' . In the case of Abelian groups the direct product may be referred to as the **direct sum** and denoted $G \oplus G'$.

So, if $a, b \in G$ and $a', b' \in G'$ then

- $(a, a'), (a, b'), (b, a'), (b, b') \in G \times G'$
- $(a, a')(b, b') = (ab, a'b')$

- the identity is $(1, 1)$ and $(a, a')^{-1} = (a^{-1}, a'^{-1})$.

Definition. The **projections** of a direct product $G \times G'$ are the maps p, p' such that,

$$p(x, x') = x, \quad p'(x, x') = x'.$$

Proposition 31. The **mapping property of direct products**: Let H be any group. The homomorphisms $\Phi : H \mapsto G \times G'$ are in bijective correspondence to pairs (ϕ, ϕ') of homomorphisms

$$\phi : H \mapsto G, \quad \phi' : H \mapsto G'.$$

The kernel of Φ is the intersection $(\ker \phi) \cap (\ker \phi')$.

Proof. Given a pair of homomorphisms (ϕ, ϕ') we can define $\Phi(x) = (\phi(x), \phi'(x))$. Then this is homomorphic because,

$$\Phi(xy) = (\phi(xy), \phi'(xy)) = (\phi(x), \phi'(x))(\phi(y), \phi'(y)) = \Phi(x)\Phi(y).$$

Conversely, given such a Φ we can recover the pair of homomorphisms with the group projections as such (outer parentheses omitted for clarity),

$$\phi(x), \phi'(x) = p(\Phi(x)), p'(\Phi(x)).$$

Since the correspondence is invertible, it is a bijection.

Clearly, also,

$$\Phi(x) = (\phi(x), \phi'(x)) = (1, 1) \iff (\phi(x) = 1) \wedge (\phi'(x) = 1)$$

so that $\ker \Phi = (\ker \phi) \cap (\ker \phi')$. □

Proposition 32. Let r, s be coprime integers. A cyclic group of order rs is isomorphic to the product of a cyclic group of order r and a cyclic group of order s .

Proof. Let $C = \{1, x, x^2, \dots, x^{rs-1}\}$, $C_1 = \{1, y, y^2, \dots, y^{r-1}\}$, $C_2 = \{1, z, z^2, \dots, z^{s-1}\}$ and define the map $\phi : C \mapsto C_1 \times C_2$ as,

$$\phi(x^i) = (y^i, z^i).$$

Then ϕ is homomorphic because it is comprised of two homomorphisms (by the mapping proper Proposition 31),

$$\phi_1(x^i) = y^i \quad \text{and} \quad \phi_2(x^i) = z^i.$$

And ϕ is injective because,

$$\phi(x^i) = (1, 1) \iff (y^i = 1) \wedge (z^i = 1) \iff (r \mid i) \wedge (s \mid i)$$

but r and s are coprime so this requires that $i = rs$ which is also the order of $x \in C$. So we have,

$$\phi(x^i) = (1, 1) \iff x^i = x^{rs} = 1.$$

Therefore $\ker \phi = \{1\}$ and, by Theorem 13, ϕ is injective.

Since ϕ is injective, its image has the same order as that of the domain C so we have,

$$|\text{im } \phi| = |C| = rs = |C \times C|$$

and ϕ is therefore surjective.

Therefore ϕ is a bijection and isomorphic. □

*Note that this is **only** the case for cyclic groups whose order is the product of two coprime numbers. For example, a cyclic group of order 4 is not isomorphic to a product of two cyclic groups of order 2 as every element in a product group $C_2 \times C_2$ has order 1 or 2. Whereas a cyclic group of order 4 has two elements of order 4 (the generating element and its inverse).*

Let $C_4 = \{1, x, x^2, x^3\}$, $C_2 = \{1, y\}$ and define the map $\phi : C_4 \mapsto C_2 \times C_2$ as $\phi(x^i) = (y^i, y^i)$. Then,

$$\phi(x^i) = (1, 1) \iff y^i = 1 \iff 2 \mid i$$

so that we have $\ker \phi = \{1, x^2\}$ and so ϕ is not injective.

2.1.8.2 Product Sets

Definition. Let A and B be subsets of the group G and denote the **product set** of A and B by

$$AB = \{x \in G \mid x = ab \text{ for some } a \in A \text{ and } b \in B\}.$$

Note that this notation is the same as one of the alternatives for the notation of the direct product so we need to be clear which is intended when we see this notation.

2.1.8.3 Relationship Between the Types of Product Groups

Proposition 33. Let H and K be subgroups of G .

- (i) If $H \cap K = \{1\}$, the product map $p : H \times K \mapsto G$ defined by $p(h, k) = hk$ is injective. Its image is the subset HK .
- (ii) If either H or K is a normal subgroup of G , then the product sets HK and KH are equal and are subgroups of G .
- (iii) If H and K are normal, $H \cap K = \{1\}$, and $HK = G$, then G is isomorphic to the direct product $H \times K$.

Proof. Proofs of each property are as follows.

- (i) If we assume that $H \cap K = \{1\}$ then, for $h, h' \in H, k, k' \in K$,

$$p(h, k) = p(h', k') \iff hk = h'k' \iff (h')^{-1}h = k'k^{-1}$$

so that $(h')^{-1}h = k'k^{-1} \in H \cap K = \{1\}$ therefore,

$$(h')^{-1}h = 1 \iff h = h', \quad k'k^{-1} = 1 \iff k' = k.$$

Therefore p is injective.

- (ii) Assume w.l.o.g. that K is a normal subgroup. Then for all $k \in K, g \in G, g^{-1}kg \in K$ and, in particular, for $h \in H, h^{-1}kh \in K$. Therefore,

$$hk \in HK \implies h(h^{-1}kh) = kh \in HK$$

and conversely, using the fact that, $h^{-1} \in H \implies hkh^{-1} \in K$,

$$kh \in KH \implies (hkh^{-1})h = hk \in KH.$$

Therefore $HK = KH$ and this implies that HK is a subgroup because, for $h, h' \in H, k, k' \in K$,

- HK is closed because

$$kh' \in KH = HK \implies kh' = h''k'' \in HK$$

for some $h'' \in H, k'' \in K$, and so,

$$hk, h'k' \in HK \implies (hk)(h'k') = h(kh')k' = h(h''k'')k' = (hh'')(k''k') \in HK.$$

- HK has inverses because for $hk \in HK$,

$$h^{-1}k^{-1} \in HK \implies k^{-1}h^{-1} \in HK.$$

- (iii) If $H \cap K = \{1\}$ then the product map p is injective and if $HK = G$ then $\text{im } p = G$ so p is surjective also and, therefore, is a bijection between $H \times K$ and G .

To show that p is a homomorphism between the direct product and the product set HK we need to show that,

$$p((h, k)(h', k')) = p((hh', kk')) = hh'kk' = hkh'k' = p((h, k))p((h', k'))$$

which will be true if $h'k = kh'$ which, in turn, will be the case if products in HK are commutative.

Now we have $H \cap K = \{1\}$ and so,

$$hk = kh \iff k^{-1}hk = h \iff k^{-1}hkh^{-1} = 1$$

implies that $H \cap K = \{k^{-1}hkh^{-1}\}$. So, if we can show that $k^{-1}hkh^{-1}$ is in both H and K then we have a homomorphism.

Since, in this case, both H and K are normal we have,

$$h, h^{-1} \in H, k \in K \implies k^{-1}hk \in H, hkh^{-1} \in K$$

which, by the group closure of H and K gives,

$$(k^{-1}hk)h \in H \quad \text{and} \quad k^{-1}(hkh^{-1}) \in K.$$

Therefore, if both H and K are normal subgroups and $H \cap K = \{1\}$, then $hk = kh$ for all $h \in H, k \in K$. This, in turn, means that the product map p is a homomorphism between the direct product $H \times K$ and the product set HK . Since p is also bijective, it is an isomorphism.

□

It is important to note that the product map of two subgroups $H \times K \mapsto HK = G$ will not be a group homomorphism unless the two subgroups commute with each other.

2.1.8.4 Examples of Product Groups

- (27) There is a group with subgroups of orders $1 \dots 12$. It is a direct product of cyclic groups of orders

$$2 \times 2 \times 2 \times 3 \times 3 \times 5 \times 7 \times 11 = |G| = 27,720$$

so it looks like,

$$G = C_2 \times C_2 \times C_2 \times C_3 \times C_3 \times C_5 \times C_7 \times C_{11}.$$

2.1.8.5 Products of Cosets

It is possible to define a law of composition on the cosets of normal subgroups. This is because,

$$aH = Ha \implies aHbH = abHH = abH$$

so that we may define a law of composition such that $aH * bH = abH$ which closes over the set of cosets of H . The identity element of this composition is $eH = H$ and the inverse of the element aH is $a^{-1}H$.

*Note that this **only** applies to normal subgroups. The reason is that if H is not normal then there exists $h \in H$ and $a \in G$ such that $aha^{-1} \notin H$ which means that $S = aHa^{-1}H$ is not in any coset.*

This last claim can be proven if we observe – remembering that cosets partition the group – that S contains $a1a^{-1}1 = 1$ which means that it has to be in H . However, S also contains $aha^{-1}1 = aha^{-1} \notin H$ so, since these are equivalence classes, S cannot be in H .

2.1.8.6 Quotient Groups

Definition. Suppose H is a normal subgroup of a group G . Then the quotient group G/N is the set of cosets of N in G with the coset product. Its order is the index of N in G , $[G : N]$.

Notation. Sometimes – when it is not necessary to specify the subgroup against which the cosets are being formed – the set of cosets in G is denoted \overline{G} and a member coset aH is denoted $\overline{a} \in \overline{G}$.

Theorem 15. Every normal subgroup of a group G is the kernel of a homomorphism.

Proof. For any normal subgroup $N \leq G$, if we define the map,

$$\pi : G \longmapsto G/N$$

then π is homomorphic because $\pi(ab) = abN = aNbN = \pi(a)\pi(b)$.

Now, Proposition 24 tells us that

$$\pi(x) = 1N = N \iff x \in N$$

which implies that $\ker \pi = N$. (We could also observe that the cosets are equivalence classes and the kernel is the equivalence class containing the identity. The coset that contains the identity is the original subgroup $N = 1N$.)

□

Theorem 16. First Isomorphism Theorem: Let $\phi : G \mapsto G'$ be a surjective group homomorphism, and let $N = \ker \phi$. Then G/N is isomorphic to G' by the map $\bar{\phi}$ which sends the coset $\bar{a} = aN$ to $\phi(a)$,

$$\bar{\phi}(\bar{a}) = \phi(a).$$

Proof. The non-empty fibres of ϕ are the cosets aN as seen in the example 17. So, G/N can be thought of either as the cosets of the kernel of ϕ or as the non-empty fibres of ϕ . Then, $\bar{\phi}$ bijectively maps the cosets in G/N with the elements of the $\text{im } \phi$ and, because ϕ is surjective, we have $\text{im } \phi = G'$ so we have a bijection $\bar{\phi} : G/N \mapsto G'$.

Also, the map $\bar{\phi}$ is homomorphic because coset multiplication is consistent with multiplication in the group,

$$\bar{\phi}(\bar{ab}) = \phi(ab) = \phi(a)\phi(b) = \bar{\phi}(\bar{a})\bar{\phi}(\bar{b}).$$

□

2.1.8.7 Examples of Quotient Groups

- (28) Let $G = (\mathbb{Z}, +)$ and $H = \{4n \mid n \in \mathbb{Z}\}$. Then the cosets of H are $\{z + 4n \mid z \in \mathbb{Z}\}$ and the product of two cosets,

$$(z_1 + H) + (z_2 + H) = (z_1 + z_2 + H).$$

- (29) Let $G = (\mathbb{R}, +)$ and $H = \{2n\pi \mid n \in \mathbb{Z}\}$. Then the cosets are the possible angles. This is an example of an infinite quotient group. The affine spaces in example 19 are another example of an infinite quotient group.

- (30) In example 16 we saw that the modulus of complex numbers is a homomorphism from complex numbers to the reals. So, its kernel is the unit circle – the set of complex numbers of modulus 1. The cosets of the unit circle are the concentric circles,

$$C_r = \{z \mid |z| = r\}.$$

Applying the product of cosets gives us $C_r C_s = C_{rs}$ which works out

because,

$$\begin{aligned} |(a + bi)(c + di)| &= |(ac - bd) + (ad + bc)i| \\ \iff |(a + bi)(c + di)| &= \sqrt{(ac - bd)^2 + (ad + bc)^2} \\ \iff |(a + bi)(c + di)| &= \sqrt{(a^2c^2 + b^2d^2 - 2abcd) + (a^2d^2 + b^2c^2 + 2abcd)} \\ \iff |(a + bi)(c + di)| &= \sqrt{(a^2 + b^2)(c^2 + d^2)} \\ \iff |(a + bi)(c + di)| &= \sqrt{(a^2 + b^2)}\sqrt{(c^2 + d^2)} \\ \iff |(a + bi)(c + di)| &= |(a + bi)| |(c + di)|. \end{aligned}$$

[TODO: Notes on Quotient Groups: Artin\[81\]](#)

2.1.8.8 Modular Arithmetic

[TODO: Describe modular arithmetic in terms of cosets: Artin\[79\]](#) [TODO: include in examples Abstract Maths ex. 14.9](#)

2.2 Fields

Definition. A **field** F is a set together with two laws of composition, addition and multiplication, satisfying the following axioms:

- (i) Addition makes F into an abelian group F^+ (or $(F, +)$). Its identity element is denoted 0 .
- (ii) Multiplication is associative and commutative and makes $(F/\{0\}, \times)$ into a group. Its identity element is denoted 1 .
- (iii) Distributive law: For all $a, b, c \in F$, $(a + b)c = ac + bc$.

The distributive law establishes a relationship between the two laws of composition such that, for $a \in \mathbb{F}$,

$$a + a = 1a + 1a = (1 + 1)a.$$

In so doing, it establishes a relationship between the two groups: the additive group and the multiplicative group.

2.2.0.1 Subfields of \mathbb{C}

Definition. A field F is a **subfield of** \mathbb{C} if the following properties hold:

- If $a, b \in F$, then $a + b \in F$.
- If $a \in F$, then $-a \in F$.
- If $a, b \in F$, then $ab \in F$.
- If $a \in F$ and $a \neq 0$, then $a^{-1} \in F$.
- $1 \in F$.

Note that using the first, second and last of these axioms we can deduce that $1 - 1 = 0$ is an element of F .

*Also notice that addition on the field makes $(F, +)$ into an abelian group and multiplication makes $(F/\{0\}, \times)$ into an abelian group also. Conversely, any subset of F for which this is also true is a **subfield**.*

[TODO: Finite Fields \(Artin\[98\]\)](#)

2.2.1 Complex Numbers

Proposition 34. *For every $\alpha \in \mathbb{C}$, there exists a unique $\beta \in \mathbb{C}$ such that $\alpha + \beta = 0$.*

Proof. By contradiction: Say there are two such elements, β, γ such that,

$$\begin{aligned}\alpha + \beta &= 0 = \alpha + \gamma \\ (\alpha + \beta) + \beta &= (\alpha + \beta) + \gamma \\ 0 + \beta &= \beta = 0 + \gamma = \gamma\end{aligned}\quad \square$$

Proposition 35. *For every $\alpha \in \mathbb{C}$ with $\alpha \neq 0$, there exists a unique $\beta \in \mathbb{C}$ such that $\alpha\beta = 1$.*

Proof. By contradiction: Say there are two such elements, β, γ then,

$$\begin{aligned}\alpha\beta &= 1 = \alpha\gamma \\ \beta &= \frac{1}{\alpha} = \gamma\end{aligned}\quad \square$$

2.2.2 Complex Numbers Problems

Find all the roots of $x^3 = 1$ for $x \in \mathbb{C}$ complex numbers Since $x^3 - 1 = (x - 1)(x^2 + x + 1)$, we have (via zero-factor theorem) possible roots from,

$$x - 1 = 0 \iff x = 1$$

$$x^2 + x + 1 = 0 \implies x = \frac{-1 \pm \sqrt{-3}}{2} = \frac{-1 \pm \sqrt{3}i}{2}$$

More generally,

$$(a + bi) + (a - bi) = 2a$$

and since also,

$$\left[\frac{-1 + \sqrt{3}i}{2} \right]^2 = \frac{-1 - \sqrt{3}i}{2}$$

as well as the reverse,

$$\left[\frac{-1 - \sqrt{3}i}{2} \right]^2 = \frac{-1 + \sqrt{3}i}{2}$$

this means that if $x = \frac{-1 \pm \sqrt{3}i}{2}$ then $x^2 + x$ is of the form $(a + bi) + (a - bi) = 2a$ and so we have that $x^2 + x = -1 \iff x^2 + x + 1 = 0$.

In addition,

$$(a + bi)(a - bi) = a^2 + b^2$$

which means that if $x = \frac{-1 \pm \sqrt{3}i}{2}$ then $x^3 = x^2 x$ is of the form $(a + bi)(a - bi) = a^2 + b^2$ so we have that $x^3 = \frac{-1}{2}^2 + \frac{\sqrt{3}}{2}^2 = \frac{1}{4} + \frac{3}{4} = 1$.

So we see that - allowing for complex x - the cubic polynomial $x^3 - 1$ has 3 roots as we should expect from the Fundamental Theorem of Algebra (*is this the correct interpretation of this?*).

2.3 Matrices

2.3.1 Basic Algebra of Matrices

Definition. Matrix **equality** is defined component-wise so that if $A = B$ then A and B must have the same dimension as well as equal values in each component.

Definition. An **identity** element e is defined as $ea = ae = a$.

The definition of an identity element above is in any context (not just for matrices). For matrices this has certain consequences.

Proposition 36. *Identity matrices must be square*

Proof. For a matrix A and an identity matrix I , $AI = IA = A$ which means that AI , IA and A must all have the same dimensions. If A is of dimension $m \times n$ then I must have dimension $n \times m$ but then AI has dimension $m \times m$ while IA has dimension $n \times n$. We conclude that $m = n$ and both matrices are square. \square

If A, B, C are matrices s.t. $AB = AC$, can we, in general, conclude that $B = C$?

The answer is no, as the following example shows:

$$A = \begin{pmatrix} 0 & 0 \\ 1 & 1 \end{pmatrix}, \quad B = \begin{pmatrix} 1 & -1 \\ 3 & 5 \end{pmatrix}, \quad C = \begin{pmatrix} 8 & 0 \\ -4 & 4 \end{pmatrix}$$

$$A = B = \begin{pmatrix} 0 & 0 \\ 4 & 4 \end{pmatrix}$$

This is because multiplication by A has no inverse (i.e. it's not a bijection and A^{-1} does not exist) as we can see by the fact that $|A| = 0$.

If A, B, C are matrices s.t. $A + 5B = A + 5C$, can we, in general, conclude that $B = C$?

The answer is yes because the matrix addition and scalar multiplication always have inverses. The inverse of $+A$ is $-A$ and the inverse of scalar multiplication by 5 is scalar multiplication by $\frac{1}{5}$. So we can say,

$$\begin{aligned}
 & A + 5B = A + 5C \\
 \iff & A + 5B - A = A + 5C - A \\
 \iff & 5B = 5C \\
 \iff & \left(\frac{1}{5}\right) 5B = \left(\frac{1}{5}\right) 5C \\
 \iff & B = C
 \end{aligned}$$

2.3.1.1 Matrix multiplication

Multiplication of matrices proceeds as a collection of dot-products of individual vectors. As a result, its properties are largely dependent on the properties of the dot-product (see: 79).

Matrix multiplication treats the two operand matrices as collections of vectors with the first matrix having the vectors as rows and the second having the vectors as columns.

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} e & f \\ g & h \end{bmatrix} = \begin{bmatrix} ae + bg & af + bh \\ ce + dg & cf + dh \end{bmatrix}$$

This difference in orientation of the vectors in the two operands results in the multiplication not being commutative - the order matters. So, the first property of the dot-product is not preserved but the others are preserved

(albeit with a slight modification for the last one).

$$\begin{aligned}\alpha \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} e & f \\ g & h \end{bmatrix} &= \begin{bmatrix} \alpha a & \alpha b \\ \alpha c & \alpha d \end{bmatrix} \begin{bmatrix} e & f \\ g & h \end{bmatrix} = \begin{bmatrix} \alpha(ae + bg) & \alpha(af + bh) \\ \alpha(ce + dg) & \alpha(cf + dh) \end{bmatrix} \\ &= \alpha \begin{bmatrix} ae + bg & af + bh \\ ce + dg & cf + dh \end{bmatrix}\end{aligned}$$

$$\begin{aligned}\left(\begin{bmatrix} a & b \\ c & d \end{bmatrix} + \begin{bmatrix} e & f \\ g & h \end{bmatrix} \right) \begin{bmatrix} i & j \\ k & l \end{bmatrix} &= \begin{bmatrix} a+e & b+f \\ c+g & d+h \end{bmatrix} \begin{bmatrix} i & j \\ k & l \end{bmatrix} \\ &= \begin{bmatrix} i(a+e) + k(b+f) & j(a+e) + l(b+f) \\ i(c+g) + k(d+h) & j(c+g) + l(d+h) \end{bmatrix} \\ &= \begin{bmatrix} ia + kb & ja + lb \\ ic + kd & jc + ld \end{bmatrix} + \begin{bmatrix} ie + kf & je + lf \\ ig + kh & jg + lh \end{bmatrix} \\ &= \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} i & j \\ k & l \end{bmatrix} + \begin{bmatrix} e & f \\ g & h \end{bmatrix} \begin{bmatrix} i & j \\ k & l \end{bmatrix}\end{aligned}$$

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} a & b \\ c & d \end{bmatrix}^T = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} a & c \\ b & d \end{bmatrix} = \begin{bmatrix} a^2 + b^2 & ac + bd \\ ac + bd & c^2 + d^2 \end{bmatrix}$$

So, to summarize:

If A, B, C are matrices and α is a scalar then,

- $\alpha AB = (\alpha A)B = A(\alpha B) = \alpha(AB)$
- $(A + B)C = C(A + B) = AC + BC$
- AA^T is a symmetric matrix with positive values along the diagonal

2.3.1.2 Block Matrix multiplication

In the description of matrix multiplication as

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} e & f \\ g & h \end{bmatrix} = \begin{bmatrix} ae + bg & af + bh \\ ce + dg & cf + dh \end{bmatrix}$$

a, b, \dots, h could also be blocks of matrices. In which case, the resulting calculations — $ae + bg$ etc. — refer to matrix multiplication ae where a and e are treated as matrices and must have compatible dimensions. This can be seen as follows.

$$\begin{aligned} \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \end{bmatrix} &= \begin{bmatrix} a_{11}b_{11} + a_{12}b_{21} & \cdots \\ a_{21}b_{11} + a_{22}b_{21} & \cdots \end{bmatrix} \\ \\ \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \\ b_{31} & b_{32} \end{bmatrix} &= \begin{bmatrix} a_{11}b_{11} + a_{12}b_{21} + a_{13}b_{31} & \cdots \\ \vdots & \end{bmatrix} \\ &= \begin{bmatrix} (a_{11}b_{11} + a_{12}b_{21}) + a_{13}b_{31} & \cdots \\ \vdots & \end{bmatrix} \\ &= \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \end{bmatrix} + \begin{bmatrix} a_{13} \\ a_{23} \end{bmatrix} \begin{bmatrix} b_{31} & b_{32} \end{bmatrix}. \end{aligned}$$

So we have, for example, working in blocks,

$$\begin{aligned} \begin{bmatrix} A & B \\ C & D \end{bmatrix} \begin{bmatrix} E & F \\ G & H \end{bmatrix} &= \begin{bmatrix} AE + BG & AF + BH \\ CE + DG & CF + DH \end{bmatrix}, \\ \begin{bmatrix} A & B \\ C & D \end{bmatrix} \begin{bmatrix} E & F \\ G & H \end{bmatrix} = \begin{bmatrix} I & 0 \\ 0 & I \end{bmatrix} &\implies AF = -BH, CE = -DG, AE + BG = CF + DH = I, \\ \begin{bmatrix} A \\ B \end{bmatrix} \begin{bmatrix} C & D \end{bmatrix} = \begin{bmatrix} I & E \\ F & G \end{bmatrix} &\implies AC = I \iff C = A^{-1}. \end{aligned}$$

2.3.1.3 Matrix transpose

Proposition 37. $(AB)^T = B^T A^T$

Proof. Denote the i th row of the matrix A as $A[i :]$ and the j th column of the matrix B as $B[:, j]$ and a matrix whose components at (i, j) are the dot-products of the i th row of the matrix A with the j th column of the matrix

B as $(\langle A[i:], B[:j] \rangle)$. Then,

$$\begin{aligned}(AB)^T &= (\langle A[i:], B[:j] \rangle)^T = (\langle A[j:], B[:i] \rangle) \\ B^T A^T &= (\langle B^T[i:], A^T[:j] \rangle) = (\langle B[:i], A[j:] \rangle)\end{aligned}$$

So, by commutativity of dot-product, $(AB)^T = B^T A^T$. \square

Proposition 38. $(A^T)^{-1} = (A^{-1})^T$

Proof.

$$\begin{aligned}I &= AA^{-1} = (AA^{-1})^T = (A^{-1})^T A^T \\ \iff I(A^T)^{-1} &= (A^{-1})^T A^T (A^T)^{-1} \\ \iff (A^T)^{-1} &= (A^{-1})^T.\end{aligned}\quad \square$$

Proposition 39. $(A + B)^T = A^T + B^T$

Proof. In $A + B$ the (i, j) element is $A_{ij} + B_{ij}$ so the (i, j) element of $(A + B)^T$ is $A_{ji} + B_{ji}$ which clearly is also the (i, j) element of $A^T + B^T$. \square

2.3.1.4 Matrix inverse

Definition. *Inverse property is: If there exists a matrix B such that $AB = BA = I$ then B is the **inverse** of A and A is the inverse of B .*

This definition is inherently bound up with the definition of the identity (\exists a matrix I s.t. $AI = IA = A$) and both define the identity and inverse elements as commutatively producing their result under matrix multiplication. Since matrix multiplication is not, in general, commutative there is no guarantee that if $AB = I$ then $BA = I$. An example of this failing is,

$$\begin{aligned}A &= \begin{bmatrix} 1 & 2 \end{bmatrix}, B = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \\ AB &= \begin{bmatrix} 1 \end{bmatrix} = I_1, BA = \begin{bmatrix} 1 & 2 \\ 0 & 0 \end{bmatrix} \neq I_2\end{aligned}$$

But we could have guessed this because Proposition 36 showed that identity matrices must be square and its product with a matrix must be defined from both the left and the right, i.e. $IA = AI = A$ meaning that the matrix A must have the same dimensions as I . So, for non-square matrices, no identity can exist. If there is no identity, then the inverse is not defined either.

Proposition 40. *If the inverses of the matrices A and B both exist then so does the inverse of the product AB and it is equal to $B^{-1}A^{-1}$.*

Proof.

$$\begin{aligned}
& (AB)(AB)^{-1} = I \\
\iff & (A^{-1}A)B(AB)^{-1} = A^{-1}I \\
\iff & (B^{-1}B)(AB)^{-1} = B^{-1}A^{-1} \\
\iff & (AB)^{-1} = B^{-1}A^{-1}
\end{aligned}$$

and since B^{-1} and A^{-1} both exist then their product exists. Furthermore, this holds for a product of any finite sequence of invertible matrices $A_1A_2 \cdots A_n$ which can easily be shown by induction on the associative product. \square

2.3.2 Basic properties of Matrices

Definition. If a_{ij} is an entry of a matrix in the i th row and j th column then the **main diagonal** of the matrix is the collection of entries a_{ij} with $i = j$.

The main diagonal is most often spoken of with respect to square matrices but the definition does not require that the matrix be square (wikipedia).

2.3.2.1 Trace

Definition. The **trace** of a matrix is the sum of the diagonal entries.

2.3.2.2 Symmetric Matrices

Definition. A **symmetric** matrix is a matrix that is invariant under transposition.

This obviously requires that the matrix be square and that the upper off-diagonal elements mirror the lower off-diagonal elements.

2.3.2.3 Diagonal Matrices

Definition. A **diagonal** matrix is a matrix whose entries outside the main diagonal are all zero.

- The definition of a diagonal matrix does **not** require that all the diagonal entries are nonzero so, for example, the zero matrix is a diagonal matrix.
- Diagonal matrices are **usually understood to be** a subset of symmetric matrices but this is not strictly necessary according to their definition (wikipedia).

Proposition 41. If $A = a_{ij}$, $B = b_{ij}$ are square diagonal matrices then multiplying them results in a square diagonal matrix $C = c_{ij}$ whose diagonal entries c_{11}, \dots, c_{nn} are the products of the corresponding diagonal entries in A and B . That's to say

$$c_{11}, \dots, c_{nn} = a_{11}b_{11}, \dots, a_{nn}b_{nn}.$$

Proof. If $A = a_{ij}$, $B = b_{ij}$ are square diagonal matrices then,

$$AB = \begin{bmatrix} a_{11} & 0 & \cdots & 0 \\ 0 & a_{22} & \cdots & 0 \\ \vdots & & & \\ 0 & 0 & \cdots & a_{nn} \end{bmatrix} \begin{bmatrix} b_{11} & 0 & \cdots & 0 \\ 0 & b_{22} & \cdots & 0 \\ \vdots & & & \\ 0 & 0 & \cdots & b_{nn} \end{bmatrix} = \begin{bmatrix} a_{11}b_{11} & 0 & \cdots & 0 \\ 0 & a_{22}b_{22} & \cdots & 0 \\ \vdots & & & \\ 0 & 0 & \cdots & a_{nn}b_{nn} \end{bmatrix}.$$

□

Corollary 11. Multiplication of square diagonal matrices is commutative.

Proof. The multiplication described in Proposition 41 is commutative because the diagonal entries of the two matrices are the same whichever way around the multiplication is performed and the multiplication of the individual entries itself is commutative. □

Corollary 12. Square diagonal matrices with nonzero diagonal entries are invertible.

Proof. A result of the multiplication described in Proposition 41 is that, if A is a square diagonal matrix with nonzero diagonal entries, then we can obtain the identity matrix by multiplication with a matrix B whose diagonal

elements are the reciprocal of the corresponding entries in A . So, if the jj th entry in A is a_{jj} , then the corresponding entry in B is $1/a_{jj}$. For example,

$$AB = \begin{bmatrix} a_{11} & 0 & \cdots & 0 \\ 0 & a_{22} & \cdots & 0 \\ \vdots & & & \\ 0 & 0 & \cdots & a_{nn} \end{bmatrix} \begin{bmatrix} 1/a_{11} & 0 & \cdots & 0 \\ 0 & 1/a_{22} & \cdots & 0 \\ \vdots & & & \\ 0 & 0 & \cdots & 1/a_{nn} \end{bmatrix} = \begin{bmatrix} 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & & & \\ 0 & 0 & \cdots & 1 \end{bmatrix}.$$

Furthermore, by 11, this multiplication is commutative so we have,

$$AB = BA = I.$$

Therefore $B = A^{-1}$. Furthermore, this matrix always exists as long as the entries of the matrix are drawn from a field (so that they have multiplicative inverses) and the diagonal entries are nonzero. \square

Proposition 42. *Square diagonal matrices with nonzero diagonal entries form an abelian group under multiplication.*

Proof. Let S be the set of $n \times n$ diagonal matrices with nonzero diagonal entries. Then,

- $I_n \in S$ so S is nonempty.
- By Proposition 41, if A, B are members of S , then AB is also a square diagonal matrix with nonzero diagonal entries and so is also a member of S .
- By 12 members of S are invertible and are members of S .

Therefore, $S \leq GL_n(\mathbb{F})$. \square

2.3.3 Matrices as linear transformations

2.3.3.1 Multiplying a vector by a matrix on the left: $A\vec{x} = \vec{y}$

Left multiplication of a matrix A of dimension $m \times n$ on a column vector \vec{x} of dimension $n \times 1$ transforms it to a column vector \vec{y} of dimension $m \times 1$.

$$\begin{bmatrix} a_{11} & \cdots & a_{1n} \\ \vdots & \vdots & \vdots \\ a_{m1} & \cdots & a_{mn} \end{bmatrix} \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix}$$

This can be thought of a function from the space of n -dimensional vectors from which \vec{x} is drawn to the space of m -dimensional vectors in which \vec{y} resides. So, for real-valued vectors, the function would be a function $f : \mathbb{R}^n \mapsto \mathbb{R}^m$ such that,

$$f(x_1, \dots, x_n) = \vec{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix}$$

Or else, this could be thought of as m n -ary functions of the form $f : \mathbb{R}^n \mapsto \mathbb{R}$,

$$\begin{aligned} f_1(x_1, \dots, x_n) &= a_{11}x_1 + \cdots + a_{1n}x_n = y_1 \\ &\vdots \\ f_m(x_1, \dots, x_n) &= a_{m1}x_1 + \cdots + a_{mn}x_n = y_m \end{aligned}$$

In this case, each row of the matrix is a real-valued function in n variables. Each of these functions is *homogenous linear* (a function of the form $a_1x_1 + \cdots + a_kx_k + c$ for scalars a_1, \dots, a_k, c and $c = 0$) and so the system of functions is called a *linear transformation*.

2.3.3.2 Multiplying a matrix of vectors by a matrix on the left: $AX = Y$


Looking at the matrix as a linear transformation from one co-ordinate space to another, consider $AX = Y$ where X is a matrix - which may be considered

a collection of vectors - transformed by the matrix A into the matrix - or collection of vectors - Y .

We transform the unit square in the source space, X in \mathbb{R}^2 , using the 2D transformation matrix A , into its image in the destination space, Y in \mathbb{R}^2 .

$$A = \begin{bmatrix} 3 & 2 \\ 1 & 4 \end{bmatrix}, X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$

$$AX = Y = \begin{bmatrix} 0 & 3 & 5 & 2 \\ 0 & 1 & 5 & 4 \end{bmatrix}$$



2.3.3.3 Change of Co-ordinates

If a matrix A has columns comprised of the axes of a co-ordinate system, then \vec{x} defined against the standard basis is, in the other system, $A^{-1}\vec{x}$. This is because the axes in A are defined relative to the standard basis axes.

Therefore, if \vec{x}_A is a vector defined against the axes in A , then the same vector against the standard basis axes \vec{x} would be $A\vec{x}_A = \vec{x}$ and so $\vec{x}_A = A^{-1}\vec{x}$.

Proposition 43. *Let B be a basis of a vector space V . The other bases are the sets of the form $B' = BP^{-1}$, where $P \in GL_n(\mathbb{F})$ is an invertible matrix.*

Proof. □

2.3.3.4 Types of Transformations

There are 3 basic types of transformation:

- **Rigid body** - preserves distances and angles.

Examples: translation and rotation.

- **Conformal** - preserves angles.

Examples: translation, rotation and uniform scaling.

- **Affine** - preserves parallelism.


Examples: translation, rotation, uniform and non-uniform scaling, shearing and reflection.

2.3.3.5 Rigid Body

Translation So as to perform the translation as multiplication by a transformation matrix we take the approach of homogeneous coordinates (see: https://en.wikipedia.org/wiki/Homogeneous_coordinates) so we form matrix with the identity in the first two columns and then a third column with the translation vector. Then, we add a row of ones to the vectors we will translate and the output vectors also have a 1 in the third row that is ignored.

$$A = \begin{bmatrix} 1 & 0 & 3 \\ 0 & 1 & 4 \\ 0 & 0 & 1 \end{bmatrix}, \quad X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{bmatrix}$$


$$AX = Y = \begin{bmatrix} 3 & 4 & 4 & 3 \\ 4 & 4 & 5 & 5 \\ 1 & 1 & 1 & 1 \end{bmatrix}$$



Rotation

$$A = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}, \quad X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$

$$AX = Y = \begin{bmatrix} 0 & 0.5253 & -0.3256 & -0.8509 \\ 0 & 0.8509 & 1.3762 & 0.5253 \end{bmatrix}$$




2.3.3.6 Conformal

Uniform Scaling is scaling by an equal amount in each dimension.

$$A = \begin{bmatrix} 3 & 0 \\ 0 & 3 \end{bmatrix}, \quad X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$

$$AX = Y = \begin{bmatrix} 0 & 3 & 3 & 0 \\ 0 & 0 & 3 & 3 \end{bmatrix}$$




2.3.3.7 Affine

Non-uniform Scaling is scaling by different amounts in the different dimensions. (The example shown here preserves the angles but for other shapes, a triangle for example, the angles would not be preserved.)

$$A = \begin{bmatrix} 2 & 0 \\ 0 & 3 \end{bmatrix}, \quad X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$


$$AX = Y = \begin{bmatrix} 0 & 2 & 2 & 0 \\ 0 & 0 & 3 & 3 \end{bmatrix}$$



Shearing


$$A = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}, \quad X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$

$$AX = Y = \begin{bmatrix} 0 & 1 & 2 & 1 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$



$$A = \begin{bmatrix} 1 & 1 \\ 0.2 & 1 \end{bmatrix}, \quad X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$


$$AX = Y = \begin{bmatrix} 0 & 1 & 2 & 1 \\ 0 & 0.2 & 1.2 & 1 \end{bmatrix}$$



Reflection

$$A = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}, \quad X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$

$$AX = Y = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & -1 & -1 \end{bmatrix}$$



2.3.4 Elementary Matrices and Row Operations

Notation. The **matrix units** - matrices with a single non-zero component whose value is 1 are traditionally named e_{ij} where i, j is the matrix co-ordinate of the 1.

An arbitrary matrix $A = (a_{ij})$ may be expressed as a sum of such unit matrices as $A = a_{11}e_{11} + \cdots + a_{nn}e_{nn}$.

$$e_{ij} = \begin{bmatrix} \cdots & \cdots & \cdots \\ \vdots & \vdots & \vdots \\ \cdots & 1 & \cdots \\ \vdots & \vdots & \vdots \\ \cdots & \cdots & \cdots \end{bmatrix}$$

So matrix units can be used to analyse matrix addition but to analyse matrix multiplication some square matrices called **elementary matrices** are more useful.

Multiplying a matrix from the left (so doing row operations), there are 3 types of elementary matrix:

Adding rows: $I + ae_{ij}$ for $i \neq j$

$$\begin{bmatrix} 1 & & & \\ & \ddots & a & \\ & & \ddots & \\ & & & 1 \end{bmatrix}$$

This adds a times some row to another row.

Swapping rows: $I + e_{ij} + e_{ji} - e_{ii} - e_{jj}$ for $i \neq j$

$$\begin{bmatrix} 1 & & & \\ & 0 & 1 & \\ & 1 & 0 & \\ & & & \ddots \end{bmatrix}$$

This swaps the rows i and j .

Scalar-multiplying a row: $I + (c - 1)e_{ii}$ for $c \neq 0$

$$\begin{bmatrix} 1 & & & \\ & \ddots & & \\ & & c & \\ & & & 1 \end{bmatrix}$$

This multiplies row i by c .

Proposition 44. *Elementary matrices are invertible and their inverses are also elementary matrices.*

Proof. Proceed by cases on the 3 elementary types of elementary matrices.

Case $I + ae_{ij}$ If R_i is row i and R_j is row j , then this matrix performs $R_i + aR_j$. Clearly this can be "undone" by performing $R_i - aR_j$. So the matrix, $I - ae_{ij}$ is the inverse and clearly this is also an elementary matrix of the same type.

Case $I - e_{ii} - e_{jj} + e_{ij} + e_{ji}$ This matrix swaps 2 rows in a permutation that is its own inverse.

Case $I + (c - 1)e_{ii}$ This matrix performs cR_i and so it is "undone" by performing $c^{-1}R_i$ (which for a real-valued matrix would be $\left(\frac{1}{c}\right)R_i$) and this inverse matrix is also an elementary matrix of the same type. \square

Proposition 45. *Suppose $AX = B$ and a series of elementary row operations on $[A \mid B]$ produces $[A' \mid B']$, then the solutions of $A'X = B'$ are the same as those of $AX = B$.*

Proof. First note that the series of elementary row operations is described as multiplication on the left by a series of elementary matrices say, E_1, E_2, \dots, E_n so that,

$$[A' \mid B'] = [(E_n \cdots E_2 E_1)A \mid (E_n \cdots E_2 E_1)B]$$

Now, let $(E_n \cdots E_2 E_1) = E$ and notice that, since each of the individual E_i is invertible the product of them is also invertible by Proposition 40 so,

$$A'X = B' \iff EAX = EB$$

and the existence of the inverse E^{-1} means that the law of cancellation is in effect so,

$$EAX = EB \iff AX = B$$

$$\therefore A'X = B' \iff AX = B$$

□

Proposition 46. *Let A be a square matrix. The following conditions are equivalent:*

- *A can be reduced to the identity by a sequence of elementary row operations.*
- *A is a product of elementary matrices.*
- *A is invertible.*
- *The system of homogeneous equations $AX = 0$ has only the trivial solution $X = 0$.*

Proof. If A can be reduced to the identity by a sequence of elementary row operations then,

$$(E_n \cdots E_2 E_1)A = I$$

and by Proposition 44 and Proposition 40 the matrix $(E_n \cdots E_2 E_1)$ is invertible so,

$$A = (E_n \cdots E_2 E_1)^{-1}I = (E_n \cdots E_2 E_1)^{-1} = E_1^{-1}E_2^{-1} \cdots E_n^{-1}$$

and, also by Proposition 44, A is a product of elementary matrices and is invertible.

Furthermore, if $AX = 0$ then $X = A^{-1}0 = 0$ - i.e. the only solution to $AX = 0$ is $X = 0$. □

2.3.5 Determinants

1×1

The determinant of a 1×1 matrix is just its unique component entry,

$$\det [a] = a$$

2×2

The determinant of a 2×2 matrix is given by the formula,


$$\det \begin{bmatrix} a & b \\ c & d \end{bmatrix} = ad - bc$$

Returning to our example of a 2d operator:

We transform the unit square in the source space, X in \mathbb{R}^2 , using the 2D transformation matrix A , into its image in the destination space, Y in \mathbb{R}^2 .

$$A = \begin{bmatrix} 3 & 2 \\ 1 & 4 \end{bmatrix}, \quad X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$

$$AX = Y = \begin{bmatrix} 0 & 3 & 5 & 2 \\ 0 & 1 & 5 & 4 \end{bmatrix}$$



We see that $\det A = 10$ and the parallelogram, Y , that is the image of the unit square, X , under the transformation represented by A has area,


$$\text{area} = b \cdot h = |\langle 3, 1 \rangle| \cdot |\langle 2 - 3, 4 - 1 \rangle| = \sqrt{10} \cdot \sqrt{10} = 10$$

And the determinant would be 0 in the case that the columns were proportional (representing co-linear vectors) and the determinant would be negative if the orientation of the output vectors were reversed w.r.t. the input vectors. So, if we swap either the columns or the rows of the transformation matrix, A , the determinant comes out -10 .

Swapping the columns:

$$A_c = \begin{bmatrix} 2 & 3 \\ 4 & 1 \end{bmatrix}, X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$

$$AX = Y = \begin{bmatrix} 0 & 2 & 5 & 3 \\ 0 & 4 & 5 & 1 \end{bmatrix}$$




Note that the result looks exactly the same - it's just that now the x-vector $\langle 1, 0 \rangle$, produces $\langle 4, 2 \rangle$ and the y-vector, $\langle 0, 1 \rangle$ produces $\langle 3, 1 \rangle$.

Swapping the rows:

$$A_r = \begin{bmatrix} 1 & 4 \\ 3 & 2 \end{bmatrix}, X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$

$$AX = Y = \begin{bmatrix} 0 & 1 & 5 & 4 \\ 0 & 3 & 5 & 2 \end{bmatrix}$$



Note that if we swap **both** the columns and the rows then we get back to a transformation with determinant 10.

$$A_{rc} = \begin{bmatrix} 4 & 1 \\ 2 & 3 \end{bmatrix}, \quad X = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$

$$AX = Y = \begin{bmatrix} 0 & 4 & 5 & 1 \\ 0 & 2 & 5 & 3 \end{bmatrix}$$

which produces the same parallelogram as the previous one but with columns reversed.

Summary So we find that,

$$\begin{bmatrix} 3 & 2 \\ 1 & 4 \end{bmatrix}, \begin{bmatrix} 4 & 1 \\ 2 & 3 \end{bmatrix} \text{ have determinant } > 0$$

$$\begin{bmatrix} 2 & 3 \\ 4 & 1 \end{bmatrix}, \begin{bmatrix} 1 & 4 \\ 3 & 2 \end{bmatrix} \text{ have determinant } < 0$$

If the product of the components on the diagonal of the matrix is greater than the components on the bottom-left to top-right diagonal then the determinant is > 0 , if the reverse is true then the determinant is < 0 , and if they are equal then the determinant $= 0$.

Note that, for the determinant to be 0 in our example, we need something like $\det A = (4 \times 3) - (4 \times 3)$ which, due to the commutativity of multiplication can be achieved by both,

$$\begin{bmatrix} 4 & 3 \\ 4 & 3 \end{bmatrix}, \begin{bmatrix} 4 & 4 \\ 3 & 3 \end{bmatrix}$$

but in both cases the columns are proportional and therefore co-linear.

$n \times n$

The determinant of an $n \times n$ matrix is defined recursively as:

- if $n = 1$ then $\det A = a_{11}$, i.e. the determinant is equal to the sole component.
- else if $n > 1$ then, defining A_{ij} as the matrix formed by leaving out the i th row and the j th column,

$$\det A = a_{11}\det A_{11} - a_{12}\det A_{12} + \cdots \pm a_{1n}\det A_{1n}$$

In the $n = 2$ case, each $\det A_{ij}$ has $n = 1$ and so is simply equal to the sole component that is neither on the same row or column as the component a_{ij} that is multiplying it. This feature of the determinant calculation continues recursively for higher dimension matrices so that the calculation is always comprised of terms that are a product of components on each of the different

columns and rows. In fact, it comprises the products of all such possible combinations of components.

For example, when $n = 2$ the only combinations are,

$$\{a_{11}, a_{22}\} \text{ and } \{a_{12}, a_{21}\}$$

so there are only 2 terms in the determinant calculation.

When $n = 3$ the possible combinations are,

$$\begin{aligned} &\{a_{11}, a_{22}, a_{33}\}, \{a_{11}, a_{32}, a_{23}\}, \\ &\{a_{12}, a_{21}, a_{33}\}, \{a_{12}, a_{31}, a_{23}\}, \\ &\{a_{13}, a_{21}, a_{32}\}, \{a_{13}, a_{31}, a_{22}\} \end{aligned}$$

so there are 6 terms in the determinant calculation. Notice that each term is generated by a different permutation of the columns while holding the rows fixed in ascending order and that the sign of each term is governed by how many permutations the permutation of columns is away from ascending order, $1, 2, \dots, n$.

$$\begin{aligned} \begin{vmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{vmatrix} &= a_{11}(a_{22}a_{33} - a_{32}a_{23}) - a_{12}(a_{21}a_{33} - a_{31}a_{23}) + a_{13}(a_{21}a_{32} - a_{31}a_{22}) \\ &= a_{11}a_{22}a_{33} - a_{11}a_{23}a_{32} - a_{12}a_{21}a_{33} + a_{12}a_{23}a_{31} + a_{13}a_{21}a_{32} - a_{13}a_{22}a_{31} \end{aligned}$$

2.3.5.1 Consequences

From this feature of the calculation we can see a number of the important properties of the determinant.

Proposition 47. *det I = 1*

Proof. Whatever the dimension of the identity matrix there will be only one combination of rows and columns that is the diagonal along which the 1s of the identity matrix reside. So, there will be a single term of the determinant calculation that is a product of 1s and all other terms will contain at 0s. In addition, the term that is along the diagonal has a positive sign in the determinant calculation. Therefore the result is 1. \square

Proposition 48. *det A is linear in the rows of the matrix*

Proof. If p and q are row vectors and we have matrices A_p, A_q, A_{pq} in which are present, respectively, the row vector p , the one q , and the row $p + q$, then linearity implies that $\det A_{pq} = \det A_p + \det A_q$. This can be seen since every term of the determinant calculation of A_{pq} will contain one of the components in the row $p + q$. So each term of the calculation will take the form,

$$(p + q)a_{ij} \cdots a_{mn} = p(a_{ij} \cdots a_{mn}) + q(a_{ij} \cdots a_{mn})$$

The other implication of linearity is that - if a row is multiplied by a scalar, c , to produce A_c then $\det A_c = c \det A$. Using a similar reasoning to the previous argument we have each term of the determinant taking the form,

$$c a_{ij} \cdots a_{mn}$$

which obviously results in the determinant being multiplied by c . \square

Proposition 49. *If two columns are exchanged in the matrix then the determinant is multiplied by -1*

Proof. If columns p and q are exchanged then the components of p and q appear in terms with signs reversed. Since the components of p and q appear in every term of the determinant calculation, every term has the sign reversed. So the determinant is multiplied by -1 . \square

Proposition 50. *$\det A = 0$ if there are two identical columns in the matrix*

Proof. If column p is identical to column q then we can swap columns p and q and we will have the same matrix so the determinant must also remain the same. But Proposition 49 proved that swapping two columns causes the determinant to be multiplied by -1 . So, if A_{pq} is the matrix A after swapping the columns,

$$\det A_{pq} = \det A = -\det A \iff \det A = 0$$

\square

Proposition 51. *Adding a multiple of one column to another leaves the determinant unchanged*

Proof. By combining Proposition 48 and Proposition 50 we find that if the columns of A are,

$$\vec{x}_1, \vec{x}_2, \cdots, \vec{x}_p, \vec{x}_q, \cdots, \vec{x}_n$$

and the columns of A_c are,

$$\vec{x}_1, \vec{x}_2, \dots, \vec{x}_p + c\vec{x}_q, \vec{x}_q, \dots, \vec{x}_n$$

then,

$$\begin{aligned} \det A_c &= \det(\vec{x}_1, \vec{x}_2, \dots, \vec{x}_p, \vec{x}_q, \dots, \vec{x}_n) + c \cdot \det(\vec{x}_1, \vec{x}_2, \dots, \vec{x}_q, \vec{x}_q, \dots, \vec{x}_n) \\ &= \det(\vec{x}_1, \vec{x}_2, \dots, \vec{x}_p, \vec{x}_q, \dots, \vec{x}_n) + c \cdot 0 \\ &= \det A \end{aligned}$$

□

2.3.5.2 Better formulation (from Rudin's Principles of Mathematical Analysis)

Let $a(i, j)$ be the component in the i th row and j th column of the matrix A and,

$$\text{sign}(x) = \begin{cases} -1 & x < 0 \\ 0 & x = 0 \\ 1 & x > 0 \end{cases}$$

$$s(j_1, \dots, j_n) = \prod_{p < q} \text{sign}(j_q - j_p)$$

Then the determinant,

$$\det A = \sum s(j_1, \dots, j_n) a(1, j_1) \cdots a(n, j_n)$$

defined over all n -tuples of n distinct values, j_1, \dots, j_n with $1 \leq j_r \leq n$ (i.e., permutations of $[1, n] \subset \mathbb{N}$) with each term being produced by a different permutation.

From this we can see that,

- **The determinant of the identity matrix is 1**

Every term of the determinant will contain at least one 0 apart from the term that traverses the main diagonal, which is all 1s. We can see that there is only one such term because the main diagonal has $i = j$ and so there is only one such j_1, \dots, j_n that satisfies this.

- **The determinant is linear in the rows or columns of the matrix, holding the others constant**

If a column, j_r , is multiplied by a scalar α and another column, j_k , is added to it, then the resulting determinant takes the form,

$$\begin{aligned} \det A &= \sum_i s(j_1, \dots, j_n) a(i, j_1) \cdots (\alpha a(i, j_r) + a(i, j_k)) \cdots a(i, j_n) \\ \iff \det A &= \alpha a(i, j_r) \sum_i s(j_1, \dots, j_n) a(i, j_1) \cdots a(i, j_n) + \\ &\quad a(i, j_k) \sum_i s(j_1, \dots, j_n) a(i, j_1) \cdots a(i, j_n) \end{aligned}$$

- **If two columns are exchanged then the determinant is multiplied by -1**

If columns p and q are exchanged then this is equivalent to swapping j_p and j_q in the n -tuple so that $s(j_1, j_2, \dots, j_n)$ changes sign and so the determinant is multiplied by -1 .

- **If two columns are equal then the determinant will be 0**

If two columns are the same then this is equivalent to a repetition of a value in the tuple j_1, \dots, j_n and so,

$$\exists p, q \text{ s.t. } \text{sign}(j_q - j_p) = 0 \implies s(j_1, \dots, j_n) = 0$$

which results in every term of the determinant being 0.

This can also be proven by using the previous property that tells us that the determinant is multiplied by -1 when we exchange the identical columns but - since the columns are identical - the resultant matrix is the same - which means that the determinant remains unchanged. Therefore, the determinant must be 0.

Proposition 52. *If A and B are $n \times n$ matrices, then*

$$\det BA = \det A \det B$$

Proof. Let the columns of A be the vectors, $\vec{x}_1, \vec{x}_2, \dots, \vec{x}_n$ so that for each column j ,

$$\vec{x}_j = \sum_i a(i, j) \vec{e}_i$$

and define,

$$\Delta_B(\vec{x}_1, \vec{x}_2, \dots, \vec{x}_n) = \Delta_B(A) = \det BA$$

so that,

$$\det(B\vec{x}_1, B\vec{x}_2, \dots, B\vec{x}_n) = \Delta_B(\vec{x}_1, \vec{x}_2, \dots, \vec{x}_n)$$

Since $B\vec{x}_j$ is linear in \vec{x}_j , $\Delta_B(\vec{x}_1, \vec{x}_2, \dots, \vec{x}_n)$ is linear in each \vec{x}_j and so,

$$\begin{aligned} \Delta_B(\vec{x}_1, \vec{x}_2, \dots, \vec{x}_n) &= \Delta_B\left(\sum_i a(i, 1)\vec{e}_i, \vec{x}_2, \dots, \vec{x}_n\right) \\ &= \sum_i a(i, 1)\Delta_B(\vec{e}_i, \vec{x}_2, \dots, \vec{x}_n) \\ &= \sum_{i_1} a(i_1, 1) \sum_{i_2} a(i_2, 2) \cdots \sum_{i_n} a(i_n, n) \Delta_B(\vec{e}_{i_1}, \vec{e}_{i_2}, \dots, \vec{e}_{i_n}) \\ &= \sum a(i_1, 1)a(i_2, 2) \cdots a(i_n, n) \Delta_B(\vec{e}_{i_1}, \vec{e}_{i_2}, \dots, \vec{e}_{i_n}) \end{aligned}$$

the sum being extended over all n-tuples, (i_1, \dots, i_n) such that $1 \leq i_j \leq n$. Also, by referring again to the properties of the determinant we see that,

$$\Delta_B(\vec{e}_{i_1}, \vec{e}_{i_2}, \dots, \vec{e}_{i_n}) = t(i_1, i_2, \dots, i_n) \Delta_B(\vec{e}_1, \vec{e}_2, \dots, \vec{e}_n)$$

where $t(i_1, i_2, \dots, i_n) = 1, 0, -1$ similar to the function s previously. So, we end up with,

$$\det BA = \sum a(i_1, 1)a(i_2, 2) \cdots a(i_n, n)t(i_1, i_2, \dots, i_n) \det B = \det A \det B$$

□

Proposition 53. *A linear operator A on \mathbb{R}^n is invertible if and only if $\det A \neq 0$*

Proof. If A is invertible then, $AA^{-1} = I$ and, using Proposition 47 and Proposition 52, we have,

$$\det AA^{-1} = \det A \cdot \det A^{-1} = 1$$

so $\det A$ cannot be 0.

Furthermore, if the columns of A are not independent then there is some linear combination of the columns that produces $\vec{0}$. Since, by Proposition 51 we know that adding multiples of columns to other columns leaves the determinant unchanged, this means that the determinant is equal to the determinant of a matrix with $\vec{0}$ as a column. Such a matrix has determinant 0, so the determinant of A is also 0. \square

Corollary 13. *For invertible matrices,*

$$\det A^{-1} = \frac{1}{\det A}$$

Corollary 14. *The determinant is the only function that has the described properties.*

Proof. Every matrix, A , can be transformed by multiplication by elementary matrices to a row-reduced form, R , which is either the identity matrix - in the case that A is invertible - or a matrix with the last row zeroes - in the case where A is not invertible. So, the determinant of the row-reduced matrix, R , is either 1 or 0. Meanwhile, the determinants of the elementary matrices are:

- Add multiple of row to another row - determinant is 1 because this operation maintains the determinant of the identity.
- Swap two rows - determinant is -1 - determinant is -1 because this operation multiplies the determinant of the identity by -1.
- Multiply a row by some scalar c - determinant is c because this operation multiplies the determinant of the identity by c .

So, we have,

$$R = E_1 E_2 \cdots E_n A \implies \det R = \det E_1 E_2 \cdots E_n \cdot \det A$$

where $\det E_1 E_2 \cdots E_n$ is a known, non-zero quantity - say d_e . Since the determinant of R is either 0 or 1 this leaves the determinant of A being either 0 or $\frac{1}{d_e}$.

So, the value of the determinant of an arbitrary matrix, A , is wholly determined by the properties described. \square

Proposition 54. *The determinant of any square matrix is equal to that of its transpose. That's to say, for an arbitrary square matrix A ,*

$$\det A^T = \det A.$$

Proof. The determinant formula from 2.3.5.2 gives us:

$$\det A = \sum s(j_1, \dots, j_n) a(1, j_1) \cdots a(n, j_n)$$

where $a(1, j)$ is the (i, j) th element of the matrix A and the summation is over all n -tuples of n distinct values, j_1, \dots, j_n with $1 \leq j_r \leq n$. So we can also deduce that,

$$\det A^T = \sum s(j_1, \dots, j_n) a(j_1, 1) \cdots a(j_n, n).$$

Clearly for the identity permutation $j_1, \dots, j_n = 1, \dots, n$ the term of the summation generated is the same: in both cases it is the product of the elements along the main diagonal $a(1, 1) \cdots a(n, n)$. On the other hand, if we take a minimal permutation, for example, where $j_1 = 2, j_2 = 1$ so that the first two elements are swapped, then we see that the term of the summation generated in the determinant of A is

$$-a(1, 2)a(2, 1)a(3, j_3) \cdots a(n, j_n)$$

while the term of the summation generated in the determinant of A^T is

$$-a(2, 1)a(1, 2)a(3, j_3) \cdots a(n, j_n)$$

so that the same permutation in the n -tuple j_1, \dots, j_n produces the exact same terms of the summation in the determinant of A and of A^T .

It's easy to see in fact, that this will happen with any permutation as the permutation of the j_1, \dots, j_n , in the determinant of A , permutes the column indices while selecting in order from the rows whereas, in the determinant of A^T , it permutes the row indices while selecting from the columns in order. So, the net result is the same in both cases. Since each permutation of the n -tuple produces equal terms of the summation, the complete summation produced by all the permutations will be the same.

Therefore $\det A^T = \det A$. \square

2.3.6 Permutation Matrices

Definition. A permutation p is a bijective map from a set S to itself. If a matrix P is the matrix associated with a permutation p then:

- the j th column of the matrix is the basis vector $e_{p(j)}$,
- P is a sum of the matrix units, $P = e_{p(1)1} + e_{p(2)2} + \cdots + e_{p(n)n} = \sum_j e_{p(j)j}$.

Proposition 55. If P, Q are permutation matrices associated with the permutations p, q then the matrix that corresponds to the permutation $p \circ q$ is PQ

Proof. $pq(i) = p(q(i))$ and $PQX = P(QX)$ □

Proposition 56. A permutation matrix P is invertible and its inverse is the transpose, $P^{-1} = P^T$

Proof. A left-multiplying permutation matrix for a permutation, p , maps each row from the input matrix using a column j in the permutation matrix, to the output row, $p(j)$. Since the permutation, by definition, is bijective, we know that this mapping is one-to-one and invertible. If we transpose the matrix P to P^T , swapping rows and columns in the permutation matrix, then the new matrix, P^T maps input rows $p(j)$ into output rows j which is clearly the inverse permutation. □

Since a permutation matrix is a the result of permuting the rows of the identity matrix, clearly, its determinant is ± 1 . A permutation is referred to as *odd* or *even* depending on whether its determinant is -1 or 1 respectively. Its determinant is called the *sign of the permutation*,

$$\text{sign } p = \det p = \pm 1$$

The determinant of an arbitrary $n \times n$ matrix can be described as,

$$\begin{aligned}
 \det A &= \sum_p \left[\det \sum_j a_{p(j)j} e_{p(j)j} \right] \\
 &= \sum_p \left[(a_{p(1)1} \cdots a_{p(n)n}) \cdot \det \sum_j e_{p(j)j} \right] \\
 &= \sum_p \left[(a_{p(1)1} \cdots a_{p(n)n}) \cdot \det P \right] \\
 &= \sum_p \left[(\text{sign } p) (a_{p(1)1} \cdots a_{p(n)n}) \right]
 \end{aligned}$$

This is the same formula as earlier and is known as the *complete expansion* of the determinant.

2.3.7 Cramer's Rule

Expansion by minors on the j th column:

$$\det A = (-1)^{j+1} a_{1j} \det A_{1j} + (-1)^{j+2} a_{2j} \det A_{2j} + \cdots + (-1)^{j+n} a_{nj} \det A_{nj}$$

Expansion by minors on the i th row:

$$\det A = (-1)^{i+1} a_{i1} \det A_{i1} + (-1)^{i+2} a_{i2} \det A_{i2} + \cdots + (-1)^{i+n} a_{in} \det A_{in}$$

Definition. If we form a matrix with elements $\alpha_{ij} = (-1)^{i+j} \det A_{ij}$ and then transpose it we get the **adjoint matrix**.

Notation. The adjoint of A is denoted $\text{adj } A$.

Following we use $[x]$ to denote a matrix as distinguished from a scalar.

Let $d = \det A$. Then, if we multiply the adjoint matrix of $[A]$ by $[A]$ we get,

$$[\text{adj } A][A] = \begin{bmatrix} d & & & \\ & d & & \\ & & \ddots & \\ & & & d \end{bmatrix}$$

The off-diagonal elements come out zero because they involve a determinant calculation that involves the same row (or column) repeated and so those determinants are zero.

Theorem 17.

$$[\text{adj } A][A] = (\det A)[I] = [A][\text{adj } A]$$

Corollary 15.

$$\frac{1}{\det A} [\text{adj } A][A] = [I] \iff [A^{-1}] = \frac{1}{\det A} [\text{adj } A]$$

This formulation of the inverse of a matrix can be used to write the solution to a system of linear equations (reverting to the normal notation) $AX = B$ as, multiplying on the left by A^{-1} ,

$$X = A^{-1}B = \frac{1}{\det A}(\text{adj } A)B$$

so that X is a vector whose components, x_j , are expressed as,

$$\begin{aligned} x_j &= \frac{1}{\det A}(b_1\alpha_{1j} + \cdots + b_n\alpha_{nj}) \\ &= \frac{1}{\det A}(b_1(-1)^{1+j} \det A_{1j} + \cdots + b_n(-1)^{n+j} \det A_{nj}) \end{aligned}$$

which is the expansion by minors of A on the j th column but with the components a_{ij} of A replaced with the components of the vector B , divided by the determinant of A .

2.3.8 Linear Algebra of Polynomials

2.3.8.1 Uniqueness of Polynomials

Proposition 57. *Three points uniquely identify a quadratic polynomial.*

Proof. Assume that there are two distinct quadratic polynomials $p(x)$, $q(x)$ that share 3 points. That is,

$$p(x_1) = q(x_1), p(x_2) = q(x_2), p(x_3) = q(x_3).$$


Then we can define another quadratic polynomial $r(x) = p(x) - q(x)$ with the property that,

$$r(x_1) = p(x_1) - q(x_1) = 0 = r(x_2) = r(x_3).$$

In other words, $r(x)$ has 3 roots: x_1, x_2, x_3 . But $r(x)$ is a quadratic polynomial and has a maximum of 2 roots. Therefore there is non such (nonzero) polynomial. \square

2.3.8.2 Lagrange Polynomials

If we look for quadratic polynomials, $p(x)$, that pass through the 3 points $(1, 3)$, $(3, 1)$ and $(5, 2)$:



Then the first has roots at $x = 1, 3$ and passes through the point $(5, 2)$. So, we have:


$$p(1) = p(3) = 0, p(5) = 2$$

meaning that $(x - 1)$ and $(x - 3)$ are factors. Therefore,

$$\begin{aligned} p(x) &= \alpha(x - 1)(x - 3) \\ &= \alpha(x^2 - 4x + 3) \end{aligned}$$

$$\begin{aligned} \implies \quad & \frac{p(5)}{\alpha(5^2 - 4(5) + 3)} = 2 \\ \iff & 8\alpha = 2 \\ \iff & \alpha = \frac{1}{4} \end{aligned}$$

$$\therefore p(x) = \frac{1}{4}(x^2 - 4x + 3)$$




The second has roots at $x = 1, 5$ and passes through the point $(3, 1)$:

$$\begin{aligned} p(x) &= \alpha(x - 1)(x - 5) \\ &= \alpha(x^2 - 6x + 5) \end{aligned}$$

$$\begin{aligned} \Rightarrow \quad & \begin{array}{lll} p(3) & = & 1 \\ \alpha(3^2 - 6(3) + 5) & = & 1 \end{array} \\ \Leftrightarrow & \begin{array}{lll} -4\alpha & = & 1 \end{array} \\ \Leftrightarrow & \begin{array}{lll} \alpha & = & -\frac{1}{4} \end{array} \end{aligned}$$

$$\therefore p(x) = -\frac{1}{4}(x^2 - 6x + 5)$$




The third has roots at $x = 3, 5$ and passes through the point $(1, 3)$:

$$\begin{aligned} p(x) &= \alpha(x - 3)(x - 5) \\ &= \alpha(x^2 - 8x + 15) \end{aligned}$$


$$\begin{aligned} & p(1) = 3 \\ \implies \alpha(1^2 - 8(1) + 15) &= 3 \\ \iff 8\alpha &= 3 \\ \iff \alpha &= \frac{3}{8} \end{aligned}$$

$$\therefore p(x) = \frac{3}{8}(x^2 - 8x + 15)$$



Adding them together we get,

$$\begin{aligned} & \frac{1}{4}(x^2 - 4x + 3) - \frac{1}{4}(x^2 - 6x + 5) + \frac{3}{8}(x^2 - 8x + 15) \\ &= \left(\frac{1}{4} - \frac{1}{4} + \frac{3}{8}\right)x^2 + \left(-1 + \frac{3}{2} - 3\right)x + \left(\frac{3}{4} - \frac{5}{4} + \frac{45}{8}\right) \\ &= \frac{3}{8}x^2 - \frac{5}{2}x + \frac{41}{8} \end{aligned}$$



2.3.8.3 Matrix approach

We are looking for the unique quadratic polynomial $p(x) = \alpha_1 x^2 + \alpha_2 x + \alpha_3$ that satisfies

$$p(1) = 3, p(3) = 1, p(5) = 2$$

so this gives us the simultaneous equations:

$$\alpha_1 + \alpha_2 + \alpha_3 = 3$$

$$9\alpha_1 + 3\alpha_2 + \alpha_3 = 1$$

$$25\alpha_1 + 5\alpha_2 + \alpha_3 = 2$$

which can be expressed as a matrix equation,

$$\begin{bmatrix} 1 & 1 & 1 \\ 9 & 3 & 1 \\ 25 & 5 & 1 \end{bmatrix} \begin{bmatrix} \alpha_1 \\ \alpha_2 \\ \alpha_3 \end{bmatrix} = \begin{bmatrix} 3 \\ 1 \\ 2 \end{bmatrix}.$$

Now,

$$\begin{bmatrix} 1 & 1 & 1 \\ 9 & 3 & 1 \\ 25 & 5 & 1 \end{bmatrix}^{-1} = \begin{bmatrix} 1/8 & -1/4 & 1/8 \\ -1 & 3/2 & -1/2 \\ 15/8 & -5/4 & 3/8 \end{bmatrix}$$

so we have,

$$\begin{bmatrix} \alpha_1 \\ \alpha_2 \\ \alpha_3 \end{bmatrix} = \begin{bmatrix} 1/8 & -1/4 & 1/8 \\ -1 & 3/2 & -1/2 \\ 15/8 & -5/4 & 3/8 \end{bmatrix} \begin{bmatrix} 3 \\ 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 3/8 - 1/4 + 1/4 \\ -3 + 3/2 - 1 \\ 45/8 - 5/4 + 3/4 \end{bmatrix} = \begin{bmatrix} 3/8 \\ -5/2 \\ 41/8 \end{bmatrix}.$$

Therefore $p(x) = (3/8)x^2 + (-5/2)x + (41/8)$.

2.3.8.4 Unified view

Consider each of the component lagrange polynomials:

- $p_1(x) = \frac{1}{4}(x^2 - 4x + 3) = (1/4)x^2 - x + (3/4)$

$$\vec{p}_1 = \begin{bmatrix} 1/4 \\ -1 \\ 3/4 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 1 & 1 & 1 \\ 9 & 3 & 1 \\ 25 & 5 & 1 \end{bmatrix} \vec{p}_1 = \begin{bmatrix} 1/4 - 1 + 3/4 \\ 9/4 - 3 + 3/4 \\ 25/4 - 5 + 3/4 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 2 \end{bmatrix}.$$

- $p_2(x) = -\frac{1}{4}(x^2 - 6x + 5)$

$$\vec{p}_2 = \begin{bmatrix} 1/4 \\ -3/2 \\ 5/4 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 1 & 1 & 1 \\ 9 & 3 & 1 \\ 25 & 5 & 1 \end{bmatrix} \vec{p}_2 = \begin{bmatrix} 1/4 - 3/2 + 5/4 \\ 9/4 - 9/2 + 5/4 \\ 25/4 - 15/2 + 5/4 \end{bmatrix} = \begin{bmatrix} 0 \\ -1 \\ 0 \end{bmatrix}.$$

- $p_3(x) = \frac{3}{8}(x^2 - 8x + 15)$

$$\vec{p}_3 = \begin{bmatrix} 3/8 \\ -3 \\ 45/8 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 1 & 1 & 1 \\ 9 & 3 & 1 \\ 25 & 5 & 1 \end{bmatrix} \vec{p}_3 = \begin{bmatrix} 3/8 - 3 + 45/8 \\ 27/8 - 9 + 45/8 \\ 75/8 - 15 + 45/8 \end{bmatrix} = \begin{bmatrix} 3 \\ 0 \\ 0 \end{bmatrix}.$$

So,

$$\begin{bmatrix} 1 & 1 & 1 \\ 9 & 3 & 1 \\ 25 & 5 & 1 \end{bmatrix} (\vec{p}_1 + \vec{p}_2 + \vec{p}_3) = \begin{bmatrix} 0 \\ 0 \\ 2 \end{bmatrix} + \begin{bmatrix} 0 \\ -1 \\ 0 \end{bmatrix} + \begin{bmatrix} 3 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 3 \\ -1 \\ 2 \end{bmatrix}.$$

2.4 Vector Spaces

2.4.1 Definition of a Vector Space

Let F denote a field which is a subfield of \mathbb{C} and V denote a vector space over F .

Definition. Addition, Scalar Multiplication

- An **addition** on a set V is a function that assigns an element $u+v \in V$ to each pair of elements $u, v \in V$.
- A **scalar multiplication** on a set V is a function that assigns an element $\lambda v \in V$ to each $\lambda \in F$ and each $v \in V$.

Note that both functions are closed over V .

A vector space is a set V along with an addition on V and a scalar multiplication on V such that the following properties hold:

commutativity $\vec{u} + \vec{v} = \vec{v} + \vec{u}$ for all $\vec{u}, \vec{v} \in V$;

associativity $(\vec{u} + \vec{v}) + \vec{w} = \vec{u} + (\vec{v} + \vec{w})$ and $(ab)\vec{v} = a(b\vec{v})$ for all $\vec{u}, \vec{v}, \vec{w} \in V$ and all $a, b \in F$;

additive identity there exists an element $\vec{0} \in V$ such that $\vec{v} + \vec{0} = \vec{v}$ for all $\vec{v} \in V$;

additive inverse for every $\vec{v} \in V$ there exists $\vec{w} \in V$ such that $\vec{v} + \vec{w} = \vec{0}$;

multiplicative identity $1\vec{v} = \vec{v}$ for all $\vec{v} \in V$;

distributive properties $a(\vec{u} + \vec{v}) = a\vec{u} + a\vec{v}$ and $(a+b)\vec{u} = a\vec{u} + b\vec{u}$ for all $a, b \in F$ and $\vec{u}, \vec{v} \in V$;

Notice that vector addition on a vector space V is an associative, commutative law of composition on the set of vectors so that $(V, +)$ **is an abelian group**.

Notice also, however, that scalar multiplication defines a law of composition between vectors in V and scalars in the field F . This is known as an **external law of composition** on the vector space. This is an important part of the definition of vectors as, it defines a relationship between the additive group of vectors V^+ and the field F in much the same way that the distributive law does for the additive and multiplicative groups inside fields (compare 2.2). Specifically the relationship takes the form,

$$(1 + 1)\vec{v} = \vec{v} + \vec{v} = 2\vec{v}.$$

Sets of objects, for example, have the property (if we were to model set union with addition) that $a + a = a$ and so cannot be modeled as vectors. They have a different sort of algebra (wikipedia).

Also, multiplication of two vectors is not a part of the definition of a vector space. Products of vectors depend on a particular co-ordinate system in which the vectors are being represented and so are considered additional structure on the vector space.

It is typically the case that when modelling some system with vectors, the system contains more structure than is represented by a vector space. Furthermore, if we view a vector space as an abelian group then the group only describes a part of the vector space structure. These abstractions of vector space and group allow us to see what is generalizable about a system and which elements are specific to the system in question.

2.4.1.1 Isomorphisms between vector spaces

Definition. An *isomorphism* $\phi : V \mapsto V'$ – where V and V' are defined over the same field F – is a bijective map compatible with the vector laws of composition. That's to say, for all $\vec{v}, \vec{v}' \in V$, $c \in F$,

$$\phi(\vec{v} + \vec{v}') = \phi(\vec{v}) + \phi(\vec{v}') \quad \text{and} \quad \phi(c\vec{v}) = c\phi(\vec{v}).$$

2.4.1.2 Examples of isomorphisms between vector spaces

- (31) The space \mathbb{F}^n of n -dimensional row vectors is isomorphic to the space of n -dimensional column vectors.
- (32) If we treat the set of complex numbers \mathbb{C} as a vector space then the map $\phi : \mathbb{R}^2 \mapsto \mathbb{C}$ defined as $\phi(a, b) = a + bi$ is an isomorphism.

2.4.2 Derived properties of a Vector Space

Proposition 58. A vector space contains a unique additive identity element.

Proof. If $\vec{0}'$ is also an additive identity then by the additive identity property,

$$\vec{0} + \vec{0}' = \vec{0}$$

but since $\vec{0}$ is also an additive identity,

$$\vec{0}' + \vec{0} = \vec{0}'$$

Then, by the commutativity of vector addition,

$$\vec{0} = \vec{0} + \vec{0}' = \vec{0}' + \vec{0} = \vec{0}' \quad \square$$

Proposition 59. A vector space contains a unique additive inverse for each element.

Proof. If \vec{v} and \vec{w} are both additive inverses of \vec{u} then, by the additive inverse property we have,

$$\vec{u} + \vec{v} = \vec{0} \text{ and also } \vec{u} + \vec{w} = \vec{0}$$

using the uniqueness of the additive identity,

$$\vec{u} + \vec{v} = \vec{0} = \vec{u} + \vec{w}$$

Then, if we add one of the additive inverses of \vec{u} to both sides,

$$\vec{u} + \vec{v} + \vec{v} = \vec{u} + \vec{w} + \vec{v}$$

and use the associativity of vector addition,

$$(\vec{u} + \vec{v}) + \vec{v} = (\vec{u} + \vec{v}) + \vec{w}$$

$$\vec{0} + \vec{v} = \vec{0} + \vec{w}$$

$$\vec{v} = \vec{w}$$

□

Because additive inverses are unique we can use the notation $-\vec{v}$ to denote the additive inverse of \vec{v} . Then we define $\vec{w} - \vec{v}$ to mean $\vec{w} + (-\vec{v})$.

Definition. *Vector Subtraction*

$$\vec{u} - \vec{v} := \vec{u} + (-\vec{v})$$

Proposition 60. $0\vec{v} = \vec{0}$ for every $\vec{v} \in V$.

Note that this proposition is asserting something about scalar multiplication and the additive identity of V . The only part of the definition of a vector space that connects scalar multiplication and vector addition is the distributive property. Therefore the distributive property must be used in this proof.

Proof. Firstly take,

$$\vec{v} + 0\vec{v} = 0\vec{v} + 1\vec{v}$$

and then use the properties of the underlying field to say

$$(0 + 1)\vec{v} = 1\vec{v} = \vec{v}$$

Now we have shown that,

$$\vec{v} + 0\vec{v} = \vec{v}$$

which, by the definition and uniqueness of the additive identity, shows that $0\vec{v} = \vec{0}$. But if we want to continue algebraically we can now add the additive inverse to both sides,

$$\begin{aligned}(\vec{v} + -\vec{v}) + 0\vec{v} &= (\vec{v} + -\vec{v}) \\ \vec{0} + 0\vec{v} &= 0\vec{v} = \vec{0}\end{aligned}\quad \square$$

Another, simpler proof exists.

Proof. Using the underlying field properties and the distributivity of scalar vector multiplication,

$$0\vec{v} = (0 + 0)\vec{v} = 0\vec{v} + 0\vec{v}$$

and then adding the additive inverse to both sides,

$$\begin{aligned}(0\vec{v} + -(0\vec{v})) &= (0\vec{v} + -(0\vec{v})) + 0\vec{v} \\ \vec{0} &= \vec{0} + 0\vec{v} = 0\vec{v}\end{aligned}\quad \square$$

Proposition 61. $a\vec{0} = \vec{0}$ for every $a \in F$.

Proof. Using the distributivity of scalar multiplication of vectors and the additive identity,

$$a\vec{0} = a(\vec{0} + \vec{0}) = a\vec{0} + a\vec{0}$$

Then, adding the additive inverse to both sides,

$$\begin{aligned}(a\vec{0} + -(a\vec{0})) &= a\vec{0} + (a\vec{0} + -(a\vec{0})) \\ \vec{0} &= a\vec{0} + \vec{0} = a\vec{0}\end{aligned}\quad \square$$

Proposition 62. $(-1)\vec{v} = -\vec{v}$ for every $\vec{v} \in V$.

Proof. Using the distributivity of scalar multiplication of vectors and the underlying field properties we have,

$$(-1)\vec{v} + \vec{v} = (-1)\vec{v} + 1\vec{v} = (-1 + 1)\vec{v} = 0\vec{v} = \vec{0}$$

Now we could add the additive inverse to both sides to show that,

$$\begin{aligned}(-1)\vec{v} + (\vec{v} + -\vec{v}) &= \vec{0} + -\vec{v} \\ (-1)\vec{v} + \vec{0} &= \vec{0} + -\vec{v} \\ (-1)\vec{v} &= -\vec{v}\end{aligned}\quad \square$$

But we already have,

$$(-1)\vec{v} + \vec{v} = \vec{0}$$

and this, by the definition of the additive inverse, proves that $(-1)\vec{v}$ is an additive inverse of \vec{v} . Since we have previously proven the uniqueness of the additive inverse in Proposition 59 we can conclude, in fact, that $(-1)\vec{v} = -\vec{v}$ the unique additive inverse of v .

2.4.3 The notation F^S

Notation. If S is a set then \mathbf{F}^S denotes the set of functions $S \mapsto F$.

Addition is defined as, for $f, g, (f + g) \in F^S$,

$$(f + g)(x) = f(x) + g(x)$$

for all $x \in S$.

Scalar multiplication is defined as, for $\lambda \in F, \lambda f \in F^S$,

$$(\lambda f)(x) = \lambda f(x)$$

for all $x \in S$.

Example: If S is the interval $[0, 1]$ and $F = \mathbb{R}$ then $\mathbb{R}^{[0,1]}$ is the set of real-valued functions on the interval $[0, 1]$. $\mathbb{R}^{[0,1]}$ is a vector space with additive identity $0 : [0, 1] \mapsto \mathbb{R}$ defined as $0(x) = 0$ and the additive inverse of some function $f \in \mathbb{R}^{[0,1]}$ is the function defined as $(-f)(x) = -f(x)$.

Any *non-empty* set S in conjunction with a subset of \mathbb{C} would similarly produce a vector space. In fact, the vector space F^n can be thought of as the space of functions from the set $\{1, 2, 3, \dots, n\}$ to F . For example, vectors in 3-dimensional space can be viewed as:

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} \equiv f : \{1, 2, 3\} \mapsto \mathbb{R} \text{ with } f(t) = \begin{cases} x & t = 1 \\ y & t = 2 \\ z & t = 3 \end{cases}$$

2.4.4 Polynomials as a vector space

A very important example involves treating a polynomial as a vector. A function $p : F \mapsto F$ is called a polynomial with coefficients in F if there exist $a_0, \dots, a_m \in F$ such that,

$$p(z) = a_0 + a_1z + a_2z^2 + \dots + a_mz^m$$

for all $z \in F$.

Then we can define a vector space, $P(F)$, to be the set of all polynomials with coefficients in F .

Addition on $P(F)$ is defined as,

$$(p + q)(z) = p(z) + q(z) \quad \text{for } p, q \in P(F), z \in F$$

whose associativity is clear from the definition and the commutativity can be shown by,

$$\begin{aligned} ((p + q) + r)(z) &= (p + q)(z) + r(z) \\ &= p(z) + q(z) + r(z) \\ &= p(z) + (q + r)(z) \\ &= (p + (q + r))(z) \end{aligned}$$

Scalar multiplication on $P(F)$ is defined as,

$$(ap)(z) = ap(z) \quad \text{for } p \in P(F), a, z \in F$$

whose associativity can be shown by substituting (ab) for a in the definition,

$$[(ab)p](z) = (ab)p(z)$$

Then, by the associativity of the multiplication of the elements of the field F we have,

$$(ab)p(z) = a[b(p(z))]$$

then we use the definition in reverse,

$$a[b(p(z))] = a[(bp)(z)] = [a(bp)](z)$$

(compare with $(ab)\vec{v} = a(b\vec{v})$)

modeling Concretely, each $p(z) \in P(F)$ is a vector that could be modeled, say, as

$$\vec{p} = \{ (a_0, a_1, \dots, a_m) \mid p(z) = a_0 + a_1z + a_2z^2 + \dots + a_mz^m \in P(F) \}$$

2.4.5 Subspaces of vector spaces

Definition. *A set U is a subspace of V if it is a subset of V and if the same addition and multiplication over U forms a vector space.*

Considering the required properties of a vector space, we can see that commutativity and associativity of the addition; associativity of the scalar multiplication; and distributivity of the scalar multiplication over the addition; will all be satisfied as we have the same addition and multiplication over a subset of the elements in V . That's to say, the vector space properties ensure that these properties hold $\forall \vec{v} \in V$ and we have $\forall \vec{u} \in U, \vec{u} \in V$.

Furthermore, the multiplicative identity also holds $\forall \vec{v} \in V$ so will also hold for every element of U .

So what remains to be proven to satisfy the requirements of a subspace?

- Existence of the additive identity
- Existence of an additive inverse for every element of U
- Closure of the addition and scalar multiplication over U

Note, however that - having proved in Proposition 62 that multiplication by -1 gives the additive inverse - closure of the scalar multiplication over U also implies the presence in U of the additive inverse of every element of U . So, actually, what we need to prove for U to be a subspace is only,

- $\vec{0} \in U$
- Closure of the addition and scalar multiplication over U

2.4.5.1 Examples of Subspaces

(33) An example of a subspace of the polynomials, $P(F)$ is,

$$\{ p \in P(F) \mid p(3) = 0 \}$$

Members of this subspace include:

- $p(z) = 3 - z$
- $p(z) = 9 - z^2$
- $p(z) = 3 - z + 3z^2 - z^3$
- $p(z) = 12z - 4z^2$
- ...etc.

To verify this we need to show that addition and multiplication are closed over this set and that $\vec{0}$ is a member of the set. It's easy to see that $\vec{0}$ is a member of the set as,

$$p(3) = 0 + 0(3) + 0(3)^2 + \cdots + 0(3)^m = 0$$

as required. Scalar multiplication is closed as,

$$ap(3) = a(0) = 0$$

whereas addition can be shown to be closed as,

$$(q + r)(3) = q(3) + r(3) = 0 + 0 = 0$$

Note that for values of $z \neq 3$, the closure of these functions is the same as for the general case of $P(F)$.

2.4.5.2 Sums and Direct Sums

Definition. If U_1, \dots, U_m are subspaces of V then their sum is defined as

$$U_1 + \cdots + U_m = \{ \vec{u}_1 + \cdots + \vec{u}_m \mid \vec{u}_1 \in U_1, \dots, \vec{u}_m \in U_m \}.$$

The sum of the subspaces of V is also a subspace of V because,

- Closure of addition

$$\begin{aligned}
& (\vec{u}_1 + \vec{u}_2 + \cdots + \vec{u}_m) + (\vec{u}'_1 + \vec{u}'_2 + \cdots + \vec{u}'_m) \\
&= (\vec{u}_1 + \vec{u}'_1) + (\vec{u}_2 + \vec{u}'_2) + \cdots + (\vec{u}_m + \vec{u}'_m) \\
&= \vec{v}_1 + \vec{v}_2 + \cdots + \vec{v}_m \quad \text{where } \vec{v}_1 \in U_1, \vec{v}_2 \in U_2, \dots, \vec{v}_m \in U_m
\end{aligned}$$

- Closure of scalar multiplication

$$\begin{aligned}
& a(\vec{u}_1 + \vec{u}_2 + \cdots + \vec{u}_m) \quad \text{where } a \in F \\
&= a\vec{u}_1 + a\vec{u}_2 + \cdots + a\vec{u}_m \\
&= \vec{v}_1 + \vec{v}_2 + \cdots + \vec{v}_m \quad \text{where } \vec{v}_1 \in U_1, \vec{v}_2 \in U_2, \dots, \vec{v}_m \in U_m
\end{aligned}$$

- Existence of $\vec{0}$

$$\begin{aligned}
& U_1, U_2, \dots, U_m \text{ are subspaces} \\
& \implies \vec{0} \in U_1, \vec{0} \in U_2, \dots, \vec{0} \in U_m \\
& \implies \vec{0} + \vec{0} + \cdots + \vec{0} \in U_1 + U_2 + \cdots + U_m
\end{aligned}$$

Note though, that this may not be the only way of producing $\vec{0}$ from the sum of vectors of these subspaces. That's to say, there could be some $(\vec{u}_1 + \vec{u}_2 + \cdots + \vec{u}_m) = \vec{0}$ and this is a key difference from direct sums.

Proposition 63. $U_1 + U_2 + \cdots + U_m$ is the smallest subspace of V containing U_1, U_2, \dots, U_m .

Proof. $U_1 + U_2 + \cdots + U_m$ is a subspace of V that contains U_1, U_2, \dots, U_m because we can obtain U_i by setting all the u_j for $j \neq i$ to $\vec{0}$.

If a subspace of V contains U_1, U_2, \dots, U_m then, by the closure of addition, it must also contain $U_1 + U_2 + \cdots + U_m$.

Therefore the smallest subspace of V that contains U_1, U_2, \dots, U_m is $U_1 + U_2 + \cdots + U_m$. \square

Definition. If U_1, \dots, U_m are subspaces of V then their **direct sum** is defined as,

$$U_1 \oplus \dots \oplus U_m = \{ \vec{u}_1 + \dots + \vec{u}_m \mid \vec{u}_1 \in U_1, \dots, \vec{u}_m \in U_m \}$$

such that,

$$\vec{u}_1 + \dots + \vec{u}_m = \vec{0} \implies \vec{u}_1 = \vec{0}, \dots, \vec{u}_m = \vec{0}.$$

In this case, the subspaces U_1, \dots, U_m are said to be **independent**.

That the unique way of obtaining $\vec{0}$ is for all of the vectors from each of the subspaces to be $\vec{0}$ is equivalent to there only being a single unique way of obtaining each resultant vector from an addition of the vectors from the individual subspaces. This can be seen as,

$$\begin{aligned} \vec{u}_1 + \vec{u}_2 + \dots + \vec{u}_m &= \vec{u}'_1 + \vec{u}'_2 + \dots + \vec{u}'_m \\ (\vec{u}_1 + \vec{u}_2 + \dots + \vec{u}_m) - (\vec{u}'_1 + \vec{u}'_2 + \dots + \vec{u}'_m) &= \vec{0} \\ (\vec{u}_1 - \vec{u}'_1) + (\vec{u}_2 - \vec{u}'_2) + \dots + (\vec{u}_m - \vec{u}'_m) &= \vec{0} \end{aligned}$$

Therefore, since vector spaces always contain $\vec{0}$ and so we will always have the representation,

$$\vec{0} + \vec{0} + \dots + \vec{0} = \vec{0}$$

if this is the unique representation of $\vec{0}$ then it follows that,

$$\begin{aligned} (\vec{u}_1 - \vec{u}'_1) = \vec{0}, (\vec{u}_2 - \vec{u}'_2) = \vec{0}, \dots, (\vec{u}_m - \vec{u}'_m) &= \vec{0} \\ \implies \vec{u}_1 = \vec{u}'_1, \vec{u}_2 = \vec{u}'_2, \dots, \vec{u}_m = \vec{u}'_m \end{aligned}$$

which means that these are the same representation. And this clearly holds in reverse also as, if there is a single way of representing each resultant vector then there must be a single way of representing $\vec{0}$ and due to the definition of a vector space we must always have the representation of all $\vec{0}$. Therefore, this is the only representation of $\vec{0}$.

Note that this is a condition on the contents of the subspaces and not on the way that the addition is performed. So, the difference between vector space sum ($U_1 + U_2$) and vector space direct sum ($U_1 \oplus U_2$) is not in the operator itself but in the operands they operate over.

For two subspaces, say, U_1, U_2 this condition on the subspaces reduces to the requirement that $U_1 \cap U_2 = \{\vec{0}\}$ which can be seen as,

$$\begin{aligned}\vec{u}_1 + \vec{u}_2 &= \vec{0} \\ \vec{u}_1 + -\vec{u}_1 + \vec{u}_2 &= \vec{0} + -\vec{u}_1 \\ \vec{u}_2 &= -\vec{u}_1 \\ \implies -\vec{u}_1 &\in U_2 \implies \vec{u}_1 \in U_2\end{aligned}$$

So, for two subspaces, obtaining $\vec{0}$ as the sum of vectors from the subspaces implies a vector in common between them. So, for $\vec{0} + \vec{0}$ to be the only way of obtaining $\vec{0}$ implies that $\vec{0}$ is the only vector in common.

However, for more than two subspaces, say U_1, U_2, U_3 , the situation is different as we could have,

$$\begin{aligned}\vec{u}_1 + \vec{u}_2 + \vec{u}_3 &= \vec{0} \\ \iff \vec{u}_1 + -\vec{u}_1 + \vec{u}_2 + -\vec{u}_2 + \vec{u}_3 &= \vec{0} + -\vec{u}_1 + -\vec{u}_2 \\ \iff \vec{u}_3 &= -\vec{u}_1 + -\vec{u}_2\end{aligned}$$

which does not imply any vectors held in common.

2.4.6 Vector Space Problems

Prove that $-(-\vec{v}) = \vec{v}$ for every $\vec{v} \in V$

$$\begin{aligned}-(-\vec{v}) &= -[(-1)\vec{v}] && \text{using Proposition 62} \\ &= (-1)[(-1)\vec{v}] && \text{using Proposition 62 again} \\ &= [(-1)(-1)]\vec{v} && \text{using associativity of scalar multiplication} \\ &= \vec{v} && \text{using field properties}\end{aligned}$$

Or, a quicker way is,

$$\begin{aligned}-\vec{v} + -(-\vec{v}) &= \vec{0} && \text{using additive identity of } -\vec{v} \\ (-\vec{v} + \vec{v}) + -(-\vec{v}) &= \vec{0} + \vec{v} && \text{adding } \vec{v} \text{ to both sides} \\ -(-\vec{v}) &= \vec{v}\end{aligned}$$

Prove that if $a \in F$, $\vec{v} \in V$, and $a\vec{v} = \vec{0}$, then $a = 0$ or $\vec{v} = \vec{0}$. We follow a proof by cases.

Case $a \neq 0$:

$$\begin{aligned}
 a\vec{v} = \vec{0}, a \neq 0 &\implies a^{-1}a\vec{v} = a^{-1}\vec{0} && \text{using field properties} \\
 &\iff 1\vec{v} = b\vec{0} && \text{where } b = a^{-1} \in F \\
 &\iff \vec{v} = \vec{0} && \text{using Proposition 61 and multiplicative identity}
 \end{aligned}$$

Case $\vec{v} \neq \vec{0}$:

$$\begin{aligned}
 a\vec{v} = \vec{0}, \vec{v} \neq \vec{0} &\implies a\vec{v} = a\vec{v} + -a\vec{v} \\
 &\iff a\vec{v} = (a + -a)\vec{v} = 0\vec{v} && \text{using field properties} \\
 \text{Wrong! } a\vec{v} = \vec{0} &\implies a\vec{v} = a\vec{v} + -a\vec{v} \\
 &\text{without need for } \vec{v} \neq \vec{0}
 \end{aligned}$$

This indicates that you are proving something that doesn't need proving. In actual fact,

Case $a = 0$: Actually, in this case there is nothing to be proven as we know from Proposition 60 that $0\vec{v} = \vec{0}$. So we have collectively exhaustive cases by looking at $a = 0$ and $a \neq 0$ and we only need to show that $a \neq 0 \implies \vec{v} = \vec{0}$ which we have already done.

Case $\vec{v} \neq \vec{0}$:

$$\begin{aligned}
 a\vec{v} = \vec{0}, \vec{v} \neq \vec{0} &\implies a\vec{v} = \vec{0} = \vec{v} + -\vec{v} \\
 \iff &a\vec{v} + (-\vec{v}) = -\vec{v} && \text{applying additive inverse} \\
 \iff &a\vec{v} + (-1)\vec{v} = -\vec{v} \\
 \iff &(a + (-1))\vec{v} = (-1)\vec{v} && \text{using distributive law} \\
 \iff &a - 1 = -1 && \text{using injectivity of scalar multiplication} \\
 \iff &a = 0. && \text{using field properties}
 \end{aligned}$$

Give an example of a nonempty subset U of \mathbb{R}^2 such that U is closed under scalar multiplication but U is not a subspace of \mathbb{R}^2 . For all $\lambda \in \mathbb{R}$ the set $\{\lambda \vec{v} \mid \vec{v} \in \{(1, 1), (-1, 1)\}\}$ is closed under scalar multiplication but not addition.




Figure 2.2: The blue arrows are vectors whose scalar multiples will all be in the same line as the blue arrows but the red arrow shows what happens if we add them; the result lies outside of both lines.

Is \mathbb{R}^2 a subspace of the complex vector space \mathbb{C}^2 ? The definition of a subspace of \mathbb{C}^2 is a set of vectors which is a subset of those in \mathbb{C}^2 and that forms a vector space under the same addition and scalar multiplication of \mathbb{C}^2 . The scalar multiplication of the vector space \mathbb{C}^2 is multiplication by scalars $\lambda \in \mathbb{C}$.

For a vector, $\vec{v} \in \mathbb{R}^2$, scaling it by a complex number, $\lambda \vec{v}$ will result in a vector that is not necessarily in \mathbb{R}^2 .

Is $\{(a, b, c) \in \mathbb{C}^3 \mid a^3 = b^3\}$ a subspace of \mathbb{C}^3 ? For $x \in \mathbb{C}$, x^3 has roots, $1, \frac{-1+\sqrt{3}i}{2}, \frac{-1-\sqrt{3}i}{2}$ so we don't have $a = b$ as we would if we were ranging over the reals.

Concretely, we can have, $(1, \frac{-1+\sqrt{3}i}{2}, 0)$ and $(1, \frac{-1-\sqrt{3}i}{2}, 0)$ but,

$$(1, \frac{-1+\sqrt{3}i}{2}, 0) + (1, \frac{-1-\sqrt{3}i}{2}, 0) = (2, -1, 0)$$

where $(2, -1, 0) \notin \{(a, b, c) \in \mathbb{C}^3 \mid a^3 = b^3\}$ meaning that addition over this set is not closed. Therefore, this is not a subspace.

Give an example of a non-empty subset U of \mathbb{R}^2 such that U is closed under addition and under taking additive inverses (meaning $-\vec{u} \in U$ whenever $\vec{u} \in U$), but U is not a subspace of \mathbb{R}^2 . First thought might be $\mathbb{R}^2 - \{\vec{0}\}$ but this is **Wrong!** If the subset is closed under addition and under taking additive inverses then it means that $\vec{u} + -\vec{u} = \vec{0} \in U$ and so the set $\mathbb{R}^2 - \{\vec{0}\}$ is not closed under addition and taking additive inverses.

The set $\{(x, y) \in \mathbb{R}^2 \mid x, y \in \mathbb{Z}\}$ however, is closed under addition because integer addition is closed and under taking additive inverses but scalar multiplication where the scalars range over the reals, will produce non-integer values for x and y .

Is the set of periodic functions over the reals a subspace of $\mathbb{R}^{\mathbb{R}}$? vector problems, periodic functions The definition of two periodic functions over the reals is

$$\begin{aligned}\exists p > 0 \in \mathbb{R} \cdot f(x) &= f(x + p) \\ \exists q > 0 \in \mathbb{R} \cdot g(x) &= g(x + q)\end{aligned}$$

Then for their sum to be periodic we need,

$$\exists \alpha, \beta \in \mathbb{Z}, m \in \mathbb{R} \cdot (m = \alpha p) \wedge (m = \beta q)$$

$$\iff \frac{q}{p} = \frac{\alpha}{\beta} \in \mathbb{Q}$$

$$\therefore (f + g)(x) = (f + g)(x + m) = f(x + m) + g(x + m)$$

$$\iff \frac{q}{p} \in \mathbb{Q}.$$

Prove that the union of two subspaces of V is a subspace of V if and only if one of the subspaces is contained within the other. Let A, B be subspaces of V and $\vec{a} \in A$, $\vec{b} \in B$ and,

$$C = A \cup B = \{ \vec{c} \mid \vec{c} \in A \vee \vec{c} \in B \}.$$

Since $\vec{a}, \vec{b} \in C$ we have (C subspace of V) $\iff \forall \alpha, \beta \in F \cdot (\alpha \vec{a} + \beta \vec{b}) \in C$.
Then,

$$\begin{aligned} \vec{b} \in A &\implies \forall \alpha, \beta \in F \cdot (\alpha \vec{a} + \beta \vec{b}) \in A \text{ (by subspace properties)} \\ &\implies (\alpha \vec{a} + \beta \vec{b}) \in C. \end{aligned}$$

A similar argument holds for $\vec{a} \in B$. Conversely,

$$\begin{aligned} \forall \alpha, \beta \in F \cdot (\alpha \vec{a} + \beta \vec{b}) \in C &\implies ((\alpha \vec{a} + \beta \vec{b}) \in A) \vee ((\alpha \vec{a} + \beta \vec{b}) \in B) \\ &\implies ((\alpha \vec{a} - \alpha \vec{a} + \beta \vec{b}) = \beta \vec{b} \in A) \vee ((\alpha \vec{a} + \beta \vec{b} - \beta \vec{b}) = \alpha \vec{a} \in B) \\ &\implies (\vec{b} \in A) \vee (\vec{a} \in B) \end{aligned}$$

$$\begin{aligned} \therefore (\text{C subspace of V}) &\iff \forall \alpha, \beta \in F \cdot (\alpha \vec{a} + \beta \vec{b}) \in C \\ &\iff (\vec{b} \in A) \vee (\vec{a} \in B) \\ &\equiv (B \subseteq A) \vee (A \subseteq B). \end{aligned}$$

Can a vector space over an infinite field be a finite union of proper subspaces? Assume that our vector space V is a finite union of proper subspaces, hence

$$V = \bigcup_{i=1}^n U_i.$$

Now, pick a non-zero vector $\vec{x} \in U_1$, and pick another vector $\vec{y} \in V \setminus U_1$.

There are infinitely many vectors $\vec{x} + k\vec{y}$, where $k \in K^*$ (K is our infinite field). Note that $\vec{x} + k\vec{y}$ is not in U_1 , hence must be contained in some U_j where $j \neq 1$.

Then since $k \in K^*$, we can have $\vec{x} + k_1\vec{y}, \vec{x} + k_2\vec{y} \in U_j$, which implies that it also contains \vec{y} and hence also \vec{x} , hence $U_1 \subset U_j$.

Explanation: There are infinitely many vectors $\vec{x} + k\vec{y}$ and only finitely many U_i so they cannot all be in different U_i so we have,

$$\begin{aligned} & \exists k_1, k_2 \in K^* \cdot \vec{x} + k_1\vec{y}, \vec{x} + k_2\vec{y} \in U_j \\ \implies & (\vec{x} + k_1\vec{y}) - (\vec{x} + k_2\vec{y}) = (k_1 - k_2)\vec{y} \in U_j \\ \implies & \vec{y} \in U_j \implies \vec{x} \in U_j \end{aligned}$$

Hence

$$V = \bigcup_{i=2}^n U_i.$$

Evidently, this can be continued, hence a contradiction arises.

Prove or give a counterexample: if U_1, U_2, W are subspaces of V such that $V = U_1 \oplus W$ and $V = U_2 \oplus W$ then $U_1 = U_2$. Counter example: $V = \mathbb{F}^2$, $U_1 = \{(x, 0) \in \mathbb{F}^2 \mid x \in F\}$, $U_2 = \{(0, x) \in \mathbb{F}^2 \mid x \in F\}$, $W = \{(x, x) \in \mathbb{F}^2 \mid x \in F\}$.

Let U_e denote the set of real-valued even functions on \mathbb{R} and let U_o denote the set of real-valued odd functions on \mathbb{R} . Show that $\mathbb{R}^{\mathbb{R}} = U_e \oplus U_o$. Every function $f \in \mathbb{R}^{\mathbb{R}}$ can be expressed as the sum of an even function and an odd function as,

$$f(x) = \frac{f(x) + f(-x)}{2} + \frac{f(x) - f(-x)}{2} = g(x) + h(x)$$

where $g(x) \in U_e$ and $h(x) \in U_o$. So, $U_e + U_o$ spans $\mathbb{R}^{\mathbb{R}}$. Furthermore,

$$\begin{aligned} f(x) \in (U_e \cap U_o) & \implies (f(-x) = f(x)) \wedge (f(-x) = -f(x)) \\ & \implies f(x) = -f(x) \\ & \implies f(x) = 0 \end{aligned}$$

Since $f(x) = 0$ is the additive identity of this space, this shows that the intersection is $\vec{0}$. So, $\mathbb{R}^{\mathbb{R}} = U_e \oplus U_o$.

2.4.7 Finite Sets of Vectors

2.4.7.1 Span and Linear Independence

Definition. The **span** of a nonempty finite set of vectors S – written $\text{span } S$ – is defined as the set of **finite** linear combinations of elements of S ,

$$\{ \alpha_1 \vec{v}_1 + \alpha_2 \vec{v}_2 + \cdots + \alpha_k \vec{v}_n \mid \vec{v}_i \in S, \alpha_j \in F \}.$$

The span of an empty set of vectors is defined to be $\{\vec{0}\}$.

The span of a set S is also sometimes known as the **subspace generated by S** .

Definition. A **linear relation** among a nonempty finite set of vectors S is any relation of a **finite** number of elements of S of the form,

$$c_1 \vec{v}_1 + \cdots + c_n \vec{v}_n = \vec{0}$$

where $c_i \in F$.

A **linearly independent** set of vectors is a nonempty set among which there is no linear relation except the trivial relation where all $c_1, \dots, c_n = 0$.

The empty set is defined to be linearly independent.

Proposition 64. If a set of vectors contains the zero vector $\vec{0}$ then it cannot be linearly independent.

Proof. Assume an arbitrary set of vectors $\vec{v}_1, \dots, \vec{v}_n$ and assume it contains some $\vec{v}_i = \vec{0}$. Then we have the linear relation $c_i \vec{v}_i = \vec{0}$ with some $c_i \neq 0$. \square

Corollary 16. If a set of vectors contains any repetition then it cannot be linearly independent.

Proof. If a set of vectors contains the same vector twice then subtracting one from the other is a non-trivial linear combination of the set of vectors equalling $\vec{0}$. \square

Corollary 17. *A set of vectors is linearly independent iff, for any vector that may be expressed as a linear combination of vectors in the set, the expression is unique.*

Proof. If the expression is not unique then the two different representations may be subtracted one from the other to produce a non-trivial linear combination resulting in $\vec{0}$ which contradicts the hypothesis that they are linearly independent. Therefore linearly independent implies unique representation.

Conversely, if the set of vectors is not linearly independent then there exists some non-trivial linear combination that results in $\vec{0}$. This, in turn, implies that there exists some linear combination of a subset of the vectors that is equal to a linear combination of the remaining vectors. Since these two linear combinations are equal, they represent two different expressions of the same resultant vector. Therefore unique representation implies linearly independent. \square

Proposition 65. *The span of a list of vectors is the smallest subspace containing those vectors.*

Note that a vector space over \mathbb{R} or \mathbb{C} is an uncountable set as - while the dimensions of the vector space may be finite - closure under scalar multiplication means that the vectors in the space are continuously valued as the field providing the scalars is continuously valued.

This means that the notion of the *smallest* subspace cannot refer to the cardinality of the set and must refer to ordering based on subset. So, the smallest subspace containing a list of vectors is a subspace that contains the list of vectors and, of which, there is no proper subset which also contains the list of vectors.

Proof.

Let $S := \text{span}(\vec{v}_1, \vec{v}_2, \dots, \vec{v}_k)$
 $:= \{ \alpha_1 \vec{v}_1 + \alpha_2 \vec{v}_2 + \dots + \alpha_k \vec{v}_k \mid \alpha_1, \alpha_2, \dots, \alpha_k \in F \}$
and let $V :=$ the smallest vector space containing $\vec{v}_1, \vec{v}_2, \dots, \vec{v}_k$.

then S contains every linear combination of $\vec{v}_1, \vec{v}_2, \dots, \vec{v}_k$ and nothing else and so is a vector space containing $\vec{v}_1, \vec{v}_2, \dots, \vec{v}_k$,

$$V \subseteq S$$

Additionally, any vector space containing the vectors $\vec{v}_1, \vec{v}_2, \dots, \vec{v}_k$ must contain all their linear combinations, $\text{span}(\vec{v}_1, \vec{v}_2, \dots, \vec{v}_k)$,

$$S \subseteq V$$

Therefore there is no proper subset of $\text{span}(\vec{v}_1, \vec{v}_2, \dots, \vec{v}_k)$ that is also a vector space containing $\vec{v}_1, \vec{v}_2, \dots, \vec{v}_k$, and so $\text{span}(\vec{v}_1, \vec{v}_2, \dots, \vec{v}_k)$ is the smallest vector space containing $\vec{v}_1, \vec{v}_2, \dots, \vec{v}_k$,

$$(V \subseteq S) \wedge (S \subseteq V) \iff V = S \quad \square$$

Proposition 66. *Let L be a linearly independent set of vectors in V and $\vec{v} \in V$. If we add \vec{v} to the set L then the resultant set L' is linearly independent iff $\vec{v} \notin \text{span } L$.*

Proof. Clearly if $\vec{v} \in \text{span } L$ then the resultant set is linearly dependent. If $v \notin \text{span } L$ however, then if we attempt to form a linear relation of the vectors in L' then we find,

$$c_1 \vec{v}_1 + \dots + c_n \vec{v}_n + b \vec{v} = \vec{0}$$

implies that $b \neq 0$ because that would leave a linear relation between the vectors of L which is not possible because L is linearly independent. Therefore,

$$\vec{v} = -(c_1/b) \vec{v}_1 + \dots + -(c_n/b) \vec{v}_n$$

which contradicts the assumption that $v \notin \text{span } L$. \square

Proposition 67. *If we add a vector $\vec{v} \in V$ to a set of vectors L in V to make a new set L' , then $\text{span } L = \text{span } L'$ iff $\vec{v} \in \text{span } L$.*

Proof. Clearly if $\vec{v} \in \text{span } L$ then adding it to the set L doesn't change its span, so $\text{span } L = \text{span } L'$.

Conversely, by construction of L' we have $\vec{v} \in L' \implies \vec{v} \in \text{span } L'$ so if $\text{span } L = \text{span } L'$ then we also have $\vec{v} \in \text{span } L$. \square

Proposition 68. *Length of every linearly independent list in a space is less than or equal to the length of a spanning list in the same space.*

Proof. Let $U = \vec{u}_1, \vec{u}_2, \dots, \vec{u}_m$ be a linearly independent list of vectors in V and $W = \vec{w}_1, \vec{w}_2, \dots, \vec{w}_n$ be a spanning list of vectors in V .

If we take \vec{u}_1 from U and add it to W then - since the other vectors in W are a spanning list - W must be linearly dependent. That's to say,

$$\begin{aligned} \exists \alpha_1, \alpha_2, \dots, \alpha_n \in \mathbb{R} \cdot \alpha_1 \vec{w}_1 + \dots + \alpha_n \vec{w}_n &= \vec{u}_1 \\ \iff \alpha_1 \vec{w}_1 + \dots + \alpha_n \vec{w}_n - \vec{u}_1 &= -\alpha_1 \vec{w}_1 \\ \iff \frac{-\alpha_1}{\alpha_i} \vec{w}_1 + \dots + \frac{-\alpha_n}{\alpha_i} \vec{w}_n + \frac{1}{\alpha_i} \vec{u}_1 &= \vec{w}_i \end{aligned}$$

So, \vec{w}_i is in the span of $\vec{u}_1, \vec{w}_2, \dots, \vec{w}_n$ and we can drop \vec{w}_i from the list, W , and it will still span the vector space.

We can keep doing this with the remaining vectors in U - each time the vector to be removed will be some \vec{w}_i because all the \vec{u}_i are linearly independent - and all the while W remains a spanning list. We continue until we have replaced (potentially) all n vectors in W , which would happen if $m > n$. At this point we would have the spanning list $W = \vec{u}_1, \vec{u}_2, \dots, \vec{u}_n$ and $(m - n)$ remaining vectors in U .

Now, since W spans the space, the $(m - n)$ vectors that remain in U will be in the span of W . But, all the vectors that originally came from U were linearly independent, so it is impossible for any vectors in U to be in the span of W (which now comprises only vectors that originally came from U). We therefore conclude that there can be no remaining vectors in U and, consequently that m cannot be greater than n , i.e. $m \leq n$. \square

Summary of Span and Linear Independence

- *The span of a set of vectors changes when adding a vector not in the existing span.*
- *A linearly independent set of vectors continues to be linearly independent when adding a vector not in the existing span.*

The fundamental point of linear independence and span is that if the span

of two sets of vectors is not completely disjoint (ignoring the zero vector) then the sets are not linearly independent. Any nonzero vector in common between the two spans implies a linear relation between the sets,

$$\begin{aligned} \vec{u}_1 + \dots + \vec{u}_n &= \vec{v} = \vec{w}_1 + \dots + \vec{w}_n \\ \iff \vec{0} &= (\vec{w}_1 + \dots + \vec{w}_n) - (\vec{u}_1 + \dots + \vec{u}_n). \end{aligned}$$

2.4.7.2 Bases

Definition. A **basis** for a vector space V is a set of vectors that is both linearly independent and spans the space V . The empty set is therefore a basis for the zero vector space $\{\vec{0}\}$.

Since a basis of V spans the space, any vector in V may be expressed as a linear combination of the vectors in the basis set and, since the basis set is linearly independent, this expression is unique.

Compare with the **generating set of a group** (2.1.1.2).

Proposition 69. A set of vectors $B = \{\vec{v}_1, \dots, \vec{v}_n\}$ in V is a basis iff every $\vec{w} \in V$ can be written in a single unique way as a linear combination of vectors in B .

Proof. If B is a basis of V then, by definition, B spans V and so every $\vec{w} \in V$ can be written as a linear combination of vectors in B . Furthermore, also by the definition of a basis, the vectors in B are linearly independent so, by corollary 17 the linear combination is unique.

Conversely if every $\vec{w} \in V$ can be written in a single unique way as a linear combination of vectors in B then B both spans the space and is linearly independent by corollary 17. \square

2.4.7.3 Examples of Bases

- (34) If we take an arbitrary finite set of vectors $S = \vec{s}_1, \dots, \vec{s}_n$ then the space $V = V(S)$ of linear combinations of elements of S is the set of all expressions of the form,

$$a_1 \vec{s}_1, \dots, a_n \vec{s}_n, \quad a_i \in \mathbb{F}.$$

In this space addition and multiplication are carried out assuming no relations among the elements of S so that,

$$(a_1 \vec{s}_1 + \dots + a_n \vec{s}_n) + (b_1 \vec{s}_1 + \dots + b_n \vec{s}_n) = (a_1 + b_1) \vec{s}_1 + \dots + (a_n + b_n) \vec{s}_n$$

and

$$c(a_1 \vec{s}_1 + \dots + a_n \vec{s}_n) = ca_1 \vec{s}_1 + \dots + ca_n \vec{s}_n.$$

Then the mapping $\phi : \mathbb{F}^n \mapsto V(S)$ defined as,

$$\phi(a_1, \dots, a_n) = a_1 \vec{s}_1, \dots, a_n \vec{s}_n$$

is an isomorphism.

Note that if the assumption of no relation between the elements of S is not valid then ϕ may fail to be isomorphic.

$V(S)$ is often referred to as *the space with basis S* or *the space of formal linear combinations of S* . If S is an infinite set then $V(S)$ is defined to be the set of all *finite* linear combinations of the elements of S .

This crops up frequently in applications, when taking weightings of different features for example; this isomorphism allows us to treat them as vectors in \mathbb{F}^n .

2.4.7.4 Finite-Dimensional Vector Spaces

Definition. A vector space is called **finite-dimensional** if there is some finite set of vectors that spans the space.

Proposition 70. *Any finite set which spans a finite-dimensional space contains a basis for the space.*

Proof. Let S be a spanning set of vectors in the space V .

If S is not linearly independent then there is some $\vec{v} \in S$ such that $\vec{v} \in \text{span}(S \setminus \{\vec{v}\})$. So we can remove \vec{v} from the set and S still spans the space. We may continue doing this until S is linearly independent, at which point we have found the minimal subset of S that spans the space. This remaining subset is a basis of the space. \square

Corollary 18. *Any finite-dimensional vector space has a basis.*

Proof. This follows from the previous proposition and the definition of a finite-dimensional vector space. \square

Proposition 71. *Any set (including infinite sets) that spans a finite-dimensional vector space, contains a finite subset which spans the space.*

Proof. Let V be a finite-dimensional vector space and S be a spanning set of V . By the definition of V as finite-dimensional there exists some finite set that spans V . Let W be a finite set of vectors that spans V . Then, since S also spans V , every vector in W may be expressed as a finite linear combination of vectors in S . The set of all the members of S that participate in the linear combinations required to produce the set W is a finite subset of S that spans V . \square

Proposition 72. *Let V be a finite-dimensional vector space. Any linearly independent set $L \subseteq V$ can be extended by adding vectors to obtain a basis of V .*

Proof. If L spans V then it is already a basis.

Assume L does not span V . Then there exists some $\vec{v} \in V$ such that $\vec{v} \notin \text{span } L$. If we add \vec{v} to L then the resulting set, say L' , continues to be linearly independent and may or may not span V . If it does then L' is a basis. If not then we can continue to repeat the same process until it does span the space at which point we have a basis. \square

Proposition 73. *For finite subsets of a vector space, any linearly independent set has cardinality less than or equal to that of any spanning set in the same space.*

Proposition 72 only tells us that, for any linearly independent set of vectors, there exists some basis whose cardinality is greater than or equal to that of the original set. What we want to prove here is the condition on the cardinality exists between any linearly independent set and spanning set in the same space.

Proof. We will show two different ways of proving this: one using an algorithm on lists and the other using simultaneous equations.

(i) **proof using lists**

Let V be a finite-dimensional vector space, S a spanning list $\vec{s}_1, \dots, \vec{s}_m$ and L a linearly independent list $\vec{l}_1, \dots, \vec{l}_n$ in V , and assume that $m < n$.

Now, since S is a spanning list, every element of L is in its span and so if we remove the first element \vec{l}_1 from L and add it to S then S will definitely contain a linear relation. If we then remove some element \vec{s}_i from the original spanning list that participates in a linear relation (i.e. $\vec{s}_i \in \text{span } S$) then we have a modified list S , of the same length as the original, but with an element replaced by \vec{l}_1 and this list continues to span the space.

We can repeat this task m times until all elements \vec{s}_i from the original spanning list have been removed and we have a spanning list $S = \vec{l}_1, \dots, \vec{l}_m$ and the remaining $n - m$ elements are still in L . But the original $L = \vec{l}_1, \dots, \vec{l}_n$ was linearly independent and no linear relation exists between them so it is impossible that the first m elements span the space because this would mean that they would participate in a linear relation with the remaining $n - m$ elements.

So we have obtained a contradiction and we therefore conclude that $m \geq n$ \square .

(ii) **proof using simultaneous equations**

Let V be a vector space, S a finite spanning set $\vec{s}_1, \dots, \vec{s}_m$ and L a finite linearly independent set $\vec{l}_1, \dots, \vec{l}_n$ in V , and assume that $m < n$. Since S spans the space, every $\vec{l}_j \in L$ can be expressed as a linear combination of vectors $\vec{s}_i \in S$ of the form,

$$\vec{l}_j = a_{1j}\vec{s}_1 + \dots + a_{mj}\vec{s}_m = \sum_{i=1}^m a_{ij}\vec{s}_i$$

for $1 \leq j \leq n$. A linear relation on the vectors of L would look like,

$$\begin{aligned} c_1\vec{l}_1 + \dots + c_n\vec{l}_n &= \vec{0} \\ \iff c_1 \sum_{i=1}^m a_{i1}\vec{s}_i + \dots + c_n \sum_{i=1}^m a_{in}\vec{s}_i &= \vec{0} \end{aligned}$$

which expands into m simultaneous equations as follows.

$$\begin{aligned} c_1a_{11}\vec{s}_1 + \dots + c_na_{1n}\vec{s}_1 &= 0 \\ &\cdot \\ &\cdot \\ c_1a_{m1}\vec{s}_m + \dots + c_na_{mn}\vec{s}_m &= 0 \end{aligned}$$

As we can see, each of the m equations has a factor \vec{s}_i in each term and so this may be factored out to give the following system.

$$\begin{aligned} \vec{s}_1(c_1a_{11} + \dots + c_na_{1n}) &= 0 \\ &\cdot \\ &\cdot \\ \vec{s}_m(c_1a_{m1} + \dots + c_na_{mn}) &= 0 \end{aligned}$$

So now we have m equations such that the i -th equation will hold if either $\vec{s}_i = \vec{0}$ or $c_1a_{i1} + \dots + c_na_{in} = 0$.

Assume that for all $\vec{s}_i \in S$, $\vec{s}_i \neq \vec{0}$. Then we end up with the following system.

$$\begin{bmatrix} a_{11} & \dots & a_{1n} \\ \cdot & & \\ \cdot & & \\ a_{m1} & \dots & a_{mn} \end{bmatrix} \begin{bmatrix} c_1 \\ \cdot \\ \cdot \\ c_n \end{bmatrix} = \begin{bmatrix} 0 \\ \cdot \\ \cdot \\ 0 \end{bmatrix}$$

It can be shown using matrix row reduction that a system such as this has non-trivial solutions if $m < n$. Therefore if $m < n$, there is a

linear relation between the vectors of L contradicting its construction as linearly independent. Therefore $m \geq n$.

However, there may be some $\vec{s}_i \in S$ s.t. $\vec{s}_i = \vec{0}$ but only one because S is a set not a list. If that is all there is (i.e. $S = \{\vec{0}\}$) then the space $V = \{\vec{0}\}$ and the only linearly independent set of vectors in the zero space is the empty set. In this case $n = 0$ and so we cannot have $m < n$. Therefore $m \geq n$ also.

On the other hand, if there are non-zero elements in S then we can remove the equation i such that $\vec{s}_i = \vec{0}$ so that we will have $m - 1$ simultaneous equations in our system. But now the linear dependence of L follows if $(m - 1) < n \iff m < (n + 1)$ and clearly $m < n \implies m < n + 1$ so once again $m < n$ implies that L is not linearly independent.

Proposition 74. *Any two bases of the same finite-dimensional vector space have the same number of elements. In other words: For a given vector space, the cardinality of bases is fixed.*

Proof. Let L, L' be finite subsets of the finite-dimensional vector space V such that both are bases. Then both L and L' are linearly independent and span the space. Therefore we have,

$$|L| \leq |L'| \quad \text{and} \quad |L| \geq |L'|$$

which implies that $|L| = |L'|$. □

2.4.7.5 Dimension of Finite-Dimensional Vector Spaces

Definition. The **dimension** of a finite-dimensional vector space v is the number of vectors in a basis. The dimension will be denoted by $\dim V$.

Theorem 18. *The dimension of a finite-dimensional vector space is an upper bound on the cardinality of a linearly independent set of vectors in the space and a lower bound on the cardinality of a spanning set of vectors in the same space.*

Proof. This theorem follows from Proposition 73. \square

Theorem 19. *Any linearly independent set of vectors in a finite-dimensional vector space V of cardinality $\dim V$ is a basis.*

Proof. Let $L \subset V$ be linearly independent set of vectors in V with $|L| = \dim V$. By Proposition 72, any linearly independent set in V may be extended with zero or more vectors to obtain a basis of V . But any basis of V has dimension $\dim V = |L|$. Therefore we extend L with zero vectors to obtain a basis. \square

Proposition 75. *If $W \subseteq V$ is a subspace of a finite-dimensional vector space then,*

$$\dim W \leq \dim V \quad \text{and} \quad (\dim W = \dim V) \iff (W = V).$$

Proof. Firstly note that W must also be finite-dimensional because by definition there is a finite set of vectors that spans V and since $W \subseteq V$, the same set must also span W .

Every basis of W is also a linearly independent set in V and so, by Proposition 73, it cannot have cardinality greater than any basis of V . Any basis in V has cardinality $\dim V$ so the cardinality of any basis of W must be less than or equal to $\dim V$. Therefore $\dim W \leq \dim V$.

If $\dim W = \dim V$ on the other hand, every basis of W has the same cardinality as every basis of V . Since every basis of W is also a linearly independent set in V with cardinality equal to $\dim V$, by Theorem 19 it is a basis of V . By similar logic in reverse, any basis of V is also a basis of W . Let B be such a basis. Then,

$$\vec{v} \in W \iff \vec{v} \in \text{span } B \iff \vec{v} \in V.$$

Therefore $W = V$. \square

Theorem 20. *Let W_1, W_2 be subspaces of a finite-dimensional vector space. Then,*

$$\dim(W_1 + W_2) = \dim W_1 + \dim W_2 - \dim(W_1 \cap W_2).$$

Proof. It is easy to show that the intersection of two subspaces is a subspace. So define a basis of $W_1 \cap W_2$ as $B = \{\vec{u}_1, \dots, \vec{u}_r\}$ where $r = \dim(W_1 \cap W_2)$. Then B is a linearly independent set in W_1 and can be extended to a basis of W_1 ,

$$B_{W_1} = \{\vec{u}_1, \dots, \vec{u}_r, \vec{v}_1, \dots, \vec{v}_{m-r}\}$$

where $m = \dim W_1$.

By the same reasoning we can extend B to a basis of W_2 ,

$$B_{W_2} = \{\vec{u}_1, \dots, \vec{u}_r, \vec{w}_1, \dots, \vec{w}_{n-r}\}$$

where $n = \dim W_2$.

Now if we can show that,

$$B' = B_{W_1} \cup B_{W_2} = \{\vec{u}_1, \dots, \vec{u}_r, \vec{v}_1, \dots, \vec{v}_{m-r}, \vec{w}_1, \dots, \vec{w}_{n-r}\}$$

is a basis of $W_1 + W_2$ then the proof will follow easily.

Clearly, B' spans $W_1 + W_2$ as B' contains a basis of both W_1 and W_2 and so is able to express any sum of two vectors chosen from the two spaces.

Linear Independence is a little more complicated however. Consider that, since B_{W_1} is linearly independent, there can be no linear relation between the vectors \vec{u}_i and \vec{v}_i and similarly there can be no such relation between the vectors \vec{u}_i and \vec{w}_i . Therefore, if there were to be a linear relation among the vectors it would have to involve the vectors \vec{v}_i with the vectors from B_{W_2} or the vectors \vec{w}_i with the vectors from B_{W_1} . If we model the linear relation as a relation between vectors that are linear combinations of the vectors \vec{u}_i, \vec{v}_i and \vec{w}_i we get,

$$\vec{u} + \vec{v} + \vec{w} = \vec{0} \iff \vec{v} = -\vec{u} - \vec{w} \in W_2$$

therefore $\vec{v} \in W_1 \cap W_2$ and so \vec{v} can be expressed as a linear combination of the vectors in B . But this implies a linear relation among the vectors of B_{W_1} which are linearly independent by construction. Therefore

$$\vec{v} = \vec{0} \implies -\vec{u} - \vec{w} = \vec{0} \iff -\vec{u} = \vec{w}$$

which implies that there is a linear relation among the vectors of B_{W_2} which are also linearly independent by construction and so $\vec{u} = \vec{w} = \vec{0}$. \square

Corollary 19. *Let W_1, \dots, W_n be a set of subspaces of a finite-dimensional vector space. Then,*

$$\dim(W_1 + \dots + W_n) \leq \dim W_1 + \dots + \dim W_n$$

with the equality case iff the spaces W_1, \dots, W_n are independent.

Proof. This follows by induction on Theorem 20.

If we take the case of $n = 2$ as our base case then,

$$\dim(W_1 + W_2) = \dim W_1 + \dim W_2 - \dim(W_1 \cap W_2) \leq \dim W_1 + \dim W_2.$$

Furthermore, if W_1, W_2 are independent then $W_1 \cap W_2 = \{\vec{0}\}$ so that,

$$\dim(W_1 + W_2) = \dim W_1 + \dim W_2 - 0 = \dim W_1 + \dim W_2.$$

For the induction step: Let $W' = W_1 + \cdots + W_{n-1}$ and take as the induction hypothesis that,

$$\dim(W') \leq \dim W_1 + \cdots + \dim W_{n-1}$$

with equality being when the spaces are independent. Then observe that, by the associativity of addition of vectors and hence of vector spaces,

$$\begin{aligned} \dim(W_1 + \cdots + W_n) &= \dim((W_1 + \cdots + W_{n-1}) + W_n) \\ &= \dim W' + \dim W_n - \dim(W' \cap W_n) && \text{by Theorem 20} \\ &\leq \dim W' + \dim W_n \\ &\leq \dim W_1 + \cdots + \dim W_{n-1} + \dim W_n && \text{by induction hypothesis.} \end{aligned}$$

Furthermore, if $\dim W' = \dim W_1 + \cdots + \dim W_{n-1}$ and if the spaces W' and W_n are independent then the intersection,

$$W' \cap W_n = \{\vec{0}\}$$

which gives,

$$\begin{aligned} \dim(W_1 + \cdots + W_n) &= \dim((W_1 + \cdots + W_{n-1}) + W_n) \\ &= \dim W' + \dim W_n - 0 && \text{by Theorem 20} \\ &= \dim W_1 + \cdots + \dim W_{n-1} + \dim W_n && \text{by induction hypothesis.} \end{aligned}$$

This shows that if the spaces are independent then

$$\dim(W_1 + \cdots + W_n) = \dim W_1 + \cdots + \dim W_n$$

but it is also easy to reverse the logic and show that if the above equality holds then, if we add a space to the sum of other spaces, then the intersection of the added space with the sum of the other spaces must have dimension 0 and therefore must be $\{\vec{0}\}$. In this way we can also prove the converse implication that the spaces must be independent. \square

Proposition 76. *Two finite-dimensional vector spaces may be isomorphic only if they have the same dimension.*

Proof. We want to prove that if $\phi : W \mapsto V$ is an isomorphism of vector spaces then it follows that $\dim W = \dim V$.

Assume for contradiction that ϕ is indeed an isomorphism between the vector spaces W and V but,

$$\dim W = m > \dim V = n.$$

Then any basis of $\vec{w}_1, \dots, \vec{w}_m \in W$ is a linearly independent set of vectors in W with the property, for $c_1, \dots, c_m \in \mathbb{F}$,

$$\begin{aligned} c_1 \vec{w}_1 + \dots + c_m \vec{w}_m = \vec{0} &\iff c_1, \dots, c_m = 0 \\ \iff \phi(c_1 \vec{w}_1 + \dots + c_m \vec{w}_m) = \phi(\vec{0}) = \vec{0} &\iff c_1, \dots, c_m = 0 \\ \iff c_1 \phi(\vec{w}_1) + \dots + c_m \phi(\vec{w}_m) = \vec{0} &\iff c_1, \dots, c_m = 0. \end{aligned}$$

But we have $\phi(\vec{w}_1), \dots, \phi(\vec{w}_m) \in V$ so that the result obtained implies that $\phi(\vec{w}_1), \dots, \phi(\vec{w}_m)$ is a linearly independent set of vectors in V of cardinality $m > n = \dim V$. By Theorem 18 this cannot be and we have obtained a contradiction. \square

[TODO:](#) example application of calculating the order of $GL_2(\mathbb{F})$ when $\mathbb{F} = \mathbb{F}_p$ is a prime field. Artin[114]

2.4.7.6 Direct Sums in Finite-Dimensional Vector Spaces

Proposition 77. *Let W_1, \dots, W_n be subspaces of a finite-dimensional vector space V , and let B_i be a basis for W_i . Then, the ordered set B obtained by listing the bases B_1, \dots, B_n in order is a basis of V iff V is the direct sum $W_1 \otimes \dots \otimes W_n$.*

Proof. Let $B = \bigcup_i B_i$ be a basis of V . Then B is a linearly independent set and so,

$$\forall \vec{b} \in B . \vec{b} \notin \text{span}(B \setminus \{\vec{b}\})$$

which also means that,

$$\forall B_i \subset B . \text{span } B_i \cap \text{span}(B \setminus B_i) = \vec{0}.$$

This implies that the sets B_i are independent spaces. Since their union is a basis of V then together they span V . Since the span of each B_i is W_i then, $V = W_1 \oplus \cdots \oplus W_n$.

Conversely, if $V = W_1 \oplus \cdots \oplus W_n$ then each subspace W_i is independent so that, for every $\vec{b} \in B_i$, $\vec{b} \notin \text{span}(B \setminus B_i)$. This means that if we begin with B_1 and add to it the elements of B_2 then the resulting set remains linearly independent and we can continue this until we have $B = \bigcup_i B_i$ as a linearly independent set. Then B contains all the basis vectors of every W_i and therefore, by $V = W_1 \oplus \cdots \oplus W_n$, B spans the space V . It is therefore a basis of V . \square

Proposition 78. *Let W be a subspace of a finite-dimensional vector space V . Then there is another subspace W' such that $W \oplus W' = V$.*

Proof. Any basis of W is a linearly independent set in V and can be extended to a basis of V by adding a set of linearly independent vectors S . Then S is a basis of a subspace W' such that $W \oplus W' = V$. \square

2.4.8 Infinite Sets of Vectors

The definition of a vector space defines what it means to add two vectors and so, by extension, arbitrarily large *finite* sums of vectors but not what it means to add an infinite number of vectors. Therefore, we define the span and linear independence of infinite sets of vectors as conditions over finite subsets of the vectors.

2.4.8.1 Span and Linear Independence of Infinite Sets of Vectors

Definition. The **span of an infinite set of vectors** S is defined to be the set of **finite** linear combinations of its elements,

$$\{ c_1 \vec{v}_1 + \cdots + c_r \vec{v}_r \mid \vec{v}_i \in S, c_i \in \mathbb{F} \}$$

where r is finite but may be arbitrarily large.

Definition. An **infinite set of vectors** is defined to be **linearly independent** if there is no linear relation among a **finite** subset of them,

$$c_1 \vec{v}_1 + \cdots + c_n \vec{v}_n = \vec{0}$$

where $c_i \in F$.

These definitions of span and linear independence are compatible with the corresponding definitions for finite sets of vectors.

2.4.8.2 Infinite-Dimensional Vector Spaces

Definition. A vector space is called **infinite-dimensional** if there is no finite set of vectors that spans the space.

2.4.8.3 Examples of Infinite-Dimensional Vector Spaces

(35) The space \mathbb{R}^∞ of infinite real vectors $(a) = (a_1, a_2, a_3, \dots)$ can also be thought of as the space of sequences $\{a_n\}$ of real numbers. It has many important subspaces:

- Convergent sequences: $C = \{(a) \in \mathbb{R}^\infty \mid \lim_{n \rightarrow \infty} a_n \text{ exists}\}$.
- Bounded sequences: $l^\infty = \{(a) \in \mathbb{R}^\infty \mid \{a_n\} \text{ is bounded}\}$.
- Absolutely convergent series: $l^1 = \{(a) \in \mathbb{R}^\infty \mid \sum_1^\infty |a_n| < \infty\}$.
- Sequences with finitely many nonzero terms:

$$Z = \{(a) \in \mathbb{R}^\infty \mid a_n = 0 \text{ for all but finitely many } n\}.$$

2.4.8.4 Bases of Infinite-Dimensional Vector Spaces

Definition. As with finite sets, a **basis of an infinite-dimensional vector space** is a linearly independent set which spans the space.

(36) Let $S = (e_1, e_2, \dots)$ be the infinite set of standard basis vectors in \mathbb{R}^∞ . S does not span \mathbb{R}^∞ because the vector $\vec{w} = (1, 1, 1, \dots)$ is not a *finite* linear combination of the elements of S . It *is*, however, a basis of the vector space Z in 35d.

It can be shown, using the Axiom of Choice, that every vector space has a basis but a basis of \mathbb{R}^∞ will be uncountably infinite and so will not be expressible as $(\vec{v}_1, \vec{v}_2, \dots)$.

There is much more to this topic. For more details see the University of Michigan Maths Dept. Linear Algebra Supplement on Infinite-Dimensional Vector Spaces.

2.4.9 Coordinate Vector Spaces

Definition. A *coordinate vector* is a representation of a vector as an ordered list of numbers that describes the vector in terms of a particular ordered basis.

Definition. Let $B = \{\vec{v}_1, \dots, \vec{v}_n\}$ be a basis of a finite-dimensional vector space V . Then for every $\vec{v} \in V$ it is possible to express \vec{v} as a linear combination of the vectors in B in the form,

$$\vec{v} = c_1\vec{v}_1 + \dots + c_n\vec{v}_n.$$

Then, the *coordinate vector of \vec{v} with respect to the basis B* is,

$$\vec{v}_B = \langle c_1, \dots, c_n \rangle.$$

Definition. The *dot product* of two coordinate vectors in an n -dimensional coordinate space is defined as,

$$\vec{v} \cdot \vec{w} = v_1w_1 + \dots + v_nw_n = \vec{v}^T \vec{w}.$$

Notation. The *dot product* may also be denoted $\langle \vec{v}, \vec{w} \rangle$ but this notation may also refer to a literal coordinate vector, i.e.

$$\begin{bmatrix} 3 \\ 7 \end{bmatrix} = \langle 3, 7 \rangle = (3, 7)^T.$$

Proposition 79. The properties of the dot product are:

- (i) $\langle \vec{x}, \vec{y} \rangle = \langle \vec{y}, \vec{x} \rangle$
- (ii) $\alpha \langle \vec{x}, \vec{y} \rangle = \langle \alpha \vec{x}, \vec{y} \rangle = \langle \vec{x}, \alpha \vec{y} \rangle$

$$(iii) \langle \vec{x} + \vec{y}, \vec{z} \rangle = \langle \vec{x}, \vec{z} \rangle + \langle \vec{y}, \vec{z} \rangle$$

$$(iv) \langle \vec{x}, \vec{x} \rangle \geq 0 \text{ and } \langle \vec{x}, \vec{x} \rangle = 0 \iff \vec{x} = 0$$

Proof. Proofs of these properties follow from the definition of the dot product.

- (i) $\langle \vec{x}, \vec{y} \rangle = \langle \vec{y}, \vec{x} \rangle$ follows from the commutativity of multiplication in a field,

$$x_i y_i = y_i x_i.$$

- (ii) $\alpha \langle \vec{x}, \vec{y} \rangle = \langle \alpha \vec{x}, \vec{y} \rangle = \langle \vec{x}, \alpha \vec{y} \rangle$ follows from the associativity of multiplication in a field,

$$\alpha x_i y_i = x_i \alpha y_i.$$

- (iii) $\langle \vec{x} + \vec{y}, \vec{z} \rangle = \langle \vec{x}, \vec{z} \rangle + \langle \vec{y}, \vec{z} \rangle$ follows from the distributivity of multiplication over addition in a field,

$$(x_i + y_i) z_i = x_i z_i + y_i z_i.$$

- (iv) $\langle \vec{x}, \vec{x} \rangle \geq 0$ and $\langle \vec{x}, \vec{x} \rangle = 0 \iff \vec{x} = 0$ follows from its form as a sum of squares,

$$x_1^2 + x_2^2 + \cdots + x_n^2.$$

□

2.4.9.1 Examples of coordinate vectors

- (37) Let $\vec{x} = \langle 3, 3 \rangle \in \mathbb{R}^2$ and B be the set $\{\langle 2, 0 \rangle, \langle 1, 3 \rangle\}$ so that B is a non-standard basis of the space \mathbb{R}^2 . Then we can also define \vec{x} with respect to the basis B as follows.

$$\vec{x} = 1 \cdot \begin{bmatrix} 2 \\ 0 \end{bmatrix} + 1 \cdot \begin{bmatrix} 1 \\ 3 \end{bmatrix}$$

Therefore, the **coordinate vector of \vec{x} with respect to the basis B** is defined as,

$$\vec{x}_B = \langle 1, 1 \rangle.$$

2.4.9.2 Bases as Matrices

Notation. If $B = \{\vec{v}_1, \dots, \vec{v}_n\}$ is a set of vectors then, when it seems necessary to clearly differentiate, the matrix whose columns are the elements of B will be denoted $[B]$. However, for convenience, when the mathematical context makes clear whether we are referring to a set of vectors or a matrix we will use B to indicate either the set or the matrix.

Proposition 80. *If \vec{x}_B is a coordinate vector of \vec{x} w.r.t. the basis B then left multiplication by the matrix $[B]$ whose columns are the elements of B converts \vec{x}_B to its standard coordinate form \vec{x} ,*

$$\vec{x} = [B]\vec{x}_B.$$

Proof. Let $B = \{\vec{v}_1, \dots, \vec{v}_n\}$ be a basis and the basis matrix

$$[B] = \begin{bmatrix} \vec{v}_1 & \dots & \vec{v}_n \end{bmatrix}$$

and $\vec{x} = \langle x_1, \dots, x_n \rangle$ is a vector in standard coordinates. Then if we calculate \vec{x} using the basis B ,

$$\vec{x} = x_{B1}\vec{v}_1 + \dots + x_{Bn}\vec{v}_n$$

we get a coordinate vector w.r.t. B ,

$$\vec{x}_B = \langle x_{B1}, \dots, x_{Bn} \rangle.$$

So, clearly, to recover the standard coordinate vector we need to apply \vec{x}_B to the basis against which it was defined,

$$[B]\vec{x}_B = \begin{bmatrix} \vec{v}_1 & \dots & \vec{v}_n \end{bmatrix} \begin{bmatrix} x_{B1} \\ \vdots \\ x_{Bn} \end{bmatrix} = x_{B1}\vec{v}_1 + \dots + x_{Bn}\vec{v}_n = \vec{x}. \quad \square$$

Corollary 20. *If \vec{x} is a coordinate vector w.r.t. the standard basis and B is an alternative basis then left multiplication by the matrix $[B]^{-1}$ converts \vec{x} to \vec{x}_B its form w.r.t. the basis B .*

Proof.

$$[B]\vec{x}_B = \vec{x} \iff \vec{x}_B = [B]^{-1}\vec{x}.$$

□

One way to think of the action of the basis matrix is that it is used to encode/decode coordinates into/from its basis coordinates.

For example, if \vec{x}_B is a coordinate vector w.r.t. the basis B , then left-multiplying it by the basis matrix $[B]$,

$$\vec{x} = [B]\vec{x}_B$$

decodes the coordinate vector into standard coordinates. Conversely we can use the inverse of the basis matrix,

$$\vec{x}_B = [B]^{-1}\vec{x}$$

to encode a standard coordinate vector into B -coordinates.

2.4.9.3 Relationship with Abstract Vector Spaces

Proposition 81. *Every vector space V of dimension n is isomorphic to the space \mathbb{F}^n of column vectors.*

Proof. Let $\phi : \mathbb{F}^n \mapsto V$ be defined as $\phi(\vec{x}) = B\vec{x}$ where B is a matrix whose columns are a basis of V . The map ϕ is surjective because the columns of B span the space V and it is injective because they are linearly independent. So ϕ is a bijection.

The structure of the vector space is preserved because,

$$\phi(\vec{x}_1 + \vec{x}_2) = B(\vec{x}_1 + \vec{x}_2) = B\vec{x}_1 + B\vec{x}_2 = \phi(\vec{x}_1) + \phi(\vec{x}_2)$$

and

$$\phi(c\vec{x}) = B(c\vec{x}) = cB\vec{x} = c\phi(\vec{x}).$$

□

Note that, by Proposition 76, \mathbb{F}^n is **not** isomorphic to \mathbb{F}^m for $m \neq n$. Every finite-dimensional vector space V is isomorphic to \mathbb{F}^n , for some uniquely determined integer n .

So, the finite-dimensional vector spaces are completely classified by Proposition 81 and any problem on finite-dimensional vector spaces may be reduced to a problem on column vectors and matrices.

It is a result of Proposition 81 that we can use coordinate spaces as vector spaces.

2.4.9.4 Using Coordinate Spaces to analyse Vectors

Span If we have a set of n vectors in \mathbb{F}^m then we can determine if a vector \vec{b} is in the span of the set of vectors by solving the system,

$$A\vec{x} = \vec{b}$$

where $\vec{x} \in \mathbb{F}^n$ and A is an $m \times n$ matrix whose columns are the set of vectors. If there is some \vec{x} that satisfies the equation then \vec{b} is in the span.

Linear Independence If we have a set of n vectors in \mathbb{F}^m then we can determine linear independence by solving a system of homogeneous linear equations,

$$A\vec{x} = \vec{0}$$

where $\vec{x} \in \mathbb{F}^n$ and A is an $m \times n$ matrix whose columns are the set of vectors. If there is a non-trivial solution — a non-zero \vec{x} for which $A\vec{x} = \vec{0}$ — then the set of vectors is not linearly independent.

2.4.9.5 Change of Basis

Definition. If we represent an abstract vector \vec{v} as a coordinate vector with respect to two different bases $B = \{\vec{b}_1, \dots, \vec{b}_n\}$ and $B' = \{\vec{b}'_1, \dots, \vec{b}'_n\}$ then we have,

$$\vec{v} = a_1 \vec{b}_1 + \dots + a_n \vec{b}_n = a'_1 \vec{b}'_1 + \dots + a'_n \vec{b}'_n.$$

Using the notation $\vec{x}_B = \langle a_1, \dots, a_n \rangle$, $\vec{x}_{B'} = \langle a'_1, \dots, a'_n \rangle$ and letting B, B' from here on refer to the matrices whose columns are the elements of B, B' , we can rewrite the previous equation with matrices as,

$$\vec{v} = B\vec{x}_B = B'\vec{x}_{B'}.$$

Then the **change of basis** from B to B' is the mapping,

$$\vec{x}_{B'} = (B')^{-1} B \vec{x}_B = P \vec{x}_B.$$

So we have,

$$(B')^{-1} B = P \iff B = B' P$$

which shows that P is the mapping between the two bases. This is known as the **matrix of change of basis**.

Note that another way to think about this is to say that \vec{x} is the coordinate vector of \vec{v} with respect to the standard basis and,

$$\vec{x} = B\vec{x}_B = B'\vec{x}_{B'}$$

so we always use the standard basis as a reference.

Proposition 82. If $B \in \mathbb{F}^{n \times n}$ is a basis of a finite vector space V then, for $P \in GL_n(\mathbb{F})$,

$$BP^{-1} = B'$$

is another basis of V .

Proof. As a member of $GL_n(\mathbb{F})$, P is invertible and so is a bijective mapping. It follows then that each of the basis vectors forming the columns of B are in the span of the columns of B' and so the columns of B' must also span the space V . Furthermore, since we must also have $B' \in \mathbb{F}^{n \times n}$, there are n columns in B' and so they are a spanning set with cardinality equal to the set of columns of B which is a basis of V . Therefore, by Theorem 19, the columns of B' are a basis of V . \square

2.4.9.6 Matrix of Change of Basis

Definition. Let B, B' be two different bases of the same space. Then, the matrix P such that,

$$B = B'P \quad \text{and} \quad \vec{x}_{B'} = P\vec{x}_B$$

is known as the **matrix of change of basis** between B and B' .

The matrix of change of basis P contains the vectors of one basis defined w.r.t. to the other basis. So,

$$B = B'P$$

tells us that if we apply the coordinate vectors in P to the basis B' we get the basis vectors of B .

Conversely, if we have a coordinate vector \vec{x}_B defined w.r.t. to B and we treat P as though it were a basis and apply the coordinate vector \vec{x}_B to it,

$$\vec{x}_{B'} = P\vec{x}_B$$

we get the same vector defined w.r.t. to the basis B' .

2.4.9.7 Euclidean Coordinate Spaces

Definition. The **Euclidean coordinate spaces** are the real coordinate spaces \mathbb{R}^n equipped with the **standard norm**,

$$|\vec{v}| = \sqrt{v_1^2 + \cdots + v_n^2}$$

which is considered to be the **length** of the vector in the space \mathbb{R}^n .

- The Euclidean norm (or standard norm) is the square root of the dot product of the vector with itself,

$$\vec{v} \cdot \vec{v} = v_1^2 + \cdots + v_n^2 = |\vec{v}|^2.$$

- In a Euclidean coordinate space the metric (wikipedia:metric), or distance function, is the length of the vector between the two points considered to be position vectors in the corresponding Euclidean vector space. So, the distance d between the point whose position vector is \vec{v} and the point whose position vector is \vec{w} is,

$$d = |\vec{v} - \vec{w}| = |\vec{w} - \vec{v}|.$$

Proposition 83. The dot product of two distinct vectors in a Euclidean space is related to their lengths by the formula,

$$\vec{v} \cdot \vec{w} = |\vec{v}| |\vec{w}| \cos \theta$$

where θ is the angle between the vectors.

Proof. The properties of the dot product (Proposition 79) tell us that,

$$\vec{v} \cdot \vec{w} = \vec{w} \cdot \vec{v} \quad \text{and} \quad (\vec{v} - \vec{w}) \cdot \vec{x} = \vec{v} \cdot \vec{x} - \vec{w} \cdot \vec{x}$$

so if we let $\vec{x} = (\vec{v} - \vec{w})$ also we obtain,

$$\begin{aligned} (\vec{v} - \vec{w}) \cdot (\vec{v} - \vec{w}) &= \vec{v} \cdot (\vec{v} - \vec{w}) - \vec{w} \cdot (\vec{v} - \vec{w}) \\ &= \vec{v} \cdot \vec{v} - \vec{v} \cdot \vec{w} - \vec{w} \cdot \vec{v} + \vec{w} \cdot \vec{w} \\ &= \vec{v} \cdot \vec{v} + \vec{w} \cdot \vec{w} - 2\vec{v} \cdot \vec{w}. \end{aligned}$$

So this gives us a formula for the square of the length of the displacement vector between \vec{v} and \vec{w} or, in other words, the square of the distance between \vec{v} and \vec{w} . Now, if we imagine \vec{v} and \vec{w} and the vector $\vec{v} - \vec{w}$ forming a triangle we can also refer to the law of cosines which tells us that,

$$c^2 = a^2 + b^2 - 2ab \cos \theta$$

where θ is the angle subtended by a and b and c is the side opposite the angle. Letting,

$$a = |\vec{v}|, b = |\vec{w}|, c = |\vec{v} - \vec{w}|$$

we obtain,

$$|\vec{v} - \vec{w}|^2 = |\vec{v}|^2 + |\vec{w}|^2 - 2 |\vec{v}| |\vec{w}| \cos \theta.$$

But the Euclidean norm is the square root of the dot product so also,

$$|\vec{v} - \vec{w}|^2 = (\vec{v} - \vec{w}) \cdot (\vec{v} - \vec{w})$$

therefore

$$\begin{aligned} & |\vec{v}|^2 + |\vec{w}|^2 - 2 |\vec{v}| |\vec{w}| \cos \theta = \vec{v} \cdot \vec{v} + \vec{w} \cdot \vec{w} - 2 \vec{v} \cdot \vec{w} \\ \iff & |\vec{v}|^2 + |\vec{w}|^2 - 2 |\vec{v}| |\vec{w}| \cos \theta = |\vec{v}|^2 + |\vec{w}|^2 - 2 \vec{v} \cdot \vec{w} \\ \iff & -2 |\vec{v}| |\vec{w}| \cos \theta = -2 \vec{v} \cdot \vec{w} \\ \iff & |\vec{v}| |\vec{w}| \cos \theta = \vec{v} \cdot \vec{w}. \quad \square \end{aligned}$$

This relationship is the foundation of analytic geometry.

2.4.9.8 Orthogonality

Definition. Two nonzero vectors \vec{v} and \vec{w} are said to be **orthogonal** if their dot product is zero, i.e.

$$\vec{v} \cdot \vec{w} = 0.$$

Theorem 21. *Two nonzero vectors \vec{v} and \vec{w} are perpendicular to each other if and only if their dot product is zero.*

Proof. If the dot product $\vec{v} \cdot \vec{w} = 0$ then, by Proposition 83,

$$\vec{v} \cdot \vec{w} = |\vec{v}| |\vec{w}| \cos \theta = 0$$

where θ is the angle between \vec{v} and \vec{w} . Since \vec{v} and \vec{w} are both nonzero, the product of their lengths $|\vec{v}| |\vec{w}| > 0$. Then we must have $\cos \theta = 0 \iff \theta \in \{\pi/2, 3\pi/2\}$.

So, $\vec{v} \cdot \vec{w} = 0$ implies that the vectors are perpendicular.

Conversely, we can use the same logic in reverse to show that, for perpendicular vectors, $\cos \theta = 0$ which means that the dot product will be 0. \square

Corollary 21. *Geometric orthogonality — that's to say perpendicularity — is equivalent to the dot product definition of orthogonality.*

2.5 Linear Transformations

2.5.1 Basic Properties of Linear Transformations

The analogue for vector spaces of a homomorphism of groups is a map,

$$T : V \longmapsto W$$

from one vector space over a field \mathbb{F} to another, which is compatible with addition and scalar multiplication:

$$T(\vec{v}_1 + \vec{v}_2) = T(\vec{v}_1) + T(\vec{v}_2) \quad \text{and} \quad T(c\vec{v}_1) = cT(\vec{v}_1),$$

for all $\vec{v}_1, \vec{v}_2 \in V$, $c \in \mathbb{F}$.

Definition. A homomorphism between two vector spaces that is also compatible with scalar multiplication is called a **linear transformation** or **linear map or mapping**.

The compatibility with addition of vectors implies that a linear transformation is a homomorphism between additive groups of vectors.

2.5.1.1 The Kernel and the Image of a Linear Transformation

Definition. Let $T : V \longmapsto W$ be any linear transformation. Then the **kernel** (or **nullspace**) of T is defined as,

$$\ker T = \{ \vec{v} \mid T(\vec{v}) = \vec{0} \}$$

and the **image** of T as,

$$\operatorname{im} T = \{ \vec{w} \in W \mid \exists \vec{v} \in V . \vec{w} = T(\vec{v}) \}.$$

Proposition 84. *The kernel of $T : V \mapsto W$ is a subspace of V and the image is a subspace of W .*

Proof. T is a homomorphism between additive groups of vectors and so the proof that the kernel and image are subspaces is the same as for general homomorphisms (see 2.1.4.1). \square

Proposition 85. *The fibres of a linear transformation are the additive cosets of the kernel.*

Proof. Let $T : V \mapsto W$ be any linear transformation with kernel $K = \ker T$. Then, for any fixed $\vec{v} \in V$,

$$\forall \vec{k} \in K . T(\vec{v} + \vec{k}) = T(\vec{v}) + T(\vec{k}) = T(\vec{v}) + \vec{0} = T(\vec{v}).$$

So every element in the additive coset $\vec{v} + K$ maps to the same value $T(\vec{v})$ in W . Therefore we have,

$$\text{im } T = \{ \vec{w} \in W \mid \exists (\vec{v} + K) \subseteq V . \{ \vec{w} \} = T(\vec{v} + K) \}.$$

We could also express this using the inverse image as in 2.1.5 but the notation would be easily confused for the inverse transformation. \square

Proposition 86. *Linear Independence is preserved across a linear transformation iff the transformation is injective.*

Proof. Let $T : V \mapsto W$ be any linear transformation with kernel $K = \ker T$. If the kernel is nontrivial then there exists a nonempty basis of the kernel $B_K = \{\vec{k}_1, \dots, \vec{k}_n\}$. Being a basis B_K is linearly independent so that,

$$\alpha_1 \vec{k}_1 + \dots + \alpha_n \vec{k}_n = \vec{0} \iff \alpha_1, \dots, \alpha_n = 0.$$

However, $T(B_K)$ the image of B_K under T is

$$\{\vec{0}, \dots, \vec{0}\}$$

which, by Proposition 64, is obviously not linearly independent.

We can also see it more directly from the definition of a linear relation.

$$\begin{aligned} & T(\alpha_1 \vec{k}_1 + \dots + \alpha_n \vec{k}_n) = \vec{0} \\ \iff & \alpha_1 T(\vec{k}_1) + \dots + \alpha_n T(\vec{k}_n) = \vec{0} \\ \iff & \alpha_1 \vec{0} + \dots + \alpha_n \vec{0} = \vec{0} \end{aligned}$$

So clearly,

$$\alpha_1 T(\vec{k}_1) + \cdots + \alpha_n T(\vec{k}_n) = \vec{0} \not\Rightarrow \alpha_1, \dots, \alpha_n = 0$$

which proves that if T is not injective then it does not preserve linear independence.

Conversely, if T does not preserve linear independence then, if $U = \{\vec{u}_1, \dots, \vec{u}_n\} \subset V$ is a linearly independent set in V , there is a nontrivial linear relation between the vectors in $T(U)$ the image of U under T . That's to say,

$$\alpha_1 T(\vec{u}_1) + \cdots + \alpha_n T(\vec{u}_n) = \vec{0} \quad \text{where} \quad \prod_{i=1}^n \alpha_i \neq 0.$$

But this means that,

$$\begin{aligned} & \alpha_1 T(\vec{u}_1) + \cdots + \alpha_n T(\vec{u}_n) = \vec{0} \\ \iff & T(\alpha_1 \vec{u}_1 + \cdots + \alpha_n \vec{u}_n) = \vec{0} \\ \iff & \alpha_1 \vec{u}_1 + \cdots + \alpha_n \vec{u}_n \in \ker T. \end{aligned}$$

Since U is linearly independent there is no nontrivial linear relation between its elements and since $\prod_{i=1}^n \alpha_i \neq 0$ we can conclude that,

$$\alpha_1 \vec{u}_1 + \cdots + \alpha_n \vec{u}_n \neq \vec{0} \implies \ker T \neq \{\vec{0}\}$$

and therefore this proves that if T does not preserve linear independence then it is not injective. \square

2.5.1.2 Examples of Linear Transformations

- (38) As previously seen in section 2.3.3, matrix multiplication on the left is a linear transformation. Let A be an $m \times n$ matrix with entries in \mathbb{F} and consider A as an operator on column vectors $A : \mathbb{F}^n \mapsto \mathbb{F}^m$. The kernel of A is the set of vectors that are solutions to $A\vec{x} = \vec{0}$ while the image (or range) is the set of vectors \vec{b} such that $A\vec{x} = \vec{b}$ has a solution. The solutions $\{\vec{x} \in \mathbb{F}^n \mid A\vec{x} = \vec{b}\}$ for some fixed $\vec{b} \in \mathbb{F}^m$ are the additive coset $\vec{v} + K$ where K is the kernel of A and $\vec{v} \in \mathbb{F}^n$ is such that $A\vec{v} = \vec{b}$. Compare with 17.

- (39) Also previously seen in 2.4.4 is that polynomials can be modeled as vectors. Let P_n be the vector space of real polynomials of degree $\leq n$. Then the derivative is a linear transformation $P_n \mapsto P_{n-1}$. The kernel of the derivative is the set of degree 0 polynomials (i.e. constant functions) and the additive cosets of the kernel are $f(x) + c$, for $f(x) \in P_n$ and constant c .

2.5.1.3 The Dimension of a Linear Transformation

Definition. The dimension of the image is called the **rank** while the dimension of the kernel is known as the **nullity**.

Dimension and rank also exist for Groups (see wikipedia) where it refers to the minimal generating set for the group.

Theorem 22. The Dimension Formula: Let $T : V \mapsto W$ be a linear transformation, and assume that V is finite dimensional. Then,

$$\dim V = \dim(\ker T) + \dim(\operatorname{im} T) = \text{rank} + \text{nullity}.$$

Proof. Let $\{\vec{k}_1, \dots, \vec{k}_m\}$ be a basis of $\ker T$. Then, by Proposition 72, it may be extended to a basis of V ,

$$B = \{\vec{k}_1, \dots, \vec{k}_m, \vec{u}_1, \dots, \vec{u}_n\}.$$

So, for any $\vec{v} \in V$, \vec{v} may be expressed as a linear combination of the vectors in B . Therefore, for any $\vec{w} \in \operatorname{im} T$,

$$\begin{aligned} \vec{w} &= T(\alpha_1 \vec{k}_1 + \dots + \alpha_m \vec{k}_m + \beta_1 \vec{u}_1 + \dots + \beta_n \vec{u}_n) \\ \iff &= T(\alpha_1 \vec{k}_1) + \dots + T(\alpha_m \vec{k}_m) + T(\beta_1 \vec{u}_1) + \dots + T(\beta_n \vec{u}_n) \\ \iff &= \vec{0} + T(\beta_1 \vec{u}_1) + \dots + T(\beta_n \vec{u}_n) \\ \iff &= \beta_1 T(\vec{u}_1) + \dots + \beta_n T(\vec{u}_n) \end{aligned}$$

This shows that $B' = \{T(\vec{u}_1), \dots, T(\vec{u}_n)\}$ spans $\text{im } T$. Furthermore, if there were a linear relation between the elements of B' then,

$$\begin{aligned}
 & \beta_1 T(\vec{u}_1) + \dots + \beta_n T(\vec{u}_n) = \vec{0} \\
 \iff & T(\beta_1 \vec{u}_1 + \dots + \beta_n \vec{u}_n) = \vec{0} \\
 \iff & \beta_1 \vec{u}_1 + \dots + \beta_n \vec{u}_n \in \ker T \\
 \iff & \beta_1 \vec{u}_1 + \dots + \beta_n \vec{u}_n = \alpha_1 \vec{k}_1 + \dots + \alpha_m \vec{k}_m
 \end{aligned}$$

where this last result implies a linear relation between the vectors of B . Since B is a basis this linear relation can only be the trivial relation and so $\beta_1, \dots, \beta_n = 0$ and B' is linearly independent also. \square

Notes about this:

- This formula bears a resemblance to Lagrange's Theorem applied to homomorphisms of finite groups (10),

$$|G| = |\ker \phi| \cdot |\text{im } \phi|.$$

The difference, however, is that the Dimension Formula of Linear Transformations is dealing with the generators of a group while Lagrange's Theorem is dealing with the orders of the groups. The orders of the groups are the number of elements in the group that are generated by the generators of the group. In the case of a real vector space, the vectors generated by the basis vectors are uncountably infinite due to scalar multiplication by real numbers and so cardinality doesn't apply in the same way.

- This formula **only applies to finite-dimensional vector spaces**. This should be clear as we simply cannot do this kind of arithmetic with ∞ . For example, if the rank is infinite then the dimension of the kernel would be $\infty - \infty = ?$.

2.5.2 Linear Transformations as Matrices

Proposition 87. *Left multiplication by a $m \times n$ matrix is a linear transformation $\mathbb{F}^n \mapsto \mathbb{F}^m$.*

Proof. Let $T : \mathbb{F}^n \mapsto \mathbb{F}^m$ be a linear transformation. Then T is a map from n -vectors to m -vectors that is compatible with the vector space operations. Let A be an $m \times n$ matrix then $A\vec{x} = \vec{b}$ where $\vec{x} \in \mathbb{F}^n$ and $\vec{b} \in \mathbb{F}^m$ showing that $T(\vec{x}) = A\vec{x} = \vec{b}$ is a map of the form $\mathbb{F}^n \mapsto \mathbb{F}^m$. Furthermore,

$$T(\vec{x}_1 + \vec{x}_2) = A(\vec{x}_1 + \vec{x}_2) = A\vec{x}_1 + A\vec{x}_2 = T(\vec{x}_1) + T(\vec{x}_2)$$

and

$$T(c\vec{x}) = A(c\vec{x}) = cA\vec{x} = cT(\vec{x})$$

which shows that left multiplication preserves vector addition and scalar multiplication so that $T(\vec{x}) = A\vec{x} = \vec{b}$ is a linear map as required. \square

Theorem 23. *Every linear transformation $\mathbb{F}^n \mapsto \mathbb{F}^m$ is left multiplication by a particular $m \times n$ matrix.*

Proof. For any $\vec{x} = \langle x_1, \dots, x_n \rangle \in \mathbb{F}^n$ we can write it as,

$$x_1\vec{e}_1 + \dots + x_n\vec{e}_n.$$

Therefore if $T : \mathbb{F}^n \mapsto \mathbb{F}^m$ then,

$$T(\vec{x}) = T(x_1\vec{e}_1 + \dots + x_n\vec{e}_n) = T(\vec{e}_1)x_1 + \dots + T(\vec{e}_n)x_n \in \mathbb{F}^m$$

and so letting $A \in \mathbb{F}^{m \times n}$ be,

$$A = \begin{bmatrix} T(\vec{e}_1) & \dots & T(\vec{e}_n) \end{bmatrix}$$

we have,

$$T(\vec{x}) = A\vec{x}. \quad \square$$

Corollary 22. *Any linear transformation between spaces isomorphic to \mathbb{F}^n and \mathbb{F}^m (refer to Proposition 81 and 34) is left multiplication by a particular $m \times n$ matrix.*

This is why linear transformations from a space to itself can be wholly characterized by what they do to the axes and also why every such linear transformation can be considered a change of basis and vice-versa.

Conceptually, a linear transformation changes the coordinates of a selection of transformed vectors. The confusion comes about because we implement the matrix of the linear transformation A by transforming the basis against which the coordinates are applied,

$$A = \begin{bmatrix} T(\vec{e}_1) & \cdots & T(\vec{e}_n) \end{bmatrix}.$$

*Whereas a change of basis transforms the basis against which coordinates are applied **and then updates the coordinates to balance out the change**. This can be seen if we deconstruct the change of basis formula:*

$$B\vec{x}_B = B'\vec{x}_{B'} \iff \vec{x}_{B'} = (B')^{-1}B\vec{x}_B$$

$$\vec{x}_{B'} = P\vec{x}_B = (B')^{-1}B\vec{x}_B$$

This can be thought of as first obtaining the coordinates against the standard basis $B\vec{x}_B$ and then applying the inverse of the target basis so as to obtain the equivalent coordinates against the target basis. But, note, we could also consider the whole thing as a linear transformation represented by P .

The biggest difference, however, is that a linear transformation can also be between different spaces — say from n -dimensional space to m -dimensional space — in which case it cannot be thought of as a change of basis as a vector $\vec{v} \in \mathbb{F}^n$ cannot be equivalently expressed using a basis of \mathbb{F}^m because the two spaces are not isomorphic.

2.5.2.1 Examples of Linear Transformations as Matrices

(40) Let $T : \mathbb{R}^2 \mapsto \mathbb{R}^2$ be a linear transformation such that,

$$T(\vec{e}_1) = \begin{bmatrix} 1 \\ 2 \end{bmatrix} \quad \text{and} \quad T(\vec{e}_2) = \begin{bmatrix} -1 \\ 0 \end{bmatrix}.$$

The transformation T has been completely described in this way because, for any $\vec{x} = \langle x_1, x_2 \rangle \in \mathbb{R}^2$ we have,

$$\begin{aligned} T(\vec{x}) &= T(x_1\vec{e}_1 + x_2\vec{e}_2) \\ \iff T(\vec{x}) &= x_1T(\vec{e}_1) + x_2T(\vec{e}_2) \\ \iff T(\vec{x}) &= x_1 \begin{bmatrix} 1 \\ 2 \end{bmatrix} + x_2 \begin{bmatrix} -1 \\ 0 \end{bmatrix} = \begin{bmatrix} x_1 - x_2 \\ 2x_1 \end{bmatrix}. \end{aligned}$$

So, T is also left multiplication by the matrix,

$$A = \begin{bmatrix} 1 & -1 \\ 2 & 0 \end{bmatrix}.$$

2.5.2.2 Linear Transformations and Change of Basis

It is often possible to achieve powerful simplifications of problems by selecting appropriate bases. In this section we will look at linear transformations represented by matrices between arbitrary bases of spaces. So here we are looking at the relationship between linear transformations and change of basis.

Definition. If the matrix A of a linear transformation $T : V \mapsto W$ is defined as, for $\vec{x} \in V$,

$$T(\vec{x}) = A\vec{x} = \vec{b} \in W$$

then **the matrix of T with respect to the bases $B \subset V$ and $B' \subset W$** is defined as the matrix A that satisfies,

$$A\vec{x}_B = \vec{b}_{B'} \in W$$

and also

$$T(\vec{x}) = [B']A\vec{x}_B = \vec{b}$$

2.5.2.3 Intuition of the Matrix of T with respect to Bases

Let $T : V \mapsto W$ be a linear transformation and let $B_V = \{\vec{v}_1, \dots, \vec{v}_n\}$ be a basis of V and $B_W = \{\vec{w}_1, \dots, \vec{w}_m\}$ be a basis of W . Then $T(\vec{v}_i) \in W$ and so there is some $m \times n$ matrix $A = (a_{ij})$ such that,

$$\begin{bmatrix} \vec{w}_1 & \cdots & \vec{w}_m \end{bmatrix} A = \begin{bmatrix} T(\vec{v}_1) & \cdots & T(\vec{v}_n) \end{bmatrix}.$$

where,

$$T(\vec{v}_j) = \begin{bmatrix} \vec{w}_1 & \cdots & \vec{w}_m \end{bmatrix} \begin{bmatrix} a_{1j} \\ \vdots \\ a_{mj} \end{bmatrix} = a_{1j}\vec{w}_1 + \cdots + a_{mj}\vec{w}_m = \sum_i a_{ij}\vec{w}_i$$

so that the j th column of the matrix A is the coordinate vector of $T(\vec{v}_j)$ with respect to the basis B_W .

Substituting $B_W = \{\vec{w}_1, \dots, \vec{w}_m\}$ we can obtain an expression for A ,

$$\begin{aligned} & \begin{bmatrix} \vec{w}_1 & \cdots & \vec{w}_m \end{bmatrix} A = \begin{bmatrix} T(\vec{v}_1) & \cdots & T(\vec{v}_n) \end{bmatrix} \\ \iff & [B_W]A = \begin{bmatrix} T(\vec{v}_1) & \cdots & T(\vec{v}_n) \end{bmatrix} \\ \iff & A = [B_W]^{-1} \begin{bmatrix} T(\vec{v}_1) & \cdots & T(\vec{v}_n) \end{bmatrix}. \end{aligned}$$

The matrix A is referred to as the **matrix of T with respect to the bases B_V and B_W** and conforms to,

$$A\vec{x}_{B_V} = \vec{b}_{B_W}.$$

This can be seen as,

$$\begin{aligned} & A\vec{x}_{B_V} = [B_W]^{-1} \begin{bmatrix} T(\vec{v}_1) & \cdots & T(\vec{v}_n) \end{bmatrix} \vec{x}_{B_V} \\ \iff & A\vec{x}_{B_V} = [B_W]^{-1} \vec{b}_{B_V} \\ \iff & A\vec{x}_{B_V} = \vec{b}_{B_W} \end{aligned}$$

so that \vec{x}_{B_V} — the coordinate vector with respect to B_V — is first transformed by applying the coordinates to the transformed version of the basis

B_V and then these coordinates are converted to B_W coordinates by left multiplication by $[B_W]^{-1}$.

If we chose different bases for the spaces we would get a different matrix. If the bases are the standard bases then the matrix is the standard matrix for the transformation.

2.5.2.4 Examples of Linear Transform Matrices w.r.t. Bases

- (41) Let $T : \mathbb{R}^2 \mapsto \mathbb{R}^2$ be a linear transform defined (against the standard basis) by,

$$T \left(\begin{bmatrix} 1 \\ 0 \end{bmatrix} \right) = \begin{bmatrix} 2 \\ 3 \end{bmatrix} \quad \text{and} \quad T \left(\begin{bmatrix} 0 \\ 1 \end{bmatrix} \right) = \begin{bmatrix} 3 \\ 2 \end{bmatrix}.$$

Then, if we define the matrix of T with respect to the standard basis only then we have,

$$A = \left[T \left(\begin{bmatrix} 1 \\ 0 \end{bmatrix} \right) \quad T \left(\begin{bmatrix} 0 \\ 1 \end{bmatrix} \right) \right] = \begin{bmatrix} 2 & 3 \\ 3 & 2 \end{bmatrix}$$

and, if we define a vector $\vec{x} = \langle 1, 1 \rangle$ against the standard basis we can see that,

$$T(\vec{x}) = 1 \cdot T \left(\begin{bmatrix} 1 \\ 0 \end{bmatrix} \right) + 1 \cdot T \left(\begin{bmatrix} 0 \\ 1 \end{bmatrix} \right) = A\vec{x} = \begin{bmatrix} 2 & 3 \\ 3 & 2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 5 \\ 5 \end{bmatrix}.$$

If we now define B , an alternative basis of \mathbb{R}^2 , as

$$B = \left\{ \begin{bmatrix} 3 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 2 \end{bmatrix} \right\}$$

then we have,

$$[B] = \begin{bmatrix} 3 & 0 \\ 0 & 2 \end{bmatrix} \quad \text{and} \quad [B]^{-1} = \frac{1}{6} \begin{bmatrix} 2 & 0 \\ 0 & 3 \end{bmatrix}$$

and the linear transform matrix **of coordinate vectors w.r.t. the basis B** is defined as,

$$A = \left[T \left(\begin{bmatrix} 3 \\ 0 \end{bmatrix} \right) \quad T \left(\begin{bmatrix} 0 \\ 2 \end{bmatrix} \right) \right] = \begin{bmatrix} 6 & 6 \\ 9 & 4 \end{bmatrix}.$$

Now *this* matrix A expects coordinate vectors w.r.t. B and so if we convert \vec{x} to basis B as follows,

$$\vec{x}_B = [B]^{-1}\vec{x} = \frac{1}{6} \begin{bmatrix} 2 & 0 \\ 0 & 3 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/2 \end{bmatrix}$$

then we find that,

$$T(\vec{x}) = A\vec{x}_B = \begin{bmatrix} 6 & 6 \\ 9 & 4 \end{bmatrix} \begin{bmatrix} 1/3 \\ 1/2 \end{bmatrix} = \begin{bmatrix} 6/3 + 6/2 \\ 9/3 + 4/2 \end{bmatrix} = \begin{bmatrix} 5 \\ 5 \end{bmatrix}.$$

If we were to define another basis of \mathbb{R}^2 called B' and construct the matrix of T with respect to B and B' then the matrix A would become,

$$A = [B']^{-1} \left[T \left(\begin{bmatrix} 3 \\ 0 \end{bmatrix} \right) \quad T \left(\begin{bmatrix} 0 \\ 2 \end{bmatrix} \right) \right]$$

and to get the result in standard coordinates we would need to apply the result to the basis vectors of B' ,

$$T(\vec{x}) = [B']A\vec{x}_B.$$

In the case where we want the result to be in the same basis as the argument vector \vec{x}_B we still need to modify the matrix A . In this case A becomes,

$$A = [B]^{-1} \left[T \left(\begin{bmatrix} 3 \\ 0 \end{bmatrix} \right) \quad T \left(\begin{bmatrix} 0 \\ 2 \end{bmatrix} \right) \right]$$

This A still expects \vec{x} to be in B coordinates but outputs a result defined in B coordinates rather than standard coordinates.

$$A\vec{x}_B = \frac{1}{6} \begin{bmatrix} 2 & 0 \\ 0 & 3 \end{bmatrix} \begin{bmatrix} 6 & 6 \\ 9 & 4 \end{bmatrix} \vec{x}_B = \begin{bmatrix} 2 & 2 \\ 9/2 & 2 \end{bmatrix} \begin{bmatrix} 1/3 \\ 1/2 \end{bmatrix} = \begin{bmatrix} 5/3 \\ 5/2 \end{bmatrix}.$$

So, to get the result in standard coordinates we need to apply the result to the basis vectors of B' ,

$$T(\vec{x}) = [B](A\vec{x}_B) = \begin{bmatrix} 3 & 0 \\ 0 & 2 \end{bmatrix} \begin{bmatrix} 5/3 \\ 5/2 \end{bmatrix} = \begin{bmatrix} 5 \\ 5 \end{bmatrix}.$$

Proposition 88. *Let A be the matrix of a linear transformation $T : V \mapsto W$ with respect to the bases B_V, B_W of dimension n and m respectively. The matrices A' which represent T with respect to other bases are those of the form,*

$$A' = QAP^{-1}$$

where $Q \in GL_m(\mathbb{F}), P \in GL_n(\mathbb{F})$.

Proof. Let $B'_V = \{\vec{v}'_1, \dots, \vec{v}'_n\}, B'_W = \{\vec{w}'_1, \dots, \vec{w}'_m\}$ be alternative bases with respect to which we want to find A' , the matrix of T . Also let,

$$B_V = B'_V P \quad \text{and} \quad B_W = B'_W Q.$$

Then for $\vec{v}'_i \in B'_V$, $T(\vec{v}') \in \text{span } B'_W$ so there exists a matrix A' such that, using the notation $T(B_V)$ to indicate the image of the set B_V under T and $[B_V]$ for the matrix whose columns are the elements of B_V ,

$$\begin{aligned} & \begin{bmatrix} \vec{w}'_1 & \cdots & \vec{w}'_m \end{bmatrix} A' = \begin{bmatrix} T(\vec{v}'_1) & \cdots & T(\vec{v}'_n) \end{bmatrix} \\ \iff & [B'_W] A' = [T(B'_V)] \\ \iff & A' = [B'_W]^{-1} [T(B'_V)] \\ \iff & A' = [B'_W]^{-1} [T(B_V P^{-1})] \\ \iff & A' = Q [B_W]^{-1} [T(B_V)] P^{-1} \quad \text{P is coefficient matrix} \\ \iff & A' = QAP^{-1}. \quad \square \end{aligned}$$

2.5.2.5 Simplification of the Matrix of a Transformation

Proposition 89. *Let $T : V \mapsto W$ be a linear transformation of rank r . Bases B_V, B_W may be chosen so that the matrix of T takes the form,*

$$A = \begin{bmatrix} I_r & \vdots \\ \cdots & 0 \end{bmatrix}$$

which is the reduced row echelon form of a bijective matrix of rank r .

Proof. Let $U = \{\vec{u}_1, \dots, \vec{u}_k\}$ be a basis of the kernel of T where $k = \dim(\ker T)$. Then U may be extended to a basis of V (Proposition 72),

$$B_V = \{\vec{v}_1, \dots, \vec{v}_r, \vec{u}_1, \dots, \vec{u}_k\}.$$

Then, let $T(\vec{v}_i) = \vec{w}_i$ so that,

$$[T(B_V)] = \begin{bmatrix} \vec{w}_1 & \dots & \vec{w}_r & \vec{0} & \dots & \vec{0} \end{bmatrix}.$$

As shown in Theorem 22, $\{\vec{w}_1, \dots, \vec{w}_r\}$ is a basis of the image of T and can also be extended to a basis of W ,

$$B_W = \{\vec{w}_1, \dots, \vec{w}_r, \vec{x}_1, \dots, \vec{x}_{m-r}\}.$$

So, the matrix A of T with respect to the bases B_V, B_W satisfies,

$$\begin{aligned} [B_W]A &= [T(B_V)] \\ \iff \begin{bmatrix} \vec{w}_1 & \dots & \vec{w}_r & \vec{x}_1 & \dots & \vec{x}_{m-r} \end{bmatrix} A &= \begin{bmatrix} \vec{w}_1 & \dots & \vec{w}_r & \vec{0} & \dots & \vec{0} \end{bmatrix}. \end{aligned}$$

If we look at the components of the matrices we see,

$$\begin{bmatrix} w_{11} \dots & w_{1r} & x_{11} \dots & x_{1(m-r)} \\ \vdots & & & \\ w_{r1} \dots & w_{rr} & x_{r1} \dots & x_{r(m-r)} \\ \vdots & & & \\ w_{m1} \dots & w_{mr} & x_{m1} \dots & x_{m(m-r)} \end{bmatrix} \begin{bmatrix} a_{11} & \dots & a_{1n} \\ \vdots & & \\ a_{r1} & \dots & a_{rn} \\ \vdots & & \\ a_{m1} & \dots & a_{mn} \end{bmatrix} = \begin{bmatrix} w_{11} \dots & w_{1r} & 0 \dots & 0 \\ \vdots & & & \\ w_{1r} \dots & w_{rr} & 0 \dots & 0 \\ \vdots & & & \\ 0 \dots & 0 & 0 \dots & 0 \end{bmatrix}$$

which makes it clear that A has the form,

$$A = \begin{bmatrix} 1 & \dots & 0 & 0 & \dots & 0 \\ \vdots & \ddots & \vdots & & & \\ 0 & \dots & 1 & 0 & \dots & 0 \\ \vdots & & & & & \\ 0 & \dots & 0 & 0 & \dots & 0 \end{bmatrix}$$

as required. □

2.5.3 Linear Operators and Eigenvectors

Definition. A **linear operator** is a linear transformation from a space to itself.

Definition. A linear operator is called **singular** if it does not have an inverse and **nonsingular** if it has an inverse.

Proposition 90. A linear operator $T : V \mapsto V$ is bijective iff it has a trivial kernel.

Proof. By Theorem 13 of homomorphisms, if T has a trivial kernel then it is injective which implies that $|\text{im } T| = |V|$ so that it is also surjective (because the domain is equal to the codomain). Therefore, it is bijective. Conversely, if it is bijective then it is injective and so it has a trivial kernel. \square

Corollary 23. A linear operator $T : V \mapsto V$ is isomorphic iff it has a trivial kernel.

Corollary 24. A linear operator is nonsingular iff it has a trivial kernel.

Note that these are not true of linear transformations in general because, for a transformation between different vector spaces of different dimensions, injectivity does not imply bijectivity and invertibility.

Proposition 91. The following conditions on a linear operator $T : V \mapsto V$ on a finite-dimensional vector space are equivalent:

- (i) $\ker T = \{0\}$.
- (ii) $\text{im } T = V$.
- (iii) If A is the matrix of the operator with respect to an arbitrary basis, then $\det A \neq 0$.

(iv) T is singular

Proof. The first two of these properties follow from the dimension formula for finite-dimensional vector spaces. The third follows since an operator with a non-trivial kernel is noninvertible (Proposition 90) and so its matrix will also be noninvertible; noninvertible matrices have determinant 0. The last property follows from the definition of singular and the fact that T is noninvertible. \square

Again it's worth noting the contrast with transformations in general: for a transformation whose domain vector space is lower dimension than its codomain, the transformation may be injective but not surjective (kernel is trivial but $\text{im } T < V$), and conversely, if the domain is higher dimension than the codomain then the transformation may be surjective while not injective ($\ker T > 0$ but image covers codomain). Furthermore, the third condition relating to the determinant, only applies to operators because the determinant is a property of square matrices.

Note that the first two properties do not hold for infinite-dimensional vector spaces. For example, let $V = \mathbb{R}^\infty$ be the space of sequences a_1, a_2, \dots . Then the "shift operator" is defined as,

$$T(a_1, a_2, \dots) = (0, a_1, a_2, \dots).$$

This is a linear operator because,

$$\begin{aligned} \alpha T(a_1, a_2, \dots) + \beta T(b_1, b_2, \dots) &= \alpha(0, a_1, a_2, \dots) + \beta(0, b_1, b_2, \dots) \\ \iff &= (0, \alpha a_1, \alpha a_2, \dots) + (0, \beta b_1, \beta b_2, \dots) \\ \iff &= (0, \alpha a_1 + \beta b_1, \alpha a_2 + \beta b_2, \dots) \\ \iff &= T(\alpha a_1 + \beta b_1, \alpha a_2 + \beta b_2, \dots). \end{aligned}$$

However, while clearly for this operator we have $\text{im } T < V$, nevertheless the kernel of this operator is the trivial kernel $\{\vec{0}\}$. This just shows further that the dimension formula (Theorem 22) for finite-dimensional vector spaces does not apply for infinite-dimensional spaces.

Proposition 92. *Let A be the matrix of a linear operator T with respect to the basis B of dimension n . The matrices A' which represent T with respect to other bases are those of the form,*

$$A' = PAP^{-1}$$

where $P \in GL_n(\mathbb{F})$.

Proof. A linear operator is a specialization of a linear transformation where the domain and codomain are the same set so we can use the definition (2.5.2.2) of the matrix of a linear transformation w.r.t. to two different bases and simply set the two bases to the same set B . This produces,

$$T(\vec{x}) = A\vec{x}_B = \vec{b}_B$$

so that,

$$[B]A = [T(\vec{b}_1) \cdots T(\vec{b}_n)] \iff A = [B]^{-1}[T(\vec{b}_1) \cdots T(\vec{b}_n)].$$

Now let there be another basis B' related to B by,

$$[B] = [B']P \iff [B]P^{-1} = [B'],$$

$$P\vec{x}_B = \vec{x}_{B'}.$$

To define a matrix A' that performs the same linear transformation as A w.r.t. to the basis B' we need,

$$T(\vec{x}) = A'\vec{x}_{B'} = \vec{b}_{B'}$$

and

$$[B']A' = [T(\vec{b}'_1) \cdots T(\vec{b}'_n)] \iff A' = [B']^{-1}[T(\vec{b}'_1) \cdots T(\vec{b}'_n)].$$

If we have the vectors of B' encoded in B -coordinates then we can use the transformed version of the basis B to produce the transformed version of the basis B' ,

$$[T(\vec{b}'_1) \cdots T(\vec{b}'_n)] = [T(\vec{b}_1) \cdots T(\vec{b}_n)][B]^{-1}[B'] = [T(\vec{b}_1) \cdots T(\vec{b}_n)]P^{-1}.$$

If we now note that,

$$A = [B]^{-1}[T(\vec{b}_1) \cdots T(\vec{b}_n)] \quad \text{and} \quad A' = [B']^{-1}[T(\vec{b}'_1) \cdots T(\vec{b}'_n)]$$

then,

$$\begin{aligned}
& A' = [B']^{-1}[T(\vec{b}_1) \cdots T(\vec{b}_n)]P^{-1} \\
\iff & A' = [B']^{-1}[B]AP^{-1} \\
\iff & A' = PAP^{-1}. \quad \square
\end{aligned}$$

2.5.3.1 Similar Matrices

Definition. If two matrices A, A' are related by,

$$A' = PAP^{-1}$$

for some $P \in GL_n(\mathbb{F})$ then they are known as **similar** matrices or **conjugates**.

Note:

- Similar matrices represent the same linear transformation expressed with respect to different bases.
- In general, linear maps are not invertible so the matrix A is not necessarily a member of a group. So, this use of the term "conjugation" only equates to the term in group theory within the general linear group $GL_n(\mathbb{F})$.
- Similar matrices have the same determinant which means that they expand or contract the space by the same amount. This is to be expected as they represent the same linear transformation defined against different bases.

2.5.3.2 Intuition of Similar Matrices

$$A' = PAP^{-1} \iff A'P = PA$$

P is the change of basis matrix such that $\vec{x}_{B'} = P\vec{x}_B$. So P is $B_{B'}$ the basis vectors of B w.r.t. to B' . Another way of looking at it: $P = [B']^{-1}[B]$ so it decodes coordinates in B -coords to standard coordinates and then encodes them into B' -coords. Meanwhile A transforms the basis vectors of B and then encodes the result in B -coords so the schematic of A is $A = [B]^{-1}[T(B)]$.

So

$$PA = P[B]^{-1}[T(B)] = [B']^{-1}[B][B]^{-1}[T(B)] = [B']^{-1}[T(B)].$$

By similar reasoning the schematic of A' is $A' = [B']^{-1}[T(B')]$. But also we have,

$$A' = PAP^{-1} = ([B']^{-1}[B])([B]^{-1}[T(B)])([B]^{-1}[B']) = [B']^{-1}[T(B)][B]^{-1}[B']$$

so that,

$$\begin{aligned} [B']^{-1}[T(B')] &= [B']^{-1}[T(B)][B]^{-1}[B'] \\ \iff [T(B')] &= [T(B)][B]^{-1}[B']. \end{aligned}$$

What this last result is saying is that if we take the basis vectors of B' and encode them in B -coordinates and then apply the result to the transformed basis of B then the result is the transformed basis of B' defined in standard coordinates.

So, we have,

$$\begin{aligned} A'P &= PA = [B']^{-1}[T(B)] \\ P^{-1}A' &= AP^{-1} = [B]^{-1}[T(B)][B]^{-1}[B'] = [B]^{-1}[T(B')]. \end{aligned}$$

Proposition 93. *Similarity of matrices is an equivalence relation.*

Proof. For any $M, N \in GL_n(\mathbb{F})$, similarity is an equivalence relation because of the following properties.

Reflexivity:

$$\begin{aligned} N &= I^{-1}NI \\ \therefore N &\sim N \end{aligned}$$

Symmetry:

$$\begin{array}{ll}
& N = P^{-1}MP \\
\Longleftrightarrow & NP^{-1} = P^{-1}M(PP^{-1}) \\
\Longleftrightarrow & NP^{-1} = P^{-1}M \\
\Longleftrightarrow & PNP^{-1} = (PP^{-1})M \\
\Longleftrightarrow & PNP^{-1} = M \\
\Longleftrightarrow & R^{-1}NR = M, \quad R \in X \\
\therefore & N \sim M \Longleftrightarrow M \sim N
\end{array}$$

Transitivity:

$$\begin{array}{ll}
& N = P^{-1}MP, \quad M = Q^{-1}AQ \\
\Rightarrow & N = P^{-1}(Q^{-1}AQ)P \\
\Longleftrightarrow & N = (P^{-1}Q^{-1})A(QP) \\
\Longleftrightarrow & N = R^{-1}AR, \quad R \in X \\
\therefore & (N \sim M) \wedge (M \sim Q) \Longleftrightarrow (N \sim Q)
\end{array}$$

□

Proposition 94. *Similar matrices have the same eigenvalues.*

Proof. Similar matrices represent the same transformation with respect to different bases and for a matrix A representing a transformation T with respect to an arbitrary basis B ,

$$T(\vec{v}) = c\vec{v} \Longleftrightarrow A\vec{v}_B = c\vec{v}_B.$$

That's to say, the eigenvalues are not dependent on the basis with respect to which a coordinate vector is defined. □

2.5.3.3 Invariant Subspaces

Definition. Let $T : V \mapsto V$ be a linear operator on a vector space. A subspace W of V is called an **invariant subspace** or a **T -invariant subspace** if it is carried to itself by the operator,

$$TW \subset W.$$

In other words, W is T -invariant if $T(\vec{w}) \in W$ for all $\vec{w} \in W$. In this case, T may be referred to as defining an operator on W that is the **restriction of T to W** .

Let $W \subseteq V$ be a T -invariant subspace of V with basis $B_W = \vec{w}_1, \dots, \vec{w}_k$. Then B_W can be extended to a basis of V ,

$$B_V = \{\vec{w}_1, \dots, \vec{w}_k, \vec{v}_1, \dots, \vec{v}_{n-k}\}.$$

If we look at the matrix A of T with respect to this basis B_V we find a characteristic pattern.

$$A = [B_V]^{-1}[T(B_V)] = [B_V]^{-1}[T(\vec{w}_1) \cdots T(\vec{w}_k) T(\vec{v}_1) \cdots T(\vec{v}_{n-k})]$$

with $T(\vec{w}_1), \dots, T(\vec{w}_k) \in W$ so that,

$$T(\vec{w}_i) = \alpha_1 \vec{w}_1 + \cdots + \alpha_k \vec{w}_k.$$

This means that when we express $T(\vec{w}_i)$ with respect to the basis B_V the coordinate vectors will take the form,

$$(\alpha_1, \dots, \alpha_k, 0, \dots, 0)^T.$$

As a result, the matrix A will take the form,

$$A = \begin{bmatrix} C & D \\ 0 & E \end{bmatrix}$$

where C is a $k \times k$ matrix that represents the restriction of T to W .

A description of a matrix of the form of the description of A here is known as a **block decomposition**.

On the other hand, if $V = W_1 \oplus W_2$ where *both* W_1 and W_2 are T -invariant subspaces then A takes the form,

$$A = \begin{bmatrix} C & 0 \\ 0 & E \end{bmatrix}$$

where, as before, C is a $k \times k$ matrix that represents the restriction of T to W_1 but, this time, also E is a $(n - k) \times (n - k)$ matrix that represents the restriction of T to W_2 .

Matrices with the form of the matrix,

$$\begin{bmatrix} C & 0 \\ 0 & E \end{bmatrix}$$

are known as **block diagonal matrices** or **diagonal block matrices**.

2.5.3.4 Eigenvectors

Definition. An **eigenvector** of a linear operator T is a nonzero vector \vec{v} with the property under T that,

$$T(\vec{v}) = c\vec{v}$$

for some constant $c \in \mathbb{F}$. The constant c is called an **eigenvalue**.

Eigenvectors may also be referred to as **characteristic vectors** and eigenvalues as **characteristic values**.

Note:

- To have eigenvectors, T must be an operator as the image of the eigen-

vector $\vec{v} \in V$ under T ,

$$T(\vec{v}) = c\vec{v} \in V$$

is, by definition, in the same space as the eigenvector itself. In fact, this is the key of the relationship between eigenvectors and invariant subspaces. The space spanned by eigenvectors of T is T -invariant. Conversely, the bases of T -invariant subspaces are eigenvectors. If W is a 2-dimensional T -invariant subspace, for example, and its image under T is one-dimensional (a line inside the plane of W) then vectors in W that are not in the plane of the line image will be in the kernel of T and so also eigenvectors but with the eigenvalue 0. It is also worth noting that they are still considered to be parallel to their image under T because, by convention, the zero vector $\vec{0}$ is considered to be parallel to all vectors.

- An eigenvector may not be $\vec{0}$ but an eigenvalue may be 0. As a consequence of this, every nonzero vector in the kernel is an eigenvector (with eigenvalue 0).
- If we speak of an eigenvector of a square matrix then we are referring to an eigenvector of **left multiplication** by the matrix, i.e. a nonzero vector \vec{x} such that,

$$A\vec{x} = c\vec{x}.$$

Clearly if \vec{x}_B is the coordinate vector of $\vec{v} \in V$ with respect to B , a basis of V and A is the matrix of T with respect to the basis B , then

$$A\vec{x}_B = c\vec{x}_B \iff T(\vec{v}) = c\vec{v}.$$

As a result, all similar matrices have the same eigenvalues.

2.5.3.5 Examples of Eigenvectors and Eigenvalues

(42) If we take the matrix,

$$\begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix}$$

we can determine its eigenvectors by finding the solutions to the equation,

$$\begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = c \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}.$$

This gives us two simultaneous equations in three unknowns: the two vector dimensions x_1, x_2 and the eigenvalue c ,

$$\begin{aligned} 3x_1 + x_2 &= cx_1 \\ 2x_2 &= cx_2 \end{aligned}$$

which imply that,

$$6x_1 = 2cx_1 - cx_2 \iff (6/c - 2)x_1 = -x_2 \iff x_2 = (2 - 6/c)x_1.$$

So, if $x_1 = 1$ then $x_2 = 2 - (6/c)$ and we have the following eigenvector/eigenvalue pairs.

$$c = 3 : \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \quad c = 1 : \begin{bmatrix} 1 \\ -4 \end{bmatrix} \text{ Wrong!}$$

Why does this not work? There is an additional constraint not expressed in the linear system here: eigenvectors are nonzero by definition. This means that we have the additional constraint,

$$x_1 x_2 \neq 0$$

which cannot be expressed in a linear system of equations. When $c \notin \{2, 3\}$ both x_1 and x_2 are zero and the vector is not an eigenvector. So, how can we systematically restrict the values of c to only those that produce valid eigenvectors? If we follow the alternative logic:

$$\begin{aligned} & \begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = c \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \\ \iff & \begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} - c \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \vec{0} \\ \iff & \left(\begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix} I - cI \right) \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \vec{0} \\ \iff & \begin{bmatrix} 3-c & 1 \\ 0 & 2-c \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \vec{0} \end{aligned}$$

Now if we let

$$A = \begin{bmatrix} 3 - c & 1 \\ 0 & 2 - c \end{bmatrix}$$

then the nullspace of A is the space of vectors $(x_1, x_2)^T$ that satisfy this equation. If A is invertible and nonsingular then this space is the trivial space $\{\vec{0}\}$ but, if A is singular however, then there exist nonzero vectors that satisfy this equation. Therefore, it is precisely the nonzero vectors that are in the nullspace of this matrix A when it is singular that are the eigenvectors we are looking for. So, if we first determine the values of c for which A is singular, we can then determine the eigenvectors. Since, by Proposition 91, A is singular if and only if $\det A = 0$, we are looking for *precisely the values of c which make $\det A = 0$.*

2.5.3.6 Matrices of Eigenvectors

If $\vec{v}_1 \in V$ is a eigenvector of a linear transformation T and we extend the set $\{\vec{v}_1\}$ (by Proposition 72) to a basis of V , say $\{\vec{v}_1, \dots, \vec{v}_n\}$, then the matrix of T will have the block form,

$$\begin{bmatrix} c & B \\ 0 & D \end{bmatrix} = \begin{bmatrix} c & \dots & \dots \\ 0 & \dots & \dots \\ \vdots & \dots & \dots \\ 0 & \dots & \dots \end{bmatrix}$$

where c is the eigenvalue of \vec{v}_1 . This is the same block decomposition as that shown for T -invariant spaces in 2.5.3.3 with the case of a 1-dimensional invariant subspace.

Proposition 95. *If A is the matrix of a linear operator T with respect to a basis B then the matrix A is diagonal iff every basis vector in B is an eigenvector of T .*

Proof. The defining property of the matrix A is that the j -th column is the coordinates of the image of the j -th basis vector in B under T ,

$$A(:, j) = T(\vec{v}_j) = a_{1j}\vec{v}_1 + \dots + a_{nj}\vec{v}_n.$$

For an eigenvector \vec{v}_j , $T(\vec{v}_j) = c\vec{v}_j = a_{jj}\vec{v}_j$ so that $a_{jj} = c$ the eigenvalue and for all a_{ij} such that $i \neq j$, $a_{ij} = 0$. \square

Corollary 25. *The matrix of a linear operator T over a vector space V is similar to a diagonal matrix iff there exists some basis of V solely comprised of eigenvectors of T .*

2.5.4 Diagonalisation

Definition. *The process of determining a diagonal matrix that is similar to a given matrix of a linear operator is known as **diagonalisation**.*

2.5.4.1 Existence of Eigenvectors

- Every linear operator on a complex vector space has at least one eigenvector and, in most cases, these form a basis.
- Linear operators over real vector spaces need not have eigenvectors (e.g. rotation of the plane \mathbb{R}^2 by an angle θ has no eigenvector unless $\theta = 0$ or π).
- Real matrices that are *positive* (having only positive components) are guaranteed to have at least one positive eigenvector.

2.5.4.2 The Effect of Multiplication by a Positive Matrix

[TODO: Artin\[134\]](#)

2.5.4.3 Determining the Eigenvectors

The process of finding eigenvectors is to first determine the eigenvalues and then calculate the eigenvectors that correspond to those eigenvalues. Let I be the identity operator. Then,

$$T(\vec{v}) = c\vec{v} \iff T(\vec{v}) - c\vec{v} = \vec{0} \iff [T - cI](\vec{v}) = \vec{0}$$

where the expression $T - cI$ is a linear combination of linear transformations and so is also a linear transformation. Furthermore, it is an operator as both its terms are operators (in fact they need to have the same dimensions in order for the expression to make sense).

Two things are clear from this expression:

- (i) The matrix of the linear operator $T - cI$ is $A - cI$ where I is the identity matrix.
- (ii) The eigenvector \vec{v} is in the kernel of $T - cI$ and so is in the nullspace of $A - cI$.

Proposition 96. *A linear operator T has a nontrivial kernel iff 0 is an eigenvalue of T .*

Proof. This follows from the fact that, if we let $\vec{v} \neq \vec{0} \in \ker T$, then

$$T(\vec{v}) = \vec{0} = 0\vec{v} = c\vec{v}$$

for $c = 0$ so that a nontrivial kernel implies that 0 is an eigenvalue. Conversely, if 0 is an eigenvalue we must have $0\vec{v} = \vec{0} = T(\vec{v})$ and, since if a vector \vec{v} is an eigenvector, by definition, $\vec{v} \neq \vec{0}$, this therefore implies that the kernel contains a nonzero vector. \square

Corollary 26. *A linear operator T has all the properties in Proposition 91 iff 0 is an eigenvalue of T .*

Proposition 97. *The eigenvalues of a linear operator T are the scalars $c \in \mathbb{F}$ such that the linear operator $[T - cI]$ is singular.*

Proof. The eigenvalues of a linear operator T are the scalars $c \in \mathbb{F}$ such that there exists a nonzero vector \vec{v} with $[T - cI](\vec{v}) = \vec{0}$. If such a vector exists then the kernel of $T - cI$ is nontrivial and so, by Proposition 91, $T - cI$ is singular. \square

Corollary 27. *If $A - cI$ is the matrix of $T - cI$, the eigenvalues of T are the scalars $c \in \mathbb{F}$ such that $\det(A - cI) = 0$.*

Corollary 28. *The eigenvalues of $A - cI$ are the same as the eigenvalues of $cI - A$.*

Proof. If A is a $n \times n$ matrix representing the operator T then,

$$\det(-A) = (-1)^n \det(A).$$

So, if $\det(A) = 0$ then $\det(-A) = 0$ also. Therefore,

$$\det(A - cI) = 0 \iff \det(cI - A) = 0. \quad \square$$

2.5.4.4 The Characteristic Polynomial

Notation. It is customary to use the variable t to denote the eigenvalue in the characteristic polynomial.

Definition. The *characteristic polynomial* of a linear operator T is the polynomial,

$$p(t) = \det(tI - A) = \sum s(j_1, \dots, j_n) a_{1j_1} \cdots a_{nj_n}$$

where the sum is defined over all permutations j_1, \dots, j_n of $\{1, \dots, n\}$ and $s(j_1, \dots, j_n)$ is the sign of the permutation.

The determinant is an expression in which every term is the product of a component in every column and row of the matrix with no column or row appearing more than once in each term,

$$tI - A = \begin{bmatrix} (t - a_{11}) & -a_{12} & \cdots & -a_{1n} \\ -a_{21} & (t - a_{22}) & \cdots & -a_{2n} \\ \vdots & & & \vdots \\ -a_{n1} & \cdots & \cdots & (t - a_{nn}) \end{bmatrix}.$$

It can be seen in the matrix of $tI - A$ that the highest power of t will be obtained in the term of the determinant that forms the product of all the diagonal terms $a_{11} \cdots a_{nn}$ which occurs when $j_1, \dots, j_n = 1, \dots, n$. This term of the determinant will be the product of precisely n terms containing the eigenvalue t . Therefore the result is a polynomial of degree n in the eigenvalue t .

Theorem 24. The eigenvalues of a linear operator are the roots of its characteristic polynomial.

Proof. If $p(t)$ is the characteristic polynomial of a linear operator T then the values of t for which $p(t) = 0$ are the values of t such that $\det(tI - A) = 0$ and these are precisely the eigenvalues. \square

Proposition 98. *The eigenvalues of an upper or lower triangular matrix are its diagonal entries.*

Proof. The determinant of a triangular matrix is equal to the product of its diagonal entries and if A is a triangular matrix then $tI - A$ is also triangular. Therefore, the characteristic polynomial is simply,

$$p(t) = (t - a_{11}) \cdots (t - a_{nn})$$

and so the eigenvalues are the diagonal entries a_{11}, \dots, a_{nn} . □

Proposition 99. *A positive matrix (a matrix whose entries are all positive) has at least one eigenvector with positive coordinates.*

An abstract vector does not have coordinates so when we refer to a "positive" vector with positive-valued coordinates, this is with respect to a particular basis. In this context, the basis in question is the basis with respect to which the matrix outputs the transformed vectors.

Proof. [TODO](#): review: this "proof" is not really a proof, more an example using a 2x2 matrix. □

Proposition 100. *The characteristic polynomial of a linear operator does not depend on the basis with respect to which the matrix of the operator is defined.*

Proof. For two similar matrices representing the same linear operator T we have,

$$A' = PAP^{-1}$$

where P is the matrix of change of basis between the bases of A and A' . If we form the characteristic polynomial of A' ,

$$\begin{aligned} tI - A' &= tI - PAP^{-1} \\ \iff &= PtIP^{-1} - PAP^{-1} & PtIP^{-1} &= tPP^{-1} = tI \\ \iff &= P(tI - A)P^{-1} & & \text{by distributivity of matrix multiplication.} \end{aligned}$$

Then,

$$\begin{aligned}
& \det(tI - A') = \det(P(tI - A)P^{-1}) \\
\iff & = \det P \cdot \det(tI - A) \cdot \det P^{-1} \\
\iff & = \det(tI - A).
\end{aligned}$$

This result, $\det(tI - A') = \det(tI - A)$, must hold for all t and therefore implies that, for p, p' characteristic polynomials of A and A' respectively,

$$\forall t \in \mathbb{F} . p(t) = p'(t).$$

This implies that the characteristic polynomials are equal. \square

Proposition 101. *The characteristic polynomial $p(t)$ of a matrix A has the form*

$$p(t) = t^n - (\text{tr } A)t^{n-1} + \cdots + (-1)^n(\det A),$$

where $\text{tr } A$ is the trace of A (see: 2.3.2.1):

$$\text{tr } A = a_{11} + \cdots + a_{nn}.$$

Proof. Calculation of the characteristic polynomial of a matrix A is calculation of $p(t) = \det(tI - A)$ the determinant of the matrix $tI - A$ which takes the form,

$$tI - A = \begin{bmatrix} (t - a_{11}) & \cdots & \cdots & \vdots \\ \vdots & (t - a_{22}) & \cdots & \vdots \\ \vdots & & & \vdots \\ \vdots & \cdots & (t - a_{(n-1)(n-1)}) & -a_{(n-1)n} \\ \vdots & \cdots & -a_{n(n-1)} & (t - a_{nn}) \end{bmatrix}.$$

This calculation proceeds with terms of products of elements from each row and a permutation of the column indices (see: 2.3.5.2) of which we examine the first two terms.

The first term is for the identity permutation $j_1, \dots, j_n = 1, \dots, n$ along the diagonal:

$$\begin{aligned} a_{11} \cdots a_{nn} &= (t - a_{11}) \cdots (t - a_{nn}) \\ &= t^n - (a_{11} + \cdots + a_{nn})t^{n-1} \\ &\quad + (a_{11}a_{22} + a_{11}a_{33} + \cdots + a_{(n-1)(n-1)}a_{nn})t^{n-2} \\ &\quad \cdots + (-1)^n(a_{11} \cdots a_{nn}) \end{aligned}$$

The second term is the permutation one swap away from identity

$$j_1, \dots, j_n = 1, \dots, n, (n-1)$$

which is an odd permutation, so the sign is -1 :

$$\begin{aligned} &(-1)a_{11} \cdots a_{(n-2)(n-2)}a_{(n-1)n}a_{n(n-1)} \\ &= (-1)(t - a_{11}) \cdots (t - a_{(n-2)(n-2)})(-a_{(n-1)n})(-a_{n(n-1)}) \\ &= -a_{(n-1)n}a_{n(n-1)}t^{n-2} + a_{(n-1)n}a_{n(n-1)}(a_{11} + \cdots + a_{(n-2)(n-2)})t^{n-3} \\ &\quad \cdots - (-1)^n(a_{11} \cdots a_{(n-2)(n-2)}a_{(n-1)n}a_{n(n-1)}) \end{aligned}$$

From these first two terms we can discern enough about the general pattern of the characteristic polynomial to see that the first two terms in t^n and t^{n-1} are produced by the first permutation and take the form

$$t^n - (a_{11} + \cdots + a_{nn})t^{n-1} = t^n - (\text{tr } A)t^{n-1}$$

as required. We can also see that the final terms of each permutation — those that involve no powers of t — are going to sum up to the value of $(-1)^n(\det A)$. Therefore the characteristic polynomial takes the form,

$$p(t) = t^n - (\text{tr } A)t^{n-1} + \cdots + (-1)^n(\det A)$$

as claimed. □

Corollary 29. *The trace of a matrix of a linear operator is independent of the basis with respect to which the matrix is defined.*

Proof. Let T be a linear operator and A be the matrix of T with respect to a basis B . Then Proposition 100 tells us that any matrix of the same linear operator defined with respect to some other basis (i.e. a similar matrix to A) has the same characteristic polynomial. Proposition 101 tells us that the coefficients of this characteristic polynomial include the trace of the matrix A . Therefore, if A' is a matrix of T defined against the basis B' and P is the change of basis matrix such that $B'P = B$ then,

$$A' = PAP^{-1} \iff p'(t) = p(t) \iff \operatorname{tr} A' = \operatorname{tr} A. \quad \square$$

As a result of this we can refer to the characteristic polynomial, determinant and trace of a linear operator T without reference to a particular matrix or basis.

Proposition 102. *Let T be a linear operator on a finite-dimensional vector space V .*

- (i) *If V has dimension n , then T has at most n eigenvalues.*
- (ii) *If \mathbb{F} is the field of complex numbers and $V \neq 0$, then T has at least one eigenvalue, and hence it has an eigenvector.*

Proof. (i) For any field \mathbb{F} , a polynomial of degree n can have at most n different roots (see Artin[373]). Since T is defined over a vector space of dimension n , the degree of the characteristic polynomial of T is n . Then, by Theorem 24 we can have a maximum of n eigenvalues.

- (ii) Every polynomial of positive degree with complex coefficients has at least one complex root. This fact is called the Fundamental Theorem of Algebra (wikipedia).

□

Proposition 103. *Let T be a linear operator on a finite-dimensional complex vector space V . There is a basis B of V such that the matrix A of T is upper triangular.*

Proof. By Proposition 102, T has at least one eigenvector. We can extend this eigenvector to a basis of V , say,

$$B' = \{\vec{v}'_1, \dots, \vec{v}'_n\}.$$

Then the first column of the matrix A' of T with respect to B' will be

$$(c_1, 0, \dots, 0)^T$$

where c_1 is the eigenvalue of \vec{v}'_1 . Therefore A' has the form

$$A' = \begin{bmatrix} c_1 & \cdots \\ 0 & D \end{bmatrix}$$

where D is a $(n-1) \times (n-1)$ matrix and, if $P = [B']^{-1}I = [B']^{-1}$ is the change of basis matrix, then

$$A' = PAP^{-1}.$$

Now we can use induction on the dimension of the matrix n and the induction hypothesis will be that there exists some upper triangular

$$Q' = QDQ^{-1}.$$

Define

$$Q_1 = \begin{bmatrix} 1 & 0 \\ 0 & Q \end{bmatrix}.$$

Then

$$(Q_1P)A(Q_1P)^{-1} = Q_1PAP^{-1}Q_1^{-1} = Q_1A'Q_1^{-1}$$

takes the form

$$\begin{bmatrix} c_1 & \cdots \\ 0 & QDQ^{-1} \end{bmatrix}$$

which is upper triangular. This proves the induction step. \square

Note that this proof is over the complex number field but the same proof would work over any field that contains all the roots of the characteristic polynomial.

Proposition 104. *Let $\vec{v}_1, \dots, \vec{v}_r \in V$ be eigenvectors for a linear operator T , with distinct eigenvalues c_1, \dots, c_r . Then the set $\{\vec{v}_1, \dots, \vec{v}_r\}$ is linearly independent.*

Proof. Assume for contradiction that there exists a linear relation between the set of eigenvectors,

$$\alpha_1 \vec{v}_1 + \dots + \alpha_r \vec{v}_r = \vec{0}.$$

Linearity of T gives us,

$$T(\alpha_1 \vec{v}_1 + \dots + \alpha_r \vec{v}_r) = \alpha_1 T(\vec{v}_1) + \dots + \alpha_r T(\vec{v}_r) = T(\vec{0}) = \vec{0}$$

while the eigenvector property gives us,

$$\alpha_1 T(\vec{v}_1) + \dots + \alpha_r T(\vec{v}_r) = \alpha_1 c_1 \vec{v}_1 + \dots + \alpha_r c_r \vec{v}_r$$

so we have the simultaneous equations,

$$\begin{aligned} \alpha_1 \vec{v}_1 + \dots + \alpha_r \vec{v}_r &= \vec{0} \\ \alpha_1 c_1 \vec{v}_1 + \dots + \alpha_r c_r \vec{v}_r &= \vec{0}. \end{aligned}$$

If we multiply the first equation by c_r and subtract the second equation from it we get,

$$\alpha_1(c_r - c_1)\vec{v}_1 + \dots + \alpha_{r-1}(c_r - c_{r-1})\vec{v}_{r-1} = \vec{0}.$$

Since all the eigenvalues are distinct, for $i \neq j$, $c_i - c_j \neq 0$ and the eigenvectors \vec{v}_i , by definition, are nonzero. So, this equation implies that either $\alpha_1, \dots, \alpha_{r-1} = 0$ or there is a linear relation between the vectors $\vec{v}_1, \dots, \vec{v}_{r-1}$. This dependence of the properties of the r -length list on the properties of the $(r-1)$ -length list signals that we can set up a proof by induction using the hypothesis that a k -length list is linearly independent.

If we use $k = 2$ as the base case, set up the linear relation and use the eigenvector property as before then this results in,

$$\alpha_1(c_2 - c_1)\vec{v}_1 = \vec{0}.$$

As before, both $c_2 - c_1$ and \vec{v}_1 are nonzero so this implies that $\alpha_1 = 0$. This, in turn, implies that $\alpha_2 = 0$ also and so, the list of length $k = 2$ is linearly independent.

Then the induction step is to assume that the list of length $k = r - 1$ is

linearly independent and show that this implies that the list of length $k = r$ is linearly independent. We have already shown that if we set up a linear relation on a list of eigenvectors of length $k = r$ then the eigenvector property implies that,

$$\alpha_1(c_r - c_1)\vec{v}_1 + \cdots \alpha_{r-1}(c_r - c_{r-1})\vec{v}_{r-1} = \vec{0}.$$

Now we can use the induction hypothesis to assert that $\vec{v}_1, \dots, \vec{v}_{r-1}$ are linearly independent implying that $\alpha_1, \dots, \alpha_{r-1} = 0$. This, in turn, implies that $\alpha_r = 0$ meaning that $\vec{v}_1, \dots, \vec{v}_r$ is linearly independent. \square

Theorem 25. *Let T be a linear operator on a vector space V of dimension n over a field F . Assume that its characteristic polynomial has n **distinct** roots in F . Then there is a basis for V with respect to which the matrix of T is diagonal.*

Proof. If the characteristic polynomial has n distinct roots then there are n distinct eigenvalues along with their associated eigenvectors. By ??, these eigenvectors form a linearly independent set. Since the dimension of the space V is n , by Theorem 19, these eigenvectors form a basis of V . Then, by Proposition 95, the matrix of T with respect to this basis is diagonal. \square

Note that:

- *The diagonal entries of the matrix of a linear operator with respect to a basis of eigenvectors are the eigenvalues. For this reason, the set of values is wholly determined by the linear operator although the order is determined on the order of the vectors in the basis set (which is not significant);*
- *If a matrix A is found to be similar to a diagonal matrix B via a change of basis expressed in the matrix P then,*

$$A^m = (P^{-1}BP)^m = P^{-1}B^mP.$$

This is because — remembering that matrices don't, in general, com-

note — A^m takes the form,

$$\begin{aligned} A^m &= (P^{-1}BP)(P^{-1}BP) \dots (P^{-1}BP)(P^{-1}BP) \\ &= P^{-1}B(P P^{-1})BP \dots P^{-1}B(P P^{-1})BP \\ &= P^{-1}B^m P. \end{aligned}$$

2.5.4.5 Jordan Canonical Form

[TODO:](#) from LSE

2.5.5 Orthonormal Bases and Orthogonal Operators

2.5.5.1 Orthonormal Bases

Definition. A basis consisting of pairwise mutually orthogonal (see: 2.4.9.8) unit vectors is called an **orthonormal basis**.

Proposition 105. Let $U = \{\vec{u}_1, \dots, \vec{u}_n\}$ be an orthonormal basis. Then, for any $\vec{u}_i, \vec{u}_j \in U$, $i \neq j$,

$$\vec{u}_i \cdot \vec{u}_j = 0 \quad \text{and} \quad \vec{u}_i \cdot \vec{u}_i = 1 = \vec{u}_j \cdot \vec{u}_j.$$

Proof. The elements of U are pairwise mutually orthogonal so — by the definition of orthogonal vectors (2.4.9.8) — $\vec{u}_i \cdot \vec{u}_j = 0$ for $i \neq j$. Furthermore,

$$\vec{u}_i \cdot \vec{u}_i = |\vec{u}_i|^2 = 1$$

because the length of unit vectors is 1. □

2.5.5.2 Orthonormal Bases in Coordinate Vectors

If we consider what an orthonormal basis would look like in coordinate vectors. Obviously, the standard basis is an orthonormal basis as,

$$\begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \cdot \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} = 0$$

and clearly for any i, j such that $i \neq j$, $\vec{e}_i \cdot \vec{e}_j = 0$ and \vec{e}_i, \vec{e}_j have unit length.

We can form other orthonormal bases in \mathbb{R}^3 though. The vectors,

$$\begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} \cdot \begin{bmatrix} -1 \\ 1 \\ 0 \end{bmatrix} = 0$$

are orthogonal but do not have unit length. We can make them unit length, though, by dividing them by $\sqrt{2}$ so that,

$$\begin{bmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \\ 0 \end{bmatrix}, \begin{bmatrix} -1/\sqrt{2} \\ 1/\sqrt{2} \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

is also an orthonormal basis of \mathbb{R}^3 .

Another orthonormal basis of \mathbb{R}^3 is

$$\begin{bmatrix} \cos \theta \\ \sin \theta \\ 0 \end{bmatrix}, \begin{bmatrix} -\sin \theta \\ \cos \theta \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

for any angle θ measured anticlockwise from the x -axis. Actually this is the general case of which the previous example is a special case when $\theta = \pi/2$.




Figure 2.3: $(\cos \pi/4, \sin \pi/4)^T = (1/\sqrt{2}, 1/\sqrt{2})^T$, $(-\sin \pi/4, \cos \pi/4)^T = (-1/\sqrt{2}, 1/\sqrt{2})^T$

But for any value of the angle θ these remain orthogonal as can be seen if we generate the basis vectors for $\theta = \pi/3$.




Figure 2.4: $(\cos \pi/3, \sin \pi/3)^T = (1/2, \sqrt{3}/2)^T$, $(-\sin \pi/3, \cos \pi/3)^T = (-\sqrt{3}/2, 1/2)^T$

2.5.5.3 Orthogonal Operators

Definition. A matrix whose columns form an orthonormal basis is called an **orthogonal matrix**.

The operation of left multiplication by such a matrix is called an **orthogonal operator**.

Proposition 106. If a matrix A is orthogonal then $A^T = A^{-1}$.

Proof. A is an orthogonal matrix so, by definition, its columns form an orthogonal basis. Then,

$$\begin{aligned}
 A^T A &= \begin{bmatrix} a_{11} & a_{21} & \cdots & a_{n1} \\ a_{12} & a_{22} & \cdots & a_{n2} \\ \vdots & & & \\ a_{1n} & a_{2n} & \cdots & a_{nn} \end{bmatrix} \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & & & \\ a_{n1} & a_{n2} & \cdots & a_{nn} \end{bmatrix} \\
 &= \begin{bmatrix} a_{11}^2 + \cdots + a_{n1}^2 & a_{11}a_{12} + \cdots + a_{n1}a_{n2} & \cdots \\ a_{12}a_{11} + \cdots + a_{n2}a_{n1} & a_{12}^2 + \cdots + a_{n2}^2 & \cdots \\ \vdots & & \\ \cdots & \cdots & \cdots & a_{1n}^2 + \cdots + a_{nn}^2 \end{bmatrix}
 \end{aligned}$$

Along the main diagonal the components take the form

$$a_{1j}^2 + a_{2j}^2 + \cdots + a_{nj}^2 = a_j \cdot a_j$$

where a_j is the j th column of the matrix A . But the columns of A are vectors in an orthonormal basis and so $a_j \cdot a_j = 1$.

Furthermore, the off-diagonal values take the form

$$a_{1j}a_{1j'} + \cdots + a_{nj}a_{nj'} = a_j \cdot a_{j'}$$

for $j \neq j'$. Since the columns are from an orthonormal basis we know that $a_j \cdot a_{j'} = 0$.

Therefore the resultant matrix looks like,

$$A^T A = \begin{bmatrix} 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & & & \\ 0 & 0 & \cdots & 1 \end{bmatrix} = I_n.$$

Clearly, the same effect would be seen for AA^T and so,

$$A^T A = AA^T = I \iff A^T = A^{-1}. \quad \square$$

Proposition 107. *Left multiplication by an orthogonal matrix preserves the dot product. In other words, for all vectors \vec{v}, \vec{w} , A is an orthogonal matrix if and only if,*

$$A\vec{v} \cdot A\vec{w} = \vec{v} \cdot \vec{w}.$$

Proof. Using Proposition 105 and the matrix formula for the dot product we can deduce that, for all vectors \vec{v}, \vec{w} ,

$$\begin{aligned} A\vec{v} \cdot A\vec{w} &= (A\vec{v})^T A\vec{w} \\ &= \vec{v}^T A^T A\vec{w} \\ &= \vec{v}^T \vec{w} && A \text{ is orthogonal so } A^T A = I \\ &= \vec{v} \cdot \vec{w} \end{aligned}$$

which is to say that an orthogonal matrix preserves the dot product.

Conversely, if we assume that A preserves the dot product then, for all vectors \vec{v}, \vec{w} ,

$$\begin{aligned}
& A\vec{v} \cdot A\vec{w} = \vec{v} \cdot \vec{w} \\
\iff & (A\vec{v})^T A\vec{w} = \vec{v}^T \vec{w} \\
\iff & \vec{v}^T A^T A\vec{w} = \vec{v}^T \vec{w} \\
\iff & \vec{v}^T A^T A\vec{w} - \vec{v}^T \vec{w} = 0 \quad \text{both terms are scalars} \\
\iff & \vec{v}^T (A^T A - I)\vec{w} = 0.
\end{aligned}$$

Now for any arbitrary matrix B ,

$$e_i^T B e_j = b_{ij}$$

where b_{ij} is the (i, j) th element of B . Then, for

$$e_i^T B e_j = 0$$

to be true *for all possible* e_i, e_j would require that,

$$\forall i, j. b_{ij} = 0 \iff B = [0]$$

where $[0]$ is the zero matrix. Therefore,

$$\forall \vec{v}, \vec{w}. \vec{v}^T (A^T A - I)\vec{w} = 0 \iff A^T A - I = [0] \iff A^T A = I$$

and so if A preserves the dot product then it is orthogonal. \square

Proposition 108. *The determinant of any orthogonal matrix is 1 or -1.*

Proof. If a matrix A is orthogonal then $A^T A = I$ which implies that,

$$\det(A^T)\det(A) = \det(I) = 1.$$

By Proposition 54, $\det(A^T) = \det(A)$ so,

$$\det(A^T)\det(A) = \det(A)^2 = 1 \iff \sqrt{\det(A)} = 1 \iff \det(A) = \pm 1. \quad \square$$

*If an orthogonal operator has determinant equal to 1 it is described as **orientation preserving** and if it is equal to -1 it is described as **orientation reversing**.*

Proposition 109. *The orthogonal matrices form a subgroup of $GL_n(\mathbb{F})$.*

Proof. Let $S = \{ A \in GL_n(\mathbb{F}) \mid A^T A = I \}$. Then,

- S is nonempty because $I \in S$.
- For $B, C \in S$,

$$(BC)^T(BC) = C^T B^T(BC) = C^T(B^T B)C = I$$

so $BC \in S$.

- For $B \in S$, by Proposition 105, $B^{-1} = B^T$ and

$$(B^T)^T B^T = BB^T = BB^{-1} = I$$

so S contains inverses.

Therefore $S \leq GL_n(\mathbb{F})$. □

*The subgroup of the general linear group formed by the orthogonal matrices is called the **orthogonal group** and is denoted O_n .*

2.5.6 Linear and Affine Transformations

2.5.6.1 Affine Spaces

Definition. An **affine space** is a generalization of a Euclidean space in which there is no particular point designated as the origin. As a result vectors can be viewed as displacements rather than points.

Let V be a vector space and P be a set of points. Then we can form an affine space over V and P by defining the vectors in V as displacements connecting members of P such that, for $Q_1, Q_2 \in P$, $\vec{v} \in V$,

$$Q_1 + \vec{v} = Q_2 \iff Q_2 - Q_1 = \vec{v} \iff Q_1 - Q_2 = -\vec{v}.$$

Definition. A **frame** of an affine space is an extension of a basis of its underlying vector space to include a point designated as an origin. If $\vec{v}_1, \dots, \vec{v}_n$ is a basis of a vector space V and Q is a point in the set of points P , then $F = (\vec{v}_1, \dots, \vec{v}_n, Q)$ is a frame of the affine space over V and P .

Definition. The **dimension** of an affine space is the dimension of the underlying vector space.

Proposition 110. Any linear combination of points in an affine space where the coefficients sum to 0 results in a vector.

Proof. Let S be a sum of n points $Q_i \in P$ in an affine space associated with a vector space V such that,

$$S = \sum_{i=0}^n \alpha_i Q_i \quad \text{and} \quad \sum_{i=0}^n \alpha_i = 0.$$

Then, if we take the partial sum of the first two points,

$$S_2 = \alpha_1 Q_1 + \alpha_2 Q_2 = \alpha_1(Q_1 - Q_2) + (\alpha_1 + \alpha_2)Q_2$$

and then the next partial sum of the first three points,

$$\begin{aligned} S_3 &= \alpha_1(Q_1 - Q_2) + (\alpha_1 + \alpha_2)Q_2 + \alpha_3Q_3 \\ &= \alpha_1(Q_1 - Q_2) + (\alpha_1 + \alpha_2)(Q_2 - Q_3) + (\alpha_1 + \alpha_2 + \alpha_3)Q_3 \end{aligned}$$

we can see that, by induction, the n th sum is,

$$\begin{aligned} S &= \alpha_1(Q_1 - Q_2) \\ &\quad + (\alpha_1 + \alpha_2)(Q_2 - Q_3) \\ &\quad + (\alpha_1 + \alpha_2 + \alpha_3)(Q_3 - Q_4) \\ &\quad \vdots \\ &\quad + (\alpha_1 + \cdots + \alpha_{n-1})(Q_{n-1} - Q_n) \\ &\quad + (\alpha_1 + \cdots + \alpha_n)Q_n. \end{aligned}$$

But we have $\alpha_1 + \cdots + \alpha_n = 0$ so the final term is 0. As a result S is a summation of terms of the form,


$$(\alpha_1 + \cdots + \alpha_{i-1} + \alpha_i)(Q_{i-1} - Q_i)$$

where $Q_{i-1} - Q_i$ is a vector. Therefore S is a linear combination of vectors in V and is therefore also a vector in V . \square

Corollary 30. *Any linear combination of points in an affine space where the coefficients sum to 1 results in a point.*

Proof. In the preceding proof if we had, instead, $\alpha_1 + \cdots + \alpha_n = 1$ then the final term would equal Q_n and the resulting summation would be a vector plus the point Q_n . Therefore the sum is a point. \square

2.5.6.2 Intuition of Affine Spaces



A translated linear subspace of a vector space like P_2 above that no longer passes through the origin is referred to as an **affine subspace**. It is not a vector space as $\vec{0} \notin P_2$ and $\vec{a}, \vec{b} \in P_2$ but $\vec{a} + \vec{b} \notin P_2$. However, if we instead consider displacements between points — e.g. $\vec{b} - \vec{a}$ — then we see the relationship between affine spaces and vector spaces: $\vec{b} - \vec{a} \in P_1$. The displacements between points in P_2 form a linear subspace. So we can define an affine space A based on the set of points in P_2 and the vectors in P_1 .

displacements between points in P_2 form a linear subspace. So we can define an affine space A based on the set of points in P_2 and the vectors in P_1 .

2.5.6.3 Affine Combinations

Definition. An **affine combination** of vectors is a combination such that the coefficients sum to 1.

The definition of an affine combination differs from that of a convex combination in that the coefficients of a convex combination are additionally required to be non-negative.

In an affine space there is no particular point designated as the origin but we can describe vector displacements between points as an ordered pair of points, for example, $(p, a) = \vec{pa}$. Affine combinations of displacements agree on the resulting point with linear combinations in the corresponding Euclidean space.

For example, imagine a point p in an affine space is at coordinates $(-1, 4)$ in the corresponding Euclidean space and similarly points a and b are at $(3, 4)$ and $(6, 1)$ respectively. Then, if we take an affine combination of the displacements to a and b the resulting point is independent of the chosen origin point.




Figure 2.5: The diagram shows the affine combination $2\vec{a} - \vec{b} = 2\vec{pa} - \vec{pb}$ where \vec{a} denotes the usual position vector in the Euclidean space.

2.5.6.4 Affine Transformations

Definition. Let $T : A_1 \mapsto A_2$ be a mapping between the affine spaces A_1 and A_2 . Then T is an **affine transformation** if:

- T maps vectors to vectors and points to points.
- T is a linear transformation over the vectors in the underlying vector space.
- $T(Q + \vec{v}) = T(Q) + T(\vec{v})$.

In an affine space a **translation** can be regarded as a change of frame in which we **change the origin point** while the vector basis may remain unchanged.

Proposition 111. *Affine transformations preserve parallelism.*

Proof. Let A be an affine space defined over a set of points P and a vector space V and let $Q_1, Q_2 \in P$ and $\vec{v} \in V$. Then, for $s, t \in \mathbb{F}$,

$$l_1 = Q_1 + t\vec{v} \quad \text{and} \quad l_2 = Q_2 + s\vec{v}$$

are parallel lines in A . Let T be an arbitrary affine transformation over A . Then,

$$l'_1 = T(l_1) = T(Q_1) + tT(\vec{v}) \quad \text{and} \quad l'_2 = T(l_2) = T(Q_2) + sT(\vec{v})$$

are also parallel. □

Proposition 112. *An affine transformation that preserves the dot product is left multiplication by an orthogonal matrix.*

Proof. Firstly, note that if a transformation m preserves the dot product and also fixes the standard basis vectors \vec{e}_i then,

$$m(\vec{e}_i) = \vec{e}_i \quad \text{and} \quad x_i = \vec{x} \cdot \vec{e}_i = m(\vec{x}) \cdot m(\vec{e}_i) = m(\vec{x}) \cdot \vec{e}_i = m(\vec{x})_i.$$

Therefore, such a transformation m is the identity transformation.

Now, assume a transformation m' preserves the dot product (but does not necessarily fix the standard basis vectors) in \mathbb{R}^n and let

$$B' = \{m'(\vec{e}_1), \dots, m'(\vec{e}_n)\}$$

be the transformed standard basis vectors. Then, if $[B']$ is the matrix whose columns are the elements of B' then, because m' preserves the dot product, $[B']$ is an orthogonal matrix as, by Proposition 108, is $[B']^{-1}$. So, composing this with m' ,

$$m'' = [B']^{-1}m'$$

is a transformation that both preserves the dot product and fixes the standard basis vectors. Therefore we have,

$$m'' = I_n = [B']^{-1}m' \iff [B'] = m'$$

so that m' is left multiplication by $[B']$, an orthogonal matrix. □

Properties of affine transformations

We can characterize affine transformations according to their form and behaviour in the Euclidean spaces \mathbb{R}^2 and \mathbb{R}^3 . Specifically, whether the transformation:

1. Fixes the origin: $T(\vec{0}) = \vec{0}$
2. Preserves the dot product: $T(\vec{v}) \cdot T(\vec{w}) = \vec{v} \cdot \vec{w}$, matrix is orthogonal
3. Preserves distances: $|T(\vec{v}) - T(\vec{w})| = |\vec{v} - \vec{w}|$
4. Preserves angles: matrix is scalar multiple of orthogonal matrix
5. Preserves orientation: transformation does not include a reflection, determinant of matrix is positive
6. Preserves parallelism: transformation exhibits affine property (linearity under affine combinations)

As we can see here, the property of preserving parallelism depends only on the affine property so clearly, all affine transformations exhibit this behaviour. As a consequence of preserving parallelism, affine transformations preserve the dimension of affine subspaces (points, lines, planes, etc.). They do not preserve distances between points however, but they do preserve ratios of distances between points lying on a straight line.

Classes of affine transformations

The most common subclassifications of affine transformations are:

- Linear: preserves the origin.
- Conformal: preserves angles.
- Isometry (also known as a congruent transformation): preserves distances in metric spaces and so also implicitly, angles. In a Euclidean space these transformations are known as a Euclidean isometries or rigid transformations.
- Rigid Motion: Euclidean isometry / rigid transformation that also preserves orientation.

2.5.6.5 Isometries

A Euclidean isometry or rigid motion, for example, carries a triangle to a congruent triangle. So, it preserves distances and angles but not necessarily orientation (a reflection flips the orientation).

The composition of two rigid motions is a rigid motion and the inverse of a rigid motion is also a rigid motion. Therefore, the rigid motions of \mathbb{R}^n form a group under composition of operations. This group is called the *group of motions* and denoted M_n .

Proposition 113. *A rigid motion that fixes the origin preserves the dot product.*

Proof. Let m be a rigid motion that fixes the origin. Then m is an isometry and so preserves distances and also m maps the origin to the origin. So we have,

$$|m(\vec{v}) - m(\vec{w})| = |\vec{v} - \vec{w}| \quad \text{and} \quad m(\vec{0}) = \vec{0}.$$

We can rewrite the isometry property as,

$$\begin{aligned} \sqrt{(m(\vec{v}) - m(\vec{w})) \cdot (m(\vec{v}) - m(\vec{w}))} &= \sqrt{(\vec{v} - \vec{w}) \cdot (\vec{v} - \vec{w})} \\ \iff (m(\vec{v}) - m(\vec{w})) \cdot (m(\vec{v}) - m(\vec{w})) &= (\vec{v} - \vec{w}) \cdot (\vec{v} - \vec{w}). \end{aligned}$$

Now, if we take the case where $\vec{w} = \vec{0} = m(\vec{w})$,

$$\begin{aligned} (m(\vec{v}) - \vec{0}) \cdot (m(\vec{v}) - \vec{0}) &= (\vec{v} - \vec{0}) \cdot (\vec{v} - \vec{0}) \\ \iff m(\vec{v}) \cdot m(\vec{v}) &= \vec{v} \cdot \vec{v}. \end{aligned}$$

and we can deduce that $m(\vec{x}) \cdot m(\vec{x}) = \vec{x} \cdot \vec{x}$ for any vector \vec{x} . Using this along with the properties of the dot product we obtain,

$$\begin{aligned} (m(\vec{v}) - m(\vec{w})) \cdot (m(\vec{v}) - m(\vec{w})) &= (\vec{v} - \vec{w}) \cdot (\vec{v} - \vec{w}) \\ \iff m(\vec{v}) \cdot m(\vec{v}) + m(\vec{w}) \cdot m(\vec{w}) - 2m(\vec{v}) \cdot m(\vec{w}) &= \vec{v} \cdot \vec{v} + \vec{w} \cdot \vec{w} - 2\vec{v} \cdot \vec{w} \\ \iff -2m(\vec{v}) \cdot m(\vec{w}) &= -2\vec{v} \cdot \vec{w} \\ \iff m(\vec{v}) \cdot m(\vec{w}) &= \vec{v} \cdot \vec{w}. \end{aligned}$$

Therefore, m preserves the dot product and, by Proposition 111, is left multiplication by an orthogonal matrix. \square

Corollary 31. *A rigid motion that fixes the origin is left multiplication by an orthogonal matrix and, therefore, also a linear operator.*

Proposition 114. *Left multiplication by any orthogonal matrix is a Euclidean isometry (a rigid motion) that fixes the origin.*

Proof. By, Proposition 106, left multiplication by an orthogonal matrix preserves the dot product. So, if m is an affine transformation such that $m(\vec{x}) = A\vec{x}$ where A is an orthogonal matrix then,

$$\begin{aligned} |m(\vec{v} - \vec{w})|^2 &= m(\vec{v} - \vec{w}) \cdot m(\vec{v} - \vec{w}) = (\vec{v} - \vec{w}) \cdot (\vec{v} - \vec{w}) = |\vec{v} - \vec{w}|^2 \\ \iff |m(\vec{v} - \vec{w})| &= |\vec{v} - \vec{w}| \end{aligned}$$

But also, left multiplication by a matrix is a linear operator so $m(\vec{v} - \vec{w}) = m(\vec{v}) - m(\vec{w})$ meaning that,

$$|m(\vec{v} - \vec{w})| = |m(\vec{v}) - m(\vec{w})| = |\vec{v} - \vec{w}|$$

which is the isometry property for m . □

Linear

Isometries that fix the origin are linear operators and rigid motions in Euclidean space.

Translation

If T is a translation then, in general, $T(\vec{0}) \neq \vec{0}$ so translations do not fix the origin and, as a result, are **not** linear transformations. However, translations preserve distances and angles (and orientation because there is no reflection) so they are rigid motions. For example:

- (43) Let $\vec{v} = (v_1, \dots, v_n)$ be any fixed vector in \mathbb{R}^n . Then translation by \vec{v} is the map,

$$t_v(\vec{x}) = \vec{x} + \vec{v} = \begin{bmatrix} x_1 + v_1 \\ \vdots \\ x_n + v_n \end{bmatrix}$$

which can be seen to be isometric (a rigid transformation) by,

$$\begin{aligned} t_v(\vec{x}) - t_v(\vec{y}) &= (\vec{x} + \vec{v}) - (\vec{y} + \vec{v}) = \vec{x} - \vec{y} \\ \implies |t_v(\vec{x}) - t_v(\vec{y})| &= |\vec{x} - \vec{y}|. \end{aligned}$$

Proposition 115. *Every rigid motion m is the composition of an orthogonal linear operator and a translation. In other words, for some orthogonal matrix A and fixed vector \vec{v} , it takes the form,*

$$m(\vec{x}) = A\vec{x} + \vec{v}.$$

Proof. Let $\vec{v} = m(\vec{0})$ and $t_v(\vec{x}) = \vec{x} + \vec{v}$ with inverse $t_{-v}(\vec{x}) = \vec{x} - \vec{v}$. Then composing this with m , the resulting transformation,

$$(t_{-v} \circ m)(\vec{x}) = m(\vec{x}) - \vec{v}$$

continues to be isometric — because it is the composition of isometric transformations — and it fixes the origin because $(t_{-v} \circ m)(\vec{0}) = m(\vec{0}) - \vec{v} = m(\vec{0}) - m(\vec{0}) = \vec{0}$. It is therefore, by Corollary 31, left multiplication by an orthogonal matrix. So we can represent it as,

$$(t_{-v} \circ m)(\vec{x}) = t_{-v}(m(\vec{x})) = A\vec{x}.$$

Since $t_{-v} = t_v^{-1}$ we can apply t_v to both sides of the equation,

$$m(\vec{x}) = t_v(A\vec{x}) = A\vec{x} + \vec{v}.$$

The obtained representation is uniquely determined by m as $\vec{v} = m(\vec{0})$ is clearly unique and then the translation t_{-v} is uniquely determined by \vec{v} and then $A = (t_{-v} \circ m)$ is unique for a given \vec{v} and m . \square

For a rigid motion $m(\vec{x}) = A\vec{x} + \vec{v}$, m is orientation-preserving if the matrix A is orientation-preserving and orientation-reversing if A is orientation-reversing.

Rotation

Rotations preserve distances, angles and orientation and so are rigid motions. Rotations also fix a vector which is known as the axis of rotation. If the axis of rotation contains the origin then they fix the origin and so are linear operators.

Theorem 26. *The rotations of \mathbb{R}^2 and \mathbb{R}^3 about the origin are the linear operators whose matrices with respect to the standard basis are orthogonal and have determinant 1.*

Proof. A rotation about the origin m involves rotating the standard basis vectors through an angle θ . It is in the definition of this rotation that the image of the standard basis vectors continue to subtend the same angle, $\pi/2$. Therefore, the rotation must preserve angles. It is also part of the definition that the image under rotation is not scaled so the rotation must preserve distances and must be a congruent transformation. Since the axis of rotation passes through the origin the origin is unchanged by this rotation and so these rotations are rigid motions that fix the origin and have the form,

$$m(\vec{x}) = A\vec{x}$$

where A is an orthogonal matrix. Additionally, rotations do not change the orientation of an shape and so their matrices have determinant 1. \square

*The rotation matrices — orthogonal matrices with determinant 1 — form a subgroup of the group O_n of orthogonal matrices called the **special orthogonal group** and denoted SO_n .*

Proposition 116. *Every member of the special orthogonal group $A \in SO_2$ is the matrix of a rotation.*

Proof. Let $A \in SO_2$. Then A is a 2×2 orthogonal matrix with determinant 1. Let \vec{v}_1 be the first column of A which, since A is orthogonal, is a unit vector. Now assume that R is the matrix of a rotation whose first column is \vec{v}_1 — which is possible because \vec{v}_1 is a unit vector so R can be orthogonal. Then the matrix

$$B = R^{-1}A$$

fixes \vec{e}_1 and also, as the composition of two orthogonal vectors, is orthogonal. Therefore the second column of B is a unit vector orthogonal to \vec{e}_1 which could be \vec{e}_2 or $-\vec{e}_2$.

However, R is an orthogonal matrix with determinant 1 and so is a member of SO_2 which means that $R^{-1}A = B$ is also in SO_2 .

This, in turn, means that B has determinant 1 which implies that the second column of B is not $-\vec{e}_2$ and is, therefore, \vec{e}_2 . So, we have obtained the result that $B = I = R^{-1}A$ which implies that $R = A$. \square

Rotating \mathbb{R}^2 about the origin

Rotating the 2-d plane about the origin means that the axis of rotation is just the origin (so the fixed vector is $\vec{0}$).

For example:

- (44) A rotation ρ_θ of the plane through an angle θ is a linear operator on \mathbb{R}^2 whose matrix with respect to the standard basis is

$$R = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}.$$

We can see that this is a rotation if we take $\vec{x} = (x_1, x_2)^T \in \mathbb{R}^2$ and write it in polar coordinates,

$$\vec{x} = (r, \alpha).$$

So, relating the polar and rectangular coordinates,

$$\vec{x} = (r \cos \alpha, r \sin \alpha)^T.$$

When we left-multiply by R ,

$$\begin{aligned} R\vec{x} &= \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \begin{bmatrix} r \cos \alpha \\ r \sin \alpha \end{bmatrix} \\ &= \begin{bmatrix} r \cos \alpha \cos \theta - r \sin \alpha \sin \theta \\ r \cos \alpha \sin \theta + r \sin \alpha \cos \theta \end{bmatrix} \\ &= \begin{bmatrix} r \cos (\alpha + \theta) \\ r \sin (\alpha + \theta) \end{bmatrix}. \end{aligned}$$

Note that R is orthogonal because

$$\begin{bmatrix} \cos \theta \\ \sin \theta \end{bmatrix} \cdot \begin{bmatrix} -\sin \theta \\ \cos \theta \end{bmatrix} = -\cos \theta \sin \theta + \sin \theta \cos \theta = 0$$

and $\det R = 1$,

$$\begin{vmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{vmatrix} = \cos^2 \theta + \sin^2 \theta = 1.$$

Rotating \mathbb{R}^3 about the origin

Definition. Define ρ as a rotation in \mathbb{R}^3 around the origin if:

- (i) ρ is a rigid motion (orientation-preserving Euclidean isometry) that fixes the origin;
- (ii) ρ also fixes a nonzero vector \vec{v} ;
- (iii) ρ operates as a rotation on the plane P orthogonal to \vec{v} .

Note that this definition could be described as selecting a 2-dimensional subspace of \mathbb{R}^3 and performing a 2-dimensional rotation on it as if it were \mathbb{R}^2 .

Condition (i) implies, by Corollary 31, that ρ is left multiplication by an orthogonal matrix. Condition (ii) states that ρ has an eigenvector \vec{v} with eigenvalue 1. Then, because ρ preserves angles, the plane P referenced in condition (iii) that is orthogonal to the eigenvector \vec{v} , must map to a plane that is orthogonal to the map of \vec{v} in the image of ρ . But \vec{v} is fixed by ρ and is unchanged in the image. Also \vec{v} uniquely identifies a plane orthogonal to it. Therefore the plane P is unchanged in the image also. In other words, P is an invariant subspace. So, condition (iii) says that the restriction of ρ to this invariant subspace is a rotation.

For example:

- (45) A rotation of \mathbb{R}^3 about the origin can be described by a pair (\vec{v}, θ) consisting of a unit vector \vec{v} , a vector of length 1, which lies in the axis of rotation, and a nonzero angle θ , the angle of rotation. The two pairs (\vec{v}, θ) and $(-\vec{v}, -\theta)$ represent the same rotation. We also consider the identity map to be a rotation, though its axis is indeterminate. The matrix representing a rotation through the angle θ about the vector \vec{e}_1 is obtained easily from the 2×2 rotation matrix. It is

$$A = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \theta & -\sin \theta \\ 0 & \sin \theta & \cos \theta \end{bmatrix}.$$

Multiplication by A fixes the first coordinate x_1 of a vector and operates by rotation on $(x_2, x_3)^T$. All rotations of \mathbb{R}^3 are linear operators, but their matrices can be fairly complicated.

Proposition 117. *Every element of SO_3 has eigenvalue 1.*

Proof. Let $A \in SO_3$. Then A is an orthogonal 3×3 matrix with determinant equal to 1. Reasoning from orthogonality of A we have,

$$\begin{aligned}
& A^T A = I \\
\iff & A^T A - A^T = I - A^T \\
\iff & A^T (A - I) = I - A^T \\
\iff & A^T (A - I) = (I - A)^T. \quad \text{by Proposition 39}
\end{aligned}$$

If we take the determinants of both sides of this equation we obtain,

$$\begin{aligned}
& \det(A^T) \cdot \det(A - I) = \det((I - A)^T) \quad \text{by Proposition 52} \\
\iff & \det(A) \cdot \det(A - I) = \det(I - A) \quad \text{by Proposition 54} \\
\iff & \det(A - I) = \det(I - A). \quad \det(A) = 1
\end{aligned}$$

But the dimension of A being 3 implies that

$$\det(-A) = (-1)^3 \det(A) = -\det(A)$$

so that,

$$\det(A - I) = \det(I - A) \iff \det(A - I) = 0.$$

Therefore A has the eigenvalue 1. □

Proposition 118. *The elements of SO_3 are precisely the rotations about the origin of \mathbb{R}^3 .*

Proof. Let $\rho : \mathbb{R}^3 \mapsto \mathbb{R}^3$ be defined as $\rho(\vec{x}) = A\vec{x}$ where $A \in SO_3$. Then,

- by Proposition 113 and orthogonality of $A \in SO_3$, left multiplication by A is a rigid motion that fixes the origin. So ρ is isometric and fixes the origin;

- Proposition 116 shows that every $A \in SO_3$ has eigenvalue 1 which implies that ρ fixes a nonzero vector;
- if we let \vec{v} be the nonzero vector fixed by ρ (i.e. its eigenvector with eigenvalue 1), then we can normalize it to find the unit vector parallel to \vec{v} , say \vec{u}_1 . Next we can find two unit vectors orthogonal to \vec{u}_1 — say \vec{u}_2 and \vec{u}_3 — and these must be a basis for the plane orthogonal to \vec{v} . Furthermore, if we select \vec{u}_2 and \vec{u}_3 to be orthogonal to each other than $B = \{\vec{u}_1, \vec{u}_2, \vec{u}_3\}$ is an orthonormal basis of \mathbb{R}^3 . Now if we define $P = [B]^{-1}$ then,

$$A' = P^{-1}AP$$

is similar to the matrix A and so has the same determinant, 1. Furthermore, because B is an orthonormal basis, the matrices $[B]$, $[B]^{-1} = P$ are orthogonal. Since both P and A are orthogonal, $P^{-1}AP = A'$ is orthogonal also. Since A' is orthogonal and has determinant equal to 1, it is a member of SO_3 .

If we examine the structure of A' , we see that the first column of A' is \vec{v}_1 — the unit vector in the direction of \vec{v} . Since \vec{v} is an eigenvector of ρ with eigenvalue 1, the first column of A' is \vec{e}_1 and since A' is orthogonal, the other columns are orthogonal to the first. So the block structure of A' looks like,

$$A' = \begin{bmatrix} 1 & 0 \\ 0 & R \end{bmatrix}$$

where R is a 2×2 matrix.

We know that the determinant of A' is 1 and this implies that the determinant of R is also 1. Furthermore, R must also be orthogonal and so $R \in SO_2$. So, by Proposition 115, R is a rotation. Therefore R represents a rotation of the plane orthogonal to \vec{v} and this implies that ρ rotates the plane orthogonal to \vec{v} as required.

□

Uniform Scaling

Uniform scaling is a scalar multiple of an orthogonal matrix and is therefore a linear operator which means that it fixes the origin. It also preserves angles but not distances between points.

2.5.6.6 Conformal Transformations

Non-uniform Scaling

Non-uniform scaling, however, its matrix is not a scalar multiple of an orthogonal matrix and, as such, it does not preserve angles. An important example is:

(46) **Mercator projection:**

This is a map projection that was designed so that rhumb lines (lines of constant bearing over the surface of the earth) are straight lines on the map. To achieve this the projection ensures that a square on the surface of the earth presents as a square on the map.

Then, modelling the earth as a sphere of radius R , if lines of latitude are horizontal grid lines across the map, then each actual line of latitude with circumference $2\pi R \cos \phi$ where ϕ is the angle of latitude will present on the map as the same length as the equator, which in reality is $2\pi R$. So they appear to be a line of length $\cos \phi$ times longer than they actually are, i.e. they are stretched by $\sec \phi$.

So, the map projection will stretch the width of a square on the surface of the earth by $\sec \phi$ and to maintain it as a square, it is necessary to stretch the height of the square by the same amount, $\sec \phi$. The actual square on the surface of the earth has height (approximately for a small square) $R\Delta\phi$ so, on the map we need a height $\Delta y \propto \Delta\phi \sec \phi$. Therefore, we have,

$$\frac{dy}{d\phi} = \sec \phi \implies y = \ln(\tan \phi + \sec \phi) + c$$

where c is a constant that we can set to 0. See <https://www.math.ubc.ca/~israel/m103/mercator/mercator.html> for a description of this

derivation. For more information on conformal map projections generally see: Map Projection, York University, Toronto and Conformal Cartographic Representations - University of Barcelona.

Reflection

Reflection usually refers to a Euclidean Isometry (rigid transformation over a Euclidean space) that fixes a hyperplane (so a line in \mathbb{R}^2 and a plane in \mathbb{R}^3) but does not preserve orientation so is not a rigid motion. However, it may also refer to a transformation that fixes an affine space of lower dimension than a hyperplane — for example, reflection in a point — in which case it does preserve orientation and is, therefore, a rigid motion (in fact, reflection in the origin in \mathbb{R}^2 is equal to rotation by π). Reflection may fix the origin or may not, depending on whether or not the origin is contained in the affine space fixed by the reflection.

2.5.6.7 Non-Rigid non-Conformal Transformations

Shear

Shear neither preserves distances nor angles. It does preserve parallelism though (as do all affine transformations) and it also fixes the origin, so it is a linear operator.

For more in-depth treatment of affine spaces and transformations see:

- *First two lectures of University of Texas - Multivariable Analysis.*
- https://www.maa.org/sites/default/files/pdf/pubs/books/meg/meg_ch12.pdf

Chapter 3

Analysis

3.1 Supremum and Infimum

3.1.1 Definitions

Definition. An upper bound on a set A is a value x such that,

$$\forall a \in A, a \leq x$$

and a lower bound is similarly defined as a value y such that,

$$\forall a \in A, a \geq y.$$

A set is said to be **upper-bounded** if there exists some upper-bound on the set and is said to be **lower-bounded** if there exists some lower bound on the set. If there exists both upper and lower bounds then the set is said to be **bounded**.

Definition. The **supremum** of a upper-bounded set A is a value σ_A such that σ_A is an upper bound on A and,

$$\sigma'_A < \sigma_A \iff \exists a \in A \text{ s.t. } a > \sigma'_A$$

which is to say that if $\sigma'_A < \sigma_A$ then σ'_A is not an upper bound on A and, if σ'_A is not an upper bound on A then it must be less than σ_A since σ_A is an upper bound on A .

An alternative, equivalent definition is,

$$\forall \epsilon > 0, \exists a \in A \text{ s.t. } a > \sigma_A - \epsilon.$$

Issue Note that there is an apparent paradox here: This second definition implies that

$$\begin{aligned} & \forall \epsilon > 0 . \exists a \in A \text{ s.t. } a + \epsilon > \sigma_A \\ \iff & \exists a \in A \text{ s.t. } a \geq \sigma_A \end{aligned}$$

which result, when combined with the upper-bound property, gives

$$\begin{aligned} & \exists a \in A \text{ s.t. } (a \geq \sigma_A) \wedge (a \leq \sigma_A) \\ \iff & \exists a \in A \text{ s.t. } a = \sigma_A \end{aligned}$$

which says that there is always an element in the bounded set that is equal to the supremum. This is not correct - the supremum may be in the set or external to it.

The initial implication is not true, however. We cannot infer that $\exists a \in A \text{ s.t. } a \geq \sigma_A$. This can be seen with another rearrangement,

$$\begin{aligned} & \forall \epsilon > 0 . \exists a \in A \text{ s.t. } a + \epsilon > \sigma_A \\ \iff & \forall \epsilon > 0 . \exists a \in A \text{ s.t. } \epsilon > \sigma_A - a \end{aligned}$$

which shows us that for any positive epsilon there needs to be an a close enough to the value of σ_A that the difference in their values is less than epsilon. Since a can approach arbitrarily close to σ_A this is achievable for any positive epsilon. This property seems to be equivalent to the fact that σ_A is a *limit point* of A but that will be covered properly in Topology.

Definition. The *infimum* of a lower-bounded set A is defined similarly to the supremum: as a value τ_A such that τ_A is a lower bound on A and,

$$\tau'_A > \tau_A \iff \exists a \in A \text{ s.t. } a < \tau'_A$$

or alternatively,

$$\forall \epsilon > 0, \exists a \in A \text{ s.t. } a < \tau_A + \epsilon.$$

Notation. The supremum of A is denoted $\sup A$ and the infimum is denoted $\inf A$.

3.1.2 Deductions using the supremum and infimum

Proposition 119. If a bounded set $A \subset \mathbb{R}$ has the property that,

$$\forall x, y \in A . |x - y| < 1$$

then it follows that,

$$(\sup A - \inf A) \leq 1.$$

Proof. Let $\sigma_A = \sup A$ and $\tau_A = \inf A$ and w.l.o.g. assume that $x > y$. By the definitions of the supremum and infimum we have,

$$\begin{aligned} & \forall \epsilon > 0 . \exists x, y \in A . (x > \sigma_A - \epsilon) \wedge (y < \tau_A + \epsilon) \\ \iff & \forall \epsilon > 0 . \exists x, y \in A . (x > \sigma_A - \epsilon) \wedge (-y > -\tau_A - \epsilon) \\ \iff & \forall \epsilon > 0 . \exists x, y \in A . (x - y) > (\sigma_A - \tau_A) - 2\epsilon \end{aligned}$$

Now suppose, for contradiction, that $(\sigma_A - \tau_A) > 1$ then we can say that,

$$\exists r > 0 . (\sigma_A - \tau_A) = 1 + r.$$

If we then constrict ϵ such that,

$$\epsilon < \frac{r}{2} \iff 2\epsilon < r \iff r - 2\epsilon > 0$$

then the previous result tells us that, for $0 < \epsilon < \frac{r}{2}$,

$$\begin{aligned} & \exists x, y \in A . (x - y) > (\sigma_A - \tau_A) - 2\epsilon \\ \iff & \exists x, y \in A . (x - y) > 1 + r - 2\epsilon > 1 \end{aligned}$$

which contradicts the set property that $\forall x, y \in A, |x - y| < 1$. So this shows that $(\sigma_A - \tau_A) \leq 1$. \square

Proposition 120. *Let $A \subset \mathbb{R}$ be a bounded set and let B be the set defined by*

$$B = \{b \mid b = f(a), a \in A\}$$

where the function f is some strictly monotonic function. Then it follows that,

$$\sup B = f(\sup A).$$

Proof. A is bounded and so $\sigma_A = \sup A$ exists. So, using the supremum properties we have,

$$\begin{aligned} & \forall a \in A . a \leq \sigma_A \\ \iff & \forall a \in A . f(a) \leq f(\sigma_A) \quad \text{by monotonicity of } f \\ \iff & \forall b \in B . b \leq f(\sigma_A) \end{aligned}$$

which is to say that $\sigma_B = f(\sigma_A)$ is an upper bound on B .

Furthermore, using the other supremum property, we have that,

$$\begin{aligned} & \sigma'_A < \sigma_A \implies \exists a \in A \text{ s.t. } a > \sigma'_A \\ \iff & f(\sigma'_A) < f(\sigma_A) \implies \exists a \in A \text{ s.t. } f(a) > f(\sigma'_A) \quad \text{by strict monotonicity of } f \\ \iff & \sigma'_B < \sigma_B \implies \exists b \in B \text{ s.t. } b > \sigma'_B. \end{aligned}$$

Therefore σ_B satisfies both requirements of the supremum and we have shown that,

$$\sup B = f(\sup A).$$

\square

3.2 Limits

3.2.1 Limits of sequences

3.2.1.1 Problems with the informal description of a limit

If we say that a sequence tends to some value L when the terms of the sequence *gets closer and closer to L* we have the following problems:

- that the sequence gets closer and closer to many numbers so that this does not specify a single specific limit.
- that the sequence can have a limit but it's not the case that every term is closer than the previous term to the limit. For example,

$$a_{2k} = 1/k, \quad a_{2k-1} = \frac{1}{k+1}$$

tends to 0 but $a_{2k} > a_{2k-1}$.

Definition. A sequence a_n is said to **tend** to L or have the **limit** L iff,

$$\forall \epsilon > 0 \in \mathbb{R}, \exists N \in \mathbb{N} \text{ s.t. } \forall n > N, |a_n - L| < \epsilon.$$

Definition. The interval $(L - \epsilon, L + \epsilon)$ is called the **ϵ -neighbourhood of L** .

Definition. A sequence a_n is said to **tend to infinity** iff,

$$\forall M > 0 \in \mathbb{R}, \exists N \in \mathbb{N} \text{ s.t. } \forall n > N, a_n > M$$

and **tend to minus-infinity** iff,

$$\forall M < 0 \in \mathbb{R}, \exists N \in \mathbb{N} \text{ s.t. } \forall n > N, a_n < M.$$

Definition. A sequence that has a limit is called **convergent** and otherwise is called **divergent**. Note that **divergent** sequences include both sequences that remain bounded but oscillate without converging and those that tend to infinity (or minus-infinity).

3.2.1.2 Examples of Convergence and Divergence

(47) Non-convergent Oscillation

The sequence $a_n = (-1)^n$ is divergent despite always remaining bounded within the interval $[-1, 1]$ as it neither converges to 1 or to -1.

(48) Limit of an Infinite Recurrence

Take the sequence given by,

$$a_1 = 1, \quad a_{n+1} = \frac{a_n}{2} + \frac{3}{2a_n} \quad (n \geq 1).$$

Assume there is an equilibrium value, a^* , then

$$\begin{aligned} a^* &= \frac{a^*}{2} + \frac{3}{2a^*} \\ \iff 2(a^*)^2 &= (a^*)^2 + 3 \\ \iff (a^*)^2 &= 3 \\ \iff a^* &= \sqrt{3} \qquad \forall n, a_n \geq 0 \end{aligned}$$

So $\sqrt{3}$ is the steady-state value that this recurrence converges to as $n \rightarrow \infty$. If the recurrence didn't converge then the assumption of an equilibrium value would result in a contradiction. Note, however, that the fact that there is an equilibrium value does *not*, by itself, prove that this sequence converges (although this sequence does).

Proposition 121. *A sequence has at most one limit. In other words, a sequence can only converge, if at all, to a single unique value.*

Proof. Let L and L' both be limits of the sequence a_n , and the constant $\alpha = L - L'$. Then,

$$\forall \epsilon > 0 \in \mathbb{R}, \exists N \in \mathbb{N} \text{ s.t. } \forall n > N, |a_n - L| < \epsilon$$

and

$$\forall \epsilon' > 0 \in \mathbb{R}, \exists N' \in \mathbb{N} \text{ s.t. } \forall n' > N', |a_{n'} - L'| < \epsilon'.$$

But also we have,

$$|a_n - L'| = |(a_n - L) + (L - L')| = |(L - L') + (a_n - L)|$$

and using the triangle inequality,

$$\begin{aligned} |L - L'| &= |(L - L') + (a_n - L) - (a_n - L)| \leq |(L - L') + (a_n - L)| + |-(a_n - L)| \\ \iff |L - L'| &\leq |(L - L') + (a_n - L)| + |a_n - L| \\ \iff |L - L'| - |a_n - L| &\leq |(L - L') + (a_n - L)| \\ \iff |\alpha| - |a_n - L| &\leq |\alpha + (a_n - L)| \end{aligned}$$

Since $\alpha = L - L'$ is constant we can consider the situation when $\epsilon = \frac{|\alpha|}{2}$ then we have that,

$$\begin{aligned} \exists N \in \mathbb{N} \text{ s.t. } \forall n > N, |a_n - L| &< \epsilon = \frac{|\alpha|}{2} \\ \iff -|a_n - L| &> -\frac{|\alpha|}{2} \\ \iff |\alpha| - |a_n - L| &> |\alpha| - \frac{|\alpha|}{2} \\ \iff |\alpha| - |a_n - L| &> \frac{|\alpha|}{2}. \end{aligned}$$

Combining this with the previous result gives, for $\forall n > N$,

$$\frac{|\alpha|}{2} < |\alpha| - |(a_n - L)| \leq |\alpha + (a_n - L)| = |a_n - L'|$$

which, by rearranging a little, is,

$$\forall n > N, |a_n - L'| > \frac{|\alpha|}{2}.$$

But this means that if we also choose $\epsilon' = \frac{|\alpha|}{2}$ then there is no N' such that $\forall n' > N', |a_{n'} - L'| < \epsilon'$ which contradicts the hypothesis that L' is also a limit of a_n . \square

Proof. Another quicker way of proving the proposition is by letting $\epsilon = \epsilon' = \frac{|\alpha|}{2}$ so that,

$$2\epsilon = |\alpha| = |L - L'| = |L - a_n + a_n - L'| \leq |L - a_n| + |a_n - L'| = |a_n - L| + |a_n - L'|$$

which gives us,

$$2\epsilon \leq |a_n - L| + |a_n - L'|.$$

But by the limit definition,

$$|a_n - L| + |a_n - L'| < \epsilon + \epsilon'$$

and since we have set $\epsilon = \epsilon'$ then,

$$|a_n - L| + |a_n - L'| < 2\epsilon$$

which contradicts $2\epsilon \leq |a_n - L| + |a_n - L'|$. \square

Definition. If a_n is a sequence and $S = \{a_n \mid n \in \mathbb{N}\}$ then a_n is said to be **bounded below** if S has a lower bound and **bounded above** if S has an upper bound, and **bounded** if it is bounded above and below.

Lemma 1. Any finite set of elements from an ordered field has a minimum and a maximum.

Proof. This can be proven quite easily using induction. Taking the base case of a set of cardinality one, clearly there is a maximum and a minimum both of which are the sole element of the set. Then, the induction step is to say, given a set S that has a maximum, s_{max} , and a minimum, s_{min} , if we add a new element e , then if e is greater than s_{max} it is the maximum of the new set and if it is less than s_{min} it is the minimum of the new set. Otherwise, the previous maximum and minimum also pertain to the new set. Therefore, adding a new element to a set that has a maximum and a minimum creates a new set with a maximum and a minimum. \square

Proposition 122. *Any convergent sequence is bounded.*

Proof. Firstly, we need to prove that any finite sequence is bounded. We can do this simply by observing that any finite set of elements from an ordered field,

$$S = \{a_1, a_2, \dots, a_n\}$$

has a minimum and a maximum.

Now, let a_n be an arbitrary convergent sequence so that,

$$\forall \epsilon > 0 . \exists N \in \mathbb{N} . \forall n > N . |a_n - L| < \epsilon$$

for some $L \in \mathbb{R}$.

Then, let S_{max} and S_{min} be the maximum and minimum respectively of the first N terms of a_n , $S = \{a_1, a_2, \dots, a_N\}$, and,

$$\exists \epsilon > 0 . \forall n > N . |a_n - L| < \epsilon$$

so that, for $n > N$, the sequence a_n is bounded in the ϵ -neighbourhood of L . So, if we define $m = \min\{S_{min}, L - \epsilon\}$ and $M = \max\{S_{max}, L + \epsilon\}$, then the whole sequence a_n for all $n \in \mathbb{N}$ is bounded below by m and bounded above by M .

Therefore a_n is bounded. \square

Definition. An *increasing* sequence is a sequence a_n such that,

$$\forall n \in \mathbb{N} . a_{n+1} \geq a_n$$

and *decreasing* if,

$$\forall n \in \mathbb{N} . a_{n+1} \leq a_n$$

and *monotonic* if either increasing or decreasing.

Proposition 123. Any increasing sequence that is bounded above has a limit.

Proof. Let a_n be an increasing sequence that is bounded above. Then,

$$\forall n \in \mathbb{N} . a_{n+1} \geq a_n$$

and let $S = \{ a_n \mid n \in \mathbb{N} \}$. Since a_n is bounded above it has a supremum. Let $\sigma = \sup S$ so that,

$$\forall a_n \in S . a_n \leq \sigma \quad \text{and} \quad \forall \epsilon > 0 . \exists a_n \in S . a_n > \sigma - \epsilon.$$

Therefore, for some arbitrary fixed $\epsilon > 0$,

$$\exists a_n \in S . a_n > \sigma - \epsilon$$

and setting $N = n$ so that $a_N > \sigma - \epsilon$, we have,

$$\forall n > N \in \mathbb{N} . a_n \geq a_N > \sigma - \epsilon$$

and so, recalling that σ is an upper bound on S ,

$$\begin{aligned} & \exists N . \forall n > N \in \mathbb{N} . (a_n > \sigma - \epsilon) \wedge (a_n \leq \sigma) \\ \iff & \exists N . \forall n > N \in \mathbb{N} . a_n \leq \sigma < a_n + \epsilon \\ \iff & \exists N . \forall n > N \in \mathbb{N} . 0 \leq \sigma - a_n < \epsilon \\ \implies & \exists N . \forall n > N \in \mathbb{N} . |\sigma - a_n| < \epsilon. \end{aligned}$$

But ϵ was an arbitrary positive value so,

$$\forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . |\sigma - a_n| < \epsilon$$

and σ is, therefore, the limit of a_n . □

Corollary 32. *Any increasing sequence that is bounded above converges to the supremum of its elements (terms, values, etc.).*

Corollary 33. *A decreasing sequence that is bounded below converges to the infimum of its elements.*

3.2.2 Algebra of limits of sequences

Proposition 124. *Let a_n and b_n be convergent sequences with limits a and b , respectively. Let C be a real number and let k be a positive integer. Then as $n \rightarrow \infty$,*

$$a) \quad Ca_n \rightarrow Ca$$

$$b) \quad |a_n| \rightarrow |a|$$

$$c) \quad a_n + b_n \rightarrow a + b$$

$$d) \quad a_n b_n \rightarrow ab$$

$$e) \quad a_n^k \rightarrow a^k$$

$$f) \quad \text{if, for all } n, b_n \neq 0 \text{ and } b \neq 0, \text{ then } \frac{1}{b_n} \rightarrow \frac{1}{b}.$$

Proof. We prove each property individually in the given order.

3.2.2.1 Proof of (a) $Ca_n \rightarrow Ca$

If $C = 0$ then $Ca_n = 0 = Ca$ for all n and the proposition holds trivially. If $C \neq 0$ then, since $a_n \rightarrow a$,

$$\forall \epsilon' > 0. \exists N. \forall n > N \in \mathbb{N}. |a_n - a| < \epsilon'.$$

Now let $\epsilon = |C| \epsilon'$. Then,

$$\begin{aligned} & \forall \epsilon > 0. \exists N. \forall n > N \in \mathbb{N}. |C| |a_n - a| < |C| \epsilon' = \epsilon \\ \iff & \forall \epsilon > 0. \exists N. \forall n > N \in \mathbb{N}. |Ca_n - Ca| < \epsilon. \end{aligned} \quad |x| |y| = |xy|$$

3.2.2.2 Proof of (b) $|a_n| \rightarrow |a|$

$$\begin{aligned} & |a_n| = |a_n - a + a| \leq |a_n - a| + |a| \quad \text{the "triangle inequality"} \\ \iff & |a_n| - |a| \leq |a_n - a|. \end{aligned} \quad (1)$$

$$\begin{aligned} & |a| = |a - a_n + a_n| \leq |a - a_n| + |a_n| \quad \text{the "triangle inequality"} \\ \iff & |a| - |a_n| \leq |a - a_n| = |a_n - a| \\ \iff & |a_n| - |a| \geq -|a_n - a|. \end{aligned} \quad (2)$$

Putting (1) and (2) together we have,

$$\begin{aligned} & -|a_n - a| \leq |a_n| - |a| \leq |a_n - a| \quad \text{the "triangle inequality"} \\ \iff & ||a_n| - |a|| \leq |a_n - a|. \end{aligned}$$

The fact that a_n converges to a implies that $|a_n - a|$ converges to zero. Since it is an upper bound on the value of $||a_n| - |a||$, the value $||a_n| - |a||$ must also converge to zero. Specifically any value of n, ϵ such that $|a_n - a| < \epsilon$ will also satisfy $||a_n| - |a|| \leq |a_n - a| < \epsilon$.

3.2.2.3 Proof of (c) $a_n + b_n \rightarrow a + b$

Using, again, the "triangle inequality",

$$|(a_n + b_n) - (a + b)| = |(a_n - a) + (b_n - b)| \leq |a_n - a| + |b_n - b|.$$

So, $|a_n - a| + |b_n - b|$ is an upper bound on the value of $|(a_n + b_n) - (a + b)|$. If we take any arbitrary $\epsilon > 0$ then,

$$\exists N_1 . \forall n > N_1 \in \mathbb{N} . |a_n - a| < \frac{\epsilon}{2}$$

and

$$\exists N_2 . \forall n > N_2 \in \mathbb{N} . |b_n - b| < \frac{\epsilon}{2}.$$

Then, if we take $N = \max\{N_1, N_2\}$, we have,

$$\forall n > N \in \mathbb{N} . |(a_n + b_n) - (a + b)| \leq |a_n - a| + |b_n - b| < \epsilon.$$

3.2.2.4 Proof of (d) $a_n b_n \rightarrow ab$

$$\begin{aligned} |a_n b_n - ab| &= |a_n b_n - ab_n + ab_n - ab| \leq |b_n(a_n - a)| + |a(b_n - b)| \quad \text{the "triangle inequality"} \\ \iff |a_n b_n - ab| &\leq |b_n| |a_n - a| + |a| |b_n - b| \end{aligned}$$

Since b_n converges, by Proposition 121, it is bounded. Therefore, $|b_n|$ has some upper bound which we shall call B . Then,

$$\forall \epsilon > 0 . \exists N_1 . \forall n > N_1 \in \mathbb{N} . |a_n - a| < \frac{\epsilon}{2B}$$

and

$$\forall \epsilon > 0 . \exists N_2 . \forall n > N_2 \in \mathbb{N} . |b_n - b| < \frac{\epsilon}{2|a|}.$$

Now, let $N = \max N_1, N_2$. Then,

$$\forall n > N \in \mathbb{N} . B |a_n - a| + |a| |b_n - b| < \epsilon.$$

3.2.2.5 Proof of (e) $a_n^k \rightarrow a^k$

Using (d) - and because k is a positive integer - we can do induction on the power k .

Base cases 0 and 1 are clearly true as $k = 0$ results in a_n being the constant 1 for all n and so trivially converges to $a = 1$; and $k = 1$ results in the same sequence as a_n .

So, we perform the induction step for $k \geq 2$. Then, by the induction hypothesis, $a_n^{k-1} \rightarrow a^{k-1}$. But $a_n^k = a_n^{k-1} a_n$ and, by (d) and the induction hypothesis, we have that $a_n^{k-1} a_n \rightarrow a^{k-1} a = a^k$. Therefore, $a_n^k \rightarrow a^k$.

3.2.2.6 Proof of (f) $\forall n . b_n, b \neq 0 \implies \frac{1}{b_n} \rightarrow \frac{1}{b}$

Again invoking Proposition 121 and letting the upper bound on the sequence b_n be B ,

$$\left| \frac{1}{b_n} - \frac{1}{b} \right| = \left| \frac{b - b_n}{b_n b} \right| = \left| \frac{b_n - b}{b_n b} \right| = \frac{|b_n - b|}{|b_n| |b|} \leq \frac{1}{B |b|} |b_n - b|.$$

Now, since $\frac{1}{B|b|}$ is a constant we can define the constant $C = \frac{1}{B|b|}$ and then we see that in (a) we have already proven that $C |b_n - b|$ converges to 0. In (a) we used that to prove that $C b_n \rightarrow C b$ but here it proves that $\frac{1}{b_n} \rightarrow \frac{1}{b}$. \square

3.2.3 Some theorems on limits of sequences

Theorem 27. *If $|a| < 1$ then $\lim_{n \rightarrow \infty} a^n = 0$.*

Proof. First of all, note that if $|a| = 0$ then $a = 0 = a^n$ for all n and so the limit holds trivially. For this reason, from here on, we will consider only the case where $a \neq 0$.

There are 3 parts to this proof:

1. $|a| < 1 \implies \lim_{n \rightarrow \infty} |a|^n = 0$,
 2. $|a|^n = |a^n|$,
 3. $\lim_{n \rightarrow \infty} |a^n| = 0 \implies \lim_{n \rightarrow \infty} a^n = 0$.
1. $|a| < 1 \implies \lim_{n \rightarrow \infty} |a|^n = 0$

It would be natural to prove this by showing that,

$$\forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . ||a|^n - 0| = |a|^n < \epsilon .$$

We can show this by "reverse engineering" the value of N from the requirement that $|a|^n$ be less than ϵ ,

$$|a|^n < \epsilon \iff n \ln |a| < \ln \epsilon \iff n > \frac{\ln \epsilon}{\ln |a|}$$

with the last step changing the direction of the inequality because we divide by $\ln |a|$ which - remembering that $|a| < 1$ - is a negative value. So, in this way, we have shown that $N = \frac{\ln \epsilon}{\ln |a|}$ is a general formula that relates a value of N with the required property with any arbitrary ϵ . However, this proof is not valid because it uses the concept of the logarithm which requires a lot of analysis that has not been proven at this stage. Since we are trying to build the fundamental basis of analysis, at this point we can only use concepts that are pre-requisites (axiomatic) in analysis or have been proven at this stage.

An alternative way to show this, using the properties of limits of sequences just proven, is as follows: Let x_n be the sequence $x_n = |a|^n$. Then, because $0 < |a| < 1$,

$$x_{n+1} = x_n \cdot |a| = |a|^{n+1} < |a|^n = x_n$$

so that x_n is a decreasing sequence. Additionally, $\forall n \in \mathbb{R} . |a|^n \geq 0$ so 0 is a lower bound on the sequence. Therefore, the sequence converges to a limit (note we haven't yet established that 0 is the limit - only that it is a candidate). Furthermore, $x_{n+1} = |a|^{n+1} = |a|^n |a|$ and, if $x_n \rightarrow L$ then $x_{n+1} \rightarrow L$ also. But, putting these two facts together, along with property (d) of limits of sequences, means that,

$$L = \lim_{n \rightarrow \infty} |a|^{n+1} = \lim_{n \rightarrow \infty} |a|^n |a| = \left(\lim_{n \rightarrow \infty} |a|^n \right) \left(\lim_{n \rightarrow \infty} |a| \right) = |a| \left(\lim_{n \rightarrow \infty} |a|^n \right) = |a| L.$$

So,

$$L = |a| L \iff L(1 - |a|) = 0$$

and, since we know that $|a| \neq 1$, therefore L must be 0.

2. $|a|^n = |a^n|$

For $a \in \mathbb{R}, n \in \mathbb{N}$ it's easy to see that $|a| |a| \dots |a| = |aa \dots a|$. In actual fact, this appears to hold even for $n \in \mathbb{Q}$, e.g.

$$\left| (-1)^{\frac{1}{2}} \right| = |i| = 1 = \left| 1^{\frac{1}{2}} \right| = |1| = 1$$

but this should be checked when studying complex numbers more thoroughly. Also, the base a , can it also be complex?

3. $\lim_{n \rightarrow \infty} |a^n| = 0 \implies \lim_{n \rightarrow \infty} a^n = 0$

This can be proved directly from the definition of the limit.

$$\begin{aligned} \lim_{n \rightarrow \infty} |a^n| = 0 &\iff \forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . ||a^n| - 0| < \epsilon \\ &\iff \forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . |a^n| < \epsilon \\ &\iff \forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . |a^n - 0| < \epsilon \\ &\iff \lim_{n \rightarrow \infty} a^n = 0. \end{aligned}$$

Bear in mind that, in general, $\lim_{n \rightarrow \infty} |x_n| = L \not\Rightarrow \lim_{n \rightarrow \infty} x_n = L$. For example, if x_n converges to $-L$ then $|x_n|$ will converge to L .

code example

Furthermore, the example 48 showed how the fact of a_n and a_{n+1} converging to the same limit produces - when the sequence is expressed as a recurrence - an equilibrium value. \square

3.2.3.1 The Sandwich Theorem

Proposition 125. *Let a_n, b_n, c_n be sequences such that,*

$$\text{for all } n, a_n \leq b_n \leq c_n \quad \text{and} \quad \lim_{n \rightarrow \infty} a_n = L = \lim_{n \rightarrow \infty} c_n.$$

Then $\lim_{n \rightarrow \infty} b_n = L$.

Proof. $\lim_{n \rightarrow \infty} a_n = L$ means that,

$$\begin{aligned} & \forall \epsilon > 0 . \exists N_1 . \forall n > N_1 \in \mathbb{N} . |a_n - L| < \epsilon \\ \iff & \forall \epsilon > 0 . \exists N_1 . \forall n > N_1 \in \mathbb{N} . -\epsilon < a_n - L < \epsilon \\ \iff & \forall \epsilon > 0 . \exists N_1 . \forall n > N_1 \in \mathbb{N} . L - \epsilon < a_n < L + \epsilon. \end{aligned}$$

By the same reasoning we also have,

$$\forall \epsilon > 0 . \exists N_2 . \forall n > N_2 \in \mathbb{N} . L - \epsilon < c_n < L + \epsilon.$$

So, if we let $N = \max\{N_1, N_2\}$ then we have,

$$\forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . L - \epsilon < a_n, c_n < L + \epsilon$$

and since we also know that $a_n \leq b_n \leq c_n$ it follows that,

$$\forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . L - \epsilon < a_n \leq b_n \leq c_n < L + \epsilon.$$

This shows that,

$$\forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . |b_n - L| < \epsilon.$$

□

3.2.3.2 An example application of the Sandwich Theorem

(49) **Prove that $|x| < 1 \implies \lim_{n \rightarrow \infty} x^n = 0$**

We have already proven this using the properties of limits in Theorem 26 but here we are going to prove it using the Sandwich Theorem.

Proof. Firstly, as we showed in 2, $|x^n - 0| = |x^n| = |x|^n$. So to show that x^n tends to zero we can show that $|x|^n$ tends to zero. So, w.l.o.g. we take $x > 0$ (since $x = 0$ makes the proposition trivially true). Then, notice that $0 < x < 1 \implies x = \frac{1}{1+h}$ for some $h > 0$. Then, we can show inductively that $(1+h)^n \geq 1 + hn$ as follows.

Base cases 0, 1: $(1 + h)^0 = 1 = 1 + h(0)$, $(1 + h)^1 = 1 + h = 1 + h(1)$.

Induction step $k > 1$

$$\begin{aligned} & (1 + h)^k \geq 1 + hk \\ \iff & (1 + h)(1 + h)^k \geq (1 + h)(1 + hk) \\ \iff & (1 + h)^{k+1} \geq 1 + hk + h + h^2k = 1 + h(k + 1) + h^2k > 1 + h(k + 1) \end{aligned}$$

This result implies that,

$$x^n = \frac{1}{(1 + h)^n} \leq \frac{1}{1 + hn}$$

so that,

$$0 < x^n \leq \frac{1}{1 + hn}.$$

Since h is some fixed value, clearly,

$$\lim_{n \rightarrow \infty} \frac{1}{1 + hn} = 0$$

and, obviously, the limit of the constant 0 is always 0 so, by the Sandwich Theorem,

$$\lim_{n \rightarrow \infty} x^n = 0.$$

□

3.2.4 Subsequences

Definition. Let $(a_n)_{n \in \mathbb{N}}$ be a sequence and consider some strictly increasing natural numbers (k_1, k_2, k_3, \dots) that is, $(k_1 < k_2 < k_3 < k_4 < \dots)$. Then the sequence $(a_{k_n})_{n \in \mathbb{N}}$ is called a **subsequence** of the sequence $(a_n)_{n \in \mathbb{N}}$. Note that a **subsequence** is always infinite (I think).

Theorem 28. If a_n is a sequence that tends to a limit, then any subsequence of it tends to the same limit.

Proof. Firstly, notice that if the n th index of some subsequence is k_n then $k_n \geq n$ (because the subsequence can only skip terms of the original - it can't add in terms). So then, if we have a sequence a_n that tends to a limit a and an arbitrary subsequence a_{k_n} then,

$$\forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . |a_n - a| < \epsilon \implies |a_{k_n} - a| < \epsilon$$

because $k_n \geq n > N$. □

Theorem 29. Every sequence has a monotonic subsequence.

Proof. Either there is an infinite number of terms that are greater than all the following terms or there is not. If there is not, then after the last such term all terms have a term that follows them that is greater than or equal to them.

In the first case we have a strict monotonic decreasing sequence and in the second we have a non-strict monotonic increasing sequence. □

If we put this theorem together with what we learned about bounded sequences in Proposition 122 and its corollaries - that monotonic bounded sequences are convergent - then we get one of the most famous results in analysis, The Bolzano-Weierstrass Theorem.

Theorem 30. The Bolzano-Weierstrass Theorem
Every bounded sequence has a convergent subsequence.

3.2.5 Examples of limits of sequences

- (50) Let $(a_n)_{n \in \mathbb{N}}$ be a sequence, and let $(b_n)_{n \in \mathbb{N}}$ be the sequence defined by $b_n = |a_n|$ for $n \in \mathbb{N}$. Which of the following two statements implies the other?

Answer: a_n converges $\implies b_n$ converges also but b_n converges $\not\Rightarrow a_n$ converges.

The first implication is because,

$$\begin{aligned} |a_n| &= |a_n - a + a| \leq |a_n - a| + |a| \\ \iff |a_n| - |a| &\leq |a_n - a| \end{aligned}$$

$$\begin{aligned} |a| &= |a - a_n + a_n| \leq |a_n - a| + |a_n| \\ \iff |a| - |a_n| &\leq |a_n - a| \\ \iff |a_n| - |a| &\geq -|a_n - a| \end{aligned}$$

which both together imply that $||a_n| - |a|| \leq |a_n - a|$.

The latter non-implication is easy to see if one thinks of a sequence that consists of two subsequences that converge to 2 and -2. Then, their absolute value would converge to 2 but their values do not converge. Remember Theorem 27, for a sequence to converge to a limit, every subsequence of it must converge to the same limit.

- (51) What is the behaviour as $n \rightarrow \infty$ of the following:

(i) $\frac{2n^3+1}{n+1} \left(\frac{3}{4}\right)^n$

$$0 < \frac{2n^3+1}{n+1} \left(\frac{3}{4}\right)^n < \frac{3n^3}{n} \left(\frac{3}{4}\right)^n = 3n^2 \left(\frac{3}{4}\right)^n \rightarrow 0$$

(ii) $\frac{2^{2n}+n}{n^3 3^n + 1}$

$$\frac{2^{2n}+n}{n^3 3^n + 1} = \frac{4^n + n}{n^3 3^n + 1} > \frac{4^n}{2n^3 3^n} = \frac{(4/3)^n}{2n^3} \rightarrow \infty$$

(52) **Let (a_n) be a sequence of non-negative numbers. Prove that if $a_n \rightarrow L$ as $n \rightarrow \infty$ then $\sqrt{a_n} \rightarrow \sqrt{L}$ as $n \rightarrow \infty$.**

Proof. We are told that $a_n \rightarrow L$ as $n \rightarrow \infty$ so we have,

$$\forall \epsilon' > 0 . \exists N . \forall n > N \in \mathbb{N} . |a_n - L| < \epsilon'.$$

We are also told that (a_n) is non-negative so $L \geq 0$.

First, consider the possibility that $L = 0$. Then $|a_n| < \epsilon'$ for $n > N$.

But,

$$|a_n| = (\sqrt{a_n})^2 < \epsilon' \iff \sqrt{a_n} < (\epsilon')^2$$

so that taking $\epsilon = (\epsilon')^2$ we obtain,

$$\forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . |\sqrt{a_n}| < \epsilon.$$

The remaining possibility is that $L > 0$. Notice that the expression we're looking to bound can be rewritten as follows:

$$|\sqrt{a_n} - \sqrt{L}| = \left| (\sqrt{a_n} - \sqrt{L}) \frac{\sqrt{a_n} + \sqrt{L}}{\sqrt{a_n} + \sqrt{L}} \right| = \left| \frac{a_n - L}{\sqrt{a_n} + \sqrt{L}} \right|.$$

Furthermore, since $a_n \rightarrow L$ and $L > 0$, clearly $0 < L/2 < L$ and there will be some $\epsilon < L/2$ such that,

$$L - L/2 < L - \epsilon < a_n < L + \epsilon < L + L/2 \iff |a_n - L| < \epsilon < L/2.$$

This means that, inside this ϵ -neighbourhood of L , we can find a constant lower bound on $\sqrt{a_n} + \sqrt{L}$ as,

$$\sqrt{a_n} + \sqrt{L} > \sqrt{L/2} + \sqrt{L} = C$$

for constant C . Now we have

$$|\sqrt{a_n} - \sqrt{L}| = \left| \frac{a_n - L}{\sqrt{a_n} + \sqrt{L}} \right| < \frac{|a_n - L|}{C} < \frac{\epsilon'}{C}$$

so we can take $\epsilon = \epsilon'/C$ to obtain $|a_n - L| < \epsilon$ as required. \square

Proof. This is an alternative proof of the $L > 0$ case using the fact that

$$\sqrt{x+y} \leq \sqrt{x} + \sqrt{y}.$$

Choose ϵ such that $0 < \epsilon \leq \sqrt{L}$ and take N so that, for $n > N$, $|a_n - L| < \epsilon^2$, or in other words $L - \epsilon^2 < a_n < L + \epsilon^2$. Then we have

$$\sqrt{L} - \epsilon \leq \sqrt{L - \epsilon^2} < \sqrt{a_n} < \sqrt{L + \epsilon^2} \leq \sqrt{L} + \epsilon$$

which places $\sqrt{a_n}$ in ϵ -neighbourhood of \sqrt{L} , so we're done.

Note that we need $\epsilon \leq \sqrt{L}$ for $\sqrt{L} - \epsilon \leq \sqrt{L - \epsilon^2}$ and also – so long as we are doing *real* analysis – this proof is only for the $L > 0$ case because, when $L = 0$, $\sqrt{L - \epsilon^2}$ will be a complex number. \square

- (53) **Let a_n be a positive decreasing sequence. Show that, if there exist numbers N and α such that**

$$0 < \frac{a_{n+1}}{a_n} < \alpha < 1 \quad \forall n > N$$

then $a_n \rightarrow 0$ as $n \rightarrow \infty$. But if, on the other hand, we have

$$0 < \frac{a_{n+1}}{a_n} < 1 \quad \forall n > N$$

then we cannot conclude that $a_n \rightarrow 0$.

The basic principle here is that, if a convergent sequence converges to a non-zero limit, then the ratio of consecutive terms must tend to 1. If the ratio of consecutive terms remains below 1 then the sequence must go to zero.

If $\frac{a_{n+1}}{a_n} < \alpha$ then α is an upper bound on the ratio of consecutive terms so we can deduce that,

$$a_{n+1} < \alpha a_n < a_n \quad \text{since } \alpha < 1$$

and that, if we let a_N be the value of the sequence at $n = N$ and consider the sequence for $n > N$,

$$a_n < \alpha^{n-N} a_N$$

which is of the form,

$$a_n < \alpha^n C$$

for constant C . So we have bound the values of the sequence below $\alpha^n C$ which clearly goes to 0 as $n \rightarrow \infty$. Since we are told that the sequence is positive so that a_n is also bounded below by 0, we can conclude, by Sandwich Theorem, that $a_n \rightarrow 0$ as $n \rightarrow \infty$. \square

An alternative way of showing the same thing is to use the algebra of limits and the fact that, in a convergent sequence, any subsequence must also converge to the same limit (Theorem 27) to deduce that if $a_n \rightarrow L$ then we must also have $a_{n+1} \rightarrow L$. Then, if $L \neq 0$ we can apply the algebra of limits to obtain,

$$\lim_{n \rightarrow \infty} \frac{a_{n+1}}{a_n} = \frac{L}{L} = 1$$

which contradicts $\frac{a_{n+1}}{a_n} < \alpha < 1$. Therefore, $L = 0$. \square

On the other hand, if $\frac{a_{n+1}}{a_n} \rightarrow 1$ as $n \rightarrow \infty$ then the sequence may converge to 0 or to some other value. For example:

(i) $a_n = \frac{1}{n} + 1$

$$\frac{a_{n+1}}{a_n} = \frac{n(n+2)}{(n+1)^2} = \frac{n^2 + 2n}{n^2 + 2n + 1} \rightarrow 1 \text{ as } n \rightarrow \infty$$

and clearly,

$$\lim_{n \rightarrow \infty} \frac{1}{n} + 1 = 1.$$

But also,

(ii) $a_n = \frac{1}{n}$

$$\frac{a_{n+1}}{a_n} = \frac{n}{n+1} \rightarrow 1 \text{ as } n \rightarrow \infty$$

even though

$$\lim_{n \rightarrow \infty} \frac{1}{n} = 0.$$

3.2.6 Limits of functions

3.2.6.1 Definition of the limit of a function

Definition. Let $f : \mathbb{R} \mapsto \mathbb{R}$ be a function. We say that L is the **limit of $f(x)$ as x approaches a** if, for each $\epsilon > 0$ there exists $\delta > 0$ such that,

$$0 < |x - a| < \delta \implies |f(x) - L| < \epsilon.$$

Definition. Let $f : \mathbb{R} \mapsto \mathbb{R}$ be a function. We say that **$f(x)$ tends to infinity as x approaches a** if, for each K there exists $\delta > 0$ such that,

$$0 < |x - a| < \delta \implies f(x) > K.$$

3.2.6.2 Examples of limits of functions

(54) Prove that if $f : \mathbb{R} \mapsto \mathbb{R}$ s.t. $f(x) = x^2 + x$ then $\lim_{x \rightarrow 2} f(x) = 6$.

Let $0 < |x - 2| < \delta$ and consider some arbitrary $\epsilon > 0$. Then we have,

$$|(x^2 + x) - 6| = |(x - 2)(x + 3)| \leq |x - 2| |x + 3| < \delta |x + 3|.$$

It's tempting at this point to say that, since we are examining the behaviour when x approaches 2 so we can assume $x \approx 2 \iff (x + 3) \approx 5$ and then we can set $\delta = \frac{\epsilon}{6}$ so that,

$$0 < |x - 2| < \delta = \frac{\epsilon}{6} \implies |f(x) - 6| < \frac{|x + 3|}{6} \epsilon < \epsilon.$$

However, there is a subtle logical problem here: We have considered any arbitrary $\epsilon > 0$, meaning that ϵ could be large. Then, when we set $\delta = \frac{\epsilon}{6}$ we linked the value of δ to that of ϵ so that δ may also be arbitrarily large. But $0 < |x - 2| < \delta$ so that $|x - 2|$ may be arbitrarily large also. This contradicts the assumption that $x \approx 2$ and so the argument we have here is only valid for small ϵ .

So, we need an alternative approach. Consider that,

$$|x + 3| = |x - 2 + 2 + 3| \leq |x - 2| + |2 + 3| = |x - 2| + 5 < \delta + 5.$$

This means that,

$$\left| (x^2 + x) - 6 \right| < \delta |x + 3| < \delta(\delta + 5).$$

But note, do **not** start trying to solve a quadratic. There is a much better way.

Now - here comes the clever bit - if we set $\delta = \min\{1, \frac{\epsilon}{6}\}$ then *both* 1 and $\frac{\epsilon}{6}$ are an upper bound on the value of δ so that we can say,

$$\left| (x^2 + x) - 6 \right| < \delta(\delta + 5) \leq 6\delta \leq 6\frac{\epsilon}{6} = \epsilon.$$

3.2.6.3 Algebra of limits of functions

Theorem 31. Let $f, g : \mathbb{R} \mapsto \mathbb{R}$ be two functions and c be any real number. Suppose that $\lim_{x \rightarrow a} f(x) = L$ and $\lim_{x \rightarrow a} g(x) = M$. Then,

1. $\lim_{x \rightarrow a} (cf)(x) = cL$
2. $\lim_{x \rightarrow a} (|f|)(x) = |L|$
3. $\lim_{x \rightarrow a} (f + g)(x) = L + M$
4. $\lim_{x \rightarrow a} (f - g)(x) = L - M$
5. $\lim_{x \rightarrow a} f(x)g(x) = LM$
6. $\lim_{x \rightarrow a} (f/g)(x) = L/M$ provided $g(x) \neq 0$ for any x in the neighbourhood of a .

Proof. see: Lamar University - Paul's Online Notes - Proofs of limit properties □

3.2.6.4 One-sided limits

Definition. Let $f : \mathbb{R} \rightarrow \mathbb{R}$ be a function. We say that L is the **limit of $f(x)$ as x approaches a from the left (or from below)**, denoted by $\lim_{x \rightarrow a^-} f(x) = L$, if for each $\epsilon > 0$, there exists $\delta > 0$ such that,

$$a - \delta < x < a \implies |f(x) - L| < \epsilon.$$

3.2.6.5 Limits at infinity

Definition. Let $f : \mathbb{R} \rightarrow \mathbb{R}$ be a function. We say that L is the **limit of $f(x)$ as x approaches ∞** , denoted by $\lim_{x \rightarrow \infty} f(x) = L$, if for each $\epsilon > 0$, there exists $M > 0$ such that,

$$x \geq M \implies |f(x) - L| < \epsilon.$$

3.2.7 Continuity

Definition. A function f is **continuous at a point a** if

- $f(a)$ is defined,
- $\lim_{x \rightarrow a} f(x) = f(a)$.

More formally, the definition of $\lim_{x \rightarrow a} f(x) = L$ is,

$$\forall \epsilon > 0 . \exists \delta \text{ s.t. } 0 < |x - a| < \delta \implies |f(x) - L| < \epsilon.$$

But if $f(a)$ is defined, then when $|x - a| = 0$ (i.e. when $x = a$) we have

$$f(x) = f(a) \implies |f(x) - f(a)| = 0 < \epsilon.$$

Therefore, if $f(a)$ is defined and $\lim_{x \rightarrow a} f(x) = f(a)$, then

$$\forall \epsilon > 0 . \exists \delta > 0 \text{ s.t. } |x - a| < \delta \implies |f(x) - f(a)| < \epsilon.$$

Definition. A function is **continuous** if it is **continuous at every point**.

Definition. A function is **continuous on the closed interval $[a, b]$** if it is

- continuous at every point in the open interval (a, b) ,
- $\lim_{x \rightarrow a^+} f(x) = f(a)$,
- $\lim_{x \rightarrow b^-} f(x) = f(b)$.

Definition. A function is **left continuous** or **continuous on the left** at a if,

$$\lim_{x \rightarrow a-} f(x) = f(a).$$

Obviously, from the other side, a function can be **right continuous**.

Theorem 32. Let $f, g : \mathbb{R} \rightarrow \mathbb{R}$ be functions that are continuous at $a \in \mathbb{R}$ and c be any real number. Then $|f|, (cf), (f - g), (f + g), (f(x)g(x))$ are all continuous at a , and (f/g) is continuous provided $g(x) \neq 0$ for any x in some neighbourhood of a .

Proof. This follows from the algebra of limits of functions given in Theorem 30. \square

Corollary 34. It follows then that every polynomial is continuous. This can be seen as the most simple polynomial is a constant - which is clearly continuous; also $f(x) = x$ is clearly continuous. Then powers of x are continuous as they are products of continuous functions and when multiplied by coefficients this is a constant multiplying a continuous function so the resultant function is continuous. Then any polynomial is a summation of such terms so the result is continuous as the sum of continuous functions is continuous.

Theorem 33. If g is a function which is continuous at a , and f is a function which is continuous at $g(a)$. Then $(f \circ g)$ is continuous at a .

Proof. Continuity of f at $g(a)$ guarantees that for any $\epsilon > 0$ there exists some $\delta' > 0$ such that $|x' - g(a)| < \delta' \implies |f(x') - f(g(a))| < \epsilon$ and continuity of g at a guarantees that for any $\delta' > 0$ there exists some $\delta > 0$ such that $|x - a| < \delta \implies |g(x) - g(a)| = |x' - g(a)| < \delta'$. \square

Corollary 35. $\lim_{x \rightarrow a} (f \circ g)(x) = \lim_{x \rightarrow a} f(g(x)) = f(\lim_{x \rightarrow a} g(x))$

3.2.7.1 Continuity of functions over sequences

We now give an alternative definition of continuity which ties in the concept of limits for sequences.

Theorem 34. A function f is continuous at a if and only if for each sequence (x_n) such that $\lim_{n \rightarrow \infty} x_n = a$ we have $\lim_{n \rightarrow \infty} f(x_n) = f(a)$.

Before the proof, an important point to note is that this theorem applies to "each sequence" with the described limit. This is important as it is possible to find an individual sequence such that the inference is not valid. For example, the constant sequence $\forall n \in \mathbb{N} . x_n = a$ clearly tends to a as $n \rightarrow \infty$ but this would not imply continuity of f as the limit of $f(x_n)$ for such a sequence would amount to saying that $f(a) = f(a)$. The definition of continuity is assertion of the equality of the limit of f over values in the neighbourhood of a (but not at a itself) with the value of f at a . So, continuity is implied by the fact that the limit of $f(x_n)$ for the constant sequence $x_n = a$ is equal to the limit of $f(x_n)$ for all other sequences x_n whose limit is a . Another way of looking at it is that f is continuous at a because the limit there equals $f(a)$ however the argument converges to a .

Proof. Breaking it down into two propositions we have,

$$(\forall x_n \text{ s.t. } \lim_{n \rightarrow \infty} x_n = a) \quad \lim_{n \rightarrow \infty} f(x_n) = f(a) \quad (P_1)$$

$$\forall \epsilon > 0 . \exists \delta > 0 . |x - a| < \delta \implies |f(x) - f(a)| < \epsilon \quad (P_2)$$

and we need to show that $P_1 \iff P_2$.

So, to begin with we'll assume show that $P_1 \implies P_2$.

Unpacking P_1 we have a function f such that

$$\forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . |f(x_n) - f(a)| < \epsilon \quad (1)$$

where x_n is a sequence such that

$$\forall \delta > 0 . \exists N' . \forall n > N' \in \mathbb{N} . |x_n - a| < \delta. \quad (2)$$

Then, if we choose a particular $\epsilon > 0$, there exists some N such that for all $n > N$ we have $|f(x_n) - f(a)| < \epsilon$. Now there will also be some $N' \leq N$ so that $\forall n > N, n > N'$ and so there is some δ such that $|x_n - a| < \delta$. Since P_1 says that this is the case *whenever* we have an x_n such as this, P_1 may be rewritten as

$$\forall \epsilon > 0 . (\exists N . \forall n > N \in \mathbb{N}) . \exists \delta > 0 . |x_n - a| < \delta \implies |f(x_n) - f(a)| < \epsilon.$$

Now, if we remove reference to the number of the term of the sequences - N, n - we have

$$\forall \epsilon > 0 . \exists \delta > 0 . |x - a| < \delta \implies |f(x) - f(a)| < \epsilon$$

which is the statement of continuity of f in P_2 .

Next we prove $P_2 \implies P_1$.

So now we begin by assuming P_2 which was that,

$$\forall \epsilon > 0 . \exists \delta > 0 . |x - a| < \delta \implies |f(x) - f(a)| < \epsilon.$$

Actually, it is quite easy to apply a similar logic as previously but in reverse to make the converse implication. We can choose any ϵ and there exists some δ such that P_2 holds. Then, as we have seen previously in (2), P_1 tells us that, for this value of δ ,

$$\exists N' . \forall n > N' \in \mathbb{N} . |x_n - a| < \delta$$

and P_2 tells us that,

$$|x_n - a| < \delta \implies |f(x_n) - f(a)| < \epsilon.$$

Putting the two together we get,

$$\forall \epsilon > 0 . \exists \delta > 0 . \exists N' . \forall n > N' \in \mathbb{N} . |x_n - a| < \delta \implies |f(x_n) - f(a)| < \epsilon$$

which implies that

$$\forall \epsilon > 0 . \exists N . \forall n > N \in \mathbb{N} . |f(x_n) - f(a)| < \epsilon.$$

So we have shown that P_2 implies that (2) implies (1) which is $P_2 \implies P_1$ as required.

For completeness, let's consider another way of proving that $P_1 \implies P_2$ using a proof by contradiction.

So, we are assuming P_1 but also assuming, for contradiction, that P_2 is false. Then we are negating the statement,

$$\forall \epsilon > 0 . \exists \delta > 0 . |x - a| < \delta \implies |f(x) - f(a)| < \epsilon$$

so we are asserting that,

$$\forall \epsilon > 0 . \nexists \delta > 0 . |x - a| < \delta \implies |f(x) - f(a)| < \epsilon$$

which is equivalent to

$$\forall \epsilon > 0 . \forall \delta > 0 . |x - a| < \delta \not\Rightarrow |f(x) - f(a)| < \epsilon$$

or alternatively,

$$\forall \epsilon > 0 . \forall \delta > 0 . \exists x . (|x - a| < \delta) \wedge (|f(x) - f(a)| \geq \epsilon).$$

So, this says that we can choose any arbitrary $\epsilon > 0$ and for all $\delta > 0$ there will be some x in the δ -neighbourhood of a such that $|f(x) - f(a)| \geq \epsilon$.

Now, if we choose a value of δ that depends on a natural number n in such a way that $\delta \rightarrow 0$ as $n \rightarrow \infty$ – for example, if $\delta = 1/n$ – then the δ -neighbourhoods around a will get smaller as $n \rightarrow \infty$. Then we can select an x from the δ -neighbourhood that corresponds to a particular value of n and call it x_n and, in this way, we create a sequence x_n that converges to a . So we have $\lim_{n \rightarrow \infty} x_n = a$.

But now, for every δ there is an x_n in the δ -neighbourhood of a with $|f(x_n) - f(a)| \geq \epsilon$ and, if we choose this value for x_n , we have a constructed a sequence that converges to a but

$$\lim_{n \rightarrow \infty} f(x_n) \neq f(a)$$

which contradicts hypothesis P_1 . □

Corollary 36. $\lim_{n \rightarrow \infty} f(x_n) = f(\lim_{n \rightarrow \infty} x_n)$

3.2.7.2 Continuous Functions on Closed Intervals

Definition. For a subset X of the domain of a function f , we say that f is bounded on X if there exists M such that $|f(x)| \leq M$ for each $x \in X$.

Definition. We define the supremum (or maximum) of f on X as $\sup \{ f(x) \mid x \in X \}$ (or $\max \{ f(x) \mid x \in X \}$ if it exists).

3.2.7.3 Examples of bounded and unbounded functions

- (55) The indicator function for rational numbers within the reals, known as the **Dirichlet Function**,

$$f(x) = \begin{cases} 0 & x \text{ is irrational} \\ 1 & x \text{ is rational} \end{cases}.$$

This function is nowhere continuous because between every two irrational numbers there is a rational number (and vice-versa) so $f(x)$ is flipping between 0 and 1 in every neighbourhood of every point however small a neighbourhood we consider. So this function never converges anywhere but *is* bounded because it only ever takes values of 1 or 0 so, clearly, its maximum is 1 and minimum is 0.

- (56) The function $f(x) = 1/x$ over the interval $(0, 1]$ is continuous at every point but unbounded as it goes to infinity as $x \rightarrow 0$. If we consider the same function over the closed interval $[0, 1]$ then we no longer have continuity over this interval as there is a singularity at $x = 0$.

3.2.7.4 Extreme Value Theorem

Theorem 35. *Let f be continuous on $[a, b]$. Then f is bounded on $[a, b]$ and it achieves its maximum; that's to say, the supremum is equal to the maximum.*

Note that, even if f is defined at every point in $[a, b]$, if it is not continuous then it may not be bounded. There do exist functions that are not continuous but bounded (for example the Dirichlet Function 55) but there also exist functions that are not continuous and unbounded such as the reciprocal function 56. So, functions that are not continuous on a closed interval may be bounded or not; and functions that are continuous on an open interval also may or may not be bounded (again, the reciprocal function is an example of a function that is continuous on an open interval but not bounded); but here we will show that functions that are continuous on a closed interval must be bounded on that interval.

Proof. It may be tempting to begin trying to prove this by reasoning as follows.

That f is continuous on $[a, b]$ means that, for any $c \in (a, b)$,

$$\forall \epsilon > 0 . \exists \delta > 0 . |x - c| < \delta \implies |f(x) - f(c)| < \epsilon.$$

This means that the value of $f(x)$ must be finite everywhere in (a, b) as, choosing any fixed point c in the open interval, $|x - c|$ is finite and, therefore, less than some δ thus implying that $|f(x) - f(c)| < \epsilon$ for some finite $\epsilon \dots$

However this is **dead wrong!** If we take the example of the reciprocal function (56) over the interval $(0, 1]$: If we take x -values approaching 0, the value of $f(x)$ grows unbounded. For any given x -value it will be less than some finite ϵ but we can always find another x -value with a greater value of $f(x)$. So, there is no maximum ϵ and so, also no maximum value of $f(x)$ on the interval.

Another tempting way to prove this is as follows.

A closed interval is an interval such that every sequence of values in the interval converges to a point in the interval.

That f is continuous on $[a, b]$ implies that for every sequence, x_n , of values in $[a, b]$ that converges to some point in $c \in [a, b]$, $\lim_{n \rightarrow \infty} f(x_n) = f(c)$. Furthermore, that the interval is closed implies that every sequence of values in the interval converges to a point in the interval. Therefore, we can conclude that $f(x)$ at every point in $[a, b]$ exists and is finite so that f is bounded and obtains a maximum on the interval.

There are two issues with this:

1. The statement about the nature of a closed interval that the proof relies upon has not been proven.
2. That $f(x)$ is defined and finite for all $x \in [a, b]$ is taken as proof that the function is bounded and obtains a maximum in the interval.

To deal with issue (1) – if we’re not going to prove the proposition about closed intervals as a pre-requisite of the proof – we need to develop a proof that doesn’t rely on this characteristic of closed intervals. To deal with (2) meanwhile, we need to explicitly prove boundedness and that f obtains a

maximum.

The proof given in LSE Abstract Mathematics course material follows.

Suppose first that f is unbounded above. For each $n \in \mathbb{N}$, let x_n be a point in $[a, b]$ such that $f(x_n) > n$. The sequence (x_n) is bounded, so has a convergent subsequence (x_{k_n}) , tending to some limit c (by Theorem 10.11). Necessarily $c \in [a, b]$. Since f is continuous at c , $f(x_{k_n}) \rightarrow f(c)$ as $n \rightarrow \infty$. But this contradicts the construction of the sequence (x_n) , since $f(x_{k_n}) > n \rightarrow \infty$. So f is bounded above. Let $M = \sup\{f(x) \mid x \in [a, b]\}$. For each $n \in \mathbb{N}$, let x_n be a point in $[a, b]$ such that $f(x_n) > M - \frac{1}{n}$. Again take a convergent subsequence (x_{k_n}) of (x_n) , tending to some limit $c \in [a, b]$. Arguing as before, we see $f(c) = M$.

This proof says: Assume that f is unbounded on the interval. Then we can construct a sequence of x -values such that, for the n th value x_n , $f(x_n) > n$. This is possible because, if f is unbounded, for any value of n , there is some subinterval of x -values such that for each of them $f(x) > n$ and any interval of the real numbers contains an infinite number of real numbers and so, a sequence x_n with $n \rightarrow \infty$. Note that, at this point, x_n is an arbitrary sequence which is not necessarily convergent (it could bounce around the interval).

Then, we notice that this sequence x_n is necessarily bounded (as it is a subinterval of $[a, b]$) and so we invoke the Bolzano-Weierstrass Theorem, Theorem 29 (called Theorem 10.11 in the quoted proof), to deduce that it has a convergent subsequence which we call x_{k_n} and call its limit c .

At this point we can use the continuity of f on the interval to deduce that $f(x_{k_n}) \rightarrow f(c)$ as $n \rightarrow \infty$ by which we obtain a contradiction to the condition we set on the values of x_{k_n} when we constructed the sequence – namely that $f(x_{k_n}) > n$. Notice that this is constructing a sequence x_{k_n} such that $f(x_{k_n})$ grows without bound as $n \rightarrow \infty$ and then saying, "but the limit of x_{k_n} as $n \rightarrow \infty$ is c which is inside the interval and so (by continuity) the limit of $f(x_{k_n})$ as $n \rightarrow \infty$ is $f(c)$ – a fixed finite value". This is the point where the fact that the interval $[a, b]$ is closed comes into play – if the interval were not closed it would be possible that c was not inside the interval and then we would not be able to invoke continuity to assert that the limit of $f(x_{k_n})$ over this sequence was finite.

So, now we have shown that if a function is continuous over a closed interval then assuming that the function is unbounded produces a contradiction and,

therefore, we can conclude that it is, in fact, bounded.

The last thing that needs to be proven is that f obtains its maximum in the interval. Having shown that the function is bounded on the interval we now know that there exists

$$M = \sup\{ f(x) \mid x \in [a, b] \}$$

and we need to show that there is such an x -value in $[a, b]$ that $f(x) = M$. In an open interval this might not be the case as the supremum of the function on the interval might occur as the limit of f over a sequence of x -values converging to a point that lies outside the interval. So, in this case, we construct a convergent sequence such that $f(x_{k_n}) \rightarrow M$ as $n \rightarrow \infty$ by selecting x_n such that $f(x_n) = M - \frac{1}{n}$. This is possible because the definition of the supremum says that, because it is the *lowest* upper bound,

$$\forall \epsilon > 0 . \exists f(x_n) . f(x_n) > M - \epsilon$$

and so we are letting $\epsilon = \frac{1}{n}$. Then, as previously, we take a convergent subsequence of this sequence and name it x_{k_n} . So, as before, we have a sequence converging on some point, we'll call it $c \in [a, b]$, at which f is continuous so that $f(x_{k_n}) \rightarrow M$ as $n \rightarrow \infty$ implies that $M = f(c)$. This, in turn, means that f obtains a maximum in the interval. \square

3.2.7.5 Intermediate Value Theorem

Theorem 36. *Let f be continuous on $[a, b]$ with $f(a) < f(b)$. Then, for all K s.t. $f(a) < K < f(b)$, there exists some $c \in (a, b)$ with $f(c) = K$.*

*Note that this theorem is **not** written for an interval such that $f(a) \leq f(b)$ because if $f(a) = f(b)$ then the only K s.t. $f(a) \leq K \leq f(b)$ is $f(a) = K = f(b)$. But now it is **not** true to say that there exists some $c \in (a, b)$ with $f(c) = K$ as there is no reason why the value of $f(x)$ at the bounds of the interval should be repeated somewhere in the interior of the interval.*

We begin by proving a special case from which the general proof will follow.

Lemma 2. *Let f be continuous on $[a, b]$ with $f(a) < 0 < f(b)$. Then there exists some $c \in (a, b)$ with $f(c) = 0$.*

Proof. A first attempt at this proof is given below.

Continuity at the interval bounds a and b means that, for some $\epsilon_1, \epsilon_2 > 0$,

$$\exists \delta_1 . 0 \leq x - a < \delta_1 \implies |f(x) - f(a)| < \epsilon_1,$$

$$\exists \delta_2 . 0 \leq b - x < \delta_2 \implies |f(x) - f(b)| < \epsilon_2.$$

So, we have a lower neighbourhood around $f(a)$ and an upper neighbourhood around $f(b)$ as follows,

$$f(a) - \epsilon_1 < f(x) < f(a) + \epsilon_1,$$

$$f(b) - \epsilon_2 < f(x) < f(b) + \epsilon_2$$

If $\epsilon_1 > |f(a)|$ and $\epsilon_2 > |f(b)|$ then both neighbourhoods contain $f(x) = 0$. Therefore, there is some interval of x such that $f(x)$ lies inside both the lower and upper neighbourhood – in the overlap of the two. In the overlap we have,

$$f(b) - \epsilon_2 < f(x) < f(a) + \epsilon_1.$$

Now if we let ϵ_1 vary freely but make ϵ_2 a function of ϵ_1 like so,

$$\epsilon_2 = f(a) + f(b) + \epsilon_1$$

then we still have $\epsilon_1 > |f(a)|$ and $\epsilon_2 > |f(b)|$ as,

$$\epsilon_1 > |f(a)| \iff -\epsilon_1 < f(a) < \epsilon_1 \iff 0 < f(a) + \epsilon_1 < 2\epsilon_1$$

so $\epsilon_2 > f(b)$ and, conversely, if we assume that $\epsilon_2 > |f(b)|$ then,

$$\epsilon_2 > |f(b)| \iff -\epsilon_2 < f(b) < \epsilon_2 \iff -\epsilon_2 - f(b) < 0 < \epsilon_2 - f(b)$$

$$\epsilon_1 = \epsilon_2 - f(a) - f(b) > -f(a) = |f(a)| \quad \text{since } f(a) < 0.$$

Therefore,

$$\epsilon_2 = f(a) + f(b) + \epsilon_1 \implies [\epsilon_1 > |f(a)| \iff \epsilon_2 > |f(b)|].$$

Now we have,

$$f(b) - \epsilon_2 = -f(a) - \epsilon_1 < f(x) < f(a) + \epsilon_1$$

which is equivalent to

$$|f(x)| < f(a) + \epsilon_1 = \epsilon_3.$$

Now, since $f(a) + \epsilon_1 > 0$ we also have $\epsilon_3 > 0$ and it can become arbitrarily small by making $|f(a)| - \epsilon_1$ arbitrarily small. So we have shown that continuity over the closed interval and $f(a) < f(b)$ imply that we can find subintervals of $[a, b]$ such that $f(x)$ is constricted to ever-decreasing neighbourhoods of 0. In other words, for some c s.t. $a < c < b$, $f(x) \rightarrow 0$ as $x \rightarrow c$ and since f is continuous on the interval this implies that $f(c) = 0$ also.

This is not bad but suffers from vagueness in a couple of areas: the overlap of the lower and upper neighbourhoods probably needs to be more precisely defined and, certainly, the final part of the proof stating that confining $f(x)$ to ever-decreasing neighbourhoods of 0 implies that there is some c such that $f(c) = 0$ needs to be drawn much more explicitly.

Here is the proof given in LSE Abstract Mathematics.

We construct a sequence of intervals $[a_n, b_n]$ such that

1. $f(a_n) < 0$, $f(b_n) > 0$ for each n
2. $[a_{n+1}, b_{n+1}] \subseteq [a_n, b_n]$ for each n .

We start by letting $[a_1, b_1] = [a, b]$. Then for each $n \geq 1$, we define $[a_{n+1}, b_{n+1}]$ as follows.

Let $c_n = (a_n + b_n)/2$, be the midpoint of the previous interval. If $f(c_n) = 0$, then the theorem is proved and so we need not continue constructing intervals!

Otherwise, if $f(c_n) < 0$, we define $a_{n+1} = c_n$ and $b_{n+1} = b_n$. And if $f(c_n) > 0$, we define $b_{n+1} = c_n$ and $a_{n+1} = a_n$. Note that the condition 1. is satisfied by choosing our intervals in this manner. Moreover, note that the $(n + 1)$ st interval is half the size of the n th interval and so $b_{n+1} - a_{n+1} \leq (b_1 - a_1)/2^n$. It follows that

$$\lim_{n \rightarrow \infty} (b_n - a_n) = 0. \tag{3}$$

Finally, note that (a_n) is increasing and bounded above (by b_1) and so it has a limit; similarly (b_n) is decreasing and bounded below and so has a limit. Thus by (3) (and algebra of limits) these limits are equal to, say, c . Thus by continuity (using Theorem 33),

$$f(c) = \lim_{n \rightarrow \infty} f(b_n) \geq 0,$$

where the last inequality follows from the fact that each $f(b_n) \geq 0$ (in fact > 0). Similarly,

$$f(c) = \lim_{n \rightarrow \infty} f(a_n) \leq 0.$$

Thus $f(c)$ must be equal to zero, and the proof is complete. □

Clearly, the general proof of the Intermediate Value Theorem follows naturally from this because,

- If $f(a) > f(b)$ then we can consider the function $g(x) = -f(x)$ so that we have $g(a) < g(b)$,
- If we have K s.t. $f(a) < K < f(b)$ with $K \neq 0$ we can consider $g(x) = f(x) - K$ so that we have $g(a) < 0 < g(b)$.

So, the general problem posed in the Intermediate Value Theorem is reducible to the case we have proved.

Corollary 37. *Suppose that the real function f is continuous on the closed interval $[a, b]$ and that f maps $[a, b]$ into $[a, b]$. Then there is $c \in [a, b]$ with $f(c) = c$.*

Proof. Let $h(x) = f(x) - x$ so that $f(c) = c$ if and only if $h(c) = 0$. Then also we have,

$$a \leq f(x) \leq b \implies h(a) \geq 0, h(b) \leq 0.$$

But this means that, either one of $h(a)$ or $h(b)$ is equal to 0 or neither are. In the case that one of them is equal to 0 then we have found our c such that $f(c) = c$. Otherwise, if neither is equal to 0, then we must have $h(a) > 0$ and $h(b) < 0$. So, we may apply the Intermediate Value Theorem to conclude that there exists $c \in (a, b)$ such that $h(c) = 0$ which is to say $f(c) = c$. □

3.2.7.6 Examples of reasoning with the Intermediate Value Theorem

- (57) *Suppose the real function f is continuous, positive and unbounded on \mathbb{R} and that $\inf\{f(x) \mid x \in \mathbb{R}\} = 0$. Use the Intermediate Value Theorem to prove that the range of f is $(0, \infty)$.*

Let $y \in (0, 1)$. We show that there is some $c \in \mathbb{R}$ such that $f(c) = y$. This shows that the range is the whole of $(0, 1)$. (The fact that it is no larger follows from the given fact that f is positive.)

Now, $\inf f(\mathbb{R}) = \inf\{f(x) \mid x \in \mathbb{R}\} = 0$, so, since $y > 0$, there must be some $y_1 \in f(\mathbb{R})$ with $y_1 < y$. This means there is some $x_1 \in \mathbb{R}$ such that $y_1 = f(x_1) < y$.

Similarly, because f is unbounded, which means $f(\mathbb{R})$ is unbounded, there must be some $y_2 \in f(\mathbb{R})$ with $y_2 > y$ and there will be some $x_2 \in \mathbb{R}$ such that $y_2 = f(x_2) > y$.

Then y lies between $f(x_1)$ and $f(x_2)$ and, since f is continuous, the Intermediate Value Theorem shows that there is some c between x_1 and x_2 with $f(c) = y$.

- (58) *Suppose the real function g is continuous on \mathbb{R} and that g maps $[a, b]$ into $[d, e]$ and maps $[d, e]$ into $[a, b]$ where $a < b$, $d < e$. By considering the function*

$$k(x) = g(g(x)),$$

prove that there are $p, q \in \mathbb{R}$ such that

$$g(p) = q, \quad g(q) = p.$$

Hence show that there is $c \in \mathbb{R}$ such that $g(c) = c$.

The function k , being a composition of continuous functions, is continuous and also maps $[a, b]$ into $[a, b]$ so that we can use Corollary 37 to deduce that there exists $c \in [a, b]$ such that $k(c) = c$. If we let $p = c$ and $q = g(p)$ then we have,

$$k(p) = p \iff g(g(p)) = p \iff g(q) = p.$$

Now we can employ the same trick again by defining

$$h(x) = g(x) - x$$

and then we have

$$h(p) = g(p) - p = q - p, \quad h(q) = g(q) - q = p - q$$

so that $h(p) = -h(q)$ and, therefore, $h(x)$ changes sign between p and q . (Note that we don't know which of p and q is the lower end and upper end of the interval but we know that between the two values the function changes sign.) Then, applying the Intermediate Value Theorem we have some c between p and q such that,

$$h(c) = 0 \iff g(c) - c = 0 \iff g(c) = c.$$

3.2.8 Relationship Between Sequences and Functions

Let $f : \mathbb{R} \mapsto \mathbb{R}$, $f(x) = y$. Now imagine that we take regular intervals on the domain, say, interval 1. We name the x -values at the upper bound of these intervals x_1, x_2, \dots for $x = 1, 2, \dots$. Then, the corresponding y -values, $y = f(x_1), f(x_2), \dots$ can be named y_1, y_2, \dots . In this way, we have defined a sequence, $y_n = f(x_n)$ where $x_n \in \mathbb{N}$ (note that this is a different notation from that used before where a sequence $x_n = f(n)$ for $n \in \mathbb{N}$). Looking at the derivative of the function,

$$\frac{dy}{dx} \approx \frac{\Delta y}{\Delta x} = \frac{f(x_{n+1}) - f(x_n)}{1} = f(x_{n+1}) - f(x_n).$$

While the ratio of consecutive terms is,

$$\frac{y_{n+1}}{y_n} = \frac{y_n + \Delta y}{y_n} = \frac{f(x_{n+1})}{f(x_n)} = \frac{f(x_n) + (f(x_{n+1}) - f(x_n))}{f(x_n)}.$$

Note, also, that

$$\frac{y_n + \Delta y}{y_n} = 1 + \frac{\Delta y}{y_n} \approx 1 + \frac{dy/dx}{y} = 1 + \frac{d}{dx} \ln y$$

which is to say that $\frac{\Delta y}{y_n}$ is the discrete form of the log derivative. Clearly, when $\Delta y > 0$, we can put this in the form

$$\frac{y_{n+1}}{y_n} = 1 + \frac{\Delta y}{y_n} = 1 + \frac{\Delta y/y_n}{1} = \frac{1 + h}{1}.$$

If $\Delta y < 0$ however,

$$\frac{y_n + \Delta y}{y_n} = \frac{1}{y_n/(y_n + \Delta y)} = \frac{1}{\frac{y_n + \Delta y - \Delta y}{y_n + \Delta y}} = \frac{1}{1 + \frac{-\Delta y}{y_n + \Delta y}} = \frac{1}{1 + h}.$$

If $\frac{\Delta y}{y_n}$ does not go to 0 then the ratio of consecutive terms stays below 1 and the sequence converges to 0. If it does go to 0 then the ratio of consecutive terms goes to 1 and the sequence may converge to 0 or to some non-zero value. These cases appear (proof?) to be distinguishable by looking at how fast $\frac{\Delta y}{y_n}$ goes to 0. For example,

$$(i) \ a_n = \frac{1}{n} + 1$$

$$\lim_{n \rightarrow \infty} \frac{1}{n} + 1 = 1$$

$$\frac{\Delta a}{a_n} = \frac{-1}{(n+1)^2}$$

$$(ii) \ a_n = \frac{1}{n}$$

$$\lim_{n \rightarrow \infty} \frac{1}{n} = 0$$

$$\frac{\Delta a}{a_n} = \frac{-1}{n+1}$$

Maybe the fact that $\frac{\Delta y}{y_n}$ goes to 0 faster as $n \rightarrow \infty$ in the first case indicates that it converges before it gets to 0?